

Metadata of the book that will be visualized in SpringerLink

Publisher Name	Springer International Publishing	
Publisher Location	Cham	
Series ID	14353	
SeriesTitle	SpringerBriefs in Probability and Mathematical Statistics	
Book ID	505524_1_En	
Book Title	Concentration of Maxima and Fundamental Limits in High-Dimensional Testing and Inference	
Book DOI	10.1007/978-3-030-80964-5	
Copyright Holder Name	The Author(s), under exclusive license to Springer Nature Switzerland AG	
Copyright Year	2021	
Corresponding Author	Family Name	Gao
	Particle	
	Given Name	Zheng
	Suffix	
	Division	Department of Statistics
	Organization	University of Michigan–Ann Arbor
	Address	Ann Arbor, MI, USA
	Email	gaozheng@umich.edu
	ORCID	https://orcid.org/0000-0003-1394-8741
Corresponding Author	Family Name	Stoev
	Particle	
	Given Name	Stilian
	Suffix	
	Division	Department of Statistics
	Organization	University of Michigan–Ann Arbor
	Address	Ann Arbor, MI, USA
	Email	sstoev@umich.edu
	ORCID	https://orcid.org/0000-0003-2108-5421

¹ SpringerBriefs in Probability and Mathematical
² Statistics

³ **Series Editors**

- ⁴ Nina Gantert, Technische Universität München, Munich, Germany
⁵ Tailen Hsing, University of Michigan, Ann Arbor, MI, USA
⁶ Richard Nickl, University of Cambridge, Cambridge, UK
⁷ Sandrine Péché, Univirsité Paris Diderot, Paris, France
⁸ Yosef Rinott, Hebrew University of Jerusalem, Jerusalem, Israel
⁹ Almut E.D. Veraart, Imperial College London, London, UK

¹⁰ **Editor-in-Chief**

- ¹¹ Gesine Reinert, University of Oxford, Oxford, UK

12 SpringerBriefs present concise summaries of cutting-edge research and practical
13 applications across a wide spectrum of fields. Featuring compact volumes of 50 to
14 125 pages, the series covers a range of content from professional to academic.
15 Briefs are characterized by fast, global electronic dissemination, standard publishing
16 contracts, standardized manuscript preparation and formatting guidelines, and
17 expedited production schedules.

18 Typical topics might include:

- 19 • A timely report of state-of-the art techniques
- 20 • A bridge between new research results, as published in journal articles, and a
21 contextual literature review
- 22 • A snapshot of a hot or emerging topic
- 23 • Lecture of seminar notes making a specialist topic accessible for non-specialist
24 readers
- 25 • SpringerBriefs in Probability and Mathematical Statistics showcase topics of
26 current relevance in the field of probability and mathematical statistics

27 Manuscripts presenting new results in a classical field, new field, or an emerging
28 topic, or bridges between new results and already published works, are encouraged.
29 This series is intended for mathematicians and other scientists with interest in
30 probability and mathematical statistics. All volumes published in this series
31 undergo a thorough refereeing process.

32 The SBPMS series is published under the auspices of the Bernoulli Society for
33 Mathematical Statistics and Probability.

34 All titles in this series are peer-reviewed to the usual standards of mathematics
35 and its applications.

36 More information about this series at <http://www.springer.com/series/14353>



³⁷ Zheng Gao · Stilian Stoev

³⁸ **Concentration of Maxima
39 and Fundamental Limits
40 in High-Dimensional Testing
41 and Inference**



Bernoulli Society
for Mathematical Statistics
and Probability

 Springer



505524_1_En_BookFrontmatter TYPESET DISK LE CP Disp.: 13/8/2021 Pages: xiv Layout: T1-Standard

43 Zheng Gao 
Department of Statistics
University of Michigan–Ann Arbor
Ann Arbor, MI, USA

Stilian Stoev 
Department of Statistics
University of Michigan–Ann Arbor
Ann Arbor, MI, USA

44 ISSN 2365-4333 ISSN 2365-4341 (electronic)
SpringerBriefs in Probability and Mathematical Statistics
ISBN 978-3-030-80963-8 ISBN 978-3-030-80964-5 (eBook)
<https://doi.org/10.1007/978-3-030-80964-5>

45 © The Author(s), under exclusive license to Springer Nature Switzerland AG 2021
46 This work is subject to copyright. All rights are solely and exclusively licensed by the Publisher, whether
47 the whole or part of the material is concerned, specifically the rights of translation, reprinting, reuse
48 of illustrations, recitation, broadcasting, reproduction on microfilms or in any other physical way, and
49 transmission or information storage and retrieval, electronic adaptation, computer software, or by similar
50 or dissimilar methodology now known or hereafter developed.
51 The use of general descriptive names, registered names, trademarks, service marks, etc. in this publication
52 does not imply, even in the absence of a specific statement, that such names are exempt from the relevant
53 protective laws and regulations and therefore free for general use.
54 The publisher, the authors and the editors are safe to assume that the advice and information in this book
55 are believed to be true and accurate at the date of publication. Neither the publisher nor the authors or
56 the editors give a warranty, expressed or implied, with respect to the material contained herein or for any
57 errors or omissions that may have been made. The publisher remains neutral with regard to jurisdictional
58 claims in published maps and institutional affiliations.

59 This Springer imprint is published by the registered company Springer Nature Switzerland AG
60 The registered company address is: Gewerbestrasse 11, 6330 Cham, Switzerland

1 *To our teachers and families with deepest
2 gratitude*

0

Preface

This text presents a collection of new results and recent developments on the phase-transition phenomena in sparse signal problems. The main theme is the study of the fundamental limits in high-dimensional testing and inference. Since the seminal works of Ingster (1998) and Donoho and Jin (2004), the subject has received a lot of attention in the literature with important contributions from Ji and Jin (2012); Genovese et al. (2012); Jin et al. (2014); Arias-Castro and Chen (2017); Butucea et al. (2018). These works, among many others, have discovered some fundamental limits in the so-called *needle in haystack* problems, where a sparse signal is observed with high-dimensional additive noise. In this setting, two archetypal problems arise—the *signal detection* and *signal support recovery*. The signal detection refers to a global hypothesis testing problem that amounts to determining the presence of non-zero signal in any of its dimensions. The support recovery, on the other hand, can be seen either as a multiple testing problem where the presence of non-zero signal is tested for each signal location of interest, or alternatively, as an inference problem that aims to estimate the signal support, i.e., the locations of the non-zero signal components. The fundamental limits of these problems are studied in the so-called high-dimensional asymptotic regime where the dimension p of the underlying signal grows to infinity, and the sample size n is either bounded or grows slowly relative to p .

From a probabilistic perspective, these aforementioned fundamental limits are stated as asymptotic zero-one type laws, as dimensionality diverges. Namely, consider a sparse signal with *support size* on the order of $p^{1-\beta}$ for some parameter $\beta \in (0, 1)$. Parameterize the non-zero *signal amplitude* by $\Delta(p^r)$, for some $r > 0$ and a suitable monotone non-decreasing function $\Delta(\cdot)$. Then, for a broad range of error distributions and statistical problems, one encounters a sharp transition between the regimes where the problem is solvable and unsolvable depending on the signal magnitude r and signal sparsity β . More precisely, there exists a boundary function $f(\beta)$ such that if the signal magnitudes are *above* the boundary, $r > f(\beta)$, then the problem can be solved with vanishing loss as $p \rightarrow \infty$, with a suitable statistical procedure. On the other hand, if the signal is below that same boundary, i.e., $r < f(\beta)$, all statistical procedures fail to provide a solution with a vanishing loss, as $p \rightarrow \infty$. Of course, depending on whether one considers the detection (testing) or support recovery (inference) problems, different loss functions quantify success and

failure. The choice of the loss functions is often guided by the applications, resulting in a rich picture of phase transitions (see, e.g., Fig. 3.2).

The contributions of this work. The fundamental limits of the classic detection problem hinge of the analysis of the discrepancy between the *null* and *alternative* hypotheses, e.g., via Hellinger distance. Thus, perhaps for technical reasons, much of the analysis in the existing literature has been done under the assumption that the additive errors are independent and/or Gaussian, or using loss functions unaffected by the dependence such as the Hamming loss. In this work, we demonstrate that the support recovery problems, especially *exact support recovery*, are best understood from the novel perspective of the *concentration of maxima* phenomenon in extreme value theory. It turns out that under a very broad range of light-tailed error distributions and under a *very* broad range of error-dependence structures, the maxima of the errors, when rescaled (but not centered!) converge in probability to a positive constant. This concentration property leads to a complete solution of the exact support recovery problem for the broad family of thresholding procedures. Most, if not all, existing support estimation procedures are **type thresholding** procedures (see Sect. 2.2). That is, the signal support estimate comprises all components exceeding a suitable (potentially data-dependent) threshold. We show, by exploiting concentration of maxima, that thresholding procedures obey a phase transition, where if the signal is above a certain boundary, asymptotically exact recovery is possible while below the boundary all thresholding procedures fail, as $p \rightarrow \infty$. Remarkably, light-tailed maxima concentrate under very broad and strong dependence. This is exemplified by our characterization of the concentration of maxima phenomenon for Gaussian triangular arrays. For example, in the special case of stationary Gaussian time series, vanishing autocovariance is necessary and sufficient for the maxima to concentrate in the same way as independent standard normal random variables. This is in stark contrast with the behavior of sums, commonly studied under short- and long-range dependence conditions (see, e.g., Dedecker et al. 2007; Pipiras and Taqqu 2017). Simply put, the notion of weak dependence that entails that the maxima of dependent variables concentrate at the same rate as in the case of independence is fundamentally weaker than the conventional mixing conditions widely used in the study of sums.

types of thresholding

Our probabilistic contributions may be of independent interest and extend classic work of Berman (1964). Concentration of maxima is a type of superconcentration phenomenon studied also in Chatterjee (2014) and Tanguy (2015a). The robustness of the concentration of maxima phenomenon to dependence can perhaps explain the universality of phase transitions in support recovery problems.

The use of concentration of maxima phenomenon highlights one core idea in our work, which allows for a first of its kind comprehensive treatment of thresholding procedures under very broad error-dependence conditions. The text involves also a full spectrum of related results such as minimax-optimality and finite-sample Bayes optimality in support estimation. Using different type of loss functions and Type I error controls, we obtain a rich picture of the exact and approximate support



77 recovery problems in high dimensions. Many of these phase-transition results have
78 not appeared in previously published literature.

79 High-dimensional support recovery problems arise in many modern applications
80 ranging from cybersecurity, theoretical computer science, to statistical genetics.
81 Genome-wide Association Studies (GWAS) in genetics are particularly natural applica-
82 tions, where the asymptotic phase-transition results help explain and quantify a
83 previously observed empirical phenomenon of the so-called *steep part of the power*
84 *curve*. In the last chapter of this work, we detail this application and highlight future
85 theoretical and practical consequences of our work.

86 **Target audience.** The original research presented in this text originates from the
87 doctoral dissertation of the first author in the Statistics Department at the University
88 of Michigan, Ann Arbor. The main goal of this text is to provide a comprehen-
89 sive treatment of the exact and approximate support recovery problems by utilizing
90 existing and newly developed probabilistic tools on concentration of maxima. The
91 text also provides a quick introduction to the state of the art in the dynamic area
92 of phase transitions in high-dimensional testing and inference. It is accessible to
93 doctoral students in Statistics with background in measure-theoretic probability and
94 statistics as well as to researchers in applied fields working with high-dimensional
95 datasets. The text can be used as a reference and a supplement to a special topics
96 course on high-dimensional inference.

Chicago, USA
97 Ann Arbor, USA
June 2021

Zheng Gao
Stilian Stoev

98 **Acknowledgements** The authors gratefully acknowledge the support of their families and all
99 colleagues from the Statistics Department at the University of Michigan, Ann Arbor. Special thanks
100 (in alphabetical order) go to Xuming He, Tailen Hsing, Michalis Kallitsis, Liza Levina, Yuanzhi
101 Li, Rodderick Little, Ya'acov Ritov, Kerby Shedden, Jinqi Shen, Jonathan Terhorst, and Gongjun
102 Xu. We would also like to thank the editorial board at the SpringerBriefs series in Probability and
103 Mathematical Statistics, and Associate Editor Richard Kruehl, who had been extremely supportive
104 throughout the preparation of the manuscript. The authors were partially supported by the NSF
105 program *Algorithms for Threat Detection*.



1 Contents

2	1 Introduction and Guiding Examples	1
3	1.1 The Additive Error Model	2
4	1.2 Genome-Wide Association Studies and the Chi-Square Model	3
5	1.3 Contents	5
6	2 Risks, Procedures, and Error Models	7
7	2.1 Statistical Risks	7
8	2.2 Statistical Procedures	10
9	2.3 Related Literature and Our Contributions	14
10	2.4 Relationships Between the Asymptotic Risks	17
11	2.5 The Asymptotic Generalized Gaussian (AGG) Models	18
12	2.6 Rapid Variation and Relative Stability	19
13	2.7 Auxiliary Facts About Gaussian Distributions	22
14	3 A Panorama of Phase Transitions	25
15	3.1 Sparse Signal Detection Problems	26
16	3.2 Sparse Signal Support Recovery Problems	31
17	3.3 The Exact Support Recovery Problem	33
18	3.4 The Approximate Support Recovery Problem	34
19	3.5 Monotonicity of the Benjamini–Hochberg Procedure	38
20	3.6 The Exact–Approximate Support Recovery Problem	40
21	3.7 The Approximate–Exact Support Recovery Problem	42
22	3.8 Asymptotic Power Analysis: A Discussion	44
23	4 Exact Support Recovery Under Dependence	47
24	4.1 Generalizations of Distributional and Dependence Assumptions	47
25	4.2 Sufficient Conditions for Exact Support Recovery	49
26	4.3 Dependence and Uniform Relative Stability	53
27	4.4 Necessary Conditions for Exact Support Recovery	54
28	4.5 Dense Signals	58
29	4.6 Numerical Illustrations for Independent Errors	59

5 Bayes and Minimax Optimality	63
5.1 Bayes Optimality in Support Recovery Problems	63
5.2 Bayes Optimality of Oracle Thresholding	65
5.3 Bayes Optimality of Likelihood Ratio Thresholding	68
5.4 Sub-optimality of Data Thresholding Procedures	70
5.5 Minimax Optimality in Exact Support Recovery	71
5.5.1 Point-Wise Minimax Optimality for Thresholding Procedures	71
5.5.2 Minimax Optimality over All Procedures	72
5.6 Optimality and Sub-optimality: A Discussion	74
6 Uniform Relative Stability for Gaussian Arrays	75
6.1 Ramsey's Theory and the Structure of Correlation Matrices	79
6.2 URS Implies UDD (Proof of the "Only If" Part of Theorem 6.1)	82
6.3 UDD Implies URS (Proof of the 'If' Part of Theorem 6.1)	84
6.3.1 Bounding the Upper Tails of AGG Maxima	85
6.3.2 Bounding the Lower Tails of Gaussian Maxima	87
6.4 Numerical Illustrations of Exact Support Recovery Under Dependence	91
7 Fundamental Statistical Limits in Genome-Wide Association Studies	97
7.1 Support Recovery Problems in Chi-Squared Models	98
7.1.1 The Exact Support Recovery Problem	99
7.1.2 The Exact–Approximate Support Recovery Problem	100
7.1.3 The Approximate Support Recovery Problem	100
7.1.4 The Approximate–Exact Support Recovery Problem	101
7.1.5 Comparison of One- Versus Two-Sided Alternatives in Additive Error Models	101
7.2 Odds Ratios and Statistical Power	102
7.3 Optimal Study Designs and Rare Variants	106
7.4 Phase Transitions in Large-Scale Association Screening Studies	108
7.5 Numerical Illustrations of the Phase Transitions in Chi-Square Models	111
7.5.1 Exact Support Recovery	111
7.5.2 Approximate, and Approximate–Exact Support Recovery	113
Appendix A: Additional Proofs	115
Appendix B: Exact Support Recovery in Non AGG Models	129
Bibliography	137



1 Acronyms

2	AGG	Asymptotically generalized Gaussian
3	BH	Benjamini–Hochberg
4	CDF	Cumulative distribution function
5	FDR	False discovery rate
6	FNR	False non-discovery rate
7	FWER	Family-wise error rate
8	FWNR	Family-wise non-discovery rate
9	GWAS	Genome-wide association studies
10	HC	Higher criticism
11	iid	Independent and identically distributed
12	LR	Likelihood ratio
13	RS	Relatively stable/relative stability
14	SNP	Single-nucleotide polymorphisms
15	URS	Uniform relatively stable/uniform relative stability

Author Queries

Chapter

Query Refs.	Details Required	Author's response
AQ1	Please check and approve the edit made in the sentence “Most, if not all, existing...” and amend if necessary.	<p style="color: red;">Please see the annotation and replace “type thresholding” by “types of thresholding”.</p>

Metadata of the chapter that will be visualized in SpringerLink

Book Title	Concentration of Maxima and Fundamental Limits in High-Dimensional Testing and Inference	
Series Title		
Chapter Title	Introduction and Guiding Examples	
Copyright Year	2021	
Copyright HolderName	The Author(s), under exclusive license to Springer Nature Switzerland AG	
Corresponding Author	Family Name Particle Given Name Prefix Suffix Role Division Organization Address Email ORCID	Gao Zheng Department of Statistics University of Michigan-Ann Arbor Ann Arbor, MI, USA gaozheng@umich.edu https://orcid.org/0000-0003-1394-8741
Author	Family Name Particle Given Name Prefix Suffix Role Division Organization Address Email ORCID	Stoev Stilian Department of Statistics University of Michigan-Ann Arbor Ann Arbor, MI, USA https://orcid.org/0000-0003-2108-5421
Abstract	This is an introductory chapter where two general problems in high-dimensional testing and inference are introduced. Namely, the signal detection and signal support estimation problems. This is done in the context of two general models—the signal-plus-noise model and the chi-square model arising in genome-wide association studies. Motivating applications of these two problems and two models are briefly discussed. The chapter concludes with a summary of the contents of the manuscript.	

Chapter 1

Introduction and Guiding Examples



0 The proliferation of information technology has enabled us to collect and consume
1 huge volumes of data at unprecedented speeds and at very low costs. This convenient
2 access to data gave rise to a fundamentally different way of pursuing scientific ques-
3 tions. In contrast with the traditional hypothesis–experiment–analysis cycle where
4 data are collected from the experiments, nowadays abundant data are often available
5 before specific questions are even formulated. Such data can be used for not just eval-
6 uating hypotheses, but also for *generating* and *selecting* the hypotheses to pursue.
7 As a result, multiple testing—where a large number of hypotheses are formulated
8 and screened for their plausibility simultaneously—has become a staple of modern
9 data-driven studies.

10 An archetypal example of multiple testing problems is genetic association studies
11 (Bush and Moore 2012). In these studies, scientists test hypotheses relating each of
12 the hundreds of thousands of genetic marker locations to phenotypic traits of interest.
13 For a phenotypic trait on which we have little prior knowledge, we cannot simply
14 test for association on one or a few specific genetic locations, as there are often not
15 enough empirical evidence or biological theory to pin point these genetic locations
16 in the first place. Rather, the goal here is to select the set of most promising genetic
17 markers from a large number of candidate locations for subsequent investigation.

18 Another example of multiple testing problems arise in cybersecurity, where mil-
19 lions of IP addresses are monitored in real time. In this engineering application,
20 statistics are collected and tests are performed for each IP address, in an attempt to
21 locate the IP addresses with anomalous network activities, so that malicious traf-
22 fic and volumetric attacks can be filtered to protect end users of network services
23 (Kallitsis et al. 2016). Similar to the genetic application above, we use data to search
24 over candidate IP addresses and identify locations of interest.

We are motivated very much by these examples to study high-dimensional multiple testing problems where a large number of hypotheses are tested simultaneously. In the rest of the introduction, we shall more review the main objectives of high-dimensional multiple testing, and elaborate on these objectives with two classes of data models in the context of various applications.

1.1 The Additive Error Model

Consider the canonical signal-plus-noise model where the observation x is a high-dimensional vector in \mathbb{R}^p ,

$$x(i) = \mu(i) + \epsilon(i), \quad i = 1, \dots, p. \quad (1.1)$$

The signal, $\mu = (\mu(i))_{i=1}^p$, is a vector with s non-zero components supported on the set $S = \{i : \mu(i) \neq 0\}$; the second term ϵ is a random error vector. The goal of high-dimensional statistics is usually twofold:

- I. *Signal detection*: To detect the presence of non-zero components in μ . That is, to test the global hypothesis $\mu = 0$.
- II. *Support recovery*: To estimate the support set S . This is also sometimes referred to as the *support estimation* or *signal identification* problem.

To illustrate, in the engineering application of cybersecurity, Internet service providers (ISP) routinely monitor a large number of network traffic streams to determine if there are abnormal surges, blackouts, or other types of anomalies. The data vector x could represent, for example, incoming traffic volumes to each server node, Internet protocol (IP) address, or port that the ISP monitors. In this case, the vector μ represents the average traffic volumes in each of the streams under normal operating conditions, and ϵ 's—the fluctuations around these normal levels of traffic. The signal detection problem in this context is then equivalent to determining if there are *any* anomalies among all data streams, and the support recovery problem is equivalent to *identifying* the streams experiencing anomalies. Similar questions of signal detection and support recovery are pursued in large-scale microarray experiments (Dudoit et al. 2003), brain imaging and fMRI analysis (Nichols and Hayasaka 2003), and numerous other anomaly detection applications.

A common theme in such applications is that the errors are *correlated*, and that the signal vectors are believed to be *sparse*: the number of non-zero (or large) components in μ is small compared to the number of tests performed. In the cybersecurity context, while a very large number of data streams are monitored, typically only just a few of them will be experiencing problems at any time, barring large-scale outages or distributed denial-of-service attacks. Under such sparsity assumptions, it is natural to ask if and when one can reliably (1) detect the signals and (2) recover the support

61 set S . In this text, we explore both the *detection* and the *support recovery* problems.
 62 More precisely, we are interested in the theoretical feasibility of both problems, and
 63 seek minimal conditions under which these problems can be consistently solved in
 64 large dimensions.

65 Model (1.1) is simple yet ubiquitous. Consider the linear model

$$66 \quad Y = X\mu + \xi,$$

67 where μ is a p -dimensional vector of regression coefficients of interest to be inferred
 68 from observations of X and Y . If the design matrix X is of full column rank,¹ then
 69 the ordinary least squares (OLS) estimator of μ can be formed

$$70 \quad \hat{\mu} = (X'X)^{-1} X'Y = \mu + \epsilon, \quad (1.2)$$

71 where $\epsilon := (X'X)^{-1} X'\xi$. Hence, we recover the generic problem (1.1). Signal detec-
 72 tion is therefore equivalent to the problem of testing the global null model, and support
 73 recovery problem corresponds to the fundamental problem of variable selection.

74 Note that the components of the observation vector x (and equivalently, the
 75 noise ϵ) in (1.1) need not be independent. In the linear regression example, even
 76 when the components of the noise term ξ are independent, those of the OLS estima-
 77 tor (1.2) need not be, except in the case of orthogonal designs. Indeed, in practice,
 78 independence is the exception rather than the rule. Therefore, a general theory of fea-
 79 sibility must address the role of the *error-dependence* structure in such testing and
 80 support estimation problems. It is also important to identify practical and/or optimal
 81 procedures that attain the performance limits in independent as well as dependent
 82 cases, as soon as the problems become theoretically feasible. We address both themes
 83 in this text.

84 1.2 Genome-Wide Association Studies and the Chi-Square 85 Model

86 The second data model we analyze is the high-dimensional chi-square model,

$$87 \quad x(i) \sim \chi_v^2(\lambda(i)), \quad i = 1, \dots, p, \quad (1.3)$$

88 where the data $x(i)$'s follow independent (non-central) chi-square distributions with
 89 v degrees of freedom and non-centrality parameter $\lambda(i)$.

¹ This, of course, requires that we have more samples than dimensions, i.e., $n > p$. Nevertheless, multiplicity of tests is still present when p itself is large—the multiple testing problem is by no means exclusive to situations where $p \gg n$.

Table 1.1 Tabulated counts of genotype-phenotype combinations in a genetic association test

# Observations	Genotype		Total by phenotype
	Variant 1	Variant 2	
Cases	O_{11}	O_{12}	n_1
Controls	O_{21}	O_{22}	n_2

Model (1.3) is motivated by large-scale categorical variable screening problems, typified by GWAS where millions of genetic factors are examined for their potential influence on phenotypic traits.

In a GWAS with a case-control design, a total of n subjects are recruited, consisting of n_1 subjects possessing some defined traits, and n_2 subjects without the traits serving as controls. The genetic compositions of the subjects are then examined for variations known as SNP at an array of p genomic marker locations, and compared between the case and the control group. These physical traits are commonly referred to as *phenotypes*, and the genetic variations are known as *genotypes* (Table 1.1).

Focusing on one specific genomic location, the counts of observed genotypes, if two variants are present, can be tabulated as follows. Researchers test for associations between the genotypes and phenotypes using, for example, the Pearson chi-square test with statistic

$$x = \sum_{j=1}^2 \sum_{k=1}^2 \frac{(O_{jk} - E_{jk})^2}{E_{jk}}, \quad (1.4)$$

where $E_{jk} = (O_{j1} + O_{j2})(O_{1k} + O_{2k})/n$.

Under the mild assumption that the counts O_{jk} 's follow a multinomial distribution (or a product-binomial distribution, if we decide to condition on one of the marginals), the statistic x in (1.4) can be shown to have an approximate $\chi^2(\lambda)$ distribution with $v = 1$ degree of freedom at large sample sizes (see, e.g., classical results in Ferguson 2017; Agresti 2018). Independence between the genotypes and phenotypes would imply a non-centrality parameter λ value of zero; if dependence exists, we would have a non-zero λ where its value depends on the underlying multinomial probabilities. More generally, if we have a J phenotypes and K genetic variants, assuming a $J \times K$ multinomial distribution, the statistic will follow approximately a $\chi^2_v(\lambda)$ distribution with $v = (J - 1)(K - 1)$ degrees of freedom, when sample sizes are large.

The same asymptotic distributional approximations also apply to the likelihood ratio statistic, and many other statistics under slightly different modeling assumptions (Gao et al. 2019). These association tests are performed at each of the p SNP marker locations throughout the whole genome, and we arrive at p statistics having approximately (non-central) chi-square distributions, $\chi^2_{v(i)}(\lambda(i))$, for $i = 1, \dots, p$, where $\lambda = (\lambda(i))_{i=1}^p$ is the p -dimensional non-centrality parameter.

Although the number of tested genomic locations p can sometimes exceed 10^5 or even 10^6 , it is often believed that only a small set of genetic locations have tangible influences on the outcome of the disease or the trait of interest. Under the stylized



assumption of sparsity, λ is assumed to have s non-zero components, with s being much smaller than the problem dimension p . The goal of researchers is again twofold: (1) to test if $\lambda(i) = 0$ for all i , and (2) to estimate the set $S = \{i : \lambda(i) \neq 0\}$. In other words, we look to first determine if there are *any* genetic variations associated with the disease, and if there are associations, we want to locate them.

The chi-square model (1.3) also plays an important role in analyzing variable screening problems under omnidirectional alternatives. A primary example is multiple testing under two-sided alternatives in the additive error model (1.1) where the errors ϵ are assumed to have standard normal distributions.

Under two-sided alternatives, unbiased test procedures call for rejecting the hypothesis $\mu(i) = 0$ at locations where observations have large absolute values or, equivalently, large squared values. Taking squares on both sides of (1.1), we arrive at Model (1.3) with non-centrality parameters $\lambda(i) = \mu^2(i)$ and degree-of-freedom parameter $v = 1$. In this case, the support recovery problem is equivalent to locating the set of observations with mean shifts, $S = \{i : \mu(i) \neq 0\}$, where the mean shifts could take place in both directions.

Therefore, a theory for the chi-square model (1.3) naturally lends itself to the study of two-sided alternatives in the Gaussian additive error model (1.1). In comparing such results with existing theory on one-sided alternatives, we will be able to quantify if, and how much of a price has to be paid for the additional uncertainty when we have no prior knowledge on the direction of the signals.

1.3 Contents

Important notions and definitions in high-dimensional testing problems are recalled in Chap. 2. We review related literature as well as key concepts and technical results used in our subsequent analyses.

In Chap. 3, we study the sparse signal detection and support recovery problems for the additive error model (1.1) when components of the noise term ϵ are independent standard Gaussian random variables. In particular, we point out several new *phase transitions* in signal detection problems, and provide a unified account of recently discovered phase transitions in support recovery problems. These results show that as the dimension $p \rightarrow \infty$, the tasks of detecting the existence of signals or identifying the support set S are either doable or impossible depending on the sparsity and signal sizes of the problems. We also identify commonly used procedures that attain the performance limits in both detection and support recovery problems.

Both the Gaussianity assumption and the independence assumption are relaxed in Chap. 4. Established are the necessary and sufficient conditions for exact support recovery in the high-dimensional asymptotic regime for the large class of thresholding procedures. This is a major theoretical contribution of our approach, which solves and expands on open problems in the recent literature (see Butucea et al. 2018; Gao and Stoev 2020). The analysis of support recovery problem is intimately related to a *concentration of maxima* phenomena in the analysis of extremes. The

latter concept is key to understanding the role played by dependence in the phase-transition phenomena of high-dimensional testing problems. In Chap. 5, we study the universality of the phase-transition phenomenon in exact support recovery. We do so by first establishing the finite-sample Bayes optimality and sub-optimality of thresholding procedures. This, combined with the results from Chap. 4, culminates in asymptotic minimax characterizations of the phase-transition phenomenon in exact support recovery across all procedures for a large class of dependence structures.

The dependence condition defined by the concentration of maxima concepts is further demystified in Chap. 6 for Gaussian errors. We offer a complete characterization of the concentration of maxima phenomenon, known as uniform relative stability, in terms of the covariance structures of the Gaussian arrays. This result may be of independent interest since it relates to the so-called *superconcentration* phenomenon coined by Chatterjee (2014). See also, Gao and Stoev (2020), Kartsioukas et al. (2019).

Chapter 7 returns to high-dimensional multiple testing problems and studies the chi-square model (1.3) inspired by the marginal association screening problems. We establish four new phase-transition-type results in the chi-square model, and illustrate their practical implications in the GWAS application. Our theory enables us to explain the long-standing empirical observation that small perturbations in the frequency and penetrance of genetic variations lead to drastic changes in the discoverability in genetic association studies.

Author Queries

Chapter 1

Query Refs.	Details Required	Author's response
AQ1	Please check and confirm if the inserted citation of Table 1.1 is correct. If not, please suggest an alternate citation. Please note that tables should be cited sequentially in the text.	Agreed. Thanks.

Metadata of the chapter that will be visualized in SpringerLink

Book Title	Concentration of Maxima and Fundamental Limits in High-Dimensional Testing and Inference	
Series Title		
Chapter Title	Risks, Procedures, and Error Models	
Copyright Year	2021	
Copyright HolderName	The Author(s), under exclusive license to Springer Nature Switzerland AG	
Corresponding Author	Family Name Particle Given Name Prefix Suffix Role Division Organization Address Email ORCID	Gao Zheng Department of Statistics University of Michigan-Ann Arbor Ann Arbor, MI, USA gaozheng@umich.edu https://orcid.org/0000-0003-1394-8741
Author	Family Name Particle Given Name Prefix Suffix Role Division Organization Address Email ORCID	Stoev Stilian Department of Statistics University of Michigan-Ann Arbor Ann Arbor, MI, USA https://orcid.org/0000-0003-2108-5421
Abstract	<p>This chapter introduces the necessary background and notation for the study of the signal detection and support estimation problems in the high-dimensional setting. First, a general notion of statistical risk is formulated and used to provide a unified treatment of a variety of signal detection and support estimation problems considered in the literature. The relationship between various choices of risk functionals and their implications on the approximate and exact support recovery are discussed. Depending on the type of control of the false discovery and false non-discovery, four different risk functionals are defined, each of which is shown to yield a different type of phase transition in the support estimation problems in Chaps. 3 and 7. The chapter proceeds with a review of several popular signal detection and support estimation procedures such as Tukey's Higher Criticism and the general class of thresholding procedures. The latter is at the center of the investigation in the subsequent chapters. The existing literature is discussed in the context of the defined statistical risks and procedures. The chapter concludes with some probabilistic background. Introduced is the class of asymptotically generalized Gaussian (AGG) distributions, which provide flexible additive error distribution models. Reviewed are the notions of rapid variation, relative stability, as well as fundamental tools for Gaussian models such as Slepian's lemma and the Sudakov–Fernique inequality.</p>	

Chapter 2

Risks, Procedures, and Error Models



0 We establish the background necessary for the study of sparse signal detection and
1 support recovery problems in this chapter. Sections 2.1 and 2.2 provide a refresher
2 on the definitions of statistical risks and some commonly used statistical procedures.
3 Section 2.3 describes the asymptotic regime under which we analyze these proce-
4 dures, and reviews the related literature in high-dimensional statistics. We discuss
5 in Sect. 2.4 the connections among the risk metrics, and point out some common
6 fallacies. The remaining sections collect the technical preparations for this text.
7 Section 2.5 defines an important class of error distributions which will be analyzed
8 in detail in later chapters. ~~And finally,~~ Sect. 2.6 introduces the concepts of concen-
9 tration of maxima, which plays a crucial role in the analysis of high-dimensional
10 support recovery problems. Finally, in Sect. 2.7, we gather well-known but indis-
11 pensable facts about Gaussian distributions.

Section

2.1 Statistical Risks

12 We define the statistical risk metrics for signal detection and signal support recovery
13 problems in this section. Formally, we denote a statistical procedure, i.e., measurable
14 function of the data, as $\mathcal{R} = \mathcal{R}(x)$. In the testing context, a procedure \mathcal{R} produces
15 a binary decision T that represents our judgment on the presence or absence of a
16 signal. In the support recovery problem, a procedure \mathcal{R} produces an index set \widehat{S} that
17 represents our estimate of the signal support. The statistical risks are then suitable
18 functionals of T and \widehat{S} in respective contexts.

20 **Signal detection.** Recall that in sparse signal detection problems, our goal is to come
 21 up with a procedure, $\mathcal{R}(x)$, such that the null hypothesis is rejected if the data x is
 22 deemed incompatible with the null. In the additive error models context (1.1), we
 23 wish to tell apart two hypotheses

$$\mathcal{H}_0 : \mu(i) = 0, \quad i = 1, \dots, p, \quad \text{v.s.} \quad \mathcal{H}_1 : \mu(i) \neq 0, \quad \text{for some } i \in \{1, \dots, p\}, \quad (2.1)$$

25 based on the p -dimensional observation x . Similarly, in the chi-square model (1.3),
 26 we look to test

$$\mathcal{H}_0 : \lambda(i) = 0, \quad i = 1, \dots, p, \quad \text{v.s.} \quad \mathcal{H}_1 : \lambda(i) \neq 0, \quad \text{for some } i \in \{1, \dots, p\}. \quad (2.2)$$

28 Since the decision is binary, we may write the outcome of the procedure in the form
 29 of an indicator function, $T(\mathcal{R}(x)) \in \{0, 1\}$, where $T = 1$ if the null is to be rejected
 30 in favor of the alternative, and 0 if we fail to reject the null. The Type I and Type
 31 II errors of the procedure, i.e., the probability of wrong decisions under the null
 32 hypothesis \mathcal{H}_0 and alternative hypothesis \mathcal{H}_1 , respectively, are defined as

$$33 \quad \alpha(\mathcal{R}) := \mathbb{P}_{\mathcal{H}_0}(T(\mathcal{R}(x)) = 1) \quad \text{and} \quad \beta(\mathcal{R}) := \mathbb{P}_{\mathcal{H}_1}(T(\mathcal{R}(x)) = 0). \quad (2.3)$$

34 The Neyman–Pearson framework of hypothesis testing then seeks tests that minimize
 35 the Type II error of the test, while controlling the Type I error of the test at low levels.
 36 We are particularly interested in the sum of the two errors,

$$37 \quad \text{risk}^D(\mathcal{R}) := \alpha(\mathcal{R}) + \beta(\mathcal{R}), \quad (2.4)$$

38 which shall be referred to as the risk of signal detection (of the procedure \mathcal{R}). It is
 39 trivial that a small risk^D would imply both small Type I and Type II errors of the
 40 procedure.

41 **Signal support recovery.** Turning to support recovery problems, our goal is to design
 42 a procedure that produces a set estimate $\widehat{S}(\mathcal{R}(x))$ of the true index set of relevant
 43 variables S . For example, in the sparse additive error model (1.1), we aim to estimate
 44 $S = \{i : \mu(i) \neq 0\}$, while in the sparse chi-square model (1.3) the goal is to esti-
 45 mate $S = \{i : \lambda(i) \neq 0\}$. For simplicity of notation, we shall write \widehat{S} for the support
 46 estimator $\widehat{S}(\mathcal{R}(x))$.

47 For a given procedure \mathcal{R} , its false discovery rate (FDR) and false non-discovery
 48 rate (FNR) are defined, respectively, as

$$49 \quad \text{FDR}(\mathcal{R}) := \mathbb{E} \left[\frac{|\widehat{S} \setminus S|}{\max\{|\widehat{S}|, 1\}} \right] \quad \text{and} \quad \text{FNR}(\mathcal{R}) := \mathbb{E} \left[\frac{|S \setminus \widehat{S}|}{\max\{|S|, 1\}} \right], \quad (2.5)$$

50 where the maxima in the denominators resolve the possible division-by-0 problem.
 51 Roughly speaking, FDR measures the expected fraction of false findings, while FNR

describes the proportion of Type II errors among the true signals, and reflects the average marginal power of the procedure.

A more stringent criterion for false discovery is the family-wise error rate (FWER), defined to be the probability of reporting at least one finding not contained in the true index set. Correspondingly, a more stringent criterion for false non-discovery is the family-wise non-discovery rate (FWNR), i.e., the probability of missing at least one signal in the true index set. That is,

$$\text{FWER}(\mathcal{R}) := 1 - \mathbb{P}[\hat{S} \subseteq S] \quad \text{and} \quad \text{FWNR}(\mathcal{R}) := 1 - \mathbb{P}[S \subseteq \hat{S}]. \quad (2.6)$$

We introduce five different statistical risk metrics, each having different asymptotic limits in the support recovery problems as we will see in Chap. 3. Following Arias-Castro and Chen (2017), we define the risk for *approximate* support recovery as

$$\text{risk}^A(\mathcal{R}) := \text{FDR}(\mathcal{R}) + \text{FNR}(\mathcal{R}). \quad (2.7)$$

Analogously, we define the risk for *exact* support recovery as

$$\text{risk}^E(\mathcal{R}) := \text{FWER}(\mathcal{R}) + \text{FWNR}(\mathcal{R}). \quad (2.8)$$

Two closely related measures of success in the exact support recovery risk are the probability of exact recovery,

$$\mathbb{P}[\hat{S} = S] = 1 - \mathbb{P}[\hat{S} \neq S], \quad (2.9)$$

and the Hamming loss

$$H(\hat{S}, S) := |\hat{S} \Delta S| = \sum_{i=1}^p |\mathbb{1}_{\hat{S}}(i) - \mathbb{1}_S(i)|, \quad (2.10)$$

which counts the number of mismatches between the estimated and true support sets.

The relationship between probability of support recovery $\mathbb{P}[\hat{S} = S]$, exact support recovery risk risk^E , and the expected Hamming loss $\mathbb{E}[H(\hat{S}, S)]$ will be discussed in Sect. 2.4.

Notice that all risk metrics introduced so far penalize false discoveries and missed signals somewhat symmetrically—the approximate support recovery risk combines proportions of errors, the exact support recovery risk combines probabilities of errors, and the Hamming loss increments the risk by one regardless of the types of errors made. In applications, however, attitudes toward false discoveries and missed signals are often asymmetric. In the example of GWAS, where the number of candidate locations p could be in the millions, and a class imbalance between the number of nulls and signals exists, researchers are typically interested in the marginal (location-wise) power of discovery, while exercising stringent (family-wise) false discovery

control. These types of asymmetric considerations, while important in applications, have not been studied theoretically. For example, the GWAS application motivates the *exact–approximate* support recovery risk, which weighs both the family-wise error rate and the marginal power of discovery:

$$\text{risk}^{\text{EA}}(\mathcal{R}) := \text{FWER}(\mathcal{R}) + \text{FNR}(\mathcal{R}). \quad (2.11)$$

The somewhat cumbersome name and notation are chosen to reflect the asymmetry in dealing with the two types of errors in support recovery. Namely, when the risk metric (2.11) vanishes, we have “exact false discovery control, and approximate false non-discovery control” asymptotically.

Analogously, we consider the *approximate–exact* support recovery risk

$$\text{risk}^{\text{AE}}(\mathcal{R}) := \text{FDR}(\mathcal{R}) + \text{FWNR}(\mathcal{R}), \quad (2.12)$$

which places more emphasis on non-discovery control over false discovery.

Theoretical limits and performance of procedures in support recovery problems will be studied in terms of the five risk metrics (2.7), (2.8), (2.9), (2.11), and (2.12), in Chaps. 3, 4, and 7. We are particularly interested in fundamental limits of signal detection and support recovery problems in terms of these metrics, as well as the optimality of commonly used procedures in high-dimensional settings.

2.2 Statistical Procedures

We review some popular procedures for signal detection and signal support recovery tasks in this section.

Signal detection. One of the commonly used statistics in sparse signal detection problems such as (2.1) and (2.2) are the L_q norms of the observations x ,

$$L_q(x) = \left(\sum_{i=1}^p |x(i)|^q \right)^{1/q}. \quad (2.13)$$

Typical choices of q include $q = 1, 2$ and ∞ , where $L_\infty(x)$ is interpreted as the limit of $L_q(x)$ norms as $q \rightarrow \infty$, and is equivalent to $\max_i |x(i)|$. Test procedures based on (2.13) may then be written as $T(\mathcal{R}(x)) = \mathbb{1}_{(t,+\infty)}(L_q(x))$, where the cutoff t can be chosen to control the Type I error at desired levels.

While (2.13) measures the deviation of the data from the origin in an omnidirectional manner, statistics that are tailored to the alternatives can be used in the hopes of power improvement if the directions of the alternatives are known. For example, in the additive error model (1.1), suppose we want to test for positive mean shifts, i.e., one-sided alternative

117 $\mathcal{H}_1 : \mu(i) > 0, \text{ for some } i \in \{1, \dots, p\}.$ (2.14)

118 Then, one might consider monitoring the sum (or equivalently, the arithmetic average)
 119 of the observations,

120 $T(x) := \sum_{i=1}^p x(i),$ (2.15)

121 or the maximum of the observations,

122 $M(x) := \max_{i=1, \dots, p} x(i).$ (2.16)

123 Other tests based on the empirical CDF are also available. Assuming the same one-
 124 sided alternative, let

125 $q(i) = 1 - \sup\{F_i(y) : y < x(i)\}, \quad i = 1, \dots, p$ (2.17)

126 be the p-values of the individual observations, where F_i is the CDF of the i -th
 127 component $x(i)$ under \mathcal{H}_0 . We define empirical CDF of the p-values as

128 $\widehat{F}_p(t) = \frac{1}{p} \sum_{i=1}^p \mathbb{1}_{[0,t]}(q(i)).$ (2.18)

129 Viewed as random elements in the space of càdlàg functions with the Skorohod J_1
 130 topology, the centered and scaled CDFs converge weakly to a Brownian bridge,

131 $\{\sqrt{p}(\widehat{F}_p(t) - t)\}_{t \in [0,1]} \implies \{\mathbb{B}(t)\}_{t \in [0,1]}, \quad \text{as } p \rightarrow \infty,$

132 under the global null \mathcal{H}_0 and mild continuity assumptions on the F_i 's (Skorokhod
 133 1956). Therefore, goodness-of-fit statistics such as Kolmogorov–Smirnov distance
 134 (Smirnov 1948), Cramer–von Mises-type statistics (Cramér 1928; Anderson and
 135 Darling 1952) that measure the departure from this limiting behavior can be used for
 136 testing \mathcal{H}_0 against \mathcal{H}_1 . Of particular interest is the higher criticism (HC) statistic,
 137 first proposed by Tukey (1976),

138 $HC(x) = \max_{0 \leq t \leq \alpha_0} \frac{\widehat{F}_p(t) - t}{\sqrt{t(1-t)/p}}.$ (2.19)

Each of the above statistics L_q , S , M , or HC gives rise to a decision rule, whereby
 the null hypothesis is rejected if the statistic exceeds a suitably calibrated threshold.
 The choice of the threshold is typically determined based on large-sample limit
 theorems. For example, as shown in Theorem 1.1 of Donoho and Jin (2004), under
 the null hypothesis

$$\frac{HC(x)}{\sqrt{2 \log \log(p)}} \longrightarrow 1, \quad \text{in probability},$$

as $p \rightarrow \infty$. Thus, one decision rule is to reject \mathcal{H}_0 , if $HC(x) > t(p, \alpha_p)$, where $t(p, \alpha_p) = \sqrt{2 \log \log(p)}(1 + o(1))$. As we will see, this yields an optimal signal detection procedure (see also Theorem 1.2 in Donoho and Jin 2004). The performance of these statistics in high-dimensional sparse signal detection problems will be reviewed in Sect. 2.3, and analyzed in Chap. 3.

Signal support recovery. In signal support recovery tasks, we shall study the performance of five procedures, all of which belong to the broad class of thresholding procedures.

Definition 2.1 (*Thresholding procedures*) A thresholding procedure for estimating the support $S := \{i : \lambda(i) \neq 0\}$ is one that takes on the form

$$\widehat{S} = \{i \mid x(i) \geq t(x)\}, \quad (2.20)$$

where the threshold $t(x)$ may depend on the data x .

Examples of thresholding procedures include ones that aim to control FWER (2.6)—Bonferroni's (Dunn 1961), Sidák's (Šidák 1967), Holm's (Holm 1979), and Hochberg's procedure (Hochberg 1988)—as well as procedures that target FDR (2.5), such as the Benjamini–Hochberg Benjamini and Hochberg (1995) and the Barber–Candès procedure (Barber and Candès 2015). Indeed, the class of thresholding procedures (2.20) is so general that it contains most (but not all) of the statistical procedures in the multiple testing literature.

Under the assumption that the data $x(i)$'s under the null have a common marginal distribution F , we review five thresholding procedures for support recovery, starting with the well-known Bonferroni's procedure which aims at controlling family-wise error rates.

Definition 2.2 (*Bonferroni's procedure*) Bonferroni's procedure with level α is the thresholding procedure that uses the threshold

$$t_p = F^\leftarrow(1 - \alpha/p), \quad (2.21)$$

where $F^\leftarrow(u) = \inf \{x : F(x) \geq u\}$ is the generalized inverse function.

The Bonferroni procedure is deterministic, i.e., non-data-dependent, and only depends on the dimension of the problem and the null distribution. A closely related procedure is Sidák's procedure (Šidák 1967), which is a more aggressive (and also deterministic) thresholding procedure that uses the threshold

$$t_p = F^\leftarrow((1 - \alpha)^{1/p}). \quad (2.22)$$

The third procedure, strictly more powerful than Bonferroni's, is the so-called Holm's procedure (Holm 1979). On observing the data x , its coordinates can be

173 ordered from largest to smallest $x(i_1) \geq x(i_2) \geq \dots \geq x(i_p)$, where (i_1, \dots, i_p) is a
 174 permutation of $\{1, \dots, p\}$. Denote these order statistics as $x_{[1]}, x_{[2]}, \dots, x_{[p]}$.

Definition 2.3 (*Holm's procedure*) Let i^* be the largest index such that

$$\bar{F}(x_{[i]}) \leq \alpha/(p - i + 1), \quad \text{for all } i \leq i^*.$$

175 Holm's procedure with level α is the thresholding procedure with threshold

$$t_p(x) = x_{[i^*]}. \quad (2.23)$$

In contrast to the Bonferroni procedure, Holm's procedure is data-dependent. A closely related, more aggressive (and also data-dependent) thresholding procedure is Hochberg's procedure (Hochberg 1988). It replaces the index i^* in Holm's procedure with the largest index such that

$$\bar{F}(x_{[i]}) \leq \alpha/(p - i + 1).$$

177 Notice that both Holm's and Hochberg's procedures compare p-values to the
 178 same thresholds $\alpha/(p - i + 1)$. However, Holm's procedure only rejects the set
 179 of hypotheses whose p-values are all smaller than their respective thresholds. On the
 180 other hand, Hochberg's procedure rejects the set of hypotheses as long as the largest
 181 of their p-values fall below its threshold, and therefore can be more powerful than
 182 Holm's procedure.

183 It can be shown that both Bonferroni's and Holm's procedures control FWER at
 184 their nominal levels, regardless of dependence in the data (Holm 1979). In contrast,
 185 Sidák's and Hochberg's procedures control FWER at nominal levels when data are
 186 independent (Šidák 1967; Hochberg 1988).

187 Last but not least, we review the BH procedure, which aims at controlling FDR
 188 in (2.5), proposed by Benjamini and Hochberg (1995).

189 Recall the order statistics of our observations are: $x_{[1]} \geq x_{[2]} \geq \dots \geq x_{[p]}$.

Definition 2.4 (*Benjamini–Hochberg's procedure*) Let i^* be the largest index such
 that

$$\bar{F}(x_{[i]}) \leq \alpha i / p.$$

190 The Benjamini–Hochberg (BH) procedure with level α is the thresholding procedure
 191 with threshold

$$t_p(x) = x_{[i^*]}. \quad (2.24)$$

193 The BH procedure is shown to control the FDR at level α when the $x(i)$'s are inde-
 194 pendent (Benjamini and Hochberg 1995). Variations of this procedure have been
 195 proposed to control the FDR under certain models of dependent observations (Ben-
 196 jamini and Yekutieli 2001).

197 The performance of these procedures in high-dimensional sparse signal support
 198 recovery problems will be reviewed in Sect. 2.3, and analyzed in Chaps. 3, 4, and 7.

199 2.3 Related Literature and Our Contributions

200 We look to derive useful asymptotic approximations for high-dimensional problems,
 201 and analyze the aforementioned procedures in the regime where the dimensionality
 202 of the observations ~~diverge~~. Throughout this text, we consider triangular arrays of ~~diverges~~.
 203 observations as described in Models (1.1) and (1.3), and study the performance of
 204 various procedures in the signal detection and support recovery tasks when

205
$$p \rightarrow \infty.$$

206 The criteria for success and failure in support recovery problems under this high-
 207 dimensional asymptotic regime are defined as follows.

208 **Definition 2.5** We say a sequence of procedures $\mathcal{R} = \mathcal{R}_p$ succeeds asymptotically
 209 in the detection problem (and, respectively, exact, exact–approximate, approximate–
 210 exact, and approximate support recovery problem) if

211
$$\text{risk}^P(\mathcal{R}) \rightarrow 0, \quad \text{as } p \rightarrow \infty, \tag{2.25}$$

212 where $P = D$ (respectively, E, EA, AE, A).

213 Conversely, we say the exact support recovery fails asymptotically in the detec-
 214 tion problem (and, respectively, exact, exact–approximate, approximate–exact, and
 215 approximate support recovery problem) if

216
$$\liminf \text{risk}^P(\mathcal{R}) \geq 1, \quad \text{as } p \rightarrow \infty, \tag{2.26}$$

217 where $P = D$ (respectively, E, EA, AE, A).

218 The choice of the constant 1 in Definition (2.26) allows us to declare failure for trivial
 219 testing procedures. For example, trivial deterministic procedures that always reject
 220 and ones that always fail to reject both have statistical risks 1 in either the detection or
 221 the support recovery problem. Similarly, a trivial randomized procedure that rejects
 222 the nulls uniformly at random also has risk of 1, and is declared as a failure in both
 223 problems.

224 **Signal detection.** The asymptotic behavior of the statistical risk in signal detection
 225 problems (2.4) in high dimensions was first studied by Yuri Izmailovich Ingster in
 226 the context of sparse additive models (1.1) with independent and Gaussian com-
 227 ponents. Specifically, Ingster (1998) considered the behavior of the theoretically
 228 optimal likelihood ratio (LR) test in the high-dimensional asymptotic regime, where
 229 the dimension p grows to infinity. Then, under certain parameterization of the size
 230 and sparsity of the signal μ , a dichotomy exists either $\text{risk}^D(\mathcal{R})$ vanishes as $p \rightarrow \infty$
 231 where \mathcal{R} is the LR test, or $\liminf_{p \rightarrow \infty} \text{risk}^D(\mathcal{R}) = 1$ for any procedure. The precise
 232 signal size and sparsity parameterizations as well as the so-called *detection boundary*
 233 discovered by Ingster are described in Chap. 3.

The LR test, unfortunately, relies on the knowledge of the signal sparsity and signal sizes which are not available in practice. The sparsity and signal size agnostic statistic HC in (2.19) was identified to attain such optimal performance limits in sparse Gaussian models in Donoho and Jin (2004). A modified goodness-of-fit test statistic in Zhang (2002) and two statistics based on thresholded- L_1 and L_2 norms in Zhong et al. (2013) were also shown to be asymptotically optimal in the detection problem. Recent studies have also focused on the behavior of detection risk (2.4) in dense and scale mixture models (Cai et al. 2011), under general distributional assumptions (Cai and Wu 2014; Arias-Castro and Wang 2017), as well as when the errors are dependent (Hall and Jin 2010). A comprehensive review focusing on the role of HC in detection problems can be found in Donoho and Jin (2015). The very recent contribution of Li and Fithian (2020) shows exciting new developments on the detection problem in a more realistic regime than the ones previously studied in the literature. It shows that the max-statistic begins to attain the optimal boundary and is on par with HC (cf Table 1, therein). Notwithstanding the extensive literature on the detection problem, the performances of simple statistics such as L_q norms (2.13) and sums (2.15), to the best of our knowledge, have only been sparingly documented. We gather relevant results in Chap. 3, and make several new contributions on the performance of several statistics commonly used in practice.

Exact support recovery. There is a wealth of literature on the so-called sparsistency (i.e., $\mathbb{P}[\widehat{S} = S] \rightarrow 1$ as $p \rightarrow \infty$) problem in the regression context. Sparsistency problems were pursued, among many others, by Zhao and Yu (2006), Wasserman and Roeder (2009) in the high-dimensional regression setting (where the number of samples $n \ll p$), and by Meinshausen and Bühlmann (2006) in graphical models. Although there have been numerous studies on the sufficient conditions for sparsistency, efforts on necessary conditions have been scarce. Notable exceptions include Wainwright (2009a, b), Comminges and Dalalyan (2012) in regression problems. We refer the reader to the recent book by Wainwright (2019) (and, in particular, the bibliographical sections of Chaps. 7 and 15 therein) for a comprehensive review.

Elaborate asymptotic minimax-optimality results under the Hamming loss were derived for methods proposed in Ji and Jin (2012), Jin et al. (2014) for regression problems. More recently, Butucea et al. (2018) also obtained similar minimax-optimality results for a specific procedure in the Gaussian additive error model (1.1) in terms of the expected Hamming loss.

Nevertheless, two important questions remained unanswered. Namely, precise phase-transition-type results for the exact support recovery risk (2.8) and for the support recovery probability (2.9) have not been established. And secondly, performance of commonly used statistical procedures reviewed in Sect. 2.2 in terms of these risk metrics have not been studied. Some of our main contributions in this text are the solutions to these problems, presented in Chaps. 3 and 4. Specifically, we show that the Bonferroni thresholding procedure (among others) is asymptotically optimal for the exact support recovery problem in (1.1) under broad classes of error distributions. Furthermore, a phase transition in the exact support recovery problem for thresholding procedures is established under broad dependence conditions on the

278 errors using the concentration of maxima phenomenon (Chap. 4). We also establish
 279 finite-sample Bayes optimality and sub-optimality results for these procedures under
 280 independence, and by extension arrive at minimax-optimality results for the exact
 281 support recovery problem (Chap. 5).

282 The landscape of the fundamental statistical limits in support estimation is yet to
 283 be fully charted. We conjecture, however, that the general concentration of maxima
 284 phenomenon will lead to its complete solution under very broad error-dependence
 285 scenarios.

286 **Approximate support recovery.** The performance limits of FDR-controlling proce-
 287 dures in the support recovery problem have been actively studied in recent years. The
 288 asymptotic optimality of the Benjamini–Hochberg procedure was analyzed under
 289 decision theoretic frameworks in Genovese and Wasserman (2002); Bogdan et al.
 290 (2011); Neuvial and Roquain (2012), with main focus on location/scale models. In
 291 particular, these papers show that the statistical risks of the procedures come close to
 292 those ~~that~~ of the oracle procedures under suitable asymptotic regimes. Strategies for deal-
 293 ing with multiple testing under general distributional assumptions can be found in,
 294 e.g., Efron (2004), Storey (2007), Sun and Cai (2007). The two-sided alternative in
 295 the additive error model was featured as the primary example in Sun and Cai (2007).

296 In the additive error model (1.1) under independent Gaussian errors and one-sided
 297 alternatives (2.14), Arias-Castro and Chen (2017) showed that a phase transition
 298 exists for the approximate support recovery risk (2.7). The BH procedure (Ben-
 299 jamini and Hochberg 1995) and the Barber–Candès procedure (Barber and Candès
 300 2015) were identified to be asymptotically optimal. However, Arias-Castro and Chen
 301 (2017), as all related work so far, assumed the non-nulls to follow a common alterna-
 302 tive distribution. We derive a new phase-transition result that relaxes this assumption
 303 on the alternatives in Chap. 3.

304 **Asymmetric statistical risks.** Although weighted sums of false discovery and non-
 305 discovery have been studied in the literature mentioned above, the case of simul-
 306 taneous family-wise error control and marginal location-wise power requirements
 307 has not been previously considered. As a result, asymmetric statistical risks such as
 308 (2.11) and (2.12) have not been investigated. As argued in Sect. 2.1, the properties
 309 of these asymmetric risks are of important practical concern in applications such as
 310 GWAS. We study the asymptotic behavior of these risks in Chaps. 3 and 7 of this
 311 text.

312 **Chi-square models and GWAS.** The high-dimensional chi-square model (1.3)
 313 seemed to have received little attention in the literature. While the sparse signal
 314 detection problem in the chi-square model has been studied Donoho and Jin (2004),
 315 to the best of our knowledge, asymptotic limits of the support recovery problems
 316 have not been studied. The chi-squared model and the motivating GWAS application
 317 are analyzed in Chap. 7. The results obtained therein help us explain a phenomenon
 318 in GWAS where statistical power decays sharply as function of sample size when the
 319 latter is in a small region known as the *steep part of the power curve*. This empirical
 320 fact has long been observed by statistical geneticists but has not been mathemati-

321 cally quantified. Gao et al. (2019) provide further details on the power and design in
322 GWAS as well as an accompanying interactive statistical software (Gao 2019).

323 2.4 Relationships Between the Asymptotic Risks

324 We now elaborate on the relationship between statistical risks, as promised in
325 Sect. 2.1. The first lemma concerns the asymptotic relationship between the proba-
326 bility of exact recovery (2.9) and the risk of exact support recovery (2.8).

327 **Lemma 2.1** *For any sequence of procedures for support recovery $\mathcal{R} = \mathcal{R}_p$, we have*

$$\mathbb{P}[\widehat{S} = S] \rightarrow 1 \iff \text{risk}^E(\mathcal{R}) \rightarrow 0, \quad (2.27)$$

329 and

$$\mathbb{P}[\widehat{S} = S] \rightarrow 0 \implies \liminf \text{risk}^E(\mathcal{R}) \geq 1, \quad (2.28)$$

331 as $p \rightarrow \infty$. Dependence on p and \mathcal{R} was suppressed for notational convenience.

332 **Proof** (Lemma 2.1) Notice that $\{\widehat{S} = S\}$ implies $\{\widehat{S} \subseteq S\} \cap \{\widehat{S} \supseteq S\}$, and therefore
333 we have for every fixed p ,

$$\text{risk}^E = 2 - \mathbb{P}[\widehat{S} \subseteq S] - \mathbb{P}[S \subseteq \widehat{S}] \leq 2 - 2\mathbb{P}[\widehat{S} = S]. \quad (2.29)$$

335 On the other hand, since $\{\widehat{S} \neq S\}$ implies $\{\widehat{S} \not\subseteq S\} \cup \{\widehat{S} \not\supseteq S\}$, we have for every
336 fixed p ,

$$1 - \mathbb{P}[\widehat{S} = S] = \mathbb{P}[\widehat{S} \neq S] \leq 2 - \mathbb{P}[\widehat{S} \subseteq S] - \mathbb{P}[S \subseteq \widehat{S}] = \text{risk}^E. \quad (2.30)$$

338 Relation (2.27) follows from (2.29) and (2.30), and Relation (2.28) from (2.30). \square

339 By virtue of Lemma 2.1, it is sufficient to study the probability of exact support
340 recovery $\mathbb{P}[\widehat{S} = S]$ in place of risk^E , if we are interested in the asymptotic properties
341 of the risk in the sense of (2.25) and (2.26).

Keen readers must have noticed the asymmetry in Relation (2.28) when we dis-
cussed the relationship between the exact support recovery risk (2.8) and the proba-
bility of exact support recovery (2.9). While a trivial procedure that never rejects and
a procedure that always rejects both have risk^E equal to 1, the converse is not true. For
example, it is possible that a procedure selects the true index set S with probability
1/2, but otherwise makes one false inclusion *and* one false omission simultaneously.
In this case, the procedure will have

$$\text{risk}^E = 1, \quad \text{and} \quad \mathbb{P}[\widehat{S} = S] = 1/2,$$

342 showing that the converse of Relation (2.28) is in fact false.

The same argument applies to risk^A: a procedure may select the true index set S with probability 1/2, but makes enough false inclusions and omissions the rest of the time, so that risk^A is equal to one. Therefore, although the class of methods with risks equal to or exceeding 1 certainly contains the trivial procedures that we mentioned, they are not necessarily “useless” as some researchers have claimed (cf. Remark 2 in Arias-Castro and Chen 2017).

Upper and lower bounds for FDR and FNR can be immediately derived by replacing the numerators in (2.5) with the Hamming loss,

$$\mathbb{E} \left[\frac{H(\widehat{S}, S)}{\max\{|\widehat{S}|, |S|, 1\}} \right] \leq \text{FDR} + \text{FNR} \leq \mathbb{E} \left[\frac{H(\widehat{S}, S)}{\max\{\min\{|\widehat{S}|, |S|\}, 1\}} \right]. \quad (2.31)$$

Therefore, it is sufficient, but not necessary, that the Hamming loss vanishes in order to have vanishing approximate support recovery risks (2.7).

Turning to the relationship between the probability of exact support recovery (2.9) and Hamming loss (2.10), we point out a natural lower bound of the former using the expectation of the latter,

$$\mathbb{P}[\widehat{S} = S] \geq 1 - \mathbb{E}[H(\widehat{S}, S)] = 1 - \sum_{i=1}^p \mathbb{E} |\mathbb{1}_{\widehat{S}}(i) - \mathbb{1}_S(i)|. \quad (2.32)$$

A key observation in Relation (2.32) is that the expected Hamming loss decouples into p terms, and the dependence of the estimates $\mathbb{1}_{\widehat{S}}(i)$ among the p locations no longer plays any role in the sum. Therefore, studying support recovery problems via the expected Hamming loss is not very informative especially under severe dependence, as the bound (2.32) may become very loose. Vanishing Hamming loss is again sufficient, but not necessary for $\mathbb{P}[\widehat{S} = S]$ or the exact support recovery risk to go to zero.

2.5 The Asymptotic Generalized Gaussian (AGG) Models

We introduce a fairly general class of distributions known as asymptotic generalized Gaussians AGG. We also state some of their tail properties which play important roles in the analysis of phase transitions of high-dimensional testing problems.

Definition 2.6 A distribution F is called asymptotic generalized Gaussian with parameter $v > 0$ (denoted AGG(v)) if

- 371 1. $F(x) \in (0, 1)$ for all $x \in \mathbb{R}$ and
- 372 2. $\log \overline{F}(x) \sim -\frac{1}{v}x^v$ and $\log F(-x) \sim -\frac{1}{v}(-x)^v$,

373 where $\overline{F}(x) = 1 - F(x)$ is the survival function, and $a(x) \sim b(x)$ is taken to mean
374 $\lim_{x \rightarrow \infty} a(x)/b(x) = 1$.



375 The AGG models include, for example, the standard Gaussian distribution ($\nu = 2$)
 376 and the Laplace distribution ($\nu = 1$) as special cases. Since the requirement is only
 377 placed on the tail behavior, this class encompasses a large variety of light-tailed
 378 models. This class is commonly used in the literature on high-dimensional testing
 379 (Cai et al. 2007; Arias-Castro and Chen 2017).

380 **Proposition 2.1** *The $(1/p)$ -th upper quantile of AGG(ν) is*

$$381 \quad u_p := F^{-1}(1 - 1/p) \sim (\nu \log p)^{1/\nu}, \quad \text{as } p \rightarrow \infty, \quad (2.33)$$

382 where $F^{-1}(q) = \inf_x \{x : F(x) \geq q\}$, $q \in (0, 1)$.

Proof (Proposition 2.1) By the definition of AGG, for any $\epsilon > 0$, there is a constant $C = C(\epsilon)$ such that for all $x \geq C$, we have

$$-\frac{1}{\nu}x^\nu(1 + \epsilon) \leq \log \bar{F}(x) \leq -\frac{1}{\nu}x^\nu(1 - \epsilon).$$

383 Therefore, for all $x < x_l := ((1 + \epsilon)^{-1}\nu \log p)^{1/\nu}$, we have

$$384 \quad -\log p = -\frac{1}{\nu}x_l^\nu(1 + \epsilon) \leq \log \bar{F}(x_l) \leq \log \bar{F}(x), \quad (2.34)$$

385 and for all $x > x_u := ((1 - \epsilon)^{-1}\nu \log p)^{1/\nu}$, we have

$$386 \quad \log \bar{F}(x) \leq \log \bar{F}(x_u) \leq -\frac{1}{\nu}x_u^\nu(1 - \epsilon) = -\log p. \quad (2.35)$$

387 **Add:**
By definition of generalized inverse,

$$388 \quad u_p := F^{-1}(1 - 1/p) = \inf\{x : \bar{F}(x) \leq 1/p\} = \inf\{x : \log \bar{F}(x) \leq -\log p\}.$$

We know from relations (2.34) and (2.35) that

$$[x_u, +\infty) \subseteq \{x : \log \bar{F}(x) \leq -\log p\} \subseteq [x_l, +\infty),$$

389 and so $x_l \leq u_p \leq x_u$, and the expression for the quantiles follows. \square

390 2.6 Rapid Variation and Relative Stability

391 The behavior of the maxima of identically distributed random variables has been
 392 studied extensively in the extreme value theory literature (see, e.g., Leadbetter et al.
 393 1983; Resnick 2013; Embrechts et al. 2013; De Haan and Ferreira 2007). The concept
 394 of rapid variation plays an important role in the light-tailed case.

395 **Definition 2.7** (*Rapid variation*) The survival function of a distribution, $\bar{F}(x) =$
 396 $1 - F(x)$, is said to be rapidly varying if

$$397 \lim_{x \rightarrow \infty} \frac{\bar{F}(tx)}{\bar{F}(x)} = \begin{cases} 0, & t > 1 \\ 1, & t = 1 \\ \infty, & 0 < t < 1 \end{cases}. \quad (2.36)$$

398 When $F(x) < 1$ for all finite x , Gnedenko (1943) showed that the distribution F
 399 has rapidly varying tails if and only if the maxima of independent observations from
 400 F are *relatively stable* in the following sense.

401 **Definition 2.8** (*Relative stability*) Let $\epsilon_p = (\epsilon_p(i))_{i=1}^p$ be a sequence of random
 402 variables with common marginal distribution F . Define the sequence $(u_p)_{p=1}^\infty$ to be
 403 the $(1 - 1/p)$ -th generalized quantile of F , i.e.,

$$404 u_p = F^\leftarrow(1 - 1/p). \quad (2.37)$$

405 The triangular array $\mathcal{E} = \{\epsilon_p, p \in \mathbb{N}\}$ is said to have relatively stable (RS) maxima
 406 if

$$407 \frac{1}{u_p} M_p := \frac{1}{u_p} \max_{i=1,\dots,p} \epsilon_p(i) \xrightarrow{\mathbb{P}} 1, \quad \text{as } p \rightarrow \infty. \quad (2.38)$$

408 In the case of independent and identically distributed $\epsilon_p(i)$'s, Barndorff-Nielsen
 409 (1963), Resnick and Tomkins (1973) obtained necessary and sufficient conditions for
 410 the *almost sure stability* of maxima, where the convergence in (2.38) holds almost
 411 surely. See also Klass (1984) for further sharp results on almost sure stability, and
 412 Naveau (2003) for almost sure stability in stationary sequences. Here, we will only
 413 need the weaker notion in (2.38) but extend our inquiry to the case of dependent
 414 $\epsilon_p(i)$'s.

415 While relative stability (and almost sure stability) is well understood in the inde-
 416 pendent case, the role of dependence has not been fully explored. We start this
 417 investigation with a small refinement of Theorem 2 in Gnedenko (1943) valid under
 418 *arbitrary dependence*.

419 **Proposition 2.2** (*Rapid variation and relative stability*) Assume that the array \mathcal{E}
 420 consists of identically distributed and possibly dependent random variables with
 421 cumulative distribution function F , where $F(x) < 1$ for all finite $x > 0$.

422 1. If F has rapidly varying right tail in the sense of (2.36), then for all $\delta > 0$,

$$423 \mathbb{P} \left[\frac{1}{u_p} M_p \leq 1 + \delta \right] \geq 1 - \frac{\bar{F}((1 + \delta)u_p)}{\bar{F}(u_p)} \rightarrow 1. \quad (2.39)$$

424 2. If the array \mathcal{E} has independent entries, then it is relatively stable if and only if F
 425 has rapidly varying tail, i.e., (2.36) holds.



426 **Proof** (Proposition 2.2) By the union bound and the fact that $p\bar{F}(u_p) \leq 1$, we have

427

$$\mathbb{P}[M_p > (1 + \delta)u_p] \leq p\bar{F}((1 + \delta)u_p) \leq \frac{\bar{F}((1 + \delta)u_p)}{\bar{F}(u_p)}. \quad (2.40)$$

428

429 In view of (2.36) (rapid variation) and the fact that $u_p \rightarrow \infty$, as $p \rightarrow \infty$, the
430 right-hand side of (2.40) vanishes as $p \rightarrow \infty$, for all $\delta > 0$. This completes the
431 proof of (2.39). Part 2 is a re-statement of a classic result dating back to
432 Gnedenko (1943). \square

433 **Remark 2.1** Part (1) of Proposition 2.2 is equivalent to

434

$$\mathbb{P}\left[\frac{1}{u_p}M_p > 1 + \delta_p\right] \longrightarrow 0, \quad \text{as } p \rightarrow \infty, \quad (2.41)$$

435 for some positive sequence $\delta_p \rightarrow 0$. Notice, on the other hand, that if M_p^* is the
436 maximum of p iid variables with distribution F , the relative stability property entails
437 $M_p^*/u_p \rightarrow 1$, in probability, as $p \rightarrow \infty$. Since the sequence $1 + \delta_p \rightarrow 1$, Relation
438 (2.41) means that the rate of growth of the maxima M_n in \mathcal{E} cannot be faster than
439 that of the independent maxima M_p^* . This somewhat surprising fact holds regardless
440 of the dependence structure of \mathcal{E} and is solely a consequence of the rapid variation
441 of F .

442 We demonstrate next that the Gaussian, Exponential, Laplace, and Gamma dis-
443 tributions all have rapidly varying tails.

444 **Example 2.1** (*Generalized AGG*) A distribution is said to have *Generalized AGG*
445 right tail, if $\log \bar{F}$ is regularly varying,

446

$$\log \bar{F}(x) = -x^\nu L(x), \quad (2.42)$$

447 where $\nu > 0$ and $L : (0, +\infty) \rightarrow (0, +\infty)$ is a slowly varying function. (A function
448 is said to be slowly varying if $\lim_{x \rightarrow \infty} L(tx)/L(x) = 1$ for all $t > 0$.) Note that the
449 AGG(ν) model corresponds to the special case where $L(x) \rightarrow 1/\nu$, as $x \rightarrow \infty$.

Relation (2.39) holds for all arrays \mathcal{E} with *generalized* AGG marginals; if the entries are independent, the maxima are relatively stable. This follows directly from Proposition 2.2, once we show that F has rapidly varying tail. Indeed, by (2.42), we have

$$\log \left(\frac{\bar{F}(tx)}{\bar{F}(x)} \right) = -L(x)x^\nu \left(t^\nu \frac{L(tx)}{L(x)} - 1 \right),$$

450 which converges to $-\infty$, 0, and $+\infty$, as $x \rightarrow \infty$, when $t > 1$, $t = 1$, and $t < 1$,
451 respectively, since $x^\nu L(x) \rightarrow \infty$ as $x \rightarrow \infty$ by definition of L .

452 The AGG class encompasses a wide variety of rapidly varying tail models such
453 as Gaussian and double exponential distributions. The larger class (2.42) is needed,
454 however, for the Gamma distribution.

455 More generally, distributions with heavier tails (e.g., log-normal) and lighter tails
 456 (e.g., Gompertz) outside the generalized AGG class (2.42) may also possess rapidly
 457 varying tails; heavy-tailed distributions like the Pareto and t-distributions, on the
 458 other hand, do not. More details on these alternative classes of models can be found
 459 in Appendix B.

460 2.7 Auxiliary Facts About Gaussian Distributions

461 We end this chapter with several facts about univariate and multivariate Gaussian
 462 distributions that will be used in the rest of the manuscript.

463 **Relative stability.** We first state the relative stability of iid standard Gaussian random
 464 variables. Since the standard Gaussian distribution falls in the class of asymptotically
 465 generalized Gaussians (AGG; see Definition 2.6), by Example 2.1, we know that the
 466 triangular array $\mathcal{E} = \{(\epsilon_p(i))_{i=1}^p, p \in \mathbb{N}\}$ has relatively stable (RS) maxima in the
 467 sense of (2.38), i.e.,

$$468 \frac{1}{u_p} \max_{i=1,\dots,p} \epsilon_p(i) \xrightarrow{\mathbb{P}} 1, \quad \text{as } p \rightarrow \infty, \quad (2.43)$$

469 where u_p is the $(1/p)$ -th upper quantile as defined in (2.33). Similarly, since the
 470 array \mathcal{E} has distributions symmetric around 0, it also has relatively stable minima

$$471 \frac{1}{u_p} \min_{i=1,\dots,p} \epsilon_p(i) \xrightarrow{\mathbb{P}} -1, \quad \text{as } p \rightarrow \infty. \quad (2.44)$$

472 The convergence in (2.43) also holds almost surely.

473 **Mill's ratio.** We give next the well-known bounds for Mill's ratio of Gaussian tails.
 474 Let Φ denote the CDF of the standard Gaussian distribution and ϕ its density. One
 475 can show that for all $x > 0$ we have

$$476 \frac{x}{1+x^2} \phi(x) \leq \bar{\Phi}(x) = 1 - \Phi(x) \leq \frac{1}{x} \phi(x), \quad (2.45)$$

477 using, e.g., integration by parts. Note that this fact may be used to verify the rapid
 478 variation of Φ , which entails the relative stability property above.

479 **Stochastic monotonicity.** The third fact is the stochastic monotonicity of the Gaussian
 480 location family. In fact, for all location families $\{F_\delta(x)\}_\delta$ where $F_\delta(x) = F(x - \delta)$, we have

$$482 F_{\delta_1}(t) \geq F_{\delta_2}(t), \quad \text{for all } t \in \mathbb{R} \quad \text{and all } \delta_1 \leq \delta_2. \quad (2.46)$$

483 Relation (2.46) holds, of course, when F is the standard Gaussian distribution.

Slepian's lemma and the Sudakov–Fernique inequality. The following two results will be instrumental in our characterization of uniform relative stability for Gaussian triangular arrays in Chap. 6. The first is the celebrated result due to Slepian (1962).

Theorem 2.1 (Slepian's lemma) *Let $\epsilon = (\epsilon(i))_{i=1}^p$ and $\eta = (\eta(i))_{i=1}^p$ be two multivariate normally distributed random vectors with zero means $\mathbb{E}[\epsilon(i)] = \mathbb{E}[\eta(i)] = 0$. If for all $i, j = 1, \dots, p$, we have*

$$\mathbb{E}[\epsilon(i)^2] = \mathbb{E}[\eta(i)^2], \quad \text{and} \quad \text{Cov}(\epsilon(i), \epsilon(j)) \leq \text{Cov}(\eta(i), \eta(j)),$$

then $\epsilon \stackrel{st}{\geq} \eta$, i.e.,

$$\mathbb{P}[\epsilon(i) \leq x_i, i = 1, \dots, p] \leq \mathbb{P}[\eta(i) \leq x_i, i = 1, \dots, p].$$

This result implies, in particular, that $M_\epsilon := \max_{i=1, \dots, p} \epsilon(i)$ dominates stochastically $M_\eta := \max_{i=1, \dots, p} \eta(i)$ in the sense that

$$\mathbb{P}[M_\eta > u] \leq \mathbb{P}[M_\epsilon > u], \quad \text{for all } u \in \mathbb{R}. \quad (2.47)$$

In this case, we shall write $M_\eta \stackrel{d}{\leq} M_\epsilon$. This result shows, for example, that the maximum of iid Gaussians is stochastically larger than the maximum of any positively correlated Gaussian vector with the same marginal distributions.

Slepian's lemma can be obtained as a corollary from the general Normal Comparison Lemma (see, e.g., Theorem 4.2.1 on page 81 in Leadbetter et al. 1983). See also Chap. 2 in Adler and Taylor (2009).

The following result, known as the Sudakov–Fernique inequality, is similar in spirit to Slepian's lemma but it does not assume that the Gaussian vectors are centered and yield a weaker conclusion—an inequality between expectations. For a proof, many insights, and, in fact, a more general result, see, e.g., Theorem 2.2.5 on page 61 in Adler and Taylor (2009).

Theorem 2.2 (The Sudakov–Fernique inequality) *Let $\epsilon = (\epsilon(i))_{i=1}^p$ and $\eta = (\eta(i))_{i=1}^p$ be two multivariate normally distributed random vectors.*

If for all $i, j = 1, \dots, p$, we have

$$\mathbb{E}[\epsilon(i)] = \mathbb{E}[\eta(i)] \quad \text{and} \quad \mathbb{E}[(\eta(i) - \eta(j))^2] \leq \mathbb{E}[(\epsilon(i) - \epsilon(j))^2],$$

then for $M_\epsilon = \max_{i=1, \dots, p} \epsilon(i)$ and $M_\eta = \max_{i=1, \dots, p} \eta(i)$, we have

$$\mathbb{E}[M_\eta] \leq \mathbb{E}[M_\epsilon].$$

With these conceptual and technical preparations, we are ready to discuss the high-dimensional signal detection and support recovery problems in the next chapter.



Author Queries

Chapter 2

Query Refs.	Details Required	Author's response
AQ1	Please check the usage of the term “to fo to” in the sentence “Vanishing Hamming loss is again...”.	This was a typo, thanks! Please replace “to fo to” by “to go to”, as indicated in the margins therein.

Metadata of the chapter that will be visualized in SpringerLink

Book Title	Concentration of Maxima and Fundamental Limits in High-Dimensional Testing and Inference	
Series Title		
Chapter Title	A Panorama of Phase Transitions	
Copyright Year	2021	
Copyright HolderName	The Author(s), under exclusive license to Springer Nature Switzerland AG	
Corresponding Author	Family Name Particle Given Name Prefix Suffix Role Division Organization Address Email ORCID	Gao Zheng Department of Statistics University of Michigan-Ann Arbor Ann Arbor, MI, USA gaozheng@umich.edu https://orcid.org/0000-0003-1394-8741
Author	Family Name Particle Given Name Prefix Suffix Role Division Organization Address Email ORCID	Stoev Stilian Department of Statistics University of Michigan-Ann Arbor Ann Arbor, MI, USA https://orcid.org/0000-0003-2108-5421
Abstract	The purpose of this chapter is to provide a unified review of a wide variety of phase-transition results in the detection and exact support recovery problems for sparse signals. In order to emphasize ideas, the focus is on the simple case of a high-dimensional sparse signal observed with additive independent Gaussian errors. The classic phase-transition result for the signal detection problem obtained first by Yuri Izmailovich Ingster as well as very recent results on approximate support recovery are reviewed. The unified approach based on the four different risk functionals introduced in Chap. 2 yields a variety of phase-transition results, many of which are new. The optimality and sub-optimality of some popular support estimation procedures are also established. The chapter provides a concise and yet complete account of the phase transitions in the context of sparse signal detection and support estimation in high dimensions.	

Chapter 3

A Panorama of Phase Transitions



0 The purpose of this chapter is to provide a unified review of the fundamental statistical
1 limits in the sparse signal detection and support recovery problems. Our goal is to
2 convey the main ideas and thus we shall focus on the simple but important setting of
3 independent Gaussian errors. Specifically, we derive the conditions under which the
4 detection and support recovery problems succeed and fail in the sense of (2.25) and
5 (2.26), in the additive error model

$$6 \quad x(i) = \mu(i) + \epsilon(i), \quad i = 1, \dots, p, \quad (3.1)$$

7 where the errors $\epsilon(i)$'s are iid standard Gaussians random variables. Once again, we
8 restrict our analysis to models with independent and identically distributed Gaussian
9 errors for the moment. Both the distributional assumption and the independence
10 assumption will be relaxed substantially in the following chapters.

As laid out in Sect. 2.3, we work under the asymptotic regime where the problem dimension p diverges to infinity. The set of non-zero entries of the signal vector $\mu = \mu_p$ will be referred to as its *support* and denoted by

$$S_p := \{i : \mu(i) \neq 0\}.$$

11 We shall assume that the size of the support is

$$12 \quad |S_p| = \lfloor p^{1-\beta} \rfloor, \quad \beta \in (0, 1], \quad (3.2)$$

13 where β parametrizes the problem sparsity. A more general parameterization of the
14 support involving a slowly varying function is considered in Chap. 4.

The closer β to 1, the sparser the support S_p . Conversely, when β is close to 0, the support is dense with many non-null signals. We consider one-sided alternatives (2.14), and parametrize the range of the non-zero (and perhaps unequal) signals with

$$\underline{\Delta} = \sqrt{2\underline{r} \log p} \leq \mu(i) \leq \bar{\Delta} = \sqrt{2\bar{r} \log p}, \quad \text{for all } i \in S_p, \quad (3.3)$$

for some constants $0 < \underline{r} \leq \bar{r} \leq +\infty$.

The parametrization of signal sparsity (3.2) and signal sizes (3.3) in the Gaussian model was first introduced in Ingster (1998), and later adopted by Hall and Jin (2010), Cai et al. (2011), Zhong et al. (2013), Cai and Wu (2014), Arias-Castro and Wang (2017), and numerous others for studying the signal detection problem in Gaussian location-scale models. Similar scalings of sparsity and signal size are also used in, e.g., Ji and Jin (2012), Jin et al. (2014), Butucea et al. (2018) to study the phase transitions of the support recovery problems under Gaussianity assumptions.

It should be noted that the “classical” setting where all signals are of equal size is not the only one that have been studied. The recent contribution of Li and Fithian (2020) investigates the signal detection problem in a more realistic setting where the signals are drawn from a general and potentially polynomial-tailed distribution. The study of such general settings in both detection and support recovery problems is an interesting new direction of research.

3.1 Sparse Signal Detection Problems

The optimality of sparse signal detection was first studied by Ingster (1998), who showed that a phase transition in the r - β plane exists for the signal detection problem. Specifically, consider the so-called *detection boundary* function:

$$f_D(\beta) = \begin{cases} \max\{0, \beta - 1/2\} & 0 < \beta \leq 3/4 \\ (1 - \sqrt{1 - \beta})^2 & 3/4 < \beta \leq 1. \end{cases} \quad \beta \in (0, 1]. \quad (3.4)$$

Assume that the non-zero signal sizes are all equal and parameterized as $\sqrt{2r \log p}$. If the signal size parameter r is *above* the detection boundary, i.e., $r > f_D(\beta)$, then the global null hypothesis $\mu(i) = 0$ for all $i = 1, \dots, p$ can be distinguished from the alternative as $p \rightarrow \infty$ in the sense of (2.25) using the likelihood ratio test. Otherwise, when the signal sizes fall below the boundary, i.e., $r < f_D(\beta)$, no test can do better than a random guess. We visualize the detection boundary in the upper panel of Fig. 3.1.

Adaptive tests such as Tukey’s HC in (2.19) (Donoho and Jin 2004) and a modified goodness-of-fit test statistic of Zhang (2002) have been identified to attain this performance limit without knowledge of the sparsity and signal sizes. It is also known that the max-statistic (2.16) is only efficient when $r > (1 + \sqrt{1 - \beta})^2$, and is therefore sub-optimal for denser signals where $1/2 \leq \beta \leq 3/4$; see Cai et al. (2011).

(Recently, Li and Fithian 2020 showed that in the more general setting where signals themselves are dispersed, the sub-optimality of the max-statistic disappears in the detection problem.) In contrast, the sum-of-square-type statistics such as L_2 was shown in Fan (1996) to be asymptotically powerless when the L_2 -norm of the signal $\|\mu\|_2^2$ is $o(\sqrt{p})$ or, equivalently, when $\beta > 1/2$ in our parametrization.

Notice that the scaling for the signal magnitude $\Delta = \sqrt{2r \log p}$ is useful for studying very sparse signals ($\beta > 1/2$), but fails to reveal the difficulties of the detection problems when the signals are relatively dense ($\beta < 1/2$). This is because $f_D(\beta) = 0$, $\beta \in (0, 1/2]$. Thus, a different scaling is needed to study the regime of small but dense signals. In this case, with slight overloading of notation, we parametrize signal sizes as

$$\Delta = p^r \leq \mu(i) \leq \bar{\Delta} = p^{\bar{r}}, \quad \text{for all } i \in S_p, \quad (3.5)$$

where r and \bar{r} are negative constants and the signal magnitude vanishes, as $p \rightarrow \infty$. In this scaling, for the so-called faint signal regime, Cai et al. (2011) established a phase-transition result characterized by the following boundary:

$$f_{D'}(\beta) = \beta - 1/2, \quad 0 < \beta \leq 1/2. \quad (3.6)$$

Specifically, if $\bar{r} < f_{D'}(\beta)$, the signal detection fails in the sense of (2.26) regardless of the procedures, while the HC statistic continues to attain asymptotically perfect detection when $r > f_{D'}(\beta)$. We visualize this boundary in the lower panel of Fig. 3.1.

To the best of our knowledge, the performance of simple statistics such as L_1 , L_2 norms, and the sum statistic T in (2.15) in this weak signal setting have not been reported in the literature. Our first theorem establishes the performance of these simple but popular statistics for detecting sparse signals in high dimensions, and summarizes the known results.

Theorem 3.1 Consider the signal detection problem in the triangular array of Gaussian error models (3.1) where the sparsity is parametrized as in (3.2).

(i) For $\beta \in (1/2, 1)$ and growing signal sizes as in (3.3), the statistics L_1 , L_2 , and T are asymptotically powerless in the sense of (2.26).

(ii) For $\beta \in (0, 1/2]$ and growing signal sizes as in (3.3), the statistics L_1 , L_2 , and T solve the detection problem in the sense of (2.25).

(iii) For dense and faint signals, i.e., $\beta \in (0, 1/2]$ under the parameterization (3.5), the sum statistic T attains the optimal detectability boundary in (3.6). That is, tests based on the sum statistic T can succeed asymptotically in the sense of (2.25) when $r > \beta - 1/2$.

(iv) In the dense and faint signal setting of (iii), the L_1 and L_2 statistics are both sub-optimal. More precisely, they succeed in the sense of (2.25) when $r > \beta/2 - 1/4$, but fail in the sense of (2.26) when $\bar{r} < \beta/2 - 1/4$.

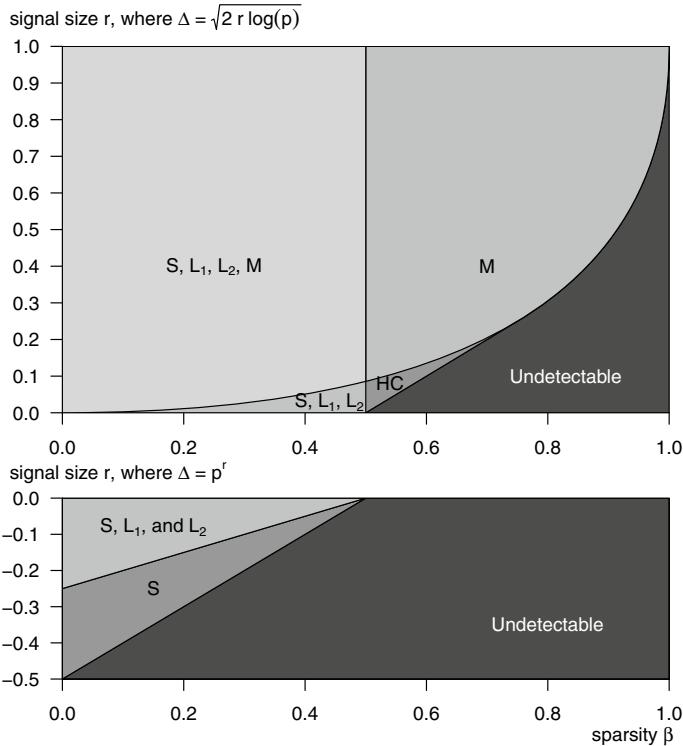


Fig. 3.1 The phase diagrams of the sparse signal detection problem. Signal size and sparsity are parametrized by r and β , respectively. The diagrams illustrate the regions where the signal detection problem can be solved asymptotically by some of the commonly used statistics: the maximum (M), the sum-of-squares (L_2), the sum-of-absolute values (L_1), and the sum (T). In each region of the diagram, the annotated statistics can make the detection risk (2.4) vanish, as dimension p diverges. Conversely, the risks has \liminf at least one. The detection problem is unsolvable for very sparse and weak signals in the undetectable regions. Notice that the L_1 and L_2 statistics are, in fact, sub-optimal for all sparsity levels. On the other hand, the max-statistic remains powerful for sparse signals ($\beta > 1/2$), and is fully efficient when the problem is very sparse ($\beta \geq 3/4$). The HC statistic can detect signals in all configurations in the detectable regions; we explicitly marked the region where signals are only detectable by HC among the statistics considered. See the text and Theorem 3.1

87 **Proof** The claims in parts (i) and (ii) about the L_1 , L_2 , and the sum statistic T in the
 88 cases of diverging signal sizes (3.3) can be found in Fan (1996), Candès (2018). We
 89 prove here the statements for the cases (iii) and (iv), where the signals are dense and
 90 small, as parametrized in (3.5).

91 For simplicity of the exposition, we will suppose that in (3.5) we have $\underline{r} = r = \bar{r}$,
 92 so that $\mu(i) = p^r$. The general case where $\underline{r} < \bar{r}$ is left as an exercise.

93 *Part (iii):* We first show that the sum statistic T or, equivalently, the simple arithmetic
 94 mean attains the sparse signal detection boundary. By the normality and inde-
 95 pendence of the summands, we have

$$96 \quad \frac{1}{\sqrt{p}} \sum_{i=1}^p x(i) \sim \begin{cases} N(0, 1), & \text{under } H_0 \\ N(p^{(r-\beta)+1/2}, 1), & \text{under } H_1. \end{cases} \quad (3.7)$$

97 It immediately follows that the two distributions can be distinguished perfectly if
 98 $p^{r-(\beta-1/2)}$ diverges, i.e., $r > \beta - 1/2$. This can be seen by simply setting the rejection
 99 region at $(p^{(r-\beta)+1/2}/2, +\infty)$ for the scaled statistic $\sum_{i=1}^p x(i)/\sqrt{p}$. According to
 100 the lower bound on the performance limit in detection problems (see Theorem 8 in
 101 Cai et al. 2011), we have shown that T attains the optimal detection boundary (3.6).

102 Part (iv): We now turn to the L_2 -norm statistic. Recall a non-central chi-square
 103 random variable $\chi_k^2(\lambda)$ has mean $(k + \lambda)$ and variance $2(k + 2\lambda)$. Since the observa-
 104 tions have distributions $N(0, 1)$ under the null and $N(p^r, 1)$ under the alternative, we
 105 have $x^2(i) \sim \chi_1^2(0)$ for $i \notin S$ and $x^2(i) \sim \chi_1^2(p^{2r})$ for $i \in S$. Therefore, the mean
 106 and variance of the (centered and scaled) L_2 statistics are

$$107 \quad \mathbb{E} \left[\frac{1}{\sqrt{p}} \sum_{i=1}^p (x(i)^2 - 1) \right] = \begin{cases} 0 & \text{under } H_0 \\ p^{1-\beta} p^{2r} p^{-1/2} = p^{1/2-\beta+2r} & \text{under } H_1, \end{cases} \quad (3.8)$$

108 and

$$109 \quad \text{Var} \left(\frac{1}{\sqrt{p}} \sum_{i=1}^p (x(i)^2 - 1) \right) = \begin{cases} \frac{1}{p} 2p = 2 & \text{under } H_0 \\ \frac{1}{p} (2p + 4p^{1-\beta+2r}) = 2(1 + 2p^{2r-\beta}) & \text{under } H_1, \end{cases} \quad (3.9)$$

110 respectively. By the central limit theorem, we have

$$111 \quad \frac{1}{\sqrt{2p}} \sum_{i=1}^p (x(i)^2 - 1) \implies N(0, 1), \quad (3.10)$$

112 under the null. On the other hand, under the alternative, since $p^{2r-\beta} \rightarrow 0$ for all
 113 $r < 0$ and $\beta > 0$, the variance in (3.9) converges to 2, as $p \rightarrow \infty$ and an application
 114 of the Lyapunov version of the CLT entails

$$115 \quad \frac{1}{\sqrt{2p}} \left(\sum_{i=1}^p (x(i)^2 - 1) - p^{1/2-\beta+2r} \right) \implies N(0, 1). \quad (3.11)$$

116 Hence, perfect detection with the L_2 -norm is possible if $p^{1/2-\beta+2r}$ diverges, i.e.,
 117 $r > \beta/2 - 1/4$. On the other hand, if $r < \beta/2 - 1/4$, the distributions of the (scaled)
 118 statistics merge under the null and the alternative.

119 The case of the L_1 -norm is treated similarly. Let $Y = |X|$ where $X \sim |N(\mu, 1)|$.
 120 Using the expressions for the mean and variance of Y (see, e.g., Tsagris et al. 2014),

$$\mu_Y = \mathbb{E}[Y] = \sqrt{\frac{2}{\pi}} e^{-\mu^2/2} + \mu(1 - \Phi(-\mu)), \quad (3.12)$$

$$\sigma_Y^2 = \text{Var}(Y) = \mu^2 + 1 - \mu_Y^2, \quad (3.13)$$

where Φ is the CDF of a standard normal random variable, we have, regardless of the value of μ ,

$$\sigma_Y^2 = \text{Var}(Y) = \mathbb{E}(Y - \mathbb{E}Y)^2 \leq \mathbb{E}(X - \mathbb{E}X)^2 = 1, \quad (3.14)$$

where the inequality holds because absolute value is a Lipschitz function with Lipschitz constant 1.

By the central limit theorem, we have

$$\frac{1}{\sqrt{p}} \left(\sum_{i=1}^p |x(i)| - \sqrt{\frac{2}{\pi}} \right) \implies N(0, 1 - 2/\pi) \quad (3.15)$$

under the null. On the other hand, when the alternative hypothesis holds, we have

$$\begin{aligned} \mathbb{E} \left[\frac{1}{\sqrt{p}} \left(\sum_{i=1}^p |x(i)| - \sqrt{\frac{2}{\pi}} \right) \right] &= \frac{p^{1-\beta}}{\sqrt{p}} \left[\left(\sqrt{\frac{2}{\pi}} e^{-\mu^2/2} + \mu(1 - 2\Phi(-\mu)) \right) - \sqrt{\frac{2}{\pi}} \right] \\ &= p^{1/2-\beta} \left[\sqrt{\frac{2}{\pi}} \left(e^{-p^{2r}/2} - 1 \right) + p^r (1 - 2\Phi(-\mu)) \right] \\ &= p^{1/2-\beta} \left[\sqrt{\frac{2}{\pi}} \left(-p^{2r}/2 - O(p^{4r}) \right) + p^r \left(\sqrt{\frac{2}{\pi}} p^r + O(p^{3r}) \right) \right] \\ &= p^{1/2-\beta} \sqrt{\frac{2}{\pi}} \left(p^{2r}/2 + O(p^{4r}) \right) \\ &= p^{1/2-\beta+2r} \sqrt{1/2\pi} + O(p^{1/2-\beta+4r}), \end{aligned}$$

and

$$\begin{aligned} \text{Var} \left(\frac{1}{\sqrt{p}} \left(\sum_{i=1}^p |x(i)| - \sqrt{\frac{2}{\pi}} \right) \right) &= \frac{1}{p} (p - p^{1-\beta})(1 - 2/\pi) + \frac{1}{p} p^{1-\beta} \sigma_Y^2 \\ &\rightarrow 1 - 2/\pi, \end{aligned}$$

by the boundedness of σ_Y^2 shown in (3.14). Again, by the Lyapunov version of the central limit theorem, we conclude asymptotic normality of the centered and scaled L_1 -norms under the alternative. In an entirely analogous argument to the L_2 -norm case, asymptotically perfect detection can be achieved if $p^{1/2-\beta+2r}$ diverges, i.e., $r > \beta/2 - 1/4$. On the other hand, when $r < \beta/2 - 1/4$, the two hypotheses cannot be told apart by the L_1 -norms since the distributions of the (scaled) statistics merge under the two hypotheses. \square



The portmanteau of results in Theorem 3.1 are visualized in Fig. 3.1. It is worth noting that the β - r parameter regions where L_1 and L_2 statistics are asymptotically powerful coincide, and these statistics are theoretically sub-optimal for both sparse regimes ($\beta > 1/2$) and relatively dense regimes ($\beta \leq 1/2$).

Ideas have been proposed to combine statistics that are powerful for different alternatives to create adaptive tests that maintain high power for all sparsity levels. Such adaptive tests can be constructed, for example, by leveraging the asymptotic independence of the sum- and supremum-type statistics (Hsing 1995). Recently, Xu et al. (2016) showed that for dependent observations under mixing and moment conditions, the sum-of-power-type statistics

$$\tilde{L}_q(x) = \sum_{i=1}^p x^q(i) \quad (3.16)$$

with distinct positive integer powers (i.e., $q = 1, 2, \dots$) are asymptotically jointly independent, and proposed an adaptive test that monitors the minimum p-value of tests constructed with \tilde{L}_q 's. This idea is further developed in Wu et al. (2019) for generalized linear models and in He et al. (2018) with U-statistics.

Optimality properties of such adaptive tests and the optimal choice of the q -combinations, however, remain open problems. Xu et al. (2016) suggested combining $q = 1, 2, 3, \dots, 6$, and $q = \infty$, based empirical evidence from numerical experiments. Theorem 3.1 here implies that, at least for detecting one-sided alternatives, the \tilde{L}_2 statistic (i.e., L_2 norm) and the L_1 norm are asymptotically dominated by the \tilde{L}_1 statistic (or, equivalently, the sum T). Therefore, it is sufficient to include only the latter in the construction of the adaptive test.

3.2 Sparse Signal Support Recovery Problems

Turning to support recovery problems in the Gaussian error model (3.1), in the rest of this chapter, we will analyze the asymptotic performance limits in terms of the risk metrics for exact, exact-approximate, approximate-exact support recovery problems (i.e., (2.8), (2.11), and (2.12), respectively), as well as the probability of support recovery (2.9). We will also review the recent result for exact support recovery risk (2.7) by Arias-Castro and Chen (2017) to reveal a rather complete landscape of support recovery problems in high-dimensional Gaussian error models.

In the rest of this chapter, we restrict our attention to the class of thresholding procedures. Specifically, the lower bounds that we develop in Theorems 3.2 through 3.5 are only meant to apply to thresholding procedures. Although it is intuitively appealing to consider only data-thresholding procedures in multiple testing problems, and such procedures are not always optimal in more general settings. The optimality of thresholding procedures and the consequences of this restriction will be treated in Chap. 5.

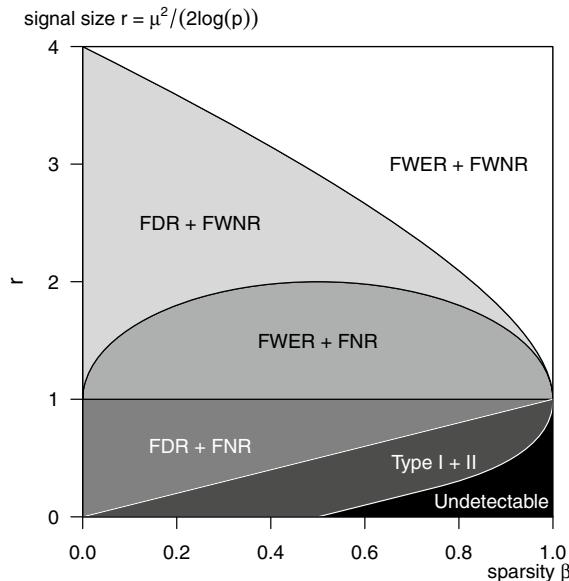


Fig. 3.2 The phase diagram of support recovery problems for the high-dimensional model (3.1), illustrating the boundaries of the exact support recovery (FWER + FWNR; top curve; Theorem 3.2), the approximate-exact support recovery (FDR + FWNR; second curve from top; Theorem 3.5), the exact-approximate support recovery (FWER + FNR; horizontal line $r = 1$; Theorem 3.4), and the approximate support recovery problems (FDR + FNR; tilted line $r = \beta$; Theorem 3.3). The signal detection problem (Type I + Type II errors of the global test; lower curve) was studied in Donoho and Jin (2004). In each region of the diagram and above, the annotated statistical risk can be made to vanish, as dimension p diverges. Conversely, the risks have \liminf at least one Add a “.”

Figure 3.2 illustrates the rich landscape of phase transitions in support recovery for the various choices of statistical risk for the family of thresholding estimators, established in the following sections. We end this brief overview with a technical notion needed in order to state our main results. We define a rate at which the nominal levels of FWER or FDR go to zero.

Definition 3.1 We say the nominal level of errors $\alpha = \alpha_p$ vanishes slowly, if

$$\alpha \rightarrow 0, \quad \text{and} \quad \alpha p^\delta \rightarrow \infty \text{ for any } \delta > 0. \quad (3.17)$$

As an example, the sequence of nominal levels $\alpha_p = 1/\log(p)$ is slowly vanishing, while the sequence $\alpha_p = 1/\sqrt{p}$ is not.



195 3.3 The Exact Support Recovery Problem

196 Our study of the exact support recovery risk (2.8) begins with a brief review of existing
 197 results for the Hamming loss (2.10). Indeed, as discussions in Sect. 2.3 suggest, the
 198 latter can be informative of the exact support recovery problems for models with
 199 independent components.

200 Inspired by the phase-transition results for the signal detection problem, Ji and
 201 Jin (2012), Genovese et al. (2012), Jin et al. (2014) derived interesting sharp results
 202 on support recovery problems in linear models under the Hamming loss $H(\widehat{S}, S)$.
 203 Specifically, these papers establish minimax-type phase-transition results in their
 204 respective settings. Under the sparsity parametrization in (3.2) and assuming equal
 205 signal sizes of $(2r \log p)^{1/2}$, Hamming losses were shown to diverge to $+\infty$ when r
 206 falls below the threshold

$$207 f_E(\beta) = (1 + (1 - \beta)^{1/2})^2, \quad (3.18)$$

208 for any method of support estimation. Conversely, under orthogonal, or near-
 209 orthogonal random designs, if $r > f_E(\beta)$, they showed that the methods they pro-
 210 posed achieve vanishing Hamming loss.

211 Very recently, Butucea et al. (2018) studied both asymptotics and non-asymptotics
 212 of support recovery problems in the additive noise model (3.1) under the assumption
 213 of equal signal sizes, using the Hamming loss. Again, the analysis of asymptotic
 214 optimality focused on a newly proposed procedure which is very specific to the
 215 Gaussian model. It is not at all clear if the optimality properties are a consequence
 216 of its mysterious construction.

217 We now show that commonly used and computationally efficient procedures can
 218 also be asymptotically optimal in the exact support recovery problem.

219 **Theorem 3.2** Consider the high-dimensional additive error model (3.1) under inde-
 220 pendent standard Gaussian errors, with signal sparsity and size as described in (3.2)
 221 and (3.3). The function (3.18) characterizes the phase transition of the exact support
 222 recovery problem. Namely, the following two results hold:

- 223 (i) If $r > f_E(\beta)$, then Bonferroni's, Sidák's, Holm's, and Hochberg's procedures
 224 with slowly vanishing nominal FWER levels (as defined in Definition 3.1) all
 225 achieve asymptotically exact support recovery in the sense of (2.25).
- 226 (ii) Conversely, if $\bar{r} < f_E(\beta)$, then for any thresholding procedure \widehat{S}_p , we have
 227 $\mathbb{P}[\widehat{S}_p = S_p] \rightarrow 0$. Therefore, in view of Lemma 2.1, exact support recovery asymp-
 228totically fails for all thresholding procedures in the sense of (2.26).

229 We illustrate this result with a β - r phase diagram in Fig. 3.2. Theorem 3.2 is, in
 230 fact, a special case of the more general Theorem 4.1, which covers dependent as
 231 well as Gaussian and non-Gaussian errors. We will study the exact support recovery
 232 problem in greater detail and generality in Chap. 4.

233 3.4 The Approximate Support Recovery Problem

234 Arias-Castro and Chen (2017) studied the performance of the Benjamini–Hochberg
 235 procedure (Benjamini and Hochberg 1995) and a stripped-down version of the
 236 Candés–Barber procedure (Barber and Candès 2015) in approximate support recov-
 237 ery problems when the components of the noise term ϵ in (3.1) have independent and
 238 symmetric distributions. A phase-transition phenomenon for the approximate sup-
 239 port recovery risk (2.7) was established in the Gaussian additive error model, where
 240 the two aforementioned methods are both shown to be asymptotically optimal.

241 The analysis therein, however, assumed equal signal sizes for the alternatives. We
 242 generalize the main results of Arias-Castro and Chen (2017) to allow for unequal
 243 signal sizes. The key to establishing this generalization is a monotonicity property
 244 of the BH procedure, presented in the following Sect. 3.5. Namely, the power of
 245 the BH procedure in terms of FNR monotonically increases for stochastically larger
 246 alternatives. This fact will be formalized in Lemma 3.2, and may be of independent
 247 interest.

248 **Theorem 3.3** *In the context of Theorem 3.2, the function*

$$249 \quad f_A(\beta) = \beta \quad (3.19)$$

the

250 characterizes the phase transition of approximate support recovery problem. Specif-
 251 cally the following two results hold:

252 (i) If $r > f_A(\beta)$, then the Benjamini–Hochberg procedure (defined in Sect. 2.2)
 253 with slowly vanishing nominal FDR levels (as defined in Definition 3.1) achieves
 254 asymptotically approximate support recovery in the sense of (2.25).

255 (ii) Conversely, if $\bar{r} < f_A(\beta)$, then approximate support recovery asymptotically
 256 fails in the sense of (2.26) for all thresholding procedures.

257 **Proof (Necessary condition in Theorem 3.3)** We first show part (ii). That is, when
 258 $\bar{r} < \beta$, no thresholding procedure is able to achieve approximate support recovery.
 259 The arguments are similar to that in Theorem 1 of Arias-Castro and Chen (2017),
 260 although we allow for unequal signal sizes.

261 Denote the distributions of $N(0, 1)$, $N(\underline{\Delta}, 1)$, and $N(\bar{\Delta}, 1)$ as F_0 , $F_{\underline{\Delta}}$, and $F_{\bar{\Delta}}$,
 262 respectively.

Recall that thresholding procedures are of the form

$$\widehat{S}_p = \{i \mid x(i) > t_p(x)\}.$$

263 Denote $\widehat{S} := \{i \mid x(i) > t_p(x)\}$, and $\widehat{S}(u) := \{i \mid x(i) > u\}$. For any threshold $u \geq t_p$,
 264 we must have $\widehat{S}(u) \subseteq \widehat{S}$, and hence

$$265 \quad \text{FDP} := \frac{|\widehat{S} \setminus S|}{|\widehat{S}|} \geq \frac{|\widehat{S} \setminus S|}{|\widehat{S} \cup S|} = \frac{|\widehat{S} \setminus S|}{|\widehat{S} \setminus S| + |S|} \geq \frac{|\widehat{S}(u) \setminus S|}{|\widehat{S}(u) \setminus S| + |S|}. \quad (3.20)$$



266 On the other hand, for any threshold $u \leq t_p$ we must have $\widehat{S}(u) \supseteq \widehat{S}$, and hence

$$267 \quad \text{NDP} := \frac{|S \setminus \widehat{S}|}{|S|} \geq \frac{|S \setminus \widehat{S}(u)|}{|S|}. \quad (3.21)$$

268 Since either $u \geq t_p$ or $u \leq t_p$ must take place, putting (3.20) and (3.21) together, we
269 have

$$270 \quad \text{FDP} + \text{NDP} \geq \frac{|\widehat{S}(u) \setminus S|}{|\widehat{S}(u) \setminus S| + |S|} \wedge \frac{|S \setminus \widehat{S}(u)|}{|S|}, \quad (3.22)$$

271 for any u . Therefore, it suffices to show that for a suitable choice of u , the RHS of
272 (3.22) converges to 1 in probability; the desired conclusion on FDR and FNR follows
273 by the dominated convergence theorem.

274 Let $t^* = \sqrt{2q \log p}$ for some fixed q , we obtain an estimate of the tail probability
275 by Mill's ratio (2.45),

$$276 \quad \overline{F}_0(t^*) \sim \frac{1}{t^*} \phi(t^*) = \frac{1}{2\sqrt{\pi q \log p}} p^{-q}, \quad (3.23)$$

where $a_p \sim b_p$ is taken to mean $a_p/b_p \rightarrow 1$. Observe that $|\widehat{S}(t^*) \setminus S|$ has distribution
Binom($p - s, \overline{F}_0(t^*)$) where $s = |S|$, denote $X = X_p := |\widehat{S}(t^*) \setminus S|/|S|$, and we
have

$$\mu := \mathbb{E}[X] = \frac{(p - s)\overline{F}_0(t^*)}{s}, \quad \text{and} \quad \text{Var}(X) = \frac{(p - s)\overline{F}_0(t^*)F_0(t^*)}{s^2} \leq \mu/s.$$

277 Therefore, for any $M > 0$, we have, by Chebyshev's inequality,

$$278 \quad \mathbb{P}[X < M] \leq \mathbb{P}[|X - \mu| > \mu - M] \leq \frac{\mu/s}{(\mu - M)^2} = \frac{1/(\mu s)}{(1 - M/\mu)^2}. \quad (3.24)$$

Now, from the expression of $\overline{F}_0(t^*)$ in (3.23), we obtain

$$\mu = (p^\beta - 1)\overline{F}_0(t^*) \sim \frac{1}{2\sqrt{\pi q \log p}} p^{\beta-q}.$$

279 Since $\bar{r} < \beta$, we can pick q such that $\bar{r} < q < \beta$. In turn, we have $\mu \rightarrow \infty$, as
280 $p \rightarrow \infty$. Therefore, the last expression in (3.24) converges to 0, and we conclude
281 that $X \rightarrow \infty$ in probability, and hence

$$282 \quad \frac{|\widehat{S}(t^*) \setminus S|}{|\widehat{S}(t^*) \setminus S| + |S|} = \frac{X}{X + 1} \rightarrow 1 \quad \text{in probability.} \quad (3.25)$$

283 On the other hand, we show that with the same choice of $u = t^*$, we have

284

$$\frac{|S \setminus \widehat{S}(t^*)|}{|S|} \rightarrow 1 \text{ in probability.} \quad (3.26)$$

285 By the stochastic monotonicity of Gaussian location family (2.46), we have the
 286 following lower bound for the probability of missed detection for each signal $\mu(i)$,
 287 $i \in S$,

288

$$\mathbb{P}[N(\mu(i), 1) \leq t^*] \geq F_{\bar{a}}(t^*). \quad (3.27)$$

289 Since $|S \setminus \widehat{S}(t^*)|$ can be written as the sum of s independent Bernoulli random
 290 variables,

291

$$|S \setminus \widehat{S}(t^*)| = \sum_{i \in S} \mathbb{1}_{(-\infty, t^*]}(x(i)),$$

292 using with (3.27), we conclude that $|S \setminus \widehat{S}(t^*)| \stackrel{d}{\geq} \text{Binom}(s, F_{\bar{a}}(t^*))$. Finally, we
 293 know that $F_{\bar{a}}(t^*)$ converges to 1 by our choice of diverging t^* , and the necessary
 294 condition is shown. \square

295 **Proof (Sufficient condition in Theorem 3.3)** We now turn to the sufficient condition,
 296 i.e., part (i). That is, when $r > \beta$, the Benjamini–Hochberg procedure with slowly
 297 vanishing FDR levels achieves asymptotic approximate support recovery.

298 The FDR vanishes by our choice of α and the FDR-controlling property of the
 299 BH procedure (Benjamini and Hochberg 1995). It only remains to show that FNR
 300 also vanishes.

301 To do so we compare the FNR under the alternative specified in Theorem 3.3 to
 302 one with all of the signal sizes equal to $\underline{\Delta}$. By Lemma 3.2, it suffices to show that
 303 the FNR under the BH procedure in this setting vanishes. Let $x(i)$ be vectors of
 304 independent observations with $p - s$ nulls having standard Gaussian distributions,
 305 and s signals having $N(\underline{\Delta}, 1)$ distributions.

306 Denote the null and the alternative distributions as F_0 and F_a , respectively. Let
 307 \widehat{G} denote the empirical survival function as in (3.36). Define the empirical survival
 308 functions for the null part and signal part

309

$$\widehat{W}_{\text{null}}(t) = \frac{1}{p-s} \sum_{i \notin S} \mathbb{1}\{x(i) \geq t\}, \quad \widehat{W}_{\text{signal}}(t) = \frac{1}{s} \sum_{i \in S} \mathbb{1}\{x(i) \geq t\}, \quad (3.28)$$

where $s = |S|$, so that

$$\widehat{G}(t) = \frac{p-s}{p} \widehat{W}_{\text{null}}(t) + \frac{s}{p} \widehat{W}_{\text{signal}}(t).$$

310 We need the following result to describe the deviations of the empirical
 311 distributions.

Lemma 3.1 (Theorem 1 of Eicker 1979) *Let Z_1, \dots, Z_k be iid with continuous
 survival function Q . Let \widehat{Q}_k denote their empirical survival function and define
 $\xi_k = \sqrt{2 \log \log(k)/k}$ for $k \geq 3$. Then*

$$\frac{1}{\xi_k} \sup_z \frac{|\widehat{Q}_k(z) - Q(z)|}{\sqrt{Q(z)(1 - Q(z))}} \rightarrow 1,$$

in probability as $k \rightarrow \infty$. In particular,

$$\widehat{Q}_k(z) = Q(z) + O_{\mathbb{P}}\left(\xi_k \sqrt{Q(z)(1 - Q(z))}\right),$$

312 uniformly in z .

313 Applying Lemma 3.1 to the two summands in \widehat{G} , we obtain $\widehat{G}(t) = G(t) + \widehat{R}(t)$,
314 where

$$G(t) = \frac{p-s}{p} \overline{F}_0(t) + \frac{s}{p} \overline{F}_a(t), \quad (3.29)$$

316 and

$$\widehat{R}(t) = O_{\mathbb{P}}\left(\xi_p \sqrt{\overline{F}_0(t) F_0(t)} + \frac{s}{p} \xi_s \sqrt{\overline{F}_a(t) F_a(t)}\right), \quad (3.30)$$

318 uniformly in t .

319 Recall (see proof of Lemma 3.2) that the BH procedure is the thresholding pro-
320 cedure with threshold set at

$$\tau = \inf\{t \mid \overline{F}_0(t) \leq \alpha \widehat{G}(t)\}. \quad (3.31)$$

The NDP may also be re-written as

$$\text{NDP} = \frac{|S \setminus \widehat{S}|}{|S|} = \frac{1}{s} \sum_{i \in S} \mathbb{1}\{x(i) < \tau\} = 1 - \widehat{W}_{\text{signal}}(\tau),$$

322 so that it suffices to show that

$$\widehat{W}_{\text{signal}}(\tau) \rightarrow 1 \quad (3.32)$$

in probability. Applying Lemma 3.1 to $\widehat{W}_{\text{signal}}$, we know that

$$\widehat{W}_{\text{signal}}(\tau) = \overline{F}_a(\tau) + O_{\mathbb{P}}\left(\xi_s \sqrt{\overline{F}_a(\tau) F_a(\tau)}\right) = \overline{F}_a(\tau) + o_{\mathbb{P}}(1).$$

324 So it suffices to show that $F_a(\tau) \rightarrow 0$ in probability. Now let $t^* = \sqrt{2q \log(p)}$ for
325 some q such that $\beta < q < \underline{r}$. We have

$$F_a(t^*) = \Phi(t^* - \Delta) = \Phi(\sqrt{2(q - \underline{r}) \log p}) \rightarrow 0. \quad (3.33)$$

327 Hence, in order to show (3.32), it suffices to show

$$\mathbb{P}[\tau \leq t^*] \rightarrow 1. \quad (3.34)$$

329 By (3.29), the mean of the empirical process \widehat{G} evaluated at t^* is

$$330 \quad G(t^*) = \frac{p-s}{p} \overline{F_0}(t^*) + \frac{s}{p} \overline{F_a}(t^*). \quad (3.35)$$

331 The first term, using Relation (3.23), is asymptotic to $p^{-q} L(p)$, where $L(p)$ is the log-
332 arithmetic term in p . The second term, since $\overline{F_a}(t^*) \rightarrow 1$ by Relation (3.33), is asymp-
333 totic to $p^{-\beta}$. Therefore, $G(t^*) \sim p^{-q} L(p) + p^{-\beta} \sim p^{-\beta}$, since $p^{\beta-q} L(p) \rightarrow 0$
334 where $q > \beta$.

335 The fluctuation of the empirical process at t^* , by Relation (3.30), is

$$336 \quad \widehat{R}(t^*) = O_{\mathbb{P}} \left(\xi_p \sqrt{\overline{F_0}(t^*) \overline{F_0}(t^*)} + \frac{s}{p} \xi_s \sqrt{\overline{F_a}(t^*) \overline{F_a}(t^*)} \right) \\ 337 \quad = O_{\mathbb{P}} \left(\xi_p \sqrt{\overline{F_0}(t^*)} \right) + o_{\mathbb{P}}(p^{-\beta}). \\ 338$$

339 By (3.23) and the expression for ξ_p , the first term is $O_{\mathbb{P}}(p^{-(q+1)/2} L(p))$ where $L(p)$
340 is a poly-logarithmic term in p . Since $\beta < \min\{q, 1\}$, we have $\beta < (q+1)/2$, and
341 hence $\widehat{R}(t^*) = o_{\mathbb{P}}(p^{-\beta})$.

Putting the mean and the fluctuation of $\widehat{G}(t^*)$ together, we obtain

$$\widehat{G}(t^*) = G(t^*) + \widehat{R}(t^*) \sim_{\mathbb{P}} G(t^*) \sim p^{-\beta},$$

and therefore, together with (3.23), we have

$$\overline{F_0}(t^*)/\widehat{G}(t^*) = p^{\beta-q} L(p)(1 + o_{\mathbb{P}}(1)),$$

which is eventually smaller than the FDR level α by Assumption (3.17) and the fact that $\beta < q$. That is,

$$\mathbb{P}[\overline{F_0}(t^*)/\widehat{G}(t^*) < \alpha] \rightarrow 1.$$

342 By definition of τ (recall (3.31)), this implies that $\tau \leq t^*$ with probability tending to
343 1, and (3.34) is shown. The proof for the sufficient condition is complete. \square

344 3.5 Monotonicity of the Benjamini–Hochberg Procedure

345 As promised in the previous section, we make a connection between power of the
346 BH procedure and the stochastic ordering of distributions under the alternative. This
347 natural result seems new.

348 **Lemma 3.2** (Monotonicity of the BH procedure) *Consider p independent obser-
349 vations $x(i)$, $i \in \{1, \dots, p\}$, where the $(p-s)$ coordinates in the null part have
350 common distribution F_0 , and the remaining s signals have alternative distributions*

351 $F_j^i, i \in S$, respectively. Compare the two alternatives $j \in \{1, 2\}$, where the distributions
 352 in Alternative 2 are stochastically larger than those in Alternative 1, i.e.,

353
$$F_2^i(t) \leq F_1^i(t), \text{ for all } t \in \mathbb{R}, \text{ and for all } i \in S.$$

354 If the BH procedure is applied at the same nominal level of FDR, then the FNR of the
 355 BH procedure under Alternative 2 is bounded above by the FNR under Alternative
 356 1. Further, the threshold of the BH procedure under Alternative 2 is stochastically
 357 smaller than that under Alternative 1.

358 Loosely put, the power of the BH procedure is monotone increasing with respect
 359 to the stochastic ordering of the alternatives, yet (the distribution of) the BH threshold
 360 is monotone decreasing in the distributions of the alternatives.

361 **Proof (Lemma 3.2)** We first re-express the BH procedure in a different form.
 362 Recall that on observing $x(i), i \in \{1, \dots, p\}$, the BH procedure is the thresholding
 363 procedure with threshold set at $x_{[i^*]}$, where $i^* := \max\{i \mid \bar{F}_0(x_{[i]}) \leq \alpha i/p\}$, and
 364 $x_{[1]} \geq \dots \geq x_{[p]}$ are the order statistics.

365 Let \widehat{G} denote the left-continuous empirical survival function

366
$$\widehat{G}(t) = \frac{1}{p} \sum_{i=1}^p \mathbb{1}\{x(i) \geq t\}. \quad (3.36)$$

367 By the definition, we know that $\widehat{G}(x_{[i]}) = i/p$. Therefore, by the definition of i^* , we
 368 have

369
$$\bar{F}_0(x_{[i]}) > \alpha \widehat{G}(x_{[i]}) = \alpha i/p \text{ for all } i > i^*.$$

370 Since \widehat{G} is constant on $(x_{[i^*+1]}, x_{[i^*]})$, the fact that $\bar{F}_0(x_{[i^*]}) \leq \alpha \widehat{G}(x_{[i^*]})$ and
 371 $\bar{F}_0(x_{[i^*+1]}) > \alpha \widehat{G}(x_{[i^*+1]})$ implies that $\alpha \widehat{G}$ and \bar{F}_0 must “intersect” on the interval by
 372 continuity of F_0 . We denote this “intersection” as

373
$$\tau = \inf\{t \mid \bar{F}_0(t) \leq \alpha \widehat{G}(t)\}. \quad (3.37)$$

374 Note that τ cannot be equal to $x_{[i^*+1]}$ since \bar{F}_0 is càdlàg. Since there is no observation
 375 in $[\tau, x_{[i^*]})$, we can write the BH procedure as the thresholding procedure with
 376 threshold set at τ .

377 Now, denote the observations under Alternatives 1 and 2 as $x_1(i)$ and $x_2(i)$. Since
 378 $x_2(i)$ stochastically dominates $x_1(i)$ for all $i \in \{1, \dots, p\}$, there exists a coupling
 379 $(\tilde{x}_1, \tilde{x}_2)$ of x_1 and x_2 such that $\tilde{x}_1(i) \leq \tilde{x}_2(i)$ almost surely for all i . We will replace
 380 \tilde{x}_1 and \tilde{x}_2 with x_1 and x_2 in what follows. Since we will compare the FNRs, i.e.,
 381 expectations with respect to the marginals of x 's in the last step, this replacement
 382 does not affect the conclusions. To simplify notation, we still write x_1 and x_2 in place
 383 of \tilde{x}_1 and \tilde{x}_2 .

384 Let \widehat{G}_k be the left-continuous empirical survival function under Alternative k , i.e.,

385
$$\widehat{G}_k(t) = \frac{1}{p} \sum_{i=1}^p \mathbb{1}\{x_k(i) \geq t\}, \quad k \in \{1, 2\}. \quad (3.38)$$

386 We define the BH thresholds τ_1 and τ_2 by replacing \widehat{G} in (3.37) with \widehat{G}_1 and \widehat{G}_2 ,
 387 respectively. Denote the set estimates of signal support $\widehat{S}_k = \{i \mid x_k(i) \geq \tau_k\}$ by the
 388 BH procedure. We claim that

389
$$\tau_2 \leq \tau_1 \quad \text{with probability 1.} \quad (3.39)$$

390 Indeed, by definition of the empirical survival function (3.38) and the fact that
 391 $x_1(i) \leq x_2(i)$ almost surely for all i , we have $\widehat{G}_1(t) \leq \widehat{G}_2(t)$ for all t . Hence, $\overline{F}_0(t) \leq$
 392 $\alpha \widehat{G}_1(t)$ implies $\overline{F}_0(t) \leq \alpha \widehat{G}_2(t)$, and Relation (3.39) follows from the definition of
 393 τ in (3.37). The claim of stochastic ordering of the BH thresholds in Lemma 3.2
 394 follows from (3.39).

395 Finally, when $\tau_2 \leq \tau_1$, we have $\tau_2 \leq \tau_1 \leq x_1(i) \leq x_2(i)$ with probability 1 for
 396 all $i \in \widehat{S}_1$. Therefore, it follows that $\widehat{S}_1 \subseteq \widehat{S}_2$ and hence $|S \setminus \widehat{S}_2| \leq |S \setminus \widehat{S}_1|$ almost
 397 surely. The first conclusion in Lemma 3.2 follows from the last inequality. \square

398 3.6 The Exact–Approximate Support Recovery Problem

399 We now derive two new asymptotic phase-transition results for the *asymmetric* statistical
 400 risks, (2.11) and (2.12), in the Gaussian error models. As discussed in Sect. 2.1,
 401 the exact–approximate support recovery risk is the natural criteria when considering
 402 the marginal power of discovery while controlling for family-wise error rates in
 403 applications such as GWAS.

404 Although there have been discussions of weighted sums of Type I and Type II
 405 errors in the literature (see, e.g., Genovese and Wasserman (Genovese and Wasserman
 406 2002) Sect. 6, where the authors sought to minimize $\text{FDR} + \lambda \text{FNR}$), asymptotic
 407 limits were not discussed. We point out that the asymptotic limits for the unequally
 408 weighted risks are no different from the equally weighted risk, so long as λ is bounded
 409 away from zero and infinity. This is because $\text{FDR} + \lambda \text{FNR}$ vanishes if and only if
 410 both FDR and FNR vanish; conversely, non-vanishing FDR and FNR are equivalent
 411 to non-vanishing weighted sums. Therefore, a different phase transition would only
 412 arise if we weight the Type I and Type II errors by combining family-wise error
 413 metrics with marginal error rates.

414 The next theorem describes the phase transition in the exact–approximate support
 415 recovery problem.

416 **Theorem 3.4** *In the context of Theorem 3.2, the function*

417
$$f_{\text{EA}}(\beta) = 1 \quad (3.40)$$

418 characterizes the phase transition of exact–approximate support recovery problem.
419 Namely, the following two results hold:

- 420 (i) If $\underline{r} > f_{EA}(\beta)$, then the procedures listed in Theorem 3.2 with slowly vanishing
421 nominal FWER levels (as defined in Definition 3.1) achieve asymptotically exact–
422 approximate support recovery in the sense of (2.25).
423 (ii) Conversely, if $\underline{r} < f_{EA}(\beta)$, then for any thresholding procedure \widehat{S} , the exact–
424 approximate support recovery fails in the sense of (2.26).

425 The phase-transition boundary (3.40) is visualized in Fig. 3.2. The proof of this
426 result uses ideas from the proof of Theorem 3.3 and is substantially shorter.

427 **Proof** (Theorem 3.4) We first show the sufficient condition. Vanishing FWER is
428 guaranteed by the properties of the procedures, and we only need to show that FNR
429 also goes to zero. Similar to the proof of Theorem 3.3, it suffices to show that

$$\text{NDP} = 1 - \widehat{W}_{\text{signal}}(t_p) \rightarrow 0, \quad (3.41)$$

431 where t_p is the threshold of Bonferroni’s procedure.

432 Since α vanishes slowly (see Definition 3.17), for any $\delta > 0$, we have $p^{-\delta} = o(\alpha)$.
433 Therefore, we have $-\log \alpha \leq \delta \log p$ for large p , and

$$\text{434} \quad 1 \leq \limsup_{p \rightarrow \infty} \frac{2 \log p - 2 \log \alpha}{2 \log p} \leq 1 + \delta,$$

for any $\delta > 0$. Therefore, by the expression for normal quantiles, we know that

$$t_p = F^{\leftarrow}(1 - \alpha/p) \sim (2 \log p - 2 \log \alpha)^{1/2} \sim (2 \log p)^{1/2}.$$

Since $\underline{r} > f_{EA}(\beta) = 1$, we can pick q such that $1 < q < \underline{r}$. Let $t^* = \sqrt{2q \log p}$, we know that $t_p < t^*$ for large p . Therefore, for large p , we have

$$\widehat{W}_{\text{signal}}(t_p) \geq \widehat{W}_{\text{signal}}(t^*) \geq \overline{F}_a(t^*) + o_{\mathbb{P}}(1),$$

where \overline{F}_a is the survival function of $N(\sqrt{2\underline{r} \log p}, 1)$; the last inequality follows from the stochastic monotonicity of the Gaussian location family (2.46), and Lemma 3.1. Indeed, by our choice of $q < \underline{r}$, we obtain

$$F_a(t^*) = \Phi\left(\sqrt{2(q - \underline{r}) \log p}\right) \rightarrow 0,$$

435 and (3.41) is shown. This completes the proof of the sufficient condition.

436 The proof of the necessary condition follows similar structure as in the proof of
437 Theorem 3.3, and uses the lower bound

438 $\text{FWER}(\mathcal{R}) + \text{FNR}(\mathcal{R}) \geq \mathbb{P} \left[\max_{i \in S^c} x(i) > u \right] \wedge \mathbb{E} \left[\frac{|S \setminus \widehat{S}(u)|}{|S|} \right], \quad (3.42)$

439 which holds for any arbitrary thresholding procedure \mathcal{R} and arbitrary real $u \in \mathbb{R}$.

440 By the assumption that $\bar{r} < f_{\text{EA}}(\beta) = 1$, we can pick q such that $\bar{r} < q < 1$ and let
441 $u = t^* = \sqrt{2q \log p}$ in (3.42). By relative stability of iid Gaussian random variables
442 (2.43), we have

443
$$\mathbb{P} \left[\frac{\max_{i \in S^c} x(i)}{\sqrt{2 \log p}} > \frac{t^*}{\sqrt{2 \log p}} \right] \rightarrow 1. \quad (3.43)$$

444 Since the first fraction in (3.43) converges to 1, while the second converges to $q < 1$.
445 Therefore, the first term on the right-hand side of (3.42) converges to 1.

446 On the other hand, by the stochastic monotonicity of Gaussian location
447 family (2.46), the probability of missed detection for each signal is lower bounded
448 by $\mathbb{P}[Z + \mu(i) \leq t^*] \geq F_{\bar{a}}(t^*)$, where Z is a standard Gaussian r.v., and $F_{\bar{a}}$ is the
449 cdf of $N(\sqrt{2\bar{r} \log p}, 1)$. Therefore, $|S \setminus \widehat{S}(t^*)| \xrightarrow{d} \text{Binom}(s, F_{\bar{a}}(t^*))$, and it suffices to
450 show that $F_{\bar{a}}(t^*)$ converges to 1. Indeed,

451
$$F_{\bar{a}}(t^*) = \Phi(\sqrt{2(q - \bar{r}) \log p}) \rightarrow 1,$$

452 by our choice of $q > \bar{r}$. Hence, both quantities in the minimum on the right-hand
453 side of (3.42) converge to 1 in the limit, and the necessary condition is shown. \square

454 **Remark 3.1** The boundary (3.40) was briefly suggested by Arias-Castro and Chen
455 (2017). Unfortunately, it was falsely claimed that the boundary characterized the
456 phase transition of the *exact* support recovery problem, and the alleged proof was
457 left as an “exercise to the reader”. This exercise was completed in Chap. 4, where the
458 correct boundary (7.4) was identified.

459 Theorem 3.4 here shows that the boundary (3.40) *does* exist, though for the slightly
460 different *exact–approximate* support recovery problem. As we will see in Sect. 7.1,
461 the boundary (3.40) also applies to the exact–approximate support recovery problem
462 in chi-square models (1.3).

463 3.7 The Approximate–Exact Support Recovery Problem

464 The last phase transition is in terms of the approximate–exact support recovery risk
465 (2.12).

466 **Theorem 3.5** *In the context of Theorem 3.2, the function*

467
$$f_{\text{AE}}(\beta) = \left(\sqrt{\beta} + \sqrt{1 - \beta} \right)^2 \quad (3.44)$$



468 characterizes the phase transition of approximate–exact support recovery problem.
 469 Namely, the following two results hold:

- 470 (i) If $\underline{r} > f_{AE}(\beta)$, then the Benjamini–Hochberg procedure with slowly vanishing
 471 nominal FDR levels (as defined in Definition 3.1) achieves asymptotically
 472 approximate–exact support recovery in the sense of (2.25).
 473 (ii) Conversely, if $\bar{r} < f_{AE}(\beta)$, then for any thresholding procedure \widehat{S} , the
 474 approximate–exact support recovery fails in the sense of (2.26).

475 The phase-transition boundary (3.44) is visualized in Fig. 3.2.

476 **Proof** (Theorem 3.5) We first show the sufficient condition (part (i)). Since FDR
 477 control is guaranteed by the BH procedure, we only need to show that the FWNR
 478 also vanishes, that is,

$$479 \mathbb{P}\left[\min_{i \in S} x(i) \geq \tau\right] \rightarrow 1, \quad (3.45)$$

480 where τ is the threshold for the BH procedure.

481 By the assumption that $\underline{r} > f_{AE}(\beta) = (\sqrt{\beta} + \sqrt{1-\beta})^2$, we have $\sqrt{\underline{r}} -$
 482 $\sqrt{1-\beta} > \sqrt{\beta}$, so we can pick $q > 0$, such that

$$483 \sqrt{\underline{r}} - \sqrt{1-\beta} > \sqrt{q} > \sqrt{\beta}. \quad (3.46)$$

484 We only need to show that with a specific choice of $t^* = \sqrt{2q \log p}$ where

$$485 \sqrt{\underline{r}} - \sqrt{1-\beta} > \sqrt{q} > \sqrt{\beta}, \quad (3.47)$$

486 we have both

$$487 \mathbb{P}[\tau \leq t^*] \rightarrow 1, \quad (3.48)$$

488 and

$$489 \mathbb{P}\left[\min_{i \in S} x(i) \geq t^*\right] \rightarrow 1, \quad (3.49)$$

490 so that

$$491 \mathbb{P}\left[\min_{i \in S} x(i) \geq \tau\right] \geq \mathbb{P}\left[\min_{i \in S} x(i) \geq t^*, t^* \geq \tau\right] \rightarrow 1.$$

492 Relation (3.48) follows in exactly the same way (3.34) did in Sect. 3.4.

AQ1

493 Dividing the left-hand side in Relation (3.49) by $\sqrt{2 \log p}$, we have

$$494 \frac{\min_{i \in S} x(i)}{\sqrt{2 \log p}} = \frac{\min_{i \in S} \mu(i) + \epsilon(i)}{\sqrt{2 \log p}} \stackrel{d}{\geq} \frac{\sqrt{2\underline{r} \log p} + \min_{i \in S} \epsilon(i)}{\sqrt{2 \log p}} \\ 495 \rightarrow -\sqrt{1-\beta} + \sqrt{\underline{r}},$$



where the last convergence follows from the relative stability of iid Gaussians minima (2.44). On the other hand, $t^*/\sqrt{2 \log p} = \sqrt{q} < \sqrt{\bar{r}} - \sqrt{1-\beta}$ by our choice of q , and Relation (3.49) follows.

The necessary condition follows from the lower bound

$$\text{FDR}(\mathcal{R}) + \text{FWNR}(\mathcal{R}) \geq \mathbb{E} \left[\frac{|\widehat{S}(u) \setminus S|}{|\widehat{S}(u) \setminus S| + |S|} \right] \wedge \mathbb{P} \left[\min_{i \in S} x(i) < u \right], \quad (3.50)$$

which holds for any thresholding procedure \mathcal{R} and for arbitrary $u \in \mathbb{R}$. In particular, we show that both terms in the minimum in (3.50) converge to 1 when we set $u = t^* = \sqrt{2q \log p}$ where

$$\sqrt{\bar{r}} - \sqrt{1-\beta} < \sqrt{q} < \sqrt{\beta}. \quad (3.51)$$

On the one hand, we have

$$\frac{\min_{i \in S} x(i)}{\sqrt{2 \log p}} \stackrel{d}{\leq} \frac{\min_{i \in S} \epsilon(i) + \sqrt{2\bar{r} \log p}}{\sqrt{2 \log p}} \rightarrow \sqrt{\bar{r}} - \sqrt{1-\beta},$$

by relative stability of iid Gaussians (2.44). On the other hand, $t^*/\sqrt{2 \log p} = \sqrt{q} > \sqrt{\bar{r}} - \sqrt{1-\beta}$ by our choice of q ; this shows that the second term on the right-hand side of (3.50) converges to 1.

Observe that $|\widehat{S}(t^*) \setminus S|$ has distribution $\text{Binom}(p-s, \bar{\Phi}(t^*))$, and define $X = X_p := |\widehat{S}(t^*) \setminus S|/|S|$, we obtain

$$\begin{aligned} \mu := \mathbb{E}[X] &= (p^\beta - 1)\bar{\Phi}(t^*) \sim (p^\beta - 1) \frac{\phi(t^*)}{t^*} \\ &\sim \frac{1}{\sqrt{2\pi}} (2q \log p)^{-1/2} p^{\beta-q} \rightarrow \infty, \end{aligned}$$

where the divergence follows from our choice of $q < \beta$. Using again Relations (3.24) and (3.25), we conclude that the first term on the right-hand side of (3.50) also converges to 1. This completes the proof of the necessary condition. \square

3.8 Asymptotic Power Analysis: A Discussion

Theorems 3.2 through 3.5 allow us to asymptotically quantify the required signal sizes in support recovery problems, as well as in the global hypothesis testing problem in the Gaussian additive error model (3.1). Specifically, these results indicate that at all sparsity levels $\beta \in (0, 1)$, the difficulties of the problems in terms of the required signal sizes have the following ordering:

$$f_D(\beta) < f_A(\beta) < f_{EA}(\beta) < f_{AE}(\beta) < f_E(\beta),$$

as previewed in Fig. 3.2. The ordering aligns with our intuition that the required signal sizes must increase as we move from detection to support recovery problems. Similarly, more stringent criteria for error control (e.g., FWER compared to FDR) require larger signals. We can now also compare $f_{EA}(\beta)$ and $f_{AE}(\beta)$, whose ordering may not be clear from this line of reasoning.

Our last comment is on the gap between FDR and FWER under sparsity assumptions. Although it is believed that FWER control is sometimes too stringent compared to, say, FDR control in support recovery problems, the fact that all five thresholds involve the same scaling indicates that the difficulties of the problems (signal detection, and the four support recovery problems) are comparable when signals are very sparse, i.e., when β is close to 1. This is illustrated with the next example.

Example 3.1 (*Power analysis for variable selection*) For Gaussian errors (AGG with $\nu = 2$), when $\beta = 3/4$, the signal detection boundary (3.4) says that signals will have to be at least of magnitude $\sqrt{(\log p)/2}$, while approximate support recovery (3.19) requires signal sizes of at least $\sqrt{3(\log p)/2}$, and exact support recovery (3.18) calls for signal sizes of at least $\sqrt{9(\log p)/2}$. The required signal sizes increases, but are within the same order of magnitude.

If m independent copies x_1, \dots, x_m of the observations were made on the same set of p locations, then by taking location-wise averages, $\bar{x}_m(j) = \frac{1}{m} \sum_{i=1}^m x_i(j)$, we can reduce error standard deviation, and hence boost the signal-to-noise ratio, by a factor of \sqrt{m} . By the simple calculations above, if m samples are needed to detect (sparse) signals of a certain magnitude, then $3m$ samples will enable approximate support recovery with false discovery and non-discovery control, and, in fact, $9m$ samples would enable exact support recovery with family-wise error rates control.

On the other hand, the gap between FDR and FWER is much larger when signals are dense. For example, if the signals are only *approximately* sparse, i.e., having a few components above (3.18) but many smaller components above (3.19), then FDR-controlling procedures will discover substantially larger proportion of signals than FWER-controlling procedures.

Indeed, as $\beta \rightarrow 0$, the required signal size for approximate support recovery (3.19) tends to 0, while the required signal size for exact support recovery (3.18) tends to 4 in the Gaussian error models. While Example 3.1 indicates that the exact support recovery is not much more stringent than approximate support recovery when signals are sparse, the gap between required signal sizes widens when signals are dense.

Author Queries

Chapter 3

Query Refs.	Details Required	Author's response
AQ1	As per Springer standard, page number cross citations are not permitted, hence the corresponding section numbers are used. Kindly check.	

Metadata of the chapter that will be visualized in SpringerLink

Book Title	Concentration of Maxima and Fundamental Limits in High-Dimensional Testing and Inference	
Series Title		
Chapter Title	Exact Support Recovery Under Dependence	
Copyright Year	2021	
Copyright HolderName	The Author(s), under exclusive license to Springer Nature Switzerland AG	
Corresponding Author	Family Name Particle Given Name Prefix Suffix Role Division Organization Address Email ORCID	Gao Zheng Department of Statistics University of Michigan-Ann Arbor Ann Arbor, MI, USA gaozheng@umich.edu https://orcid.org/0000-0003-1394-8741
Author	Family Name Particle Given Name Prefix Suffix Role Division Organization Address Email ORCID	Stoev Stilian Department of Statistics University of Michigan-Ann Arbor Ann Arbor, MI, USA https://orcid.org/0000-0003-2108-5421
Abstract	This chapter describes some of the core contributions of this work. It focuses on the exact support recovery problem for a sparse signal observed with additive noise. The noise is assumed to be light tailed with potentially arbitrary dependence structure and the general class of thresholding estimators for the signal support is considered. In this setting, a universal phase-transition result for the exact support recovery is established when the dimension of the signal diverges to infinity. Namely, when the (suitably parameterized) signal strength is above the so-called classification boundary (a certain function of the signal sparsity parameter), then thresholding estimators are shown to recover the signal support exactly with probability converging to one, as the dimension goes to infinity. This is so regardless of the dependence of the errors. Conversely, when the signal is below the classification boundary function, all thresholding estimators fail to recover the support with probability converging to one. This is so provided that the error array is uniformly relatively stable—a type of concentration of maxima condition. The URS condition is optimal and characterized further in Chap. 6.	

Chapter 4

Exact Support Recovery Under Dependence



We focus on exact support recovery problems in this chapter. Recall from Lemma 2.1 that in order to study the asymptotic behaviors of risk^E, it is sufficient to establish minimal conditions under which the support sets can be consistently estimated, i.e.,

$$\mathbb{P}[\widehat{S}_p = S_p] \longrightarrow 1 \text{ as } p \rightarrow \infty, \quad (4.1)$$

where \widehat{S}_p is an estimate of the true support set S_p of a high-dimensional signal vector μ_p .

We will establish minimal conditions such that (4.1) holds, by generalizing the results we obtained in Chap. 3 to additive error models with relaxed distributional and dependence assumptions on the additive error array.

4.1 Generalizations of Distributional and Dependence Assumptions

Consider the additive error model (1.1) with the triangular array of errors,

$$\mathcal{E} = \{(\epsilon_p(i))_{i=1}^p, p = 1, 2, \dots\}, \quad (4.2)$$

where the $\epsilon_p(i)$'s have common cumulative distribution function $F(x) = \mathbb{P}[\epsilon_p(i) \leq x]$. In contrast to the assumptions in Chap. 3, we only require the errors to have common marginal distributions.

Although our method of analysis applies to all light-tailed error distributions with rapidly varying tails (see Definition 2.7), to be concrete and better convey the main



ideas, we will focus on the class of AGG(ν) laws (see Definition 2.6). Extensions of the results to other classes of error models are presented in Sect. B.1.

This generalized distributional assumption on the errors call for a suitable generalization of the signal size parametrization in order to analyze the problem as we did in the previous chapter. As before, we assume the signals in model (1.1) to be a sparse vector $\mu_p = (\mu_p(i))_{i=1}^p$, with support $S_p := \{i : \mu_p(i) \neq 0\}$. The sparsity of μ_p , with a few exceptions which will be explicitly stated, is parametrized in terms of a fixed regularly varying sequence $\{s_p^\dagger\}$ as follows:

$$|S_p| = \lfloor s_p^\dagger \rfloor, \quad \text{where } s_p^\dagger := \ell(p)p^{1-\beta}, \quad (4.3)$$

for some fixed slowly varying function ℓ . Recall that a function ℓ is slowly varying if $\ell(\lambda t)/\ell(t) \rightarrow 1$, as $t \rightarrow \infty$, for all $\lambda > 0$. As before, the exponent

$$0 < \beta \leq 1$$

controls the sparsity.

We assume that the non-zero entries of μ are positive and take values in the interval $[\underline{\Delta}, \bar{\Delta}] \subset (0, \infty)$. That is, $0 < \underline{\Delta} \leq \mu(i) \leq \bar{\Delta} \leq +\infty$, for all $i \in S_p$. The lower and upper bounds on the signal sizes $\underline{\Delta}$ and $\bar{\Delta}$ are parametrized as

$$\underline{\Delta} = \underline{\Delta}(p) = (\nu r \log p)^{1/\nu} \quad \text{and} \quad \bar{\Delta} = \bar{\Delta}(p) = (\nu \bar{r} \log p)^{1/\nu}, \quad (4.4)$$

with parameters $0 < \underline{r} \leq \bar{r} \leq +\infty$.

We now turn to the dependence conditions. Several authors have studied the support recovery problem in terms of the Hamming loss and obtained minimax-optimality results (see, e.g., Ji and Jin 2012; Genovese et al. 2012; Jin et al. 2014; Butucea et al. 2018). In the special case of Gaussian marginals, Butucea et al. (2018) showed that the boundary (3.18) exists in a minimax sense. That is, when the errors are *independent* Gaussians, the Hamming loss cannot be made to vanish if the signal sizes are sufficiently small by any procedure. Conversely, if signal size falls below, the Hamming loss can be made to vanish for some thresholding procedure. However, as pointed out in Sect. 2.4, vanishing Hamming loss is only sufficient, not necessary for support recovery (4.1), and results on the former do not carry over directly to the study of the exact support recovery problem. More importantly, since Hamming loss decomposes into expectations on individual terms that are not affected by dependence, Hamming loss-minimax studies do not reveal the difference in probability of support recovery between independent and dependent observations. This prevents one from fully exploring the phase-transition phenomena under other dependence conditions. As a result, so far in the literature, the role of dependence in model (1.1) has remained largely unexplored.

We take a different approach in this text. In particular, we study the exact support recovery problem (4.1) directly, and show that for thresholding procedures the phase-transition phenomena exist universally in a large class of dependence structures, and not just in a minimax sense.

In a first step, we show that in the AGG model under *arbitrary* dependence, under the scaling described in (4.3) and (4.4), the function

$$f_E(\beta) = f_{E,v}(\beta) = (1 + (1 - \beta)^{1/v})^v, \quad v > 0 \quad (4.5)$$

demarcates the region of possibility for the exact support recovery problem. That is, if the signal sizes are above the boundary (i.e., $\underline{r} > f_E(\beta)$), then FWER-controlling procedures with appropriately calibrated levels achieve exact support recovery (Theorem 4.1). We refer to (4.5) as the *strong classification boundary*.

Conversely, we show that for a surprisingly large class of dependence structures characterized by the concept of *uniform relative stability* (URS, see Definition 4.1), when the signal size is below the boundary (i.e., $\underline{r} < f_E(\beta)$), no thresholding procedure can achieve the asymptotically perfect support recovery. In fact,

$$\mathbb{P}[\widehat{S}_p = S_p] \rightarrow 0, \quad \text{as } p \rightarrow \infty, \quad (4.6)$$

for all thresholding procedures (Theorem 4.2). These two results show that the thresholding procedures obey a phase-transition phenomenon in a strong, *point-wise* sense over the class of URS dependence structures, and over the class of AGG(v), $v > 0$ error distributions.

4.2 Sufficient Conditions for Exact Support Recovery

Following Butucea et al. (2018), we define the parameter space for the signals μ as

$$\begin{aligned} \Theta_p^+(\beta, \underline{r}) = \{\mu \in \mathbb{R}^p : \text{there exists a set } S_p \subseteq \{1, \dots, p\} \text{ such that } |S_p| \leq s_p^\dagger, \\ \mu(i) \geq (\nu \underline{r} \log p)^{1/v} \text{ for all } i \in S_p, \text{ and } \mu(i) = 0 \text{ for all } i \notin S_p\}, \end{aligned} \quad (4.7)$$

where s_p^\dagger is as in (4.3). Our first result states that, when $F \in \text{AGG}(v)$ with $v > 0$, regardless of the error-dependence structure, (asymptotic) perfect support recovery is achieved by applying Bonferroni's procedure with appropriately calibrated FWER, as long as the minimum signal size \underline{r} is above the strong classification boundary (4.5).

Theorem 4.1 *Let the errors have common marginal distribution $F \in \text{AGG}(v)$ with $v > 0$. Let \widehat{S}_p be Bonferroni's procedure (2.21) with vanishing FWER $\alpha = \alpha(p) \rightarrow 0$, such that $\alpha p^\delta \rightarrow \infty$ for every $\delta > 0$. If*

$$\underline{r} > f_E(\beta) = (1 + (1 - \beta)^{1/v})^v, \quad (4.8)$$



85 then we have

$$86 \quad \lim_{p \rightarrow \infty} \sup_{\mu \in \Theta_p^+(\beta, \underline{r})} \mathbb{P}[\widehat{S}_p \neq S_p] = 0. \quad (4.9)$$

87 **Proof** Throughout the proof, the dependence on p will be suppressed to simplify
88 notations when such omissions do not lead to ambiguity.

89 Under the AGG(v) model, it is easy to see from Eq. (2.33) that the thresholds in
90 Bonferroni's procedure are

$$91 \quad t_p = F^\leftarrow(1 - \alpha/p) = (v \log(p/\alpha))^{1/v}(1 + o(1)). \quad (4.10)$$

92 It is known that Bonferroni's procedure $\widehat{S}_p = \{i : x(i) > t_p\}$ controls the FWER.
93 Indeed,

$$94 \quad \mathbb{P}[\widehat{S} \subseteq S] = 1 - \mathbb{P}\left[\max_{i \in S^c} x(i) > t_p\right] = 1 - \mathbb{P}\left[\max_{i \in S^c} \epsilon(i) > t_p\right] \\ 95 \quad \geq 1 - \sum_{i=1}^p \mathbb{P}[\epsilon(i) > t_p] \geq 1 - \alpha(p) \rightarrow 1, \quad (4.11)$$

97 where we used the union bound in the first inequality. Notice that the lower bound
98 (4.11) is independent of the parameter μ (as well as the dependence structures), and
99 hence holds uniformly over the parameter space, i.e.,

$$100 \quad \lim_{p \rightarrow \infty} \inf_{\mu \in \Theta_p^+(\beta, \underline{r})} \mathbb{P}[\widehat{S}_p \subseteq S_p] = 1. \quad (4.12)$$

101 On the other hand, for the probability of no missed detection, we have

$$102 \quad \mathbb{P}[\widehat{S} \supseteq S] = \mathbb{P}\left[\min_{i \in S} x(i) > t_p\right] = \mathbb{P}\left[\min_{i \in S} x(i) - (v\underline{r} \log p)^{1/v} > t_p - (v\underline{r} \log p)^{1/v}\right].$$

103 Since the signal sizes are no smaller than $(v\underline{r} \log p)^{1/v}$, we have

$$104 \quad x(i) - (v\underline{r} \log p)^{1/v} \geq \epsilon(i), \quad \text{for all } i \in S,$$

105 and hence we obtain

$$106 \quad \mathbb{P}[\widehat{S} \supseteq S] \geq \mathbb{P}\left[\min_{i \in S} \epsilon(i) > (v \log(p/\alpha))^{1/v}(1 + o(1)) - (v\underline{r} \log p)^{1/v}\right], \quad (4.13)$$

107 where we plugged in the expression for t_p in (4.10). Now, since the minimum signal
108 size is bounded below by $\underline{r} > (1 + (1 - \beta)^{1/v})^v$, we have $\underline{r}^{1/v} - (1 - \beta)^{1/v} > 1$,
109 and so we can pick a $\delta > 0$ such that

$$110 \quad \delta < (\underline{r}^{1/v} - (1 - \beta)^{1/v})^v - 1. \quad (4.14)$$



Since by assumption, for all $\delta > 0$, we have $p^{-\delta} = o(\alpha(p))$, there is an $M = M(\delta)$ such that $p/\alpha(p) < p^{1+\delta}$ for all $p \geq M$. Thus, from (4.13), we further conclude that for $p \geq M$ we have

$$\begin{aligned} \mathbb{P}\left[\widehat{S} \supseteq S\right] &\geq \mathbb{P}\left[\min_{i \in S} \epsilon(i) > ((1 + \delta)\nu \log p)^{1/\nu} (1 + o(1)) - (\nu \underline{\ell} \log p)^{1/\nu}\right] \\ &= \mathbb{P}\left[\max_{i \in S} (-\epsilon(i)) < \underbrace{(\underline{\ell}^{1/\nu} - (1 + \delta)^{1/\nu})(\nu \log p)^{1/\nu}(1 + o(1))}_{=: A}\right] \\ &\geq 1 - \ell(p)p^{1-\beta} \times \overline{F}_-(A), \end{aligned} \quad (4.15)$$

where $\overline{F}_-(x) = \mathbb{P}[-\epsilon(i) > x]$ is the survival function of the $(-\epsilon(i))$'s. Notice that (4.15) follows from the union bound and the assumption that $|S_p| \leq s_p^\dagger = \ell(p)p^{1-\beta}$ in (4.7). Therefore, the lower bound does not depend on μ (nor on the error-dependence structure), and holds uniformly in the parameter space. In turn, we obtain

$$\inf_{\mu \in \Theta_p^+(\beta, \underline{\ell})} \mathbb{P}[\widehat{S}_p \supseteq S_p] \geq 1 - \ell(p)p^{1-\beta} \times \overline{F}_-(A). \quad (4.16)$$

We first show that the right-hand side of (4.16) converges to 1 when $\beta = 1$. Indeed, since $F \in \text{AGG}(\mu)$, we have, for sufficiently large p ,

$$\overline{F}_-(A) \leq \overline{F}_-(c(\nu \log(p))^{1/\nu}) = O(p^{-c'}),$$

for some $c > c' > 0$. On the other hand, the celebrated Potter bounds for slowly varying functions (see, e.g., Bingham et al. 1987) entail $\ell(p) = o(p^{c'})$, for every $c' > 0$ and hence $\ell(p)\overline{F}_-(A) \rightarrow 0$, as $p \rightarrow \infty$.

Let now $\beta \in (0, 1)$ and $u_p^- := F_-^\leftarrow(1 - 1/p)$. The fact that $p\overline{F}_-(u_p^-) \leq 1$ implies

$$s_p^\dagger \times \overline{F}_-(A) \leq \frac{\overline{F}_-\left(B \times u_{s_p^\dagger}^-\right)}{\overline{F}_-\left(u_{s_p^\dagger}\right)}, \quad (4.17)$$

where $B := A/u_{s_p^\dagger}^-$.

Notice that by assumption, the $-\epsilon(i)$'s are also $\text{AGG}(\nu)$ distributed, and by Proposition 2.1, $u_p^- := F_-^\leftarrow(1 - 1/p) \sim (\nu \log(p))^{1/\nu}$, as $p \rightarrow \infty$. Therefore, we have

$$u_{s_p^\dagger}^- \equiv u_{\ell(p)p^{1-\beta}}^- \sim (\nu(1 - \beta) \log p)^{1/\nu}, \quad (4.18)$$

where we used the fact that $\log(\ell(p)) = o(\log(p))$. Hence,

$$B = \frac{A}{u_{s_p^\dagger}^-} = \frac{\underline{\ell}^{1/\nu} - (1 + \delta)^{1/\nu}}{(1 - \beta)^{1/\nu}} (1 + o(1)) \rightarrow c > 1$$



138 as $p \rightarrow \infty$, by our choice of δ in (4.14).

139 Finally, since the distribution F_- has *rapidly varying tails* (by Definition 2.7 and
 140 Example 2.1), applying Proposition 2.2, we conclude that (4.17) vanishes. Conse-
 141 quently, the lower bound on the right-hand side of (4.16) converges to 1. This,
 142 combined with (4.12), entails $\lim_{p \rightarrow \infty} \inf_{\mu \in \Theta_p^+(\beta, r)} \mathbb{P}[\hat{S}_p = S_p] = 1$, and hence the
 143 desired Conclusion (4.9), which completes the proof. \square

144 We end this section with several comments and applications of Theorem 4.1.

145 **Corollary 4.1** (Classes of procedures attaining the boundary) *Relation (4.9) holds*
 146 *for any FWER-controlling procedure that is strictly more powerful than Bonferroni's*
 147 *procedure. This includes Holm's procedure (Holm 1979), and in the case of inde-
 148 *pendent errors, Hochberg's procedure (Hochberg 1988), and the Šidák procedure*
 149 *(Šidák 1967).**

150 **Example 4.1** Under Gaussian errors, the particular choice of the thresholding
 151 at $t_p = \sqrt{2 \log p}$ in (2.21) corresponds to a Bonferroni's procedure with FWER
 152 decreasing at a rate of $O((\log p)^{-1/2})$, and hence Theorem 4.1 applies. By Corol-
 153 lary 4.1, Holm's procedure—and when the errors are independent, the Šidák, and
 154 Hochberg procedures—with FWER controlled at $(\log p)^{-1/2}$ all achieve perfect sup-
 155 port recovery provided that $r > f_E(\beta)$.

Proof (Example 4.1) By Mill's ratio for the standard Gaussian distribution,

$$\frac{t_p \mathbb{P}[Z > t_p]}{\phi(t_p)} \rightarrow 1, \quad \text{as } t_p \rightarrow \infty,$$

where $Z \sim N(0, 1)$. Using the expression for $t_p = \sqrt{2 \log p}$, we have

$$p \mathbb{P}[Z > t_p] \sim \sqrt{2\pi}^{-1} (2 \log p)^{-1/2} \rightarrow 0,$$

156 as desired. The rest of the claims follow from Corollary 4.1. \square

157 The statements in Theorem 4.1 can be strengthened, to prepare us for a minimax
 158 result given in Sect. 5.5.

159 **Remark 4.1** In the proof of Theorem 4.1, both (4.11) and (4.15) hold uniformly over
 160 all error-dependence structures. Therefore, (4.12) and (4.16) may be strengthened to
 161 yield

$$\lim_{p \rightarrow \infty} \sup_{\substack{\mu \in \Theta_p^+(\beta, r) \\ \mathcal{E} \in D(F)}} P[\hat{S}_p \neq S_p] = 0, \quad (4.19)$$

163 for $r > f_E(\beta)$, where $D(F)$ is the collection of all arrays with common marginal F ,
 164 i.e.,

$$D(F) = \{\mathcal{E} = (\epsilon_p(i))_p : \epsilon_p(i) \sim F \text{ for all } i = 1, \dots, p, \text{ and } p = 1, 2, \dots\}. \quad (4.20)$$



166 **Remark 4.2** We emphasize that Theorem 4.1 holds for errors with *arbitrary* depen-
167 dence structures. Intuitively, this is because the maxima of the errors grow at their
168 fastest in the case of independence (recall Remark 2.1). Formally, the light-tailed
169 nature of the error distribution allowed us to obtain sharp tail estimates via simple
170 union bounds, valid under arbitrary dependence.

171 4.3 Dependence and Uniform Relative Stability

172 An important ingredient needed for a converse of Theorem 4.1 is an appropriate char-
173 acterization of the error-dependence structure under which the strong classification
174 boundary (4.5) is tight. The notion of *uniform relative stability* turns out to be the
175 key.

176 **Definition 4.1** (*Uniform Relative Stability*) Under the notations established in Def-
177 inition 2.8, the triangular array \mathcal{E} is said to have uniform relatively stable (URS)
178 maxima if for every sequence of subsets $S_p \subseteq \{1, \dots, p\}$ such that $|S_p| \rightarrow \infty$, we
179 have

$$\frac{1}{u_{|S_p|}} M_{S_p} := \frac{1}{u_{|S_p|}} \max_{i \in S_p} \epsilon_p(i) \xrightarrow{\mathbb{P}} 1, \quad (4.21)$$

181 as $p \rightarrow \infty$, where u_q , $q \in \{1, \dots, p\}$ is the generalized quantile in (2.37). The
182 collection of arrays $\mathcal{E} = \{\epsilon_p(i)\}$ with URS maxima is denoted $U(F)$.

183 Uniform relative stability is, as its name suggests, a stronger requirement on
184 dependence than relative stability (recall Definition 2.8). Proposition 2.2 states that
185 an array with iid components sharing a marginal distribution F with rapidly varying
186 tails (Definition 2.7) has relatively stable maxima; it is easy to see that URS also
187 follows, by independence of the entries.

188 **Corollary 4.2** *An independent array \mathcal{E} with common marginals $F \in AGG(\nu)$, $\nu >$*
 0 , *is URS; in this case, URS holds with $u_{|S_p|} \sim (\nu \log |S_p|)^{1/\nu}$.*

190 On the other hand, RS and URS hold under much broader dependence structures
191 than just independent errors. These conditions are extremely mild and can be shown
192 to hold for many classes of error models. In Chap. 6, we will focus extensively on the
193 Gaussian case, which is of great interest in applications and is rather challenging. We
194 will provide simple necessary and sufficient condition for uniform relative stability
195 in terms of the covariance structures.

196 The relative stability concepts are important because they characterize the depen-
197 dence structures under which the maxima of error sequences *concentrate* around the
198 quantiles (2.37) in the sense of (2.38). This concentration of maxima phenomena, in
199 turn, is the key to establishing the necessary conditions of the phase-transition results
200 in support recovery problems.

201 4.4 Necessary Conditions for Exact Support Recovery

202 With the preparations from Sect. 4.3, we are ready to state the necessary conditions
 203 for exact support recovery (4.1) by thresholding procedures. It turns out that the
 204 strong classification boundary (4.5) is tight, under the general dependence structure
 205 characterized by URS (Definition 4.1).

206 Formally, we define the parameter space for the signals μ to be

$$207 \Theta_p^-(\beta, \bar{r}) = \{\mu \in \mathbb{R}^p : \text{there exists a set } S_p \subseteq \{1, \dots, p\} \text{ such that } |S_p| = \lfloor s_p^\dagger \rfloor, \\ 208 0 < \mu(i) \leq (\sqrt{r} \log p)^{1/\nu} \text{ for all } i \in S_p, \text{ and } \mu(i) = 0 \text{ for all } i \notin S_p\}, \\ 209 \quad (4.22)$$

210 where $s_p^\dagger = \ell(p)p^{1-\beta}$ is as in (4.3).

211 **Theorem 4.2** Let \mathcal{E} be a triangular array with common $\text{AGG}(v)$ marginal F , $v > 0$.
 212 Assume further that the errors \mathcal{E} have uniform relatively stable maxima and minima,
 213 i.e., $\mathcal{E} \in U(F)$, and $(-\mathcal{E}) = \{-\epsilon_p(i)\} \in U(F)$. If

$$214 \bar{r} < f_E(\beta) = (1 + (1 - \beta)^{1/\nu})^v, \quad (4.23)$$

215 then

$$216 \lim_{p \rightarrow \infty} \inf_{\widehat{S}_p \in \mathcal{T}} \inf_{\mu \in \Theta_p^-(\beta, \bar{r})} \mathbb{P}[\widehat{S}_p \neq S_p] = 1, \quad (4.24)$$

217 where \mathcal{T} is the class of all thresholding procedures (2.20).

218 **Proof** To avoid cumbersome double subscript notations, we will sometimes suppress
 219 dependence on p of the set sequences \widehat{S}_p and S_p in the proof.

220 Since the estimator $\widehat{S}_p = \{x(i) \geq t_p(x)\}$ is thresholding, exact support recovery
 221 takes place if and only if the threshold separates the signals and null part, i.e.,

$$222 \mathbb{P}[\widehat{S}_p = S_p] = \mathbb{P}\left[\max_{i \in S^c} x(i) < t_p(x) \leq \min_{i \in S} x(i)\right] \leq \mathbb{P}\left[\max_{i \in S^c} x(i) < \min_{i \in S} x(i)\right].$$

223 Since the right-hand side does not depend on the procedure \widehat{S}_p , we also have

$$224 \sup_{\widehat{S}_p \in \mathcal{T}} \mathbb{P}[\widehat{S}_p = S_p] \leq \mathbb{P}\left[\max_{i \in S^c} x(i) < \min_{i \in S} x(i)\right] \leq \mathbb{P}\left[\max_{i \in S^c} \epsilon(i) < \bar{\Delta} + \min_{i \in S} \epsilon(i)\right], \\ (4.25)$$

where we used the assumption that the signal sizes are no greater than $\bar{\Delta}$. Let $S^* = S_p^*$
 be a sequence of support sets that maximize the right-hand side of (4.25), i.e., let

$$S_p^* = \arg \max_{S \subseteq \{1, \dots, p\}: |S| = \lfloor s_p^\dagger \rfloor} \mathbb{P}\left[\max_{i \in S^c} \epsilon(i) < \bar{\Delta} + \min_{i \in S} \epsilon(i)\right],$$

where $s_p^\dagger = \ell(p)p^{1-\beta}$ is the size of the true support set, and ties are broken lexicographically if multiple maximizers exist. Then, we obtain the following bound which only depends on \bar{r} and the distribution of \mathcal{E} ,

$$\begin{aligned} \sup_{\widehat{S}_p \in \mathcal{T}} \sup_{\mu \in \Theta_p^-(\beta, \bar{r})} \mathbb{P}[\widehat{S}_p = S_p] &\leq \mathbb{P} \left[\max_{i \in S^{*c}} \epsilon(i) < \bar{\Delta} + \min_{i \in S^*} \epsilon(i) \right] \\ &= \mathbb{P} \left[\frac{M_{S^{*c}}}{u_p} < \frac{\bar{\Delta} - m_{S^*}}{u_p} \right], \end{aligned} \quad (4.26)$$

where $M_{S^{*c}} = \max_{i \in S^{*c}} \epsilon(i)$ and $m_{S^*} = \max_{i \in S^*} (-\epsilon(i))$. Since the error arrays \mathcal{E} and $(-\mathcal{E})$ are URS by assumption, using the expression for the AGG quantiles (2.33), we have

$$\frac{M_{S^{*c}}}{u_p} = \frac{M_{S^{*c}}}{u_{|S^{*c}|}} \frac{u_{|S^{*c}|}}{u_p} \xrightarrow{\mathbb{P}} 1, \quad \text{and} \quad \frac{m_{S^*}}{u_p} = \frac{m_{S^*}}{u_{|S^*|}} \frac{u_{|S^*|}}{u_p} \xrightarrow{\mathbb{P}} (1 - \beta)^{1/\nu}, \quad (4.27)$$

so that the two random terms in probability (4.26) converge to constants. Notice that the second relation in (4.27) holds by URS for any $\beta \in (0, 1)$. When $\beta = 1$, the relation holds because $\{m_{S^*}/u_{|S^*|}\}$ is tight, while $0 \leq u_{|S^*|}/u_p \leq u_{\ell(p)}/u_p \rightarrow 0$ since $\ell(p) = o(p)$ by the Potter bounds for slowly varying functions (see, e.g., Bingham et al. 1987).

Since signal sizes are bounded above by $\bar{r} < (1 + (1 - \beta)^{1/\nu})^\nu$, we can write $\bar{r}^{1/\nu} = 1 + (1 - \beta)^{1/\nu} - d$ for some $d > 0$. By our parametrization of $\bar{\Delta}$, we have

$$\frac{\bar{\Delta}}{u_p} = (1 + (1 - \beta)^{1/\nu} - d)(1 + o(1)). \quad (4.28)$$

Combining (4.27) and (4.28), we conclude that the right-hand side of the probability (4.26) converges in probability to a constant strictly less than 1, that is,

$$\frac{\bar{\Delta} - m_{S^*}}{u_p} \xrightarrow{\mathbb{P}} 1 - d, \quad (4.29)$$

while $M_{S^{*c}}/u_p \xrightarrow{\mathbb{P}} 1$. Therefore, the probability in (4.26) must go to 0. \square

We end this section with several remarks on the scope and consequences of our results. Our first comment is on the signal sizes and, in particular, on the gap between the sufficient conditions (Theorem 4.1) and the necessary conditions (Theorem 4.2).

Remark 4.3 (*Minding the gap*) The sufficient condition in Theorem 4.1 requires that *all* signals be larger than the strong classification boundary $f_E(\beta)$ in order to achieve exact support recovery (4.1), while Theorem 4.2 states that exact support recovery fails (in the sense of (4.6)) when *all* signal sizes are below the boundary—the two conditions are *not* complements of each other. This gap between the sufficient

and necessary conditions on signal sizes, however, may be difficult to bridge. Indeed, in general, when signal sizes straddle the boundary $f_E(\beta)$, either outcome is possible, as we demonstrate in Example 4.2.

Example 4.2 (*Signals straddling the boundary*) Let the signal μ have $|S_p| = \lfloor p^{(1-\beta)} \rfloor$ non-zero entries composed of two disjoint sets $S_p = S_p^{(1)} \cup S_p^{(2)}$. Let also the magnitude of the signals be equal within the two sets, i.e., $\mu(i) = \sqrt{2r^{(k)}} \log p$ if $i \in S_p^{(k)}$ for some constants $r^{(k)} > 0$ for $k = 1, 2$. For simplicity, assume that the errors are iid standard Gaussians.

Consider two scenarios

- 264 1. $r^{(1)} = (1 + \delta)f_E(\beta)$, $r^{(2)} = (1 + \delta)$, with $|S_p^{(1)}| = |S_p| - 1$, $|S_p^{(2)}| = 1$,
- 265 2. $r^{(1)} = (1 + \delta)f_E(\beta)$, $r^{(2)} = (1 - \delta)f_E(\beta)$, with $|S_p^{(1)}| = \lfloor |S_p|/2 \rfloor$, $|S_p^{(2)}| =$
- 266 $|S_p| - |S_p^{(1)}|$

267 for some constants $0 < \delta < 1 - \beta < 1$. In both cases, signals in $S_p^{(1)}$ (respectively, 268 $S_p^{(2)}$) are above (respectively, below) the strong classification boundary (4.5). How- 269 ever, in the first scenario, we have $\mathbb{P}[\widehat{S}_p^{\text{Bonf}} = S_p] \rightarrow 1$ where $\widehat{S}_p^{\text{Bonf}}$ is Bonfer- 270 roni's procedure described in Theorem 4.1, while in the second scenario, we have 271 $\mathbb{P}[\widehat{S}_p = S_p] \rightarrow 0$ for all thresholding procedures \widehat{S}_p .

272 **Proof** (Example 4.2) In the first scenario, signal sizes in $S_p^{(1)}$ are by definition above 273 the strong classification boundary (4.5). The signal in $S_p^{(2)}$ has size parameter $1 + \delta <$ 274 $2 - \beta < (1 + \sqrt{1 - \beta})^2$, and therefore falls below the boundary.

It remains to show that $\mathbb{P}[\widehat{S}_p^{\text{Bonf}} = S_p] \rightarrow 1$. To do so, we define two new arrays

$$\mathcal{Y}^{(k)} = \{y_p^{(k)}(j), j = 1, 2, \dots, p\}, \quad k \in \{1, 2\}_p,$$

where $y_p^{(k)}(j) = x_p(j)$ if $j \notin S_p^{(k)}$, and $y_p^{(k)}(j) = \tilde{\epsilon}_p(j)$ if $j \in S_p^{(k)}$, using an independent error array $\{\tilde{\epsilon}_p(j), j = 1, \dots, p\}$ with iid standard Gaussian elements. That is, we replace the elements in $S_p^{(1)}$ and $S_p^{(2)}$ with iid standard Gaussian noise. Notice both arrays $\mathcal{Y}^{(1)}$ and $\mathcal{Y}^{(2)}$ satisfy the conditions in Theorem 4.1 (with sparsity parameter equal to β and 1, respectively). Hence, we have

$$\mathbb{P}[\widehat{S}_p^{\text{Bonf}} \subseteq S_p] = \mathbb{P}\left[\max_{j \in S^c} x(j) \leq t_p\right] \leq \mathbb{P}\left[\max_{j \in S^c} y^{(1)}(j) \leq t_p\right] \rightarrow 0,$$

275 and

$$\begin{aligned} 276 \quad \mathbb{P}[\widehat{S}_p^{\text{Bonf}} \supseteq S_p] &= \mathbb{P}\left[\min_{j \in S} x(j) > t_p\right] \geq 1 - \mathbb{P}\left[\min_{j \in S^{(1)}} x(j) \leq t_p\right] - \mathbb{P}\left[\min_{j \in S^{(2)}} x(j) \leq t_p\right] \\ 277 &\geq 1 - \mathbb{P}\left[\min_{j \in S^{(1)}} y_p^{(2)}(j) \leq t_p\right] - \mathbb{P}\left[\min_{j \in S^{(2)}} y_p^{(1)}(j) \leq t_p\right] \rightarrow 1, \end{aligned}$$

279 where t_p is the threshold in Bonferroni's procedure. The conclusion follows.

In the second scenario, the signal sizes in $S^{(2)}$ by definition fall below the strong classification boundary (4.5). To see that no thresholding procedure succeeds, we adapt the proof of Theorem 4.2. In particular, we obtain

$$\mathbb{P}[\widehat{S}_p = S_p] \leq \mathbb{P}\left[\max_{j \in S^c} x(j) \leq t_p < \min_{j \in S} x(j)\right] \leq \mathbb{P}\left[\max_{j \in S^c} x(j) < \min_{j \in S^{(2)}} x(j)\right].$$

280 By the assumption that signals in $S^{(2)}$ have size parameter $(1 - \delta)f_E(\beta)$, we have

$$\mathbb{P}\left[\max_{j \in S^c} x(j) < \min_{j \in S^{(2)}} x(j)\right] = \mathbb{P}\left[\frac{M_{S^c}}{u_p} < \frac{\sqrt{2(1 - \delta)f_E(\beta) \log p} - m_{S^{(2)}}}{u_p}\right], \quad (4.30)$$

281 where $M_{S^c} = \max_{j \in S^c} \epsilon(j)$ and $m_{S^{(2)}} = \max_{j \in S^{(2)}} (-\epsilon(j))$. The ratio on the left-hand
282 side of the inequality converges to 1 as in (4.27) in the main text, whereas the term
283 on the right-hand side
284

$$\begin{aligned} 285 \frac{\sqrt{2(1 - \delta)f_E(\beta) \log p} - m_{S^{(2)}}}{u_p} &= \sqrt{(1 - \delta)f_E(\beta)} - \frac{m_{S^{(2)}}}{u_{|S^{(2)}|}} \frac{u_{|S^{(2)}|}}{u_p} \\ 286 &\xrightarrow{\mathbb{P}} \sqrt{(1 - \delta)} + \sqrt{1 - \beta}(\sqrt{(1 - \delta)} - 1) < 1, \end{aligned}$$

where we used the URS of the error arrays, and that

$$u_{|S^{(2)}|} \sim \sqrt{2 \log(p^{1-\beta}/2)} = \sqrt{2((1 - \beta) \log p - \log 2)} \sim \sqrt{2(1 - \beta) \log p}$$

288 to conclude the convergence in probability. □

289 Our second remark is on the restriction to thresholding procedures.

290 **Remark 4.4** Since the sharp phase-transition result just established apply only to
291 the general class of thresholding procedures, it is natural to ask if other good proce-
292 dures have left out by this restriction. We will establish later in Chap. 5 that in many
293 cases the optimal procedures are, in fact, thresholding procedures. In general, how-
294 ever, thresholding procedures can be sub-optimal, e.g., when the errors have heavy
295 (regularly varying) tails. We will also demonstrate the absence of a phase-transition
296 phenomenon in exact support recovery by thresholding, in Supplement Sect. B.2.

297 Our final comment is on the interplay between thresholding procedures and the
298 dependence class characterized by URS.

299 **Remark 4.5** Paraphrasing Theorems 4.1 and 4.2: if we consider only thresholding
300 procedures, then for a very large class of dependence structures, we cannot improve
301 upon the Bonferroni procedure $\widehat{S}_p^{\text{Bonf}}$. Specifically, for all $\mathcal{E} \in U(F)$ and $-\mathcal{E} \in U(F)$,
302 and for all $S_p \in \mathcal{S}$, where $\mathcal{S} = \{S \subseteq \{1, \dots, p\}; |S| = \lfloor \ell(p)p^{1-\beta} \rfloor\}$, we have

$$\lim_{p \rightarrow \infty} \mathbb{P}[\widehat{S}_p^{\text{Bonf}} \neq S_p] = \begin{cases} \limsup_{p \rightarrow \infty} \inf_{\widehat{S}_p \in \mathcal{T}} \mathbb{P}[\widehat{S}_p \neq S_p] = 0, & \text{if } \underline{r} > f_E(\beta), \\ \liminf_{p \rightarrow \infty} \inf_{\widehat{S}_p \in \mathcal{T}} \mathbb{P}[\widehat{S}_p \neq S_p] = 1, & \text{if } \bar{r} < f_E(\beta), \end{cases} \quad (4.31)$$

where \mathcal{T} is the set of all thresholding procedures (2.20).

Theorem 4.2 answers a question raised in Butucea et al. (2018). In particular, the authors of (Butucea et al. 2018) commented that independent error is the “least favorable model” in the problem of support recovery, and conjectured that the support recovery problem may be easier to solve under dependence, similar to how the problem of signal detection is easier under dependent errors (Hall and Jin 2010). Surprisingly, our results here state that asymptotically *all* error-dependence structures in the large URS class are equally difficult for thresholding procedures. Therefore, the phase-transition behavior is universal in the class of dependence structures characterized by URS.

We emphasize the restriction to the URS dependence class in Theorem 4.2 is not an assumption of convenience. The dependence condition characterized by uniform relative stability is, in fact, one of the weakest in the literature. We will characterize the class URS dependence class in Chap. 6.

4.5 Dense Signals

We treat briefly the case of dense signals, where the size of the support set is proportional to the problem dimension, i.e., $s \sim cp$ for some constant $c \in (0, 1)$. We show that in this case, a phase-transition-type result still holds, independently of the value of c . Analogous to the setup of Theorems 4.1 and 4.2, let

$$\Theta_p^{d+}(c, \underline{r}) = \{\mu \in \mathbb{R}^p : \text{there exists a set } S_p \subseteq \{1, \dots, p\} \text{ such that } |S_p| \leq \lfloor cp \rfloor, \\ \mu(i) \geq (\nu \underline{r} \log p)^{1/\nu} \text{ for all } i \in S_p, \text{ and } \mu(i) = 0 \text{ for all } i \notin S_p\}, \quad (4.32)$$

where “d” in the notation Θ_p^{d+} stands for “dense”. Similarly, define

$$\Theta_p^{d-}(c, \bar{r}) = \{\mu \in \mathbb{R}^p : \text{there exists a set } S_p \subseteq \{1, \dots, p\} \text{ such that } |S_p| = \lfloor cp \rfloor, \\ 0 < \mu(i) \leq (\nu \bar{r} \log p)^{1/\nu} \text{ for all } i \in S_p, \text{ and } \mu(i) = 0 \text{ for all } i \notin S_p\}. \quad (4.33)$$

Theorem 4.3 Let $c \in (0, 1)$ be a fixed constant, and let $\widehat{S} = \widehat{S}_p^{\text{Bonf}}$ denote Bonferroni’s procedure as described in Theorem 4.1. In the context of Theorem 4.1, if $\underline{r} > 1$, then we have

$$\lim_{p \rightarrow \infty} \sup_{\mu \in \Theta_p^{d+}(c, \underline{r})} \mathbb{P}[\widehat{S}_p \neq S_p] = 0. \quad (4.34)$$



334 While in the context of Theorem 4.2, if $\bar{r} < 1$, then

$$\lim_{p \rightarrow \infty} \inf_{\widehat{S}_p \in \mathcal{T}} \inf_{\mu \in \Theta_p^{d-}(c, \bar{r})} \mathbb{P}[\widehat{S}_p \neq S_p] = 1, \quad (4.35)$$

336 where \mathcal{T} is the class of all thresholding procedures (2.20).

337 **Remark 4.6** Notice that the boundary for the signal size parameter is identically 1
338 in this dense regime. Therefore, if we interpret $\beta = 0$ of the parametrization (4.3)
339 as $s \sim cp$, where $c \in (0, 1)$, then the strong classification boundary (4.5) may be
340 continuously extended to the left-end point where $f_E(0) = 1$.

Proof (Theorem 4.3) The proof is entirely analogous to that of Theorems 4.1 and 4.2. Specifically, (4.34) follows by replacing $\lfloor p^{1-\beta} \rfloor$ with $\lfloor cp \rfloor$ in Relation (4.15) onward, and replacing (4.18) with

$$u_s^- \sim (\nu \log cp)^{1/\nu} \sim (\nu \log p)^{1/\nu}$$

in the proof of Theorem 4.1. Similarly, (4.35) follows the proof of Theorem 4.2. Indeed, by using the fact that

$$\frac{u_{|S^*|}}{u_p} \sim \frac{(\nu \log (1-c)p)^{1/\nu}}{(\nu \log p)^{1/\nu}} \rightarrow 1$$

341 and $u_{|S^*|}/u_p \rightarrow 1$ for all $c \in (0, 1)$, we see that Relation (4.27) holds with $\beta = 0$,
342 and the rest of Theorem 4.2 applies. \square

343 4.6 Numerical Illustrations for Independent Errors

We examine numerically Boundaries (4.5) under several error tail assumptions for independence errors in this section. Numerical experiments for dependent errors will be deferred until we characterize the URS conditions in Chap. 6.

To demonstrate the phase-transition phenomenon under different error tail densities, we simulate from the additive error model (1.1) with

- 349 • Gaussian errors, where the density is given by $f(x) = \frac{1}{\sqrt{2\pi}} \exp \{-x^2/2\}$.
- 350 • Laplace errors, where the density is given by $f(x) = \frac{1}{2} \exp \{-|x|\}$.
- 351 • Generalized Gaussian $\nu = 1/2$, with density $f(x) = \frac{1}{2} \exp \{ -2|x|^{1/2} \}$.

The sparsity and signal size of the sparse mean vector are parametrized as in Eqs. (4.3) and (4.4), respectively. The support set S is estimated with $\widetilde{S} = \{i : x(i) > \sqrt{2 \log p}\}$ under the Gaussian errors, $\widetilde{S} = \{i : x(i) > \log p + (\log \log p)/2\}$ under the Laplace errors, and with $\widetilde{S} = \{i : x(i) > \frac{1}{4} (W(-c/(ep \log p)) + 1)^2\}$ under the generalized Gaussian ($\nu = 1/2$) errors. Here W is the Lambert W function, i.e., $W = f^{-1}$ where

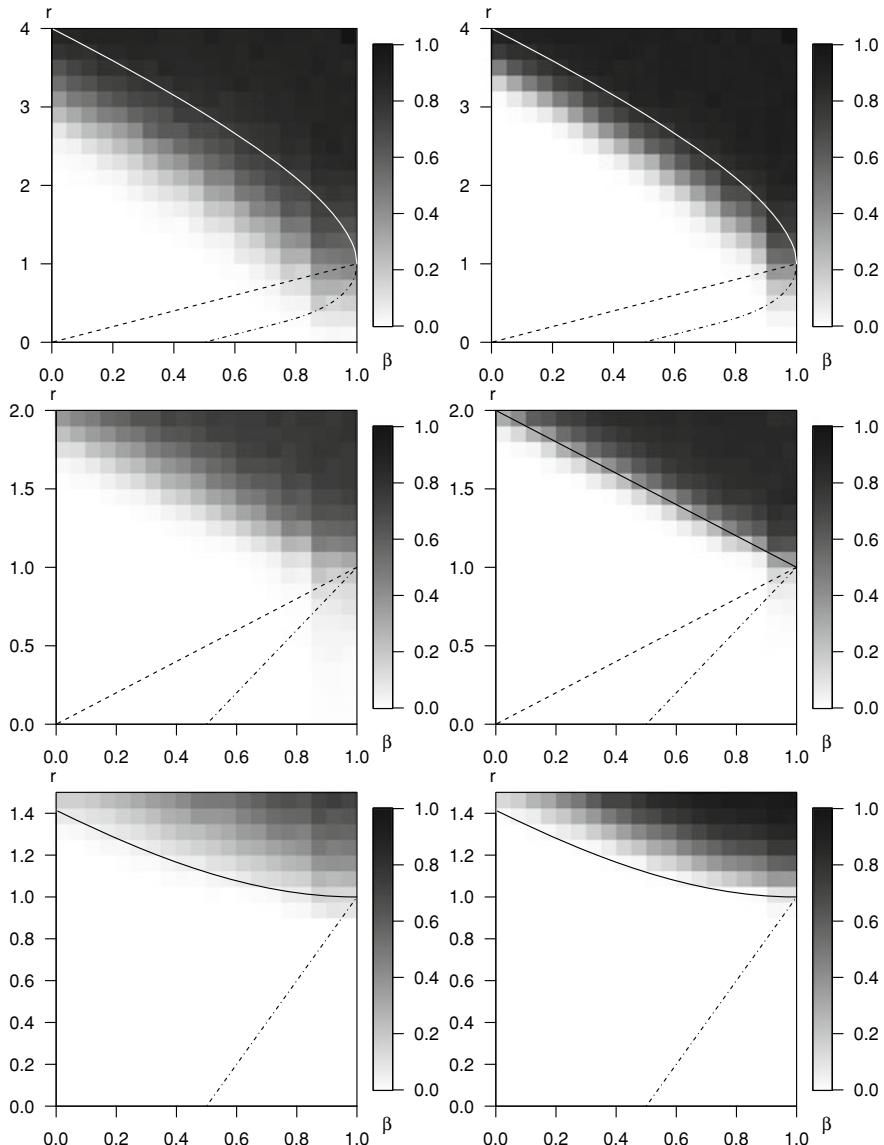


Fig. 4.1 The empirical probability of exact support recovery from numerical experiments, as a function of sparsity level β and signal sizes r , from Gaussian error models (upper panels), Laplace error models (middle panels), and generalized Gaussian with $\nu = 1/2$ (lower panels); darker color indicates higher probability of exact support recovery. The experiments were repeated 1000 times for each sparsity–signal size combination, and for dimensions $p = 100$ (left panels) and $p = 10000$ (right panels). Numerical results agree with the boundaries described in Theorem 4.1; convergence is noticeably slower for under generalized Gaussian ($\nu = 1/2$) errors. For reference, the dashed and dash-dotted lines represent the weak classification and detection boundaries (see Chap. 3)

357 $f(x) = x \exp(x)$. The choices of thresholds correspond to Bonferroni's procedures
358 with FWER decreasing at a rate of $1/\sqrt{\log p}$, therefore satisfying the assumptions in
359 Theorem 4.1. Experiments were repeated 1000 times under each sparsity and signal
360 size combination.

361 The results of the numerical experiments are shown in Fig. 4.1. The numerical
362 results illustrate that the predicted boundaries are not only accurate in high dimen-
363 sions ($p = 10000$, right panels of Fig. 4.1), but also practically meaningful even at
364 moderate dimensions ($p = 100$, left panels of Fig. 4.1).

Metadata of the chapter that will be visualized in SpringerLink

Book Title	Concentration of Maxima and Fundamental Limits in High-Dimensional Testing and Inference	
Series Title		
Chapter Title	Bayes and Minimax Optimality	
Copyright Year	2021	
Copyright HolderName	The Author(s), under exclusive license to Springer Nature Switzerland AG	
Corresponding Author	Family Name Particle Given Name Prefix Suffix Role Division Organization Address Email ORCID	Gao Zheng Department of Statistics University of Michigan-Ann Arbor Ann Arbor, MI, USA gaozheng@umich.edu https://orcid.org/0000-0003-1394-8741
Author	Family Name Particle Given Name Prefix Suffix Role Division Organization Address Email ORCID	Stoev Stilian Department of Statistics University of Michigan-Ann Arbor Ann Arbor, MI, USA https://orcid.org/0000-0003-2108-5421
Abstract	The finite-sample Bayes optimality of support estimators for the signal plus noise model is investigated in this chapter. In general, it is shown that likelihood ratio thresholding procedures are always Bayes optimal. Consequently, the Oracle threshold estimation procedures are optimal for the class of log-concave error models and, otherwise, they are typically sub-optimal. These results lead to a slew of asymptotic minimax characterizations of the phase-transition phenomenon in exact support recovery. In particular, the previously established phase transition in the exact support recovery problem for thresholding estimators is shown to hold for all support estimators in a minimax sense, provided the class of error arrays includes independent log-concave models.	

Chapter 5

Bayes and Minimax Optimality



0 In this chapter, we investigate the universality of the phase-transition results on exact
1 support recovery established in Chap. 4. Specifically, we would like to know to what
2 extent the strong classification boundary applies to all support estimators and not just
3 thresholding ones. The answer to this question will complete the characterization of
4 the fundamental limits in exact support recovery.

5 To this end, we begin by characterizing the finite-sample Bayes optimality of the
6 thresholding procedures. As we will see, the so-called oracle thresholding estimators
7 are in fact finite-sample optimal for many classes of models. These optimality results
8 allow us to establish a minimax formulation of the exact support recovery phase-
9 transition phenomenon that covers arbitrary procedures.

10 Perhaps surprisingly, thresholding estimators can be sub-optimal. This is so, for
11 example, in the additive noise model when the error tail densities are heavier than
12 exponential. In this case, we will see that *likelihood ratio thresholding* rather than
13 *data thresholding* are the optimal support estimators.

14 5.1 Bayes Optimality in Support Recovery Problems

15 In studying support recovery problems, restrictions to the thresholding procedures
16 are sometimes justified by arguing that such procedures are the “reasonable” choice
17 for estimating the support set (see, e.g., Arias-Castro and Chen 2017). We show in this
18 chapter that, perhaps surprisingly, for general error models, thresholding procedures
19 are not always optimal, even when the observations are independent.

20 We shall identify the optimal procedure for support recovery problems under a
21 Bayesian setting with general distributional assumptions (including but not limited
22 to additive models (1.1)). Specifically, we assume that there is an ordered set $P =$
23 (i_1, \dots, i_s) , $i_i \in \{1, \dots, p\}$, and s not necessarily equal densities f_1, \dots, f_s , such
24 that the observations indexed by set P have corresponding densities. That is,

25 $x(i_j) \sim f_j, \quad j = 1, \dots, s.$ (5.1)

26 Let also the rest $(p - s)$ observations have common density f_0 , i.e., $x(i) \sim f_0$ for
 27 $i \notin S$. We further assume that the observations x are mutually independent.

28 We adopt here a Bayesian framework to measure statistical risks. Let the ordered
 29 support $P = (i_1, \dots, i_s)$ have prior

30 $\pi((i_1, \dots, i_s)) = (p - s)!/p!,$ (5.2)

31 for all distinct $1 \leq i_1 < \dots < i_s \leq p$. Consequently, the unordered support $S =$
 32 $\{i_1, \dots, i_s\}$ is distributed uniformly in the collection of all set of size s , with
 33 the unordered uniform distribution π^u . That is, for all $S \in \mathcal{S} := \{S \subseteq \{1, \dots, p\};$
 34 $|S| = s\}$, we have

35 $\pi^u(\{i_1, \dots, i_s\}) = \sum_{\sigma} \pi((i_{\sigma(1)}, \dots, i_{\sigma(s)})) = (p - s)!s!/p!,$ (5.3)

36 where the sum is taken over all permutations of $\{1, 2, \dots, s\}$.

For any fixed configuration P , consider the loss function,

$$\ell(\hat{S}, S) := \mathbb{P}[\hat{S} \neq S] = \mathbb{P}_P[\hat{S} \neq S],$$

37 where the probability is taken over the randomness in the observations x only. The
 38 Bayes optimal procedures, by definitions, should minimize

39 $\mathbb{E}_{\pi} \mathbb{P}[\hat{S} \neq S],$ (5.4)

40 where the expectation is taken over the random configurations P , with a uniform
 41 distribution π as specified in (5.2).

42 If, however, the sparsity $s = |S|$ of the problem is known, then a “natural” esti-
 43 mator for S would be based on the set of top s -order statistics. Such estimators will
 44 be referred to as oracle thresholding estimators and formally defined next.

For any collection of numbers $\{a_i, i = 1, \dots, s\}$, let

$$\langle a_1, \dots, a_s \rangle := (a_{[1]}, \dots, a_{[s]})$$

45 denote the vector of a_i ’s arranged in a non-increasing order.

Definition 5.1 (Oracle data thresholding) Let $x_{[1]} \geq \dots \geq x_{[p]}$ be the order
 statistics of the data vector x . Any estimator $\hat{S}^* := \{i_1, \dots, i_s\}$, where

$$\langle x(i_1), \dots, x(i_s) \rangle = (x_{[1]}, \dots, x_{[s]})$$

46 will be referred to as an *oracle thresholding estimator*.

Simply put, the oracle thresholding estimators are comprised of the indices corresponding to the s largest values in the data. Note that, in the absence of ties among the largest $s + 1$ data values, the oracle thresholding estimator is unique. For concreteness, one can break possible ties lexicographically. In many cases, the oracle thresholding estimators will be almost surely unique.

5.2 Bayes Optimality of Oracle Thresholding

In this section, we study the Bayes optimality of the oracle thresholding procedures. The following *monotone likelihood ratio* (MLR) property will play a key role.

Definition 5.2 (*Monotone Likelihood Ratio*) A family of positive densities on \mathbb{R} , $\{f_\delta, \delta \in U\}$, is said to have the MLR property if, for all $\delta_0, \delta_1 \in U \subseteq \mathbb{R}$ such that $\delta_0 < \delta_1$, the likelihood ratio $(f_{\delta_1}(x)/f_{\delta_0}(x))$ is an increasing function of x .

The next result provides a general criterion for the finite-sample Bayes optimality of the oracle thresholding procedure \widehat{S}^* .

Theorem 5.1 *Let the observations $x(i)$, $i = 1, \dots, p$ be as prescribed as in (5.1) through (5.2). If each of the pairs $\{f_0, f_1\}, \dots, \{f_0, f_s\}$ forms an MLR family, then every oracle data thresholding procedure \widehat{S}^* is finite-sample optimal in terms of Bayes risk $\mathbb{E}_\pi \mathbb{P}[\widehat{S} \neq S]$. That is,*

$$\widehat{S}^* \in \arg \min_{\widehat{S}} \mathbb{E}_\pi \mathbb{P}[\widehat{S} \neq S]. \quad (5.5)$$

for all s and p .

Proof The problem of support recovery can be equivalently stated as a classification problem, where the discrete parameter space is $\mathcal{S} = \{S \subseteq \{1, \dots, p\} : |S| = s\}$, and the observation $x \in \mathbb{R}^p$ has likelihood $f(x|S)$ indexed by the support set S .

By the optimality of the Bayes classifier (see, e.g., Domingos and Pazzani 1997), a set estimator that maximizes the probability of support recovery is one such that

$$\widehat{S} \in \arg \max_{S \in \mathcal{S}} f(x|S)\pi(S).$$

Since we know from (5.3) that $\pi(\cdot)$ is uniform, the problem in our context reduces to showing that $f(x|\widehat{S}^*) = f(x|\widehat{S})$, where $f(x|S)$ is the conditional distribution of data given the unordered support S ,

$$f(x|S) = \sum_{P \in \sigma(S)} f(x|P)\pi^{\text{ord}}(P|S) = \frac{1}{s!} \left(\sum_{P \in \sigma(S)} \prod_{i=1}^s f_i(x(P(i))) \right) \prod_{k \notin S} f_0(x(k)),$$

where $\sigma(S)$ is the set of all permutations of the indices in the support set S .

70 Suppose that \widehat{S} is *not* an oracle thresholding estimator, then there must be indices
 71 $j \in \widehat{S}$ and $j' \notin \widehat{S}$ such that $x(j) < x(j')$. We exchange the classifications of $x(j)$
 72 and $x(j')$, and form a new estimate $\widehat{S}' = (\widehat{S} \setminus \{j\}) \cup \{j'\}$. Comparing the likelihoods
 73 under \widehat{S} and \widehat{S}' , we have

$$\begin{aligned}
 74 \quad f(x|\widehat{S}) - f(x|\widehat{S}') &= \frac{1}{s!} \sum_{P \in \sigma(\widehat{S})} \prod_{i=1}^s f_i(x(P(i))) f_0(x(j')) \prod_{k \notin \widehat{S} \cup \{j'\}} f_0(x(k)) - \\
 75 \quad &\quad - \frac{1}{s!} \sum_{P' \in \sigma(\widehat{S}')} \prod_{i=1}^s f_i(x(P'(i))) f_0(x(j)) \prod_{k \notin \widehat{S}' \cup \{j\}} f_0(x(k)) \\
 76 \quad &= \frac{1}{s!} \left(\sum_{i=1}^s a_i \left(f_i(x(j)) f_0(x(j')) - f_i(x(j')) f_0(x(j)) \right) \right) \prod_{k \notin \widehat{S} \cup \{j'\}} f_0(x(k)),
 \end{aligned} \tag{5.6}$$

77

78 where the last equality follows by first summing over all permutations fixing $P(i) = j$
 79 and $P'(i) = j'$, and setting $a_i = \sum_{P \in \sigma(\widehat{S} \setminus \{j\})} \prod_{i' \neq i} f_{i'}(x(P(i')))$. Notice that the a_i 's
 80 are non-negative.

Since $x(j) < x(j')$, and since each of $\{f_0, f_i\}$ is an MLR family, we have

$$\frac{f_i(x(j))}{f_0(x(j))} - \frac{f_i(x(j'))}{f_0(x(j'))} \leq 0 \implies f_i(x(j)) f_0(x(j')) - f_i(x(j')) f_0(x(j)) \leq 0.$$

81 Using Relation (5.6), we conclude that $f(x|\widehat{S}) \leq f(x|\widehat{S}')$. Continuing this way,
 82 we can successively improve the likelihood of every estimator until we arrive at
 83 an oracle thresholding estimator, proving the desired optimality. Note that with the
 84 same argument, we obtain that any two oracle thresholding estimators have the same
 85 likelihood. \square

86 We emphasize that under the MLR conditions in Theorem 5.1, the oracle thresh-
 87 olding procedures are in fact *finite-sample optimal* in the above Bayesian context. Fur-
 88 ther, our setup allows for different alternative distributions, and relaxes the assump-
 89 tions of Butucea et al. (2018) when studying distributional generalizations, where
 90 the alternatives are assumed to be identically distributed.

91 It remains to understand when the key MLR property holds. We elaborate on this
 92 question next. Returning to the more concrete signal-plus-noise model (1.1), it turns
 93 out that the error tail behavior is what determines the optimality of data thresholding
 94 procedures. In this setting, log-concavity of the error densities is *equivalent* to the
 95 MLR property (Lemma 5.1). This, in turn, yields the finite-sample optimality of data
 96 thresholding procedures (Corollary 5.1).

97 **Lemma 5.1** *Let δ be the magnitude of the non-zero signals in the signal-plus-noise
 98 model (1.1) with positive error density f_0 , and let $f_\delta(x) = f_0(x - \delta)$. The family
 99 $\{f_\delta, \delta \in \mathbb{R}\}$ has the MLR property if and only if the error density f_0 is log-concave.*



100 **Proof** Suppose MLR holds, we will show that $f_0(t) = \exp\{\phi(t)\}$ for some concave
 101 function ϕ . By the assumption of MLR, for any $x_1 < x_2$, setting $\delta_0 = 0$, and $\delta_1 =$
 102 $(x_2 - x_1)/2 > 0$, we have

$$103 \log \frac{f_{\delta_1}(x_2)}{f_{\delta_0}(x_2)} = \phi\left(\frac{(x_1 + x_2)}{2}\right) - \phi(x_2) \geq \phi(x_1) - \phi\left(\frac{(x_1 + x_2)}{2}\right) = \log \frac{f_{\delta_1}(x_1)}{f_{\delta_0}(x_1)}.$$

104 This implies that the log-density $\phi(t)$ is midpoint-concave, i.e., for all x_1 and x_2 , we
 105 have

$$106 \phi\left(\frac{(x_1 + x_2)}{2}\right) \geq \frac{1}{2}\phi(x_1) + \frac{1}{2}\phi(x_2). \quad (5.7)$$

107 For Lebesgue measurable functions, midpoint concavity is equivalent to concavity
 108 by the Sierpinski Theorem see, (e.g., Sec I.3 of Donoghue 2014). This proves the
 109 “only-if” part.

110 For the “if” part, when $\phi(t) = \log(f_0(t))$ is log-concave, then for any $\delta_0 < \delta_1$,
 111 and any $x < y$, we have

$$112 \log \frac{f_{\delta_1}(y)}{f_{\delta_0}(y)} - \log \frac{f_{\delta_1}(x)}{f_{\delta_0}(x)} = \phi(y - \delta_1) - \phi(y - \delta_0) - \phi(x - \delta_1) + \phi(x - \delta_0) \geq 0, \quad (5.8)$$

113 where the last inequality is a simple consequence of concavity (see Lemma 5.2
 114 below). This proves the “if” part. \square

Lemma 5.2 Let ϕ be any concave function on \mathbb{R} . For any $x < y \in \mathbb{R}$, and $\delta > 0$,
 we have

$$\phi(x) + \phi(y + \delta) \leq \phi(y) + \phi(x + \delta).$$

115 **Proof** Pick $\lambda = \delta/(y - x + \delta)$, by concavity of f , we have

$$116 \lambda\phi(x) + (1 - \lambda)\phi(y + \delta) \leq \phi(\lambda x + (1 - \lambda)(y + \delta)) = \phi(y), \quad (5.9)$$

117 and

$$118 (1 - \lambda)\phi(x) + \lambda\phi(y + \delta) \leq \phi((1 - \lambda)x + \lambda(y + \delta)) = \phi(x + \delta). \quad (5.10)$$

119 Summing up (5.9) and (5.10) and we arrive at the conclusion as desired. \square

120 Theorem 5.1 and Lemma 5.1 yield immediately the following.

121 **Corollary 5.1** Consider the additive error model (1.1), where the $\epsilon(i)$'s are inde-
 122 pendent with common distribution F . Let the signal μ have s positive entries with
 123 magnitudes $0 < \delta_1 \leq \dots \leq \delta_s$, located on $\{1, \dots, p\}$ as prescribed in (5.2).



If F has a positive, log-concave density f , then the support estimator

$$\widehat{S}^* := \{i : x(i) \geq x_{[s]}\}$$

124 is finite-sample optimal in terms of Bayes risk in the sense of (5.5).

125 **Proof** The independence and the fact that the observations have densities implies
126 the absence of ties among the order statistics $\{x_{[i]}\}$, with probability one. Thus, the
127 oracle thresholding procedure is a.s. unique and given by $\widehat{S}^* = \{i : x(i) \geq x_{[s]}\}$. The
128 result then follows from Theorem 5.1 and Lemma 5.1. \square

129 **Remark 5.1** Theorem 5.1 and Corollary 5.1 show that under MLR (or equivalently,
130 log-concavity of the errors in additive models), the oracle thresholding procedures
131 are finite-sample optimal even in the case where the signals have different (positive)
132 sizes. This fascinating property perhaps explains the success of the thresholding
133 estimators.

134 The assumption of log-concavity of the densities is compatible with the AGG
135 model when $v \geq 1$, as demonstrated in the next example.

136 **Example 5.1** The generalized Gaussian density $f(x) \propto \exp\{-|x|^v/v\}$ is
137 log-concave for all $v \geq 1$. Therefore in the additive error model (1.1), according
138 to Corollary 5.1, the oracle thresholding procedure is Bayes optimal in the sense
139 of (5.5).

140 5.3 Bayes Optimality of Likelihood Ratio Thresholding

141 When the MLR condition in Theorem 5.1 is violated, the oracle thresholding proce-
142 dures can in fact be sub-optimal (see Example 5.2 and Sect. 5.4, below).

143 In this section, we demonstrate that thresholding the *likelihood ratio* rather than
144 signal values yields the finite-sample Bayes optimal procedures. We consider a spe-
145 cial but sufficiently general case of signal models with equal densities.

Namely, let the observations $x(i)$, $i = 1, \dots, p$ have s signals as prescribed in (5.2) with *common* “signal” density f_a , and let the remaining $(p - s)$ locations have common “error” density f_0 . Define the likelihood ratios

$$L(i) := f_a(x(i)) / f_0(x(i)),$$

146 and let $L_{[1]} \geq L_{[2]} \geq \dots \geq L_{[p]}$ be the order statistics of the $L(i)$ ’s.

Definition 5.3 (Oracle likelihood ratio thresholding) Recall that $\langle a_1, \dots, a_s \rangle$ denotes the vector of a_i ’s arranged in a non-increasing order. Any estimator $\widehat{S} = \{i_1, \dots, i_s\}$ such that

$$\langle L(i_1), \dots, L(i_s) \rangle = (L_{[1]}, \dots, L_{[s]}),$$

¹⁴⁷ will be referred to as an *oracle likelihood thresholding* estimator of the support S .

¹⁴⁸ **Theorem 5.2** Any oracle likelihood ratio thresholding procedure \widehat{S}_{LRT} is finite-
¹⁴⁹ sample optimal in terms of Bayes risk. That is,

$$\widehat{S}_{\text{LRT}} \in \arg \min_{\widehat{S} \in \mathcal{S}} \mathbb{E}_\pi \mathbb{P}[\widehat{S} \neq S] \quad (5.11)$$

¹⁵¹ for all s and p , where the infimum on \widehat{S} is taken over all support estimators of size s .

Proof The proof is analogous to that of Theorem 5.1. We need to show that $\widehat{S}_{\text{LRT}} \in \arg \max_{S \in \mathcal{S}} f(x|S)\pi(S)$. Since the distribution π of the support S is uniform (recall (5.3)), it is equivalent to prove that

$$f(x|\widehat{S}_{\text{LRT}}) = \max_{S \in \mathcal{S}} f(x|S),$$

¹⁵² where $f(x|S)$ is the conditional distribution of the data given the unordered support
¹⁵³ S ,

$$f(x|S) = \sum_P f(x|P)\pi^{\text{ord}}(P|S) = \prod_{j \in S} f_a(x(j)) \prod_{j \notin S} f_0(x(j)). \quad (5.12)$$

¹⁵⁵ Suppose $\widehat{S} \in \mathcal{S}$ is *not* an oracle likelihood thresholding estimator. Then from the
¹⁵⁶ definition of the likelihood ratio thresholding procedure, there must be indices $j \in \widehat{S}$
¹⁵⁷ and $j' \notin \widehat{S}$ such that $L(j) < L(j')$. If we exchange the labels of $L(j)$ and $L(j')$,
¹⁵⁸ that is, we form a new estimate $\widehat{S}' = (\widehat{S} \setminus \{j\}) \cup \{j'\}$, comparing the log-likelihoods
¹⁵⁹ under \widehat{S} and \widehat{S}' , we have

$$\log f(x|\widehat{S}) - \log f(x|\widehat{S}') = \log f_a(x(j)) + \log f_0(x(j')) - \log f_a(x(j')) - \log f_0(x(j)).$$

¹⁶¹ By the definition of $L(j)$'s, and the order relations, we obtain

$$\log f(x|\widehat{S}) - \log f(x|\widehat{S}') = \log L(j) - \log L(j') > 0$$

¹⁶³ This shows that \widehat{S} cannot be Bayes optimal unless it is a likelihood thresholding
¹⁶⁴ estimator. Note that with the same argument for every two likelihood thresholding
¹⁶⁵ estimators \widehat{S}' and \widehat{S}'' , we have $f(x|\widehat{S}') = f(x|\widehat{S}'')$, proving the desired
¹⁶⁶ optimality. \square

¹⁶⁷ The characterization of optimal likelihood ratio thresholding procedures in Theorem 5.2
¹⁶⁸ may not always yield practical estimators, as the density of the alternatives,
¹⁶⁹ and the number of signals s are typically unknown. Still, some insights can be gained
¹⁷⁰ by virtue of Theorem 5.2. In particular, when MLR fails (for example, when the errors
¹⁷¹ in model (1.1) do not have log-concave densities), data thresholding is sub-optimal.

Example 5.2 (*Sub-optimality of data thresholding*) Let the errors have iid generalized Gaussian density with $\nu = 1/2$, i.e., $\log f_0(x) \propto -x^{1/2}$. Let dimension $p = 2$, sparsity $s = 1$ with uniform prior, and signal size $\delta = 1$. That is, $\mathbb{P}[\mu = (0, 1)^T] = \mathbb{P}[\mu = (1, 0)^T] = 1/2$. If the observations take on values $x = (x_1, x_2)^T = (1, 2)^T$, we see from a comparison of the likelihoods (and hence, the posteriors),

$$\log \frac{f(x|\{1\})}{f(x|\{2\})} = 2x_1^{1/2} + 2(x_2 - 1)^{1/2} - 2x_2^{1/2} - 2(x_1 - 1)^{1/2} = 4 - 2\sqrt{2} > 0,$$

that even though $x_1 < x_2$, the set $\{1\}$ is a better estimate of support than $\{2\}$, i.e., $\mathbb{P}[S = \{1\} | x] > \mathbb{P}[S = \{2\} | x]$.

This simple example shows that, in the case when the errors have super-exponential tails, the optimal procedures are in general *not* data thresholding. A slightly more general conclusion can be found in Corollary 5.2.

5.4 Sub-optimality of Data Thresholding Procedures

We provide a slightly more general result on the sub-optimality of data thresholding procedures.

Corollary 5.2 Consider the additive error model (1.1). Let the errors ϵ be independent with common distribution F . Let each of the s signals be located on $\{1, \dots, p\}$ uniformly at random with equal magnitude $0 < \delta < \infty$. Assume the errors $\epsilon(i)$'s are iid with density f that is log-convex on $[K, +\infty)$, for some $K > 0$.

If \widehat{S}_{opt} is the Bayes optimal (i.e., the oracle likelihood thresholding estimator), then, whenever $j \in \widehat{S}_{opt}$ for some $x(j) > K + \delta$, we must necessarily have $j' \in \widehat{S}_{opt}$ for all j' such that $K + \delta \leq x(j') < x(j)$.

Specifically, if there are m observations exceeding $K + \delta$, with $m > s$, then the top $m - s$ observations will *not* be included in the optimal estimator \widehat{S}_{opt} . This shows that, in the case when the errors have super-exponential tails, the optimal procedures are in general *not* data thresholding.

Proof (Corollary 5.2) Since the density of the alternatives $f_a(t) = f(t - \delta)$ is log-convex on $[K + \delta, \infty)$, by Relation (5.8) in the proof of Lemma 5.1 and appealing to log-convexity (rather than log-concavity), the likelihood ratio

$$L(j) := \frac{f_a(x(j))}{f_0(x(j))}$$

is decreasing in $x(j)$ on $[K + \delta, \infty)$. The claim follows from Theorem 5.2. □



192 **Remark 5.2** As we have seen, the thresholding estimators are no longer optimal
193 in the additive model with error densities heavier than exponential. Thanks to The-
194 orem 5.2, the oracle likelihood thresholding procedures are promising alternatives
195 that can lead us to practical support estimators.

196 In the case where the signals have different sizes, however, the argument in the
197 proof of Theorem 5.2 breaks down since the signal densities need to be identical for
198 Relation (5.12) to hold. In such cases, the characterization of the optimal procedure
199 is an open problem.

200 5.5 Minimax Optimality in Exact Support Recovery

201 We establish in this section minimax versions of our results from Chap. 4. Specifically,
202 if we restrict ourselves to *the class of thresholding procedures* \mathcal{T} (defined in (2.20)),
203 then Bonferroni's procedure is minimax optimal, for *any* fixed dependence structure
204 in the URS class. This is formalized in Corollary 5.3 below. We refer to this result as
205 *point-wise* minimax, to emphasize the fact that this optimality holds for every *fixed*
206 URS array.

207 Meanwhile, if we search over *all procedures*, but expand the model space to
208 include *all* dependence structures, then a different minimax optimality statement
209 holds for Bonferroni's procedure. This result, formally stated in Sect. 5.5.2, is a con-
210 sequence of our characterization of the finite-sample Bayes optimality of thresholding
211 procedures in Sect. 5.2.

212 5.5.1 Point-Wise Minimax Optimality for Thresholding 213 Procedures

214 Theorems 4.1 and 4.2 can be cast in the form of an asymptotic minimax statement.

215 **Corollary 5.3** (Point-wise minimax) Let $\widehat{S}^{\text{Bonf}}$ be the sequence of Bonferroni's
216 procedure described in Theorem 4.1. Let also the errors have common $\text{AGG}(v)$ dis-
217 tribution F with parameter $v > 0$, and $\Theta_p^+(\beta, \underline{r})$ be as defined in (4.7). If $\underline{r} > f_E(\beta)$,
218 then we have

$$\limsup_{p \rightarrow \infty} \sup_{\mu \in \Theta_p^+(\beta, \underline{r})} \mathbb{P}(\widehat{S}_p^{\text{Bonf}} \neq S_p) = 0, \quad (5.13)$$

220 for arbitrary dependence structure of the error array $\mathcal{E} = \{\epsilon_p(i)\}_p$. Let \mathcal{T} be the
221 class of thresholding procedures (2.20). If $\underline{r} < f_E(\beta)$, then we have

$$\liminf_{p \rightarrow \infty} \inf_{\widehat{S}_p \in \mathcal{T}} \sup_{\mu \in \Theta_p^+(\beta, \underline{r})} \mathbb{P}(\widehat{S}_p \neq S_p) = 1, \quad (5.14)$$

223 for any error-dependence structure such that $\mathcal{E} \in U(F)$ and $(-\mathcal{E}) \in U(F)$.

Proof The first conclusion (5.13) is a restatement of Theorem 4.1.

For the second statement (5.14), since $\underline{r} < f_E(\beta)$, we can pick a sequence $\mu^* \in \Theta_p^+(\beta, \underline{r})$ such that $|S_p| = \lfloor \ell(p)p^{1-\beta} \rfloor$, with signals having the same signal size $\mu(i) = (2r \log p)^{1/\nu}$ for all $i \in S_p$, where $\underline{r} < r < f_E(\beta)$. For this particular choice of μ^* , we have $\mu^* \in \Theta_p^-(\beta, \bar{r})$ (recall (4.22)), where $r < \bar{r} < f_E(\beta)$, and according to Theorem 4.2, we obtain $\lim_{p \rightarrow \infty} \inf_{\widehat{S}_p \in \mathcal{T}} \mathbb{P}[\widehat{S}_p \neq S_p] = 1$, for all dependence structures in the URS class. \square

Remark 5.3 Theorem 4.2 is a stronger result than the traditional minimax claim in Relation (5.14). Indeed, (4.24) involves an infimum (over the class Θ_p^-), while (5.14) has a supremum (over the class Θ_p^+).

On the other hand, Corollary 5.3 is more informative than many minimax-type statements, since it applies “point-wise” to any fixed error-dependence structure in the URS class.

Remark 5.4 Corollary 5.3 echoes Remark 4.5: for a very large class of dependence structures, we cannot improve upon Bonferroni’s procedure in exact support recovery problems (asymptotically), unless we look beyond thresholding procedures.

5.5.2 Minimax Optimality over All Procedures

Consider the asymptotic Bayes risk as defined in (5.4). The statement for the necessary condition of support recovery in Theorem 4.2, with the help of Corollary 5.1, can be strengthened to include all procedures (in the Bayesian context), regardless of whether they are thresholding or not.

Theorem 5.3 Consider the additive model (1.1) where the $\epsilon_p(i)$ ’s are independent and identically distributed with log-concave densities in the AGG class. Let the signals be as prescribed in Corollary 5.1. If the signal sizes fall below the strong classification boundary (4.5), i.e. $\bar{r} < f_E(\beta)$, then we have

$$\liminf_{p \rightarrow \infty} \inf_{\widehat{S}_p} \mathbb{E}_{\pi} \mathbb{P}[\widehat{S}_p \neq S_p] = 1, \quad (5.15)$$

where the infimum on \widehat{S}_p is taken over all procedures.

Proof When the errors are independent with log-concave density, the oracle thresholding procedure \widehat{S}_p^* , by Corollary 5.1, minimizes the Bayes risk (5.4) among all procedures. That is,

$$\liminf_{p \rightarrow \infty} \inf_{\widehat{S}_p} \mathbb{E}_{\pi} \mathbb{P}[\widehat{S}_p \neq S_p] \geq \liminf_{p \rightarrow \infty} \mathbb{E}_{\pi} \mathbb{P}[\widehat{S}_p^* \neq S_p].$$

Since \widehat{S}_p^* belongs to the class of all thresholding procedures, we have

$$\begin{aligned} \liminf_{p \rightarrow \infty} \mathbb{E}_\pi \mathbb{P}[\widehat{S}_p^* \neq S_p] &\geq \liminf_{p \rightarrow \infty} \inf_{\widehat{S}_p \in \mathcal{T}} \mathbb{E}_\pi \mathbb{P}[\widehat{S}_p \neq S_p] \\ &\geq \liminf_{p \rightarrow \infty} \inf_{\widehat{S}_p \in \mathcal{T}} \inf_{S_p} \mathbb{P}[\widehat{S}_p \neq S_p] = 1, \end{aligned}$$

when $\bar{r} < f_E(\beta)$, where the last line follows from Theorem 4.2. \square

Theorem 5.3 allows us to state another minimax conclusion—one in which we search over *all procedures*, by allowing the supremum in the minimax statement to be taken over the dependence structures.

Corollary 5.4 *Let $D(F)$ be the collection of error arrays with common marginal F as defined in (4.20) where F is an AGG(v) distribution. Let also \widehat{S}_p^{Bonf} be Bonferroni's procedure as described in Theorem 4.1. If $\underline{r} > f_E(\beta)$, then we have*

$$\limsup_{p \rightarrow \infty} \sup_{\substack{\mu \in \Theta_p^+(\beta, \underline{r}) \\ \mathcal{E} \in D(F)}} \mathbb{P}(\widehat{S}_p^{Bonf} \neq S_p) = 0. \quad (5.16)$$

Further, when $\underline{r} < f_E(\beta)$, and F has a positive log-concave density f , we have

$$\liminf_{p \rightarrow \infty} \inf_{\widehat{S}_p} \sup_{\substack{\mu \in \Theta_p^+(\beta, \underline{r}) \\ \mathcal{E} \in D(F)}} \mathbb{P}(\widehat{S}_p \neq S_p) = 1, \quad (5.17)$$

where the infimum on \widehat{S}_p is taken over all procedures.

Proof Relation (5.16) is a re-statement of Remark 4.1.

For any distribution π (with a slight abuse of notation) over the parameter space $\Theta_p^+ \times D(F)$, we have

$$\liminf_{p \rightarrow \infty} \inf_{\widehat{S}_p} \sup_{\substack{\mu \in \Theta_p^+(\beta, \underline{r}) \\ \mathcal{E} \in D(F)}} \mathbb{P}(\widehat{S}_p \neq S_p) \geq \liminf_{p \rightarrow \infty} \inf_{\widehat{S}_p} \mathbb{E}_\pi \mathbb{P}(\widehat{S}_p \neq S_p), \quad (5.18)$$

since the supremum is bounded from below by expectations. In particular, define π to be the uniform distribution over the configurations $\Theta_p^* \times I(f)$, where

$$\begin{aligned} \Theta_p^* = \{\mu \in \mathbb{R}^d : |S_p| = \lfloor \ell(p) p^{1-\beta} \rfloor, \mu(i) = 0 \text{ for all } i \notin S, \text{ and} \\ \mu(i) = (\nu r \log p)^{1/\nu} \text{ for all } i \in S, \text{ where } \underline{r} < r < f_E(\beta)\}, \end{aligned}$$

and

$$I(f) = \{\mathcal{E} = (\epsilon_p(i))_p : \epsilon_p(i) \text{ iid with density } f(x) \propto \exp\{-|x|^\nu/\nu\}\}.$$

Since the density f of F is log-concave, the distribution of the signal configurations satisfies the conditions in Theorem 5.3. Thus, the desired conclusion (5.17) follows from Theorem 5.3 and (5.18). \square

Remark 5.5 Since the class $\text{AGG}(\nu)$, $\nu \geq 1$ contains distributions with log-concave densities (Example 5.1), the minimax statement (5.17) continues to hold if the supremum is taken over the entire class $F \in \text{AGG}(\nu)$, $\nu \geq 1$. We opted for a more informative formulation which emphasizes the log-concavity condition on the density of F .

Remark 5.6 Corollary 5.4 is no stronger than Corollary 5.3. In Corollary 5.3 we search over only the class of thresholding procedures, but offer a tight, point-wise lower bound on the asymptotic risk over the class of URS-dependence structures. On the other hand, Corollary 5.4 provides a uniform lower bound for the asymptotic risk over all dependence structures, which may not be tight except in the case of independent errors.

5.6 Optimality and Sub-optimality: A Discussion

We conclude with a brief summary on the optimality and sub-optimality of the thresholding procedures in the problem of exact support estimation. For clarity, we focus on the model (1.1) with *independent* errors.

Theorem 5.3 and Corollary 5.4 provide a nearly complete picture of the difficulty in the exact support recovery problem, in the regime when the thresholding estimators are optimal. Specifically, in such cases the signal classification boundary is universal. On the other hand, Theorem 5.2, and indeed, Example 5.2 demonstrate that thresholding procedures are *sub-optimal* for $\text{AGG}(\nu)$ models with $\nu < 1$. Therefore, the optimality of thresholding procedures (specifically, Bonferroni's procedure) only applies to $\text{AGG}(\nu)$ models with $\nu \geq 1$.

If we restrict the space of methods to only thresholding procedures, then the results in Sect. 5.5.1 state that the phase transition phenomenon—the 0–1 law in the sense of Corollary 5.3—is universal in all error models with rapidly varying tails. This includes $\text{AGG}(\nu)$ models for all $\nu > 0$. In contrast, models with heavy (regularly varying) tailed errors do not exhibit this phenomenon (form more details, see Theorem B.3). We summarize the properties of thresholding procedures in Table 5.1.

Table 5.1 Properties of thresholding procedures under different error distributions when the errors are independent. Properties of the error distributions are listed in brackets

Thresholding procedure (Error distributions)	Bayes optimality (Log-concave density)	Phase transition (Rapidly-varying tails)
$\text{AGG}(\nu)$, $\nu \geq 1$	Yes (Yes)	Yes (Yes)
$\text{AGG}(\nu)$, $0 < \nu < 1$	No (No)	Yes (Yes)
Power laws	No (No)	No (No)



Author Queries

Chapter 5

Query Refs.	Details Required	Author's response
AQ1	Kindly note that the citation “Theorem 5.1” has been changed to “Corollary 5.1”. Kindly check and confirm.	

Metadata of the chapter that will be visualized in SpringerLink

Book Title	Concentration of Maxima and Fundamental Limits in High-Dimensional Testing and Inference	
Series Title		
Chapter Title	Uniform Relative Stability for Gaussian Arrays	
Copyright Year	2021	
Copyright HolderName	The Author(s), under exclusive license to Springer Nature Switzerland AG	
Corresponding Author	Family Name Particle Given Name Prefix Suffix Role Division Organization Address Email ORCID	Gao Zheng Department of Statistics University of Michigan-Ann Arbor Ann Arbor, MI, USA gaozheng@umich.edu https://orcid.org/0000-0003-1394-8741
Author	Family Name Particle Given Name Prefix Suffix Role Division Organization Address Email ORCID	Stoev Stilian Department of Statistics University of Michigan-Ann Arbor Ann Arbor, MI, USA https://orcid.org/0000-0003-2108-5421
Abstract	The general phase-transition results for the exact support recovery investigated in Chap. 4 exploit a concentration of maxima phenomenon for dependent arrays known as the uniform relative stability (URS) property. This property holds for a wide range of models. The purpose of this chapter is to provide a complete characterization of URS for Gaussian triangular arrays. The main result shows that the Gaussian arrays are URS if and only if they are uniformly decreasingly dependent (UDD), which is a simple condition expressed in terms of the covariance structure of the array. In particular, classic results of Simeon Berman are recovered and extended. The proofs utilize tools such as Slepian's lemma, the Sudakov–Fernique inequality, as well as, a curious result on the structure of correlation matrices derived using elementary Ramsey's theory.	

Chapter 6

Uniform Relative Stability for Gaussian Arrays



0 The notion of uniform relative stability (URS) in Definition 4.1 is the key to the
1 necessary conditions for exact support recovery established in Theorem 4.2. In this
2 chapter, we provide a complete characterization of the class of URS Gaussian arrays
3 in terms of a simple condition on their covariance structure. The condition is as
4 follows.

Definition 6.1 (*Uniformly decreasing dependence (UDD)*) Consider a triangular array of jointly Gaussian distributed errors $\mathcal{E} = \{(\epsilon_p(i))_{i=1}^p, p = 1, 2, \dots\}$ with unit variances,

$$\epsilon_p \sim N(0, \Sigma_p), \quad p = 1, 2, \dots$$

5 The array \mathcal{E} is said to be uniform decreasingly dependent (UDD) if for every $\delta > 0$
6 there exists a finite $N(\delta) < \infty$, such that for every $i \in \{1, \dots, p\}$, and $p \in \mathbb{N}$, we
7 have

$$8 \quad \left| \{k \in \{1, \dots, p\} : \Sigma_p(i, k) > \delta\} \right| \leq N(\delta) \quad \text{for all } \delta > 0. \quad (6.1)$$

9 That is, for every coordinate i , the number of elements which are more than δ -
10 correlated with $\epsilon_p(i)$ does not exceed $N(\delta)$.

11 Note that the bound in (6.1) holds uniformly in i and p , and only depends on δ .
12 Also observe that on the left-hand side of (6.1), we merely count in each row of Σ_p
13 the number of exceedances of covariances (not their absolute values!) over level δ .

Remark 6.1 Without loss of generality, we may require that $N(\delta)$ be a monotone non-increasing function of δ , for we can take

$$N(\delta) = \sup_{p,i} \left| \{k : \Sigma_p(i, k) > \delta\} \right|,$$

14 which is non-increasing in δ . Definition 6.1 therefore states that the array is UDD
15 when $N(\delta) < \infty$ for all $\delta > 0$.



16 Observe that the UDD condition does not depend on the order of the coordinates
 17 in the error vector $\epsilon_p = (\epsilon_p(i))_{i=1}^p$. Often times, however, the errors are thought of
 18 coming from a stochastic process indexed by time or space. To illustrate the generality
 19 of the UDD condition, we formulate next a simple sufficient condition (UDD') that
 20 is easier to check in a time-series context.

21 **Definition 6.2 (UDD')** For $\epsilon_p \sim N(0, \Sigma_p)$ with unit variances, an array
 22 $\mathcal{E} = (\epsilon_p(i))_{i=1}^p$ is said to satisfy the UDD' condition if there exist:

- 23 (i) permutations l_p of $\{1, \dots, p\}$, for all $p \in \mathbb{N}$, and
 24 (ii) a non-negative sequence $(r_n)_{n=1}^\infty$ converging to zero $r_n \rightarrow 0$, as $n \rightarrow \infty$,

25 such that

$$26 \quad \sup_{p \in \mathbb{N}} |\Sigma_p(i', j')| \leq r_{|i-j|}. \quad (6.2)$$

27 where $i' = l_p(i)$, $j' = l_p(j)$, for all $i, j \in \{1, \dots, p\}$.

28 **Remark 6.2** Without loss of generality, we may also require that r_n be non-
 29 increasing in n , for we can replace r_n with the non-increasing sequence $r'_n =$
 30 $\sup_{m \geq n} r_m$.

31 **Proposition 6.1** *UDD' implies UDD.*

32 **Proof** Since $r_n \rightarrow 0$, for any $\delta > 0$, there exists an integer $M = M(\delta) < \infty$ such
 33 that $r_n \leq \delta$, for all $n \geq M$. Thus, by (6.2), for every fixed $j' \in \{1, \dots, p\}$, we can
 34 have $|\text{Cov}(\epsilon_p(k'), \epsilon_p(j'))| > \delta$, only if k' belongs to the set:

$$\{k' \in \{1, \dots, p\} : j - M \leq k := l_p^{-1}(k') \leq j + M\},$$

35 where $j := l_p^{-1}(j')$. That is, there are at most $2M + 1 < \infty$ indices $k' \in \{1, \dots, p\}$,
 36 whose covariances with $\epsilon_p(j')$ may exceed δ . Since this holds uniformly in $j' \in$
 37 $\{1, \dots, p\}$, Condition UDD follows with $N(\delta) = 2M + 1$. \square

38 We now state the main result of this chapter. It states that a Gaussian array is URS
 39 if and only if it is UDD. The URS condition essentially requires that the dependencies
 40 decay in a uniform fashion, the rate at which dependence decay does *not* matter.

38 **Theorem 6.1** *Let \mathcal{E} be a Gaussian triangular array with standard normal marginals.
 39 The array \mathcal{E} has uniformly relatively stable (URS) maxima if and only if it is uniformly
 40 decreasing dependent (UDD).*

41 Specifically, for stationary Gaussian arrays, we have the following corollary.

42 **Corollary 6.1** *Let $\mathcal{E} = \{\epsilon_p(i) = Z(i)\}$ for a stationary Gaussian time series $Z =$
 43 $\{Z(i)\}$. Then \mathcal{E} is URS if and only if the autocovariance function $\text{Cov}(Z(k), Z(0)) \rightarrow$
 44 0, as $k \rightarrow \infty$.*

45 Corollary 6.1 follows by Theorem 6.1 and the observation that UDD is equivalent
 46 to vanishing autocovariance of \mathcal{Z} . A slightly weaker form of the “if” part was
 47 established in Theorem 3 of Berman (1964).

48 Returning again to the study of support recovery problems, Theorem 6.1 and the
 49 necessary condition for exact support recovery in Theorem 4.2 yield the following
 50 result.

51 **Corollary 6.2** *For UDD Gaussian errors, the result in Theorem 4.2 holds.*

52 One may ask, whether the UDD (equivalently, URS) condition can be relaxed
 53 further for the phase-transition result in Theorem 4.2 to hold. As a counterpart to
 54 Remark 4.5, we demonstrate next that the dependence conditions in Theorem 4.2
 55 are nearly optimal. Specifically, we show that if the URS-dependence condition is
 56 violated, then it may be possible to recover the support of weaker signals, falling
 57 below the boundary. The main idea is to use the equivalence of URS and UDD
 58 to construct a Gaussian error array, whose correlations do not decay in a uniform
 59 fashion (UDD fails). As we will see, in such a case one can do substantially better
 60 in terms of support recovery. This shows that the URS condition is nearly optimal in
 61 the Gaussian setting. Numerical simulations illustrating this example can be found
 62 in Sect. 4.6, below.

Example 6.1 (*On the tightness of the URS condition for exact support recovery*)
 Suppose $\mathcal{E} = (\epsilon_p(i))_{i=1}^p$ is Gaussian, and is comprised of $\lfloor p^{1-\beta} \rfloor$ blocks, each of
 size at least $\lfloor p^\beta \rfloor$. Let the elements within each block have correlation 1, and let the
 elements from different blocks be independent. If $r \geq 4(1 - \beta)$, then the procedure

$$\widehat{S} = \{i : x(i) > \sqrt{2(1 - \beta) \log p}\}$$

63 yields exact support recovery, i.e., $\mathbb{P}[\widehat{S} = S] \rightarrow 1$, as $p \rightarrow \infty$. This requirement on
 64 the signal size is strictly weaker than that of the strong classification boundary, since
 65 $4(1 - \beta) < (1 + \sqrt{1 - \beta})^2$ on $\beta \in (0, 1)$.

66 **Proof** (Example 6.1) Let $t_p^* = \sqrt{2(1 - \beta) \log p}$ and observe that $\widehat{S} = \{j : x(j) >
 67 t_p^*\}$. Analogous to (4.11) in the proof of Theorem 4.1, we have

$$\begin{aligned} 68 \quad \mathbb{P}[\widehat{S} \subseteq S] &= 1 - \mathbb{P}\left[\max_{j \in S^c} x(j) > t_p^*\right] = 1 - \mathbb{P}\left[\max_{j \in S^c} \epsilon(j) > t_p^*\right] \\ 69 \quad &\geq 1 - \mathbb{P}\left[\max_{j \in \{1, \dots, p\}} \epsilon(j) > t_p^*\right] \geq 1 - \mathbb{P}\left[\max_{j \in \{1, \dots, \lfloor p^{1-\beta} \rfloor\}} \widetilde{\epsilon}(j) > t_p^*\right] \end{aligned}$$

71 where $(\widetilde{\epsilon})_{j=1}^{\lfloor p^{1-\beta} \rfloor}$'s are independent Gaussian errors; in the last inequality we used the
 72 assumption that there are at most $\lfloor p^{1-\beta} \rfloor$ independently distributed Gaussian errors
 73 in $(\epsilon_p(j))_{j=1}^p$. By Example 4.1 (with $\lfloor p^{1-\beta} \rfloor$ taking the role of p), we know that the
 74 FWER goes to 0 at a rate of $(2 \log \lfloor p^{1-\beta} \rfloor)^{-1/2}$. Therefore, the probability of no false
 75 inclusion converges to 1.

76 On the other hand, since the signal sizes are no smaller than $(v\underline{r} \log p)^{1/v} =$
 77 $\sqrt{2\underline{r} \log p}$ (for $v = 2$), similar to (4.13), we obtain

$$\begin{aligned} 78 \quad \mathbb{P}[\widehat{S} \supseteq S] &\geq \mathbb{P}\left[\min_{j \in S} \epsilon(j) > \sqrt{2(1-\beta) \log p} - \sqrt{2\underline{r} \log p}\right] \\ 79 \quad &= \mathbb{P}\left[\max_{j \in S} (-\epsilon(j)) < \sqrt{2 \log p} (\sqrt{\underline{r}} - \sqrt{1-\beta})\right] \\ 80 \quad &= \mathbb{P}\left[\frac{\max_{j \in S} (-\epsilon(j))}{u_{|S|}} < \frac{\sqrt{\underline{r}} - \sqrt{1-\beta}}{\sqrt{1-\beta}} (1 + o(1))\right], \end{aligned} \quad (6.3)$$

82 where in the last line we used the quantiles (2.33). Since the minimum signal size
 83 is bounded below by $\underline{r} > 4(1-\beta)$, the right-hand side of the inequality in (6.3)
 84 converges to a constant strictly larger than 1. While the left-hand side, by Slepian's
 85 lemma (recall Theorem 2.1 and Relation 2.47), is stochastically smaller than a r.v.
 86 going to 1. Namely, we have

$$87 \quad \frac{1}{u_{|S|}} \max_{j \in S} (-\epsilon(j)) \stackrel{d}{\leq} \frac{1}{u_{|S|}} \max_{j \in S} \epsilon^*(j) \xrightarrow{\mathbb{P}} 1, \quad (6.4)$$

88 where $(\epsilon^*)_{j=1}^{\lfloor p^{1-\beta} \rfloor}$'s are independent Gaussian errors. Therefore the probability in (6.3)
 89 must also converge to 1. \square

90 Before proceeding to the proof of Theorem 6.1, we will briefly discuss the relationships
 91 between UDD and other dependence conditions in the context of extreme
 92 value theory. The main idea we would like to convey is that UDD (and equivalently
 93 URS) is an exceptionally mild condition on the dependence of the array.

94 **The Berman and UDD conditions.** Suppose that the array of errors \mathcal{E} comes
 95 from a stationary Gaussian time series $\epsilon(i)$, $i \in \mathbb{N}$, with autocovariance $r_p =$
 96 $\text{Cov}(\epsilon(i+p), \epsilon(i))$. One is interested in the asymptotic behavior of the maxima
 97 $M_p := \max_{i=1, \dots, p} \epsilon(i)$.

98 In this setting, Berman's condition, introduced in Berman (1964), requires that

$$99 \quad r_p \log p \rightarrow 0, \quad \text{as } p \rightarrow \infty. \quad (6.5)$$

100 This condition entails that

$$101 \quad a_p(M_p - b_p) \xrightarrow{d} Z, \quad \text{as } p \rightarrow \infty, \quad (6.6)$$

with the Gumbel limit distribution $\mathbb{P}[Z \leq x] = \exp\{-e^{-x}\}$, $x \in \mathbb{R}$, where

$$a_p = \sqrt{2 \log p}, \quad b_p = \sqrt{2 \log p} - \frac{1}{2} \left(\sqrt{2 \log p} \right)^{-1} (\log \log(p) + \log(4\pi)),$$



are *the same* centering and normalization sequences as in the case of iid $\epsilon(i)$'s. Berman's condition is one of the weakest dependence conditions in the literature for which the convergence in (6.6) holds. See, e.g., Theorem 4.4.8 in Embrechts et al. (2013), where (6.5) is described as "very weak".

Instances where the dependence in the time series is so strong that Berman's condition (6.5) fails have also been studied. In such cases, one may continue to have (6.6) but typically the sequences of normalizing and centering constants will be *different* from the iid case, and the corresponding limit is usually no longer Gumbel; see, for example, Theorems 6.5.1 and 6.6.4 in Leadbetter et al. (1983) and McCormick and Mittal (1976).

In our high-dimensional support estimation context, the notion of relative stability is sufficient and more natural than the finer notions of distributional convergence. If one is merely interested in the asymptotic relative stability of the Gaussian maxima, then Berman's condition can be relaxed significantly (see also, Theorem 4.1 of Berman 1964). Observe that by Proposition 6.1, the Berman condition (6.5) implies UDD and hence relative stability (Theorem 6.1), i.e.,

$$\frac{1}{b_p} M_p \xrightarrow{\mathbb{P}} 1, \quad \text{as } p \rightarrow \infty. \quad (6.7)$$

This *concentration of maxima* property can be readily deduced from (6.6), since $a_p b_p \sim 2 \log(p) \rightarrow \infty$ as $p \rightarrow \infty$. Theorem 6.1 shows that (6.7) holds if the much weaker uniform dependence condition UDD holds. Note that our condition is coordinate free—neither monotonicity of the sequence r_p nor stationarity of the underlying array is required. This makes it substantially broader than the time series setting in the seminal work Berman (1964).

The rest of this chapter is devoted to the proof of the main result, i.e., Theorem 6.1. We first introduce a key lemma regarding the structure of an *arbitrary* correlation matrix of high-dimensional random variables. The proof uses a surprising, yet elegant application of Ramsey's Theorem from the study of combinatorics. The "only if" part of Theorem 6.1 follows from this lemma, in Sect. 6.2.

The proof of the "if" part is detailed in Sect. 6.3. The arguments there have been recently extended to establish bounds on the rate of concentration of maxima in Kartsioskas et al. (2019); see also, Tanguy (2015b) and the related notion of super-concentration of Chatterjee (2014).

6.1 Ramsey's Theory and the Structure of Correlation Matrices

Given any integer $k \geq 1$, there is always an integer $R(k, k)$ called the *Ramsey number*:

$$k \leq R(k, k) \leq \binom{2k-2}{k-1} \quad (6.8)$$



such that the following property holds: every undirected graph with at least $R(k, k)$ vertices will contain *either* a clique of size k , or an *independent set* of k nodes. Recall that a clique is a complete sub-graph where all pairs of nodes are connected, and an independent set is a set of nodes where no two nodes are connected.

This result is a consequence of the celebrated work of Ramsey (2009), which gave birth to Ramsey Theory (see e.g., Conlon et al. 2015). The Ramsey Theorem and the upper bound (6.8) (established first in Erdős and Szekeres 1935) are at the heart of the proof of the following result. A recent improvement on the upper bound is given by Sah (2020).

Proposition 6.2 Fix $\gamma \in (0, 1)$ and let $P = (\rho(i, j))_{n \times n}$ be an arbitrary correlation matrix. If

$$k := \lfloor \log_2(n)/2 \rfloor \geq \lceil 1/\gamma \rceil + 1, \quad (6.9)$$

then there is a set of k indices $K = \{l_1, \dots, l_k\} \subseteq \{1, \dots, n\}$ such that

$$\rho(i, j) \geq -\gamma, \text{ for all } i, j \in K. \quad (6.10)$$

Proof By using (6.8) and a refinement of Stirling's formula, we will show at the end of the proof that for $k \leq \log_2(n)/2$, we have

$$R(k, k) \leq n, \quad (6.11)$$

where $R(k, k)$ is the Ramsey number.

Now, construct a graph with vertices $\{1, \dots, n\}$ such that there is an edge between nodes i and j if and only if $\rho(i, j) \geq -\gamma$. In view of (6.11) and Ramsey's theorem (see e.g., Theorem 1 in Fox (2009) or Conlon et al. (2015) for a recent survey on Ramsey theory), there is a subset of k nodes $K = \{l_1, \dots, l_k\}$, which is either a *complete graph* or an *independent set*. Recall that in a complete graph, every two nodes are connected with an edge; while in independent sets, no two nodes are connected.

If K is a complete graph, then by our construction of the graph, Relation (6.10) holds.

Now, suppose that K is a set of independent nodes. This means, again by the construction of our graph, that

$$\rho(i, j) < -\gamma, \text{ for all } i \neq j \in K.$$

Let Z_i , $i \in K$ be zero-mean random variables such that $\rho(i, j) = \mathbb{E}[Z_i Z_j]$. Observe that

$$\text{Var}\left(\sum_{i \in K} Z_i\right) = \sum_{i \in K} \text{Var}(Z_i) + \sum_{\substack{i \neq j \\ i, j \in K}} \text{Cov}(Z_i, Z_j) < k - k(k-1)\gamma, \quad (6.12)$$



168 since $\text{Var}(Z_i) = 1$ and $\rho(i, j) < -\gamma$ for $i \neq j$. By our assumption, $k \geq (\lceil 1/\gamma \rceil + 1)$,
 169 or equivalently, $(k - 1) \geq 1/\gamma$, the variance in (6.12) is negative. This is a contra-
 170 diction showing that there are no independent sets K with cardinality k .

To complete the proof, it remains to show that Relation (6.11) holds. In view of the upper bound on the Ramsey numbers (6.8), it is enough to show that $k \leq \log_2(\sqrt{n})$ implies

$$\binom{2k-2}{k-1} \leq n.$$

This follows from a refinement of the Stirling formula, due to Robbins (1955):

$$\sqrt{2\pi}m^{m+1/2}e^{-m}e^{\frac{1}{(12m+1)}} \leq m! \leq \sqrt{2\pi}m^{m+1/2}e^{-m}e^{\frac{1}{12m}}.$$

171 Indeed, letting $\tilde{k} := k - 1$, and applying the above upper and lower bounds to the
 172 terms $(2\tilde{k})!$ and $\tilde{k}!$, respectively, we obtain

$$\binom{2k-2}{k-1} \equiv \frac{(2\tilde{k})!}{(\tilde{k}!)^2} \leq \frac{2^{2\tilde{k}}}{\sqrt{\pi\tilde{k}}} \exp\left\{\frac{1}{24\tilde{k}} - \frac{2}{12\tilde{k}+1}\right\} < 2^{2k}$$

173 where the last two inequalities follow by simply dropping positive factors less than
 174 1. Since $2k \leq \log_2(n)$, the above bound implies Relation (6.11) and the proof is
 175 complete. \square

176 Using Proposition 6.2, we establish the key lemma used in the proof of
 177 Theorem 6.1.

178 **Lemma 6.1** Let $c \in (0, 1)$, and $P = (\rho(i, j))_{(n+1) \times (n+1)}$ be a correlation matrix
 179 such that

$$\rho(1, j) > c \quad \text{for all } j = 1, \dots, n+1. \quad (6.13)$$

180 If $n \geq 2^{\lceil 2/c^2 \rceil + 4}$, then there is a set of indices $K = \{l_1, \dots, l_k\} \subseteq \{2, \dots, n+1\}$ of
 181 cardinality $k = |K| = \lfloor \log_2 \sqrt{n} \rfloor$, such that

$$\rho(i, j) > \frac{c^2}{2} \quad \text{for all } i, j \in K. \quad (6.14)$$

182 That is, all entries of the $k \times k$ sub-correlation matrix $P_K := (\rho(i, j))_{i,j \in K}$ are larger
 183 than $c^2/2$.

184 **Proof (Lemma 6.1)** Let Z_1, \dots, Z_{n+1} be random variables with covariance matrix
 185 P . Denote $\rho_j = \rho(1, j)$ and define

$$R_j = \begin{cases} \frac{1}{\sqrt{1-\rho_j^2}} (Z_j - \rho_j Z_1), & \text{if } \rho_j < 1, \\ R^* & \text{if } \rho_j = 1, \end{cases} \quad (6.15)$$



where R^* is an arbitrary zero-mean, unit-variance random variable. It is easy to see that $\text{Var}(R_j) = 1$, and

$$\begin{aligned} \text{Cov}(Z_i, Z_j) &= \text{Cov}\left(\rho_i Z_1 + \sqrt{1 - \rho_i^2} R_i, \rho_j Z_1 + \sqrt{1 - \rho_j^2} R_j\right) \\ &= \rho_i \rho_j + \sqrt{1 - \rho_i^2} \sqrt{1 - \rho_j^2} \text{Cov}(R_i, R_j) \\ &> c^2 + \min\{\text{Cov}(R_i, R_j), 0\}. \end{aligned}$$

Therefore, Relation (6.14) would hold if we can find a set of indices $K = \{l_1, \dots, l_k\}$ such that $\text{Cov}(R_i, R_j) \geq -c^2/2$ for all $i, j \in K$, where $k = |K| = \lfloor \log_2 \sqrt{n} \rfloor$. This, however, follows from Proposition 6.2 applied to $(R_j)_{j=2}^{n+1}$ with $\gamma = c^2/2$, provided that

$$k = \lfloor \log_2 \sqrt{n} \rfloor \geq \lceil 2/c^2 \rceil + 1.$$

The last inequality indeed follows from the assumption that $n \geq 2^{2\lceil 2/c^2 \rceil + 4}$. \square

6.2 URS Implies UDD (Proof of the “Only If” Part of Theorem 6.1)

In view of Remark 6.1, UDD is equivalent to the requirement that $N(\delta) := 1 + \sup_p N_p(\delta) < \infty$ for all $\delta \in (0, 1)$, where

$$N_p(\delta) := \max_{j \in \{1, \dots, p\}} |\{i : i \neq j, \Sigma_p(j, i) > \delta\}|. \quad (6.16)$$

Therefore, if \mathcal{E} is not UDD, then there must exist a constant $c \in (0, 1)$ for which $N(c)$ is infinite, i.e., there is a subsequence $\tilde{p} \rightarrow \infty$ such that $N_{\tilde{p}}(c) \rightarrow \infty$. Without loss of generality, we may assume that $\tilde{p} = p$.

Let $j_p(c)$ be the maximizers of (6.16), and let

$$S_p(c) := \{i \in \{1, \dots, p\} : \Sigma_p(j_p(c), i) > c\}. \quad (6.17)$$

Observe that $|S_p(c)| = N_p(c) + 1 \rightarrow \infty$, as $p \rightarrow \infty$ (note $j_p(c) \in S_p(c)$).

Applying Lemma 6.1 to the set of random variables indexed by $S_p(c)$, we conclude, for $N_p(c) \geq 2^{2\lceil 2/c^2 \rceil + 4}$, there must be a further subset

$$K_p(c) \subseteq S_p(c), \quad (6.18)$$



212 of cardinality

$$213 \quad k_p(c) := |K_p(c)| \geq \log_2 \sqrt{N_p(c)}, \quad (6.19)$$

214 such that all pairwise correlations of the random variables indexed by $K_p(c)$ are
 215 greater than $c^2/2$. Since the sequence $N_p(c) \rightarrow \infty$, by (6.19), we have $k_p(c) \rightarrow \infty$
 216 as $p \rightarrow \infty$.

217 Therefore, we have identified a sequence of subsets $K_p(c) \subseteq \{1, \dots, p\}$ with the
 218 following two properties:

- 219 1. $k_p(c) := |K_p(c)| \rightarrow \infty$, as $p \rightarrow \infty$, and
 220 2. For all $i, j \in K_p(c)$, we have

$$221 \quad \Sigma_p(i, j) > c^2/2. \quad (6.20)$$

222 Without loss of generality, we may assume $K_p(c) = \{1, \dots, k_p(c)\} \subseteq \{1, \dots, p\}$,
 223 upon re-labeling of the coordinates.

Now consider a Gaussian sequence $\epsilon^* = \{\epsilon^*(j), j = 1, 2, \dots\}$, independent of \mathcal{E} , defined as follows:

$$\epsilon^*(j) := Z \left(c/\sqrt{2} \right) + Z(j) \sqrt{1 - c^2/2}, \quad j = 1, 2, \dots,$$

224 where Z and $Z(j)$, $j = 1, 2, \dots$ are independent standard normal random variables.
 225 Hence,

$$226 \quad \text{Var}(\epsilon^*(j)) = 1 = \text{Var}(\epsilon_p(j)), \quad (6.21)$$

227 and

$$228 \quad \text{Cov}(\epsilon^*(i), \epsilon^*(j)) = \frac{c^2}{2} \leq \text{Cov}(\epsilon_p(i), \epsilon_p(j)), \quad (6.22)$$

229 for all p , and all $i \neq j$, $i, j \in K_p(c)$. Thus, we have, as $p \rightarrow \infty$,

$$230 \quad \frac{1}{u_{k_p(c)}} \max_{j \in K_p(c)} \epsilon^*(j) = \frac{c/\sqrt{2}}{u_{k_p(c)}} Z + \frac{\sqrt{1 - c^2/2}}{u_{k_p(c)}} \max_{j \in K_p(c)} Z(j) \xrightarrow{\mathbb{P}} \sqrt{1 - \frac{c^2}{2}}, \quad (6.23)$$

231 where the convergence in probability follows from Proposition 2.2 part 2.

232 Relations (6.21) and (6.22), by Slepian’s Lemma (recall Theorem 2.1), also imply,

$$233 \quad \frac{1}{u_{k_p(c)}} \max_{j \in K_p(c)} \epsilon^*(j) \stackrel{d}{\geq} \frac{1}{u_{k_p(c)}} \max_{j \in K_p(c)} \epsilon_p(j). \quad (6.24)$$

Therefore, by (6.24) and (6.23), for all $\sqrt{1 - c^2/2} \leq \delta < 1$, we have

$$232 \quad \mathbb{P} \left[\frac{1}{u_{k_p(c)}} \max_{j \in K_p(c)} \epsilon_p(j) < \delta \right] \rightarrow 1 \quad \text{as } p \rightarrow \infty.$$



234 This contradicts the definition of URS (with the particular choice of $S_p := K_p(c)$),
 235 and the proof of the “only if” part of Theorem 6.1 is complete.

236 **6.3 UDD Implies URS (Proof of the ‘If’ Part
 237 of Theorem 6.1)**

238 Recall that our objective is to show (4.21). We will do so in two stages; namely, we
 239 will prove that for all $\delta > 0$, we have

240
$$\mathbb{P} \left[\frac{M_{S_p}}{u_{|S_p|}} > 1 + \delta \right] \rightarrow 0, \quad (6.25)$$

241 and

242
$$\mathbb{P} \left[\frac{M_{S_p}}{u_{|S_p|}} < 1 - \delta \right] \rightarrow 0, \quad (6.26)$$

243 for any sequence of subsets S_p such that $|S_p| \rightarrow \infty$. Although the first step (6.25) was
 244 already shown in Proposition 2.2, regardless of the dependence structure, we provide
 245 in this section a more refined result. Specifically, the following result states that for
 246 the AGG model, the constant δ in Proposition 2.2 can be replaced by a vanishing
 247 sequence $c_p \rightarrow 0$.

248 **Lemma 6.2** (Upper tails of AGG maxima) *Let \mathcal{E} be an array with marginal distri-
 249 bution $F \in \text{AGG}(\nu)$, $\nu > 0$. If we pick*

250
$$c_p = \frac{u_p \log p}{u_p} - 1, \quad (6.27)$$

251 where $u_p = F^\leftarrow(1 - 1/p)$, then we have $c_p > 0$, $c_p \rightarrow 0$, and

252
$$\mathbb{P} \left[\frac{M_p}{u_p} - (1 + c_p) > 0 \right] \rightarrow 0. \quad (6.28)$$

253 The proof can be found in Sect. 6.3.1 below.

254 Since Lemma 6.2 holds regardless of the dependence structure, the same conclu-
 255 sions hold if one replaces M_p by $M_{S_p} = \max_{j \in S_p} \epsilon(j)$ and p by $q = q(p) = |S_p|$,
 256 where S_p is any sequence of sets such that $q \equiv |S_p| \rightarrow \infty$. This entails (6.25).

257 On the other hand, the proof of (6.26) uses a more elaborate argument based on
 258 the Sudakov–Fernique bound. We proceed by first bounding the probability by an
 259 expectation. For all $\delta > 0$, we have

$$\begin{aligned}
 260 \quad & \mathbb{P} \left[\frac{M_{S_p}}{u_q} < 1 - \delta \right] = \mathbb{P} \left[- \left(\frac{M_{S_p}}{u_q} - (1 + c_q) \right) > \delta + c_q \right] \\
 261 \quad & \leq \mathbb{P} \left[\left(\frac{M_{S_p}}{u_q} - (1 + c_q) \right)_- > \delta + c_q \right] \\
 262 \quad & \leq \frac{1}{\delta + c_q} \mathbb{E} \left[\left(\frac{M_{S_p}}{u_q} - (1 + c_q) \right)_- \right], \tag{6.29} \\
 263
 \end{aligned}$$

264 where $(x)_- := \max\{-x, 0\}$ and the last line follows from the Markov inequality.
 265 The next result shows that the upper bound in (6.29) vanishes.

266 **Lemma 6.3** *Let \mathcal{E} be a Gaussian UDD array and $S_p \subseteq \{1, \dots, p\}$ be an arbitrary
 267 sequence of sets such that $q = q(p) = |S_p| \rightarrow \infty$. Then, for $M_{S_p} := \max_{j \in S_p} \epsilon_p(j)$
 268 and c_q as in (6.27), we have*

$$269 \quad \mathbb{E} \left[\left(\frac{M_{S_p}}{u_q} - (1 + c_q) \right)_- \right] \rightarrow 0, \quad \text{as } p \rightarrow \infty. \tag{6.30}$$

270 The proof of the lemma is given in Sect. 6.3.2 below.

271 Going back to the proof of Theorem 6.1, we observe that Relations (6.29) and
 272 (6.30) imply (6.26), which completes the proof of the ‘if’ part of Theorem 6.1. \square

273 **Remark 6.3** Only the Sudakov–Fernique minorization argument used in the proof
 274 of Lemma 6.3, relies on the Gaussian assumption. We expect the techniques and
 275 results here to be useful in extending Theorem 6.1 to more general class of distribu-
 276 tions, say, the AGG model.

277 6.3.1 Bounding the Upper Tails of AGG Maxima

278 **Proof** (Lemma 6.2) Recall by (2.33) that

$$279 \quad u_q \sim (\nu \log q)^{1/\nu}, \quad q \rightarrow \infty,$$

280 so that

$$281 \quad c_p = \frac{u_{p \log p}}{u_p} - 1 = \left(\frac{\log p + \log \log p}{\log p} \right)^{1/\nu} (1 + o(1)) - 1 \rightarrow 0 \quad \text{as } p \rightarrow \infty. \tag{6.31}$$

282 By the union bound, we have

$$\begin{aligned}
 283 \quad & \mathbb{P} \left[\frac{M_p}{u_p} > 1 + c_p \right] \leq \sum_{j=1}^p \mathbb{P} \left[\frac{\epsilon_p(j)}{u_p} > 1 + c_p \right] = p \bar{F}(u_p \log p) \quad (6.32) \\
 284 \quad & = p \bar{F} \left(F^{-1} \left(1 - \frac{1}{p \log p} \right) \right) \leq \frac{1}{\log p} \rightarrow 0. \\
 285
 \end{aligned}$$

286 where the last inequality follows from the fact that $F(F^{-1}(u)) \geq u$ for all
 287 $u \in [0, 1]$. \square

288 In addition to Lemma 6.2, which says the upper tail vanishes in probability, we
 289 will also prepare a result which states that the upper tail also vanishes in expectation.

Lemma 6.4 *Let M_p and c_p be as in Lemma 6.2, and denote*

$$\xi_p := \frac{M_p}{(1 + c_p)u_p}.$$

290 Then there exist $p_0, t_0 > 0$, and an absolute constant $C > 0$ such that

$$291 \quad \mathbb{P} [\xi_p > t] \leq \exp \{-Ct^\nu\}, \quad \text{for all } p > p_0, t > t_0. \quad (6.33)$$

292 In particular, the set of random variables $\{(\xi_p)_+, p \in \mathbb{N}\}$ is uniformly integrable.

293 **Proof (Lemma 6.4)** Recalling that $(1 + c_p)u_p = u_p \log p$, and by applying the union
 294 bound as in (6.32), we have

$$\begin{aligned}
 295 \quad & \log \mathbb{P} [\xi_p > t] \leq \log p + \log \bar{F}(u_p \log p t) \\
 296 \quad & \leq \log p - \frac{1}{\nu} (u_p \log p t)^\nu (1 - \delta). \quad (6.34)
 \end{aligned}$$

297 for $t > t_0(\delta) > 0$, where $\delta \in (0, 1)$ is an arbitrarily small number fixed in advance.
 298 This follows from the assumption that $F \in \text{AGG}(\nu)$ and Definition 2.6 of the AGG
 299 distribution. Using in (6.34), the explicit expressions for the quantiles in (2.33), we
 300 obtain
 301

$$\begin{aligned}
 302 \quad & \log \mathbb{P} [\xi_p > t] \leq \log p - \underbrace{(1 + o(1))(1 - \delta)t^\nu}_{\text{greater than 1 for large } t} \log p - t^\nu \underbrace{\log \log p (1 + o(1))(1 - \delta)}_{\text{greater than } C \text{ for large } p}. \quad (6.35)
 \end{aligned}$$

303 For large t , we have $(1 + o(1))(1 - \delta)t^\nu > 1$ so that the sum of the first two terms
 304 on the right-hand side of (6.35) is negative. Also, for p larger than some constant
 305 $p_0(\delta)$, we have $\log \log p (1 + o(1))(1 - \delta) > C$ for some constant C that does not
 306 depend on p . Therefore (6.33) holds for $t > t_0(\delta)$ and $p > p_0(\delta)$, and the proof is
 307 complete. \square



308 **Corollary 6.3** *The upper tails of AGG maxima vanish in expectation, i.e.,*

309
$$\mathbb{E} \left[\left(\frac{M_p}{u_p} - (1 + c_p) \right)_+ \right] \rightarrow 0 \text{ as } p \rightarrow \infty, \quad (6.36)$$

310 where $(a)_+ := \max\{a, 0\}$.

311 **Proof (Corollary 6.3)** Since $c_p \geq 0$ is a sequence converging to 0, we have $c_p < 1$
312 for $p \geq p_0$. Hence, for any $t > 0$, we have

313
$$\begin{aligned} \mathbb{P} \left[\left(\frac{M_p}{u_p} - (1 + c_p) \right)_+ > t \right] &= \mathbb{P} \left[(1 + c_p) (\xi_p - 1)_+ > t \right] \\ \text{314} \quad &\leq \mathbb{P} \left[(\xi_p - 1)_+ > t/2 \right] \leq \mathbb{P} [\xi_p > t/2]. \end{aligned} \quad (6.37)$$

315

316 By Lemma 6.4, $\{(\xi_p)_+\}$ is u.i., therefore by Relation (6.37), $\{(M_p/u_p - (1 + c_p))_+,$
317 $p \in \mathbb{N}\}$ is u.i. as well. Since by Lemma 6.2, $(M_p/u_p - (1 + c_p))_+ \rightarrow 0$ in
318 probability, Relation (6.36) follows from the established uniform integrability (see,
319 e.g., Theorem 6.6.1 in Resnick 2014). \square

320 6.3.2 Bounding the Lower Tails of Gaussian Maxima

321 The main goal of this section is to establish the following result.

322 **Proposition 6.3** *For every UDD Gaussian array \mathcal{E} , and any sequence of subsets
323 $S_p \subseteq \{1, \dots, p\}$ such that $q = q(p) = |S_p| \rightarrow \infty$, we have*

324
$$\liminf_{p \rightarrow \infty} \mathbb{E} \left[\frac{M_{S_p}}{u_q} \right] \geq 1, \quad (6.38)$$

325 where $M_S = \max_{j \in S} \epsilon(j)$.

326 We will first show that Lemma 6.3, which is the key to the proof of the ‘if’ part
327 of Theorem 6.1, follows immediately from this proposition.

Proof (Lemma 6.3) We start with the identity

$$\mathbb{E} \left[\frac{M_{S_p}}{u_q} - (1 + c_q) \right] = \mathbb{E} \left[\left(\frac{M_{S_p}}{u_q} - (1 + c_q) \right)_+ \right] - \mathbb{E} \left[\left(\frac{M_{S_p}}{u_q} - (1 + c_q) \right)_- \right].$$

328 By re-arranging terms and taking limsup/liminf, we obtain

$$329 \quad 0 \leq \limsup_{p \rightarrow \infty} \mathbb{E} \left[\left(\frac{M_{S_p}}{u_q} - (1 + c_q) \right)_- \right]$$

$$330 \quad \leq \limsup_{p \rightarrow \infty} \mathbb{E} \left[\left(\frac{M_{S_p}}{u_q} - (1 + c_q) \right)_+ \right] - \liminf_{p \rightarrow \infty} \mathbb{E} \left[\frac{M_{S_p}}{u_q} - (1 + c_q) \right] \quad (6.39)$$

$$331 \quad = - \liminf_{p \rightarrow \infty} \mathbb{E} \left[\frac{M_{S_p}}{u_q} - (1 + c_q) \right], \quad (6.40)$$

332 where the last equality follows from the fact that the lim-sup in (6.39) vanishes by Corollary 6.3. On the other hand, since $c_q \rightarrow 0$, we have

$$\liminf_{p \rightarrow \infty} \mathbb{E} \left[\frac{M_{S_p}}{u_q} - (1 + c_q) \right] = \liminf_{p \rightarrow \infty} \mathbb{E} \left[\frac{M_{S_p}}{u_q} - 1 \right] \geq 0,$$

333 where the last inequality follows from Proposition 6.3. This shows that the right-hand
334 side of (6.40) is non-positive and hence (6.30) holds. \square

335 The following interesting fact about the relationship between the upper quantiles
336 and the expectation of iid maxima will be needed for the proof of Proposition 6.3.

Lemma 6.5 *Let $(X_i)_{i=1}^p$ be p iid random variables with distribution F such that $\mathbb{E}[(X_i)_-]$ exists, i.e.,*

$$\mathbb{E}[\max\{-X_i, 0\}] < \infty.$$

Let $M_p = \max_{i=1, \dots, p} X_i$. Assume that F has a density f , which is eventually decreasing. More precisely, we suppose there exists a C_0 such that $0 < F(C_0) < 1$, and $f(x_1) \geq f(x_2)$ whenever $C_0 < x_1 \leq x_2$. Under these assumptions, we have,

$$\liminf_{p \rightarrow \infty} \frac{\mathbb{E} M_p}{u_{p+1}} \geq 1,$$

337 where $u_{p+1} = F^\leftarrow(1 - 1/(p+1))$.

Proof The idea comes from an argument in the monograph of Boucheron et al. (2013). Write

$$X_i = F^\leftarrow(U_i)$$

338 where U_i are iid uniform random variables on $(0, 1)$. Denote M_p^U as the maximum
339 of the U_i 's, we have $\mathbb{E} M_p = \mathbb{E}[F^\leftarrow(M_p^U)]$, and by conditioning, we obtain

$$340 \quad \mathbb{E} M_p = \mathbb{E} [F^\leftarrow(M_p^U) \mid M_p^U \geq F(C_0)] \mathbb{P}[M_p^U \geq F(C_0)] + \\ 341 \quad + \mathbb{E} [F^\leftarrow(M_p^U) \mid M_p^U < F(C_0)] \mathbb{P}[M_p^U < F(C_0)]. \quad (6.41)$$



Focus on the first term in the summation. Since f is decreasing beyond C_0 , F is concave on (C_0, ∞) , and F^\leftarrow is convex on $(F(C_0), 1)$. By Jensen’s inequality, we have

$$\mathbb{E}[F^\leftarrow(M_p^U) \mid M_p^U \geq F(C_0)] \geq F^\leftarrow(\mathbb{E}[M_p^U \mid M_p^U \geq F(C_0)]).$$

Using the fact that M_p^U is the maximum of iid Uniform(0, 1) random variables, with a direct calculation one can show that

$$F^\leftarrow(\mathbb{E}[M_p^U \mid M_p^U \geq F(C_0)]) = F^\leftarrow\left(\left(1 - \frac{1}{p+1}\right)\left(\frac{1-F(C_0)^{p+1}}{1-F(C_0)^p}\right)\right),$$

and hence

$$\begin{aligned} \mathbb{E}[F^\leftarrow(M_p^U) \mid M_p^U \geq F(C_0)] &\geq F^\leftarrow\left(\left(1 - \frac{1}{p+1}\right)\left(\frac{1-F(C_0)^{p+1}}{1-F(C_0)^p}\right)\right) \\ &\geq F^\leftarrow\left(1 - \frac{1}{p+1}\right) = u_{p+1}. \end{aligned} \quad (6.42)$$

Now, focus on the second term in (6.41). Since $\mathbb{P}[M_p^U \leq m \mid M_p^U < F(C_0)] = (m/F(C_0))^p \leq m/F(C_0)$ for $m \leq F(C_0)$, we have

$$(M_p^U \mid M_p^U < F(C_0)) \stackrel{\text{d}}{\geq} (U_1 \mid U_1 < F(C_0)),$$

where and the latter is the uniform distribution on $(0, F(C_0))$. Therefore, for the second term of the sum in (6.41), by the monotonicity of F^\leftarrow , we obtain

$$\begin{aligned} \mathbb{E}[F^\leftarrow(M_p^U) \mid M_p^U < F(C_0)] &\geq \mathbb{E}[F^\leftarrow(U_1) \mid U_1 < F(C_0)] \\ &= \mathbb{E}[X_1 \mid X_1 < C_0]. \end{aligned} \quad (6.43)$$

Finally, since $\mathbb{P}[M_p^U < F(C_0)] = F(C_0)^p = 1 - \mathbb{P}[M_p^U \geq F(C_0)]$, by (6.42) and (6.43), we have

$$\frac{\mathbb{E}M_p}{u_{p+1}} \geq (1 - F(C_0)^p) + \frac{\mathbb{E}[X_1 \mid X_1 < C_0]}{u_{p+1}} F(C_0)^p.$$

The conclusion follows since the right-hand side of the last inequality converges to 1. \square

We are now ready to prove Proposition 6.3. This is where the UDD dependence assumption is used.

Proof (Proposition 6.3) Recall that $\mathcal{E} = \{\epsilon_p(i), i = 1, \dots, p, p \in \mathbb{N}\}$ is a Gaussian array with standardized marginals. Define the canonical (pseudo) metric on $S_p \subset \{1, \dots, p\}$,

$$d(i, j) = \sqrt{\mathbb{E}[(\epsilon(i) - \epsilon(j))^2]}.$$

It can be easily checked that the canonical metric takes values between 0 and 2. For an arbitrary $\delta \in (0, 1)$, take $\gamma = \sqrt{2(1 - \delta)}$, and let \mathcal{N} be a γ -packing of S_p . That is, let \mathcal{N} be a subset of S_p , such that for any $i, j \in \mathcal{N}, i \neq j$, we have $d(i, j) \geq \gamma$, i.e.,

$$d(i, j) = \sqrt{2(1 - \Sigma_p(i, j))} \geq \gamma = \sqrt{2(1 - \delta)}, \quad (6.44)$$

or equivalently, $\Sigma_p(i, j) \leq \delta$. We claim that we can find a γ -packing \mathcal{N} whose number of elements is at least

$$|\mathcal{N}| \geq q/N(\delta). \quad (6.45)$$

Indeed, \mathcal{N} can be constructed iteratively as follows:

1: Set $S_p^{(1)} := S_p$ and $\mathcal{N} := \{j_1\}$, where $j_1 \in S_p^{(1)}$ is an arbitrary element. Set $k := 1$.

2: Set $S_p^{(k+1)} := S_p^{(k)} \setminus B_\gamma(j_k)$, where

$$B_\gamma(j_k) := \{i \in S_p : d(i, j_k) < \gamma \equiv \sqrt{2(1 - \delta)}\}.$$

3: If $S_p^{(k)} \neq \emptyset$, pick an arbitrary $j_{k+1} \in S_p^{(k)}$, set $\mathcal{N} := \mathcal{N} \cup \{j_{k+1}\}$, and $k := k + 1$, go to step 2; otherwise, stop.

By the definition of UDD (see Definition 6.1), there are at most $N(\delta)$ coordinates whose covariance with $\epsilon(j)$ exceed δ . Therefore at each iteration, $|B_\gamma(j_k)| \leq N(\delta)$, and hence

$$|S_p^{(k+1)}| \geq |S_p^{(k)}| - |B_\gamma(j_k)| \geq q - kN(\delta).$$

The construction can continue for at least $q/N(\delta)$ iterations, and we have $|\mathcal{N}| \geq \lfloor q/N(\delta) \rfloor$ as desired.

Now we define on this γ -packing \mathcal{N} an independent Gaussian process $(\eta(j))_{j \in \mathcal{N}}$,

$$\eta(j) = \frac{\gamma}{\sqrt{2}} Z(j) \quad j \in \mathcal{N},$$

where $Z(j)$'s are iid standard normal random variables. Observe that by the definition of γ -packing in (6.44), the increments of the new process are smaller than those of the original process in the following sense,

$$\mathbb{E}[(\eta(i) - \eta(j))^2] = \gamma^2 \leq d^2(i, j) = \mathbb{E}[(\epsilon(i) - \epsilon(j))^2]$$

382 for all $i \neq j, i, j \in \mathcal{N}$. Applying the Sudakov–Fernique inequality (see Theorem 2.2)
 383 to $(\eta(j))_{j \in \mathcal{N}}$ and $(\epsilon(j))_{j \in \mathcal{N}}$, we have

$$384 \quad \mathbb{E} \left[\max_{j \in \mathcal{N}} \eta(j) \right] \leq \mathbb{E} \left[\max_{j \in \mathcal{N}} \epsilon(j) \right] \leq \mathbb{E} \left[\max_{j \in S_p} \epsilon(j) \right]. \quad (6.46)$$

385 Since the $(\eta(j))_{j \in \mathcal{N}}$ are independent Gaussians, Lemma 6.5 yields the lower bound,

$$386 \quad \liminf_{p \rightarrow \infty} \mathbb{E} \left[\frac{\max_{j \in \mathcal{N}} \eta(j)}{u_{|\mathcal{N}|}} \right] \geq \frac{\gamma}{\sqrt{2}} = \sqrt{1 - \delta}. \quad (6.47)$$

387 Using (6.45) and the expressions (2.33) for the quantiles of AGG models (with $v = 2$
 388 here), we have

$$389 \quad \frac{u_{|\mathcal{N}|}}{u_q} \geq \left(\frac{\log q - \log N(\delta)}{\log q} \right)^{1/2} (1 + o(1)) \rightarrow 1, \quad (6.48)$$

390 since $N(\delta)$ does not depend on $q = q(p) \rightarrow \infty$.

391 By combining (6.46), (6.47) and (6.48), we conclude that

$$\begin{aligned} 392 \quad \liminf_{p \rightarrow \infty} \mathbb{E} \left[\frac{\max_{j \in S_p} \epsilon(j)}{u_q} \right] &\geq \liminf_{p \rightarrow \infty} \mathbb{E} \left[\frac{\max_{j \in \mathcal{N}} \eta(j)}{u_q} \right] && \text{by (6.46)} \\ 393 \quad &\geq \liminf_{p \rightarrow \infty} \mathbb{E} \left[\frac{\max_{j \in \mathcal{N}} \eta(j)}{u_{|\mathcal{N}|}} \right] && \text{by (6.48)} \\ 394 \quad &\geq \sqrt{1 - \delta}. && \text{by (6.47)} \end{aligned}$$

396 Since $\delta > 0$ is arbitrary, (6.38) follows as desired. □

397 6.4 Numerical Illustrations of Exact Support Recovery 398 Under Dependence

399 The characterization of URS with the UDD condition allows us to simulate Gaussian
 400 errors and illustrate the effect of dependence on the phase transition behavior in finite
 401 dimensions. We shall compare the performance of the Bonferroni’s procedure, which
 402 is agnostic to both sparsity and signal size, with the oracle procedure which picks
 403 the top-s observations.

404 The first set of experiments explores short-range dependent errors from an auto-
 405 regressive (AR) models.

- 406 • AR(1) Gaussian errors with parameter $\rho = -0.5$, $\rho = 0.5$, and $\rho = 0.9$, where
 407 the autocovariance functions decay exponentially, $\rho_k = \rho^k$.



We again apply both the sparsity- and signal-size agnostic Bonferroni's procedure, i.e., $\widehat{S} = \{i : x(i) > \sqrt{2 \log p}\}$, as well as the oracle procedure $\widehat{S}^* = \{i : x(i) \geq x_{[s]}\}$, $s = |S|$, to all settings. Results of the numerical experiments for the AR models are shown in Fig. 6.1.

As was commented in the main text, for dependent errors the oracle procedures is able to recover support of signals with higher probability than the Bonferroni procedures in finite dimensions; compare left and right columns of Fig. 6.1. Short range dependent observations, however, there is not a pronounced difference. The results of the experiments are very similar to that of the independent Gaussian case.

The second set of experiments explores exact support recovery in additive error models in the cases of long-range dependent but UDD, as well as non-UDD errors. In particular we simulate

- Fractional Gaussian noise (fGn) with Hurst parameter $H = 0.75$ and $H = 0.9$. The autocovariance functions are

$$\rho_k \sim 0.75k^{-0.6} \quad \text{and} \quad \rho_k \sim 1.44k^{-0.2},$$

as $k \rightarrow \infty$. Both fGn models represent the regime of long-range dependence, where covariances decay very slowly to zero, so that $\sum |\rho_k| = \infty$; see, e.g., Taqqu (2003). Observe that every stationary Gaussian process with vanishing autocovariance gives rise to an UDD array as concluded in Corollary 6.1.

- The non-UDD Gaussian errors described in Example 6.1.

We will apply both the sparsity-and-signal-size-agnostic Bonferroni's procedure, i.e., $\widetilde{S} = \{i : x(i) > \sqrt{2 \log p}\}$, as well as the oracle procedure $\widehat{S}^* = \{i : x(i) \geq x_{[s]}\}$, $s = |S|$, to all settings. Results of the numerical experiments for the fGn and non-UDD models are shown in Fig. 6.2.

Notice that the oracle procedure sets its thresholds more aggressively (at roughly $\sqrt{2 \log s}$) than the Bonferroni procedure (at $\sqrt{2 \log p}$). Although this difference vanishes as $p \rightarrow \infty$, in finite dimensions ($p = 10\,000$) the advantage can be felt. Indeed, in all our experiments the oracle procedure is able to recover support of signals with higher probability than the Bonferroni procedures; compare left and right columns of Fig. 6.2. Notice also that there is an increase in probability of recovery near $\beta = 0$ for oracle procedures. This is an artifact in finite dimensions due to the fact that $s = \lfloor p^{1-\beta} \rfloor < p/2$, and there are more signals than nulls. The oracle procedures is able to adjust to this reversal by lowering its threshold accordingly.

For UDD errors, Theorem 4.2 predicts that exact recovery of the support is impossible when signal sizes are below the boundary (4.5), even with oracle procedures. However, the rate of this convergence (i.e., $\mathbb{P}[\widehat{S}^* = S] \rightarrow 0$ or 1) can be very slow when the errors are heavily dependent, even though all AR and fGn models demonstrate qualitatively the same behavior in line with the predicted boundary (4.5). In finite dimensions ($p = 10\,000$), as dependence in the errors increases (fGN(H = 0.75) to fGN(H = 0.9)), the oracle procedure becomes more powerful at recovering signal support with high probability for weaker signals.

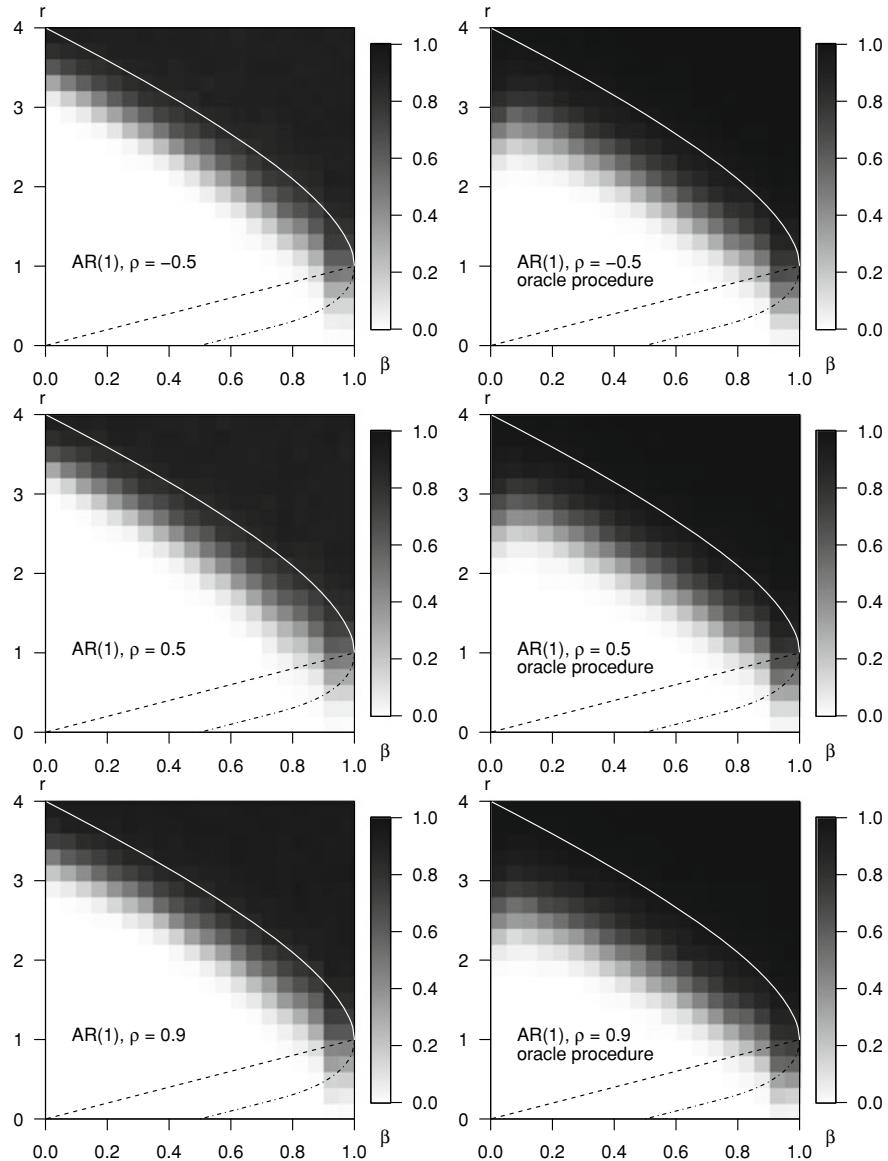


Fig. 6.1 The empirical probability of exact support recovery from numerical experiments, as a function of sparsity level β and signal sizes r . Darker colors indicate higher probability of exact support recovery. Three AR(1) models with autocorrelation functions $(-0.5)^k$ (upper), 0.5^k (middle), and 0.9^k (lower) are simulated. The experiments were repeated 1000 times for each sparsity-signal size combination. In finite dimensions ($p = 10000$), the Bonferroni procedures (left) suffers small loss of power compared to the oracle procedures (right). A phase transition in agreement with the predicted boundary (4.5) can be seen in the AR models. The boundaries (solid, dashed, and dash-dotted lines) are as in Fig. 4.1

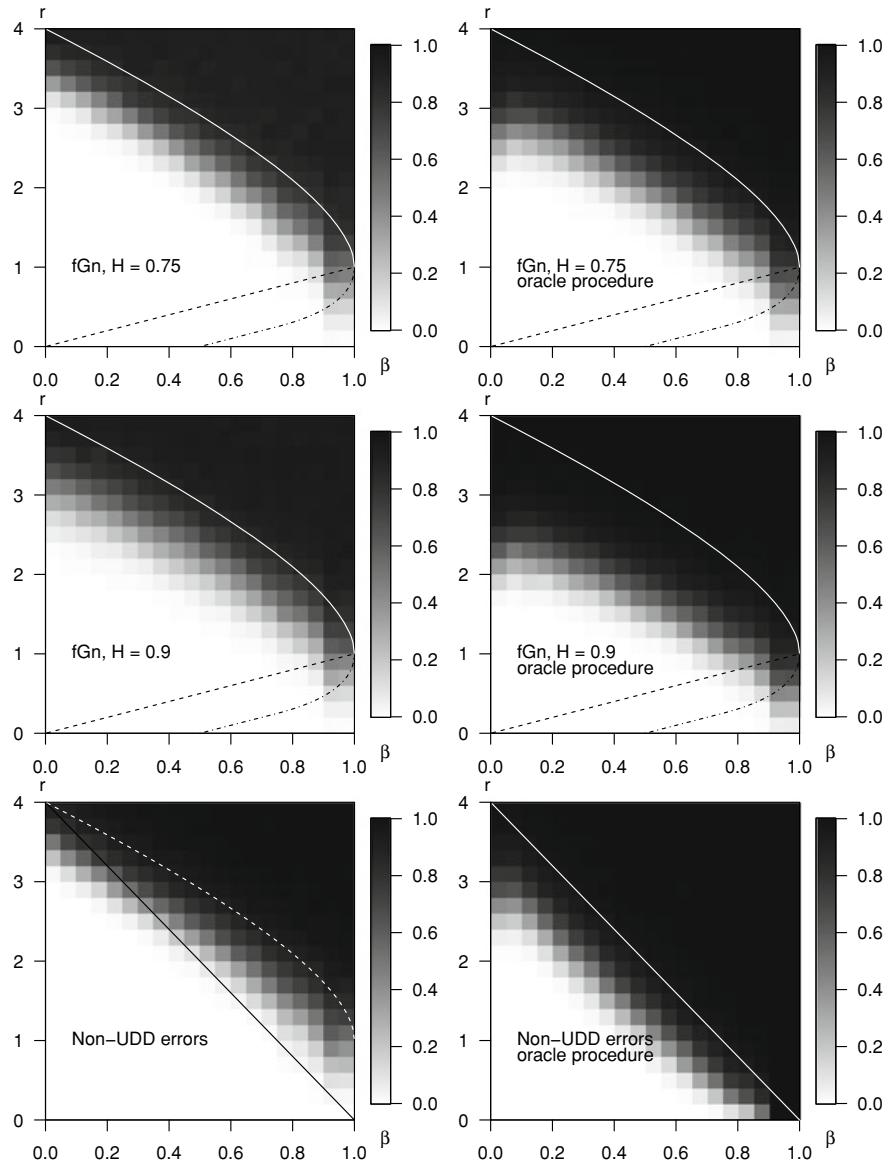


Fig. 6.2 The empirical probability of exact support recovery from numerical experiments, as a function of sparsity level β and signal sizes r . Darker colors indicate higher probability of exact support recovery. Two fGn models with Hurst parameter $H = 0.75$ (upper), $H = 0.9$ (middle), and the non-UDD errors in Example 6.1 (lower) are simulated. The experiments were repeated 1000 times for each sparsity-signal size combination. In finite dimensions ($p = 10000$), the oracle procedures (right) is able to recover support for weaker signals than the Bonferroni procedures (left) when errors are heavily dependent, although they have the same phase transition limit. The non-UDD errors demonstrate qualitatively different behavior, enabling support recovery for strictly weaker signals. The boundaries (solid, dashed, and dash-dotted lines) are as in Fig. 4.1. In the non-UDD example, dashed lines represent the limit attained by Bonferroni's procedures. See text for additional comments



446 On the other hand, as demonstrated in Example 6.1, non-UDD errors yield qual-
447 itatively different behavior; exact support recovery is possible for signal sizes strictly
448 weaker than that in the UDD case. Lower-right panel of Fig. 6.2 demonstrates in this
449 example that the signal support can be recovered as long as the signal sizes are larger
450 than $4(1 - \beta)$.

Metadata of the chapter that will be visualized in SpringerLink

Book Title	Concentration of Maxima and Fundamental Limits in High-Dimensional Testing and Inference	
Series Title		
Chapter Title	Fundamental Statistical Limits in Genome-Wide Association Studies	
Copyright Year	2021	
Copyright HolderName	The Author(s), under exclusive license to Springer Nature Switzerland AG	
Corresponding Author	Family Name Particle Given Name Prefix Suffix Role Division Organization Address Email ORCID	Gao Zheng Department of Statistics University of Michigan-Ann Arbor Ann Arbor, MI, USA gaozheng@umich.edu https://orcid.org/0000-0003-1394-8741
Author	Family Name Particle Given Name Prefix Suffix Role Division Organization Address Email ORCID	Stoev Stilian Department of Statistics University of Michigan-Ann Arbor Ann Arbor, MI, USA https://orcid.org/0000-0003-2108-5421
Abstract	<p>This chapter focuses on an important application of high-dimensional testing and inference to statistical genetics. Genome-wide association studies (GWAS) aim to discover the set of gene variants associated with a disease. In the simple case of single-gene marginal association studies, one typically considers 2-by-2 contingency tables and computes non-central chi-square-type statistics (one for each gene), where the non-centrality parameter plays the role of the signal. In this context, the support of the signal corresponds to the set of genes associated with the disease. It is shown that all phase-transition results for the signal-plus-noise model discussed in Chap. 3 have their counterparts for the chi-square model. This leads to a first of its kind explanation of the observed power dichotomy in GWAS known as the “steep part of the power curve”. To clarify further the relationship between signal size and power, the chapter develops canonical parameterizations in terms of the marginal frequencies, odds ratio, and sample sizes. This leads to a simplified characterization of the power in GWAS and to further insights into optimal study design. These results are accompanied by an online R Shiny web application for power analysis in association studies.</p>	

Chapter 7

Fundamental Statistical Limits in Genome-Wide Association Studies



0 The process of scientific discovery, as explained by Richard Feynman, usually starts
1 with guesses. The consequences of such guesses are then computed and compared
2 with experimental results. If the predictions disagree with the experiment then our
3 guesses are wrong. “That is all there is to it” (Feynman 2017).

4 In the previous chapters, we delved deep into the theoretical underpinnings of
5 the phase-transition phenomena in high-dimensional multiple testing problems. The
6 results are interesting in their own right. However, we have not discovered any scientific
7 law in the spirit of Feynman, but merely worked out mathematical consequences
8 of our postulated models. In this chapter, our goal is to relate these predictions to
9 real experimental data from the field of genetics, where large-scale simultaneous
10 hypotheses testing problems often arise. From such comparisons, we will demonstrate
11 that the phase transition laws are indeed reasonable predictions of some curious
12 phenomena in that field. The accuracy of our predictions will lend credibility to the
13 application of these “laws of large dimensions” in actual applications.

14 In our case, the experimental data used as the measuring stick come from genome-
15 wide association studies (GWAS), introduced in Sect. 1.2. Recall that in GWAS, a
16 large number of marginal association tests are conducted simultaneously, resulting
17 in statistics that can be approximated by

$$18 \quad x(i) \sim \chi_v^2(\lambda(i)), \quad i = 1, \dots, p, \quad (7.1)$$

19 where $\chi_v^2(\lambda(i))$ is a chi-square distributed random variable with $v > 0$ degrees of
20 freedom¹ and non-centrality parameter $\lambda(i)$.

21 We establish our theoretical predictions in two steps. In Sect. 7.1 below, we shall
first establish the phase transitions of the model (7.1). In parallel to results in Chap. 3,

¹ The parameter v here should not be confused with the shape parameter of the AGG(v) distributions, which will not appear in this chapter.

we show that several commonly used family-wise error rate-control procedures—including Bonferroni’s procedure—are asymptotically optimal for the exact, and exact–approximate support recovery problems (recall Definition 2.5) in idealized chi-square models with independent components. Analogously, the BH procedure is asymptotically optimal for the approximate, and approximate–exact support recovery problems. Under appropriate parametrization of the signal sizes and sparsity, we establish the phase transitions of support recovery problems in the chi-square model. Remarkably, the degree-of-freedom parameter does not affect the asymptotic boundaries in any of the four support recovery problems.

In the second step, we translate the canonical signal size and sparsity parametrizations into the vernacular of statistical geneticists in Sect. 7.2. We do so by characterizing the relationship between the signal size λ and the marginal frequencies, odds ratio, and sample sizes for association tests on 2-by-2 contingency tables. This is important because these parameters are often estimated and reported in GWAS, while we have never seen the elusive signal size parameter λ reported. As a bonus, we point out the implications of this relationship on statistically optimal study designs for association studies in Sect. 7.3: perhaps surprisingly, balanced designs with equal number of cases and controls are often statistically inefficient.

Armed with the results on phase transitions in the chi-square model, and a translation from the language of high-dimensional statistics to the patois of association screening studies, we finally present in Sect. 7.4 the consequences of the phase transitions in GWAS, and compare against real experimental data to evaluate the success of our predictions.

The phase transitions in the chi-square models are demonstrated with numerical simulations in Sect. 7.5. The proofs, which are closely resemble those of the results in Chap. 3, are collected in Appendix A.

7.1 Support Recovery Problems in Chi-Squared Models

Similar to the analysis of additive error models in Chap. 3, we will work with triangular arrays of chi-square models (7.1) indexed by p . We adopt the same parametrization for the sparsity of the non-centrality parameter vectors $\lambda = \lambda_p$,

$$|S_p| = \lfloor p^{1-\beta} \rfloor, \quad \beta \in (0, 1] \quad (7.2)$$

where $S_p := \{i : \lambda(i) > 0\}$ is the *signal support* set and β parametrizes the problem *sparsity*. More general parameterizations of the support size are possible as in (4.3). Here, however, we drop the slowly varying term $\ell(\cdot)$ for simplicity. The closer β is to 1, the sparser the support S_p ; conversely, when β is close to 0, the support is dense with many non-null signals.



We parametrize the range of the non-zero and perhaps unequal signals in the chi-square model with

$$\underline{\Delta} = 2\underline{r} \log p \leq \lambda(i) \leq \bar{\Delta} = 2\bar{r} \log p, \quad \text{for all } i \in S_p, \quad (7.3)$$

for some constants $0 < \underline{r} \leq \bar{r} \leq +\infty$.

7.1.1 The Exact Support Recovery Problem

The first main result characterizes the phase transition phenomenon in the exact support recovery problem under the chi-square model. It parallels Theorem 3.2.

Theorem 7.1 Consider the high-dimensional chi-squared model (7.1) with signal sparsity and size as described in (7.2) and (7.3). The function

$$f_E(\beta) = \left(1 + \sqrt{1 - \beta}\right)^2 \quad (7.4)$$

characterizes the phase transition of exact support recovery problem. Namely, the following two results hold.

(i) If $\underline{r} > f_E(\beta)$, then Bonferroni's, Sidák's, Holm's, and Hochberg's procedures with slowly vanishing (see Definition 3.1) nominal FWER levels all achieve asymptotically exact support recovery in the sense of (2.25).

(ii) Conversely, if $\bar{r} < f_E(\beta)$, then for any thresholding procedure \widehat{S}_p , we have $\mathbb{P}[\widehat{S}_p = S_p] \rightarrow 0$. Therefore, in view of Lemma 2.1, exact support recovery asymptotically fails for all thresholding procedures in the sense of (2.26).

The procedures listed in Theorem 7.1 were reviewed in Sect. 2.2. The proof of the theorem can be found in Sect. A.2.

It is evident that the exact support recovery boundary (7.4) coincides with that in parallel results for the Gaussian additive error models (1.1) in Chap. 3. Implications of these results will be discussed in Sect. 7.1.5 below.

Remark 7.1 Theorem 7.1 predicts that the asymptotic boundaries are the same for all values of the degrees of freedom parameter v . In simulations (Sect. 7.5), we find this asymptotic prediction to be quite accurate for $v \leq 3$ even in moderate dimensions ($p = 100$). For $v > 3$, the phase transitions take place somewhat above the boundary g . The behavior is qualitatively similar for the other three phase transitions (see Theorems 7.2, 7.3, and 7.4 below).

87 7.1.2 The Exact–Approximate Support Recovery Problem

88 The next theorem describes the phase transition in the exact–approximate support
 89 recovery problem. Recall also Theorem 3.4.

90 **Theorem 7.2** *In the context of Theorem 7.1, the function*

$$91 \quad f_{EA}(\beta) = 1 \quad (7.5)$$

92 characterizes the phase transition of exact–approximate support recovery problem.
 93 Namely, the following two results hold.

- 94 (i) If $\underline{r} > f_{EA}(\beta)$, then the procedures listed in Theorem 7.1 with slowly vanishing
 95 nominal FWER levels achieve asymptotically exact–approximate support recovery
 96 in the sense of (2.25).
- 97 (ii) Conversely, if $\bar{r} < f_{EA}(\beta)$, then for any thresholding procedure \widehat{S}_p , the exact–
 98 approximate support recovery fails in the sense of (2.26).

99 Theorem 7.2 is proved in Sect. A.4.

100 7.1.3 The Approximate Support Recovery Problem

101 Our third asymptotic result characterizes the phase-transition phenomenon in the
 102 approximate support recovery problem in the chi-square model. It closely parallels
 103 Theorem 3.3 for the additive errors model.

104 **Theorem 7.3** *Consider the high-dimensional chi-squared model (7.1) with signal
 105 sparsity and size as described in (7.2) and (7.3). The function*

$$106 \quad f_A(\beta) = \beta \quad (7.6)$$

107 characterizes the phase transition of approximate support recovery problem. Specifically,
 108 the following two results hold.

- 109 (i) If $\underline{r} > f_A(\beta)$, then the BH procedure \widehat{S}_p (defined in Sect. 2.2) with slowly vanishing
 110 (see Definition 3.1) nominal FDR levels achieves asymptotically approximate
 111 support recovery in the sense of (2.25).
- 112 (ii) Conversely, if $\bar{r} < f_A(\beta)$, then approximate support recovery asymptotically
 113 fails in the sense of (2.26) for all thresholding procedures.

114 Theorem 7.3 is proved in Sect. A.4 below.



¹¹⁵ **7.1.4 The Approximate–Exact Support Recovery Problem**

¹¹⁶ A counterpart of Theorem 3.5 also holds in the chi-square models.

¹¹⁷ **Theorem 7.4** *In the context of Theorem 7.3, the function*

$$\sup_{118} f_{AE}(\beta) = \left(\sqrt{\beta} + \sqrt{1 - \beta} \right)^2 \quad (7.7)$$

¹¹⁹ characterizes the phase transition of approximate–exact support recovery problem.
¹²⁰ Namely, the following two results hold.

¹²¹ (i) If $\underline{r} > f_{AE}(\beta)$, then the Benjamini–Hochberg procedure with slowly vanishing
¹²² nominal FDR levels achieves asymptotically approximate–exact support recovery
¹²³ in the sense of (2.25).

¹²⁴ (ii) Conversely, if $\bar{r} < f_{AE}(\beta)$, then for any thresholding procedure \widehat{S}_p , the
¹²⁵ approximate–exact support recovery fails in the sense of (2.26).

¹²⁶ Theorem 7.4 is proved in Sect. A.2.

¹²⁷ Notice that all phase-transitions boundaries are identical to those in the Gaussian
¹²⁸ additive error model (1.1) under one-sided alternative. We refer readers to Fig. 3.2
¹²⁹ in Sect. 3.2 for a visualization of the results in Theorems 7.1 through 7.4.

¹³⁰ All four phase transitions results in Theorems 7.1 through 7.4 focus only on the
¹³¹ idealized models (7.1) where the statistics are *independent*. Support recovery prob-
¹³² lems under dependent observations remain to be explored. Recall in Chap. 4 we
¹³³ showed that the boundary for the exact support recovery problem in the additive
¹³⁴ error model (1.1) continues to hold even under *severe dependence* and general distri-
¹³⁵ buctional assumptions. We conjecture that the concentration of maxima phenomenon,
¹³⁶ which is at the heart of the results in Chap. 4, will play a role and all of the above
¹³⁷ phase-transition results will continue to hold, under broad dependence conditions
¹³⁸ in the chi-square models. As an example, in the GWAS application, dependence
¹³⁹ among the genetic markers at different locations (known as linkage disequilibrium)
¹⁴⁰ decay as a function of their physical distances on the genome (Bush and Moore
¹⁴¹ 2012), resulting in locally dependent test statistics. It would be of great interest to
¹⁴² extend the current theory to cover important dependence structures that arise in such
¹⁴³ applications.

¹⁴⁴ **7.1.5 Comparison of One- Versus Two-Sided Alternatives
¹⁴⁵ in Additive Error Models**

¹⁴⁶ As alluded to in Sect. 1.2 in the introduction, we draw explicit comparisons between
¹⁴⁷ the one-sided and two-sided alternatives in Gaussian additive error models (1.1).

¹⁴⁸ The exact support recovery problem in the dependent Gaussian additive error
¹⁴⁹ model (1.1) was studied in Chap. 3, with parametrization of sparsity identical to that

150 in (7.2), whereas the range of the non-zero (and perhaps unequal) mean shifts $\mu(i)$
 151 was parametrized as

$$152 \quad \underline{\Delta} = \sqrt{2r \log p} \leq \mu(i) \leq \bar{\Delta} = \sqrt{2\bar{r} \log p}, \quad \text{for all } i \in S_p,$$

153 for some constants $0 < r \leq \bar{r} \leq +\infty$. Under this one-sided alternative, a phase transi-
 154 tion in the r - β plane was described, where the boundary was found to be identical
 155 to (7.4) in Theorem 7.1 for the chi-square models (7.1).

156 As discussed in Sect. 1.2, support recovery problems in the chi-square model
 157 with $\nu = 1$ correspond to the support recovery problems in the additive model under
 158 two-sided alternatives. This implies that the asymptotic signal size requirements are
 159 identical between the two-sided alternative and its one-sided counterpart, in order
 160 to achieve exact support recovery. As we shall see in numerical experiments (in
 161 Sect. 7.5 below), the difference is not very pronounced even in moderate dimensions,
 162 and vanishes as $p \rightarrow \infty$, in accordance with Theorem 7.1.

163 Comparisons can also be drawn in the approximate, approximate-exact, and
 164 exact-approximate support recovery problems between the two types of alternatives.

165 Specifically, the approximate support recovery problem in the Gaussian additive
 166 error model (1.1) under one-sided alternatives exhibits a phase transition phe-
 167 nomenon characterized by a boundary that coincides with (7.6) in Theorem 7.3. Sim-
 168 ilar to the exact support recovery problem, this indicates vanishing difference in the
 169 difficulties of the two types alternatives in approximate support recovery problems.

170 Comparing Theorems 7.2 and 3.4 as well as Theorems 7.4 and 3.5, we see that the
 171 phase transition boundaries under the two types of alternatives are also identical in
 172 the exact-approximate and approximate-exact support recovery problems.

173 To complete the comparisons, we point out that the phase-transition boundaries
 174 for the sparse signal detection problem in the two types of alternatives are both
 175 identical to (3.4). This was analyzed in Donoho and Jin (2004).

176 Therefore, all phase-transition boundaries coincide with those in the additive error
 177 models obtained in Chap. 3 under their respective parametrizations. This indicates
 178 vanishing differences between the difficulties of the one-sided and two-sided alter-
 179 natives in the Gaussian additive error model (1.1). The additional uncertainty in
 180 the two-sided alternatives does not call for larger signal sizes in these problems,
 181 asymptotically.

182 7.2 Odds Ratios and Statistical Power

183 We return to the application of association screenings for categorical variables, and
 184 put the results in the previous section to use. In particular, we focus on the exact-
 185 approximate support recovery problem, and demonstrate the consequences of its
 186 phase transition (Theorem 7.2) in genetic association studies.

187 In order to do so, we must first connect the concept of statistical signal size
 188 λ with some key quantities in association tests. While the term “signal size” likely

Table 7.1 Probabilities of the multinomial distribution in a genetic association test. (Compare and contrast with Table 1.1. We have $\mathbb{E}[O_{ij}] = n\mu_{ij}$, $i, j = 1, 2$, where $n = \sum_{i,j} O_{ij}$.)

Probabilities	Genotype		Total by phenotype
	Variant 1	Variant 2	
Cases	μ_{11}	μ_{12}	ϕ_1
Controls	μ_{21}	μ_{22}	ϕ_2
Total by genotype	θ_1	θ_2	1

sounds foreign to most practitioners, it is intimately linked with the concept of “effect sizes”—or odds ratios—in association studies, which are frequently estimated and reported in GWAS catalogs. Effect sizes, on the other hand, may be alien to some statisticians. In this section, we aim to bridge the two languages by characterizing the relationship between “signal size” and “odds-ratio” parameterizations in the special, but fairly common case of association tests on 2-by-2 contingency tables.

Recall the general setup of genetic association testing in Sect. 1.2, where one wants to detect the association between genetic variations at a specific location and the occurrence of a disease. An individual randomly drawn from the target population will have two (random) characteristics: a phenotype indicating whether the individual has the condition or is healthy (i.e., belonging to the Case group or the Control group), and a genotype that encodes the genetic variation in question. Table 1.1 in the introduction summarizes the *counts* for all phenotype–genotype combinations for the individuals in a given study sampled from the population. These counts may be assumed to follow a multinomial distribution, with probabilities given in Table 7.1 below.

Consider a 2-by-2 multinomial distribution with marginal probabilities of phenotypes (ϕ_1, ϕ_2) and genotypes (θ_1, θ_2). The *probability* table (as opposed to the table of multinomial *counts* in the introduction) is as follows.

The odds ratio (i.e., “effect size”) is defined as the ratio of the phenotype frequencies between the two genotype variants,

$$R := \frac{\mu_{11}}{\mu_{21}} / \frac{\mu_{12}}{\mu_{22}} = \frac{\mu_{11}\mu_{22}}{\mu_{12}\mu_{21}}. \quad (7.8)$$

The multinomial distribution is fully parametrized by the trio (θ_1, ϕ_1, R) . Odds ratios further away from 1 indicate greater contrasts between the probability of outcomes. Independence between the genotypes and phenotypes would imply an odds ratio of one, and hence $\mu_{jk} = \phi_j\theta_k$, for all $j, k \in \{1, 2\}$.

For a sequence of local alternatives $\mu^{(1)}, \mu^{(2)}, \dots$, such that $\sqrt{n}(\mu_{jk}^{(n)} - \phi_j\theta_k)$ converges to a constant table $\delta = (\delta_{jk})$, the chi-square test statistics converge in distribution to the non-central chi-squared distribution with non-centrality parameter



$$\lambda = \sum_{j=1}^2 \sum_{k=1}^2 \delta_{jk}^2 / (\phi_j \theta_k).$$

See, e.g., Ferguson (2017). Hence, for large samples from a fixed distribution (μ_{ij}) , the statistic is well approximated by a $\chi_1^2(\lambda)$ distribution, where

$$\lambda = n \sum_{j=1}^2 \sum_{k=1}^2 \frac{(\mu_{jk} - \phi_j \theta_k)^2}{\phi_j \theta_k}. \quad (7.9)$$

Power calculations therefore only depend on the μ_{jk} 's through $\lambda = nw^2$, where we define

$$w^2 := \lambda/n \quad (7.10)$$

to be the *signal size per sample*. Statistical power would be increasing in w^2 for fixed sample sizes.

The next proposition states that the statistical signal size per sample can be parametrized by the odds ratio and the marginals in the probability table.

Proposition 7.1 Consider a 2-by-2 multinomial distribution with marginal distributions $(\phi_1, \phi_2 = 1 - \phi_1)$ and $(\theta_1, \theta_2 = 1 - \theta_1)$. Let signal size w^2 be defined as in (7.10), and odds ratio R be defined as in (7.8). If $R = 1$, we have $w^2 = 0$; if $R \in (0, 1) \cup (1, +\infty)$, then we have

$$w^2(R) = \frac{1}{4A(R-1)^2} \left(B + CR - \sqrt{(B+CR)^2 - 4A(R-1)^2} \right)^2, \quad (7.11)$$

where $A = \phi_1 \theta_1 \phi_2 \theta_2$, $B = \phi_1 \theta_1 + \phi_2 \theta_2$, and $C = \phi_1 \theta_2 + \phi_2 \theta_1$.

Proof We parametrize the 2-by-2 multinomial distribution with the parameter δ ,

$$\mu_{11} = \phi_1 \theta_1 + \delta, \quad \mu_{12} = \phi_1 \theta_2 - \delta, \quad \mu_{21} = \phi_2 \theta_1 - \delta, \quad \mu_{22} = \phi_2 \theta_2 + \delta. \quad (7.12)$$

By relabeling of categories, we may assume $0 < \theta_1, \phi_1 \leq 1/2$ without loss of generality. Note that δ must lie within the range $[\delta_{\min}, \delta_{\max}]$, where

$$\delta_{\min} := \max\{-\phi_1 \theta_1, -\phi_2 \theta_2, \phi_1 \theta_2 - 1, \phi_2 \theta_1 - 1\} = -\phi_1 \theta_1,$$

and

$$\delta_{\max} := \min\{1 - \phi_1 \theta_1, 1 - \phi_2 \theta_2, \phi_1 \theta_2, \phi_2 \theta_1\} = \min\{\phi_1 \theta_2, \phi_2 \theta_1\},$$

in order for $\mu_{ij} \geq 0$ for all $i, j \in \{1, 2\}$. Under this parametrization, Relation (7.8) then becomes

$$R = \frac{\mu_{11}\mu_{22}}{\mu_{12}\mu_{21}} = \frac{\phi_1 \theta_1 \phi_2 \theta_2 + \delta(\phi_1 \theta_1 + \phi_2 \theta_2) + \delta^2}{\phi_1 \theta_1 \phi_2 \theta_2 - \delta(\phi_1 \theta_2 + \phi_2 \theta_1) + \delta^2} = \frac{A + \delta B + \delta^2}{A - \delta C + \delta^2}, \quad (7.13)$$

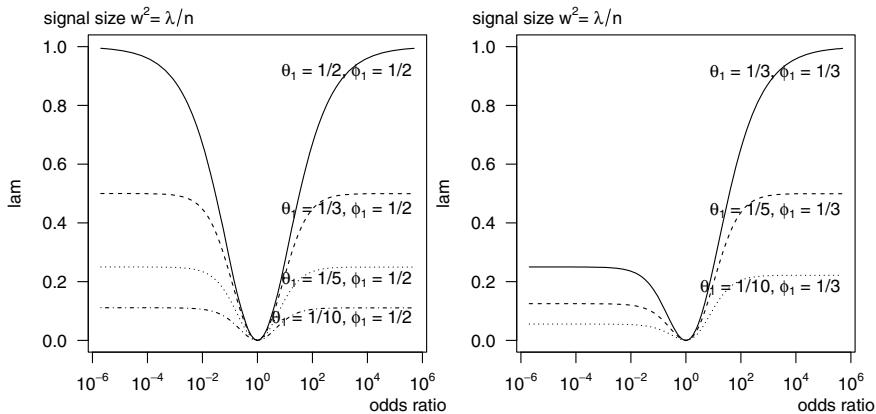


Fig. 7.1 Signal sizes per sample w^2 as functions of odds ratios in 2-by-2 multinomial distributions for selected genotype marginals in balanced (left) and unbalanced (right) designs; see Relation (7.11) in Proposition 7.1. For given marginal distributions, extreme odds ratios imply stronger statistical signals at a given sample size. However, the signal sizes are bounded above by constants that depend on the marginal distributions; see Relations (7.15) and (7.16)

which is one-to-one and increasing in δ on $(\delta_{\min}, \delta_{\max})$. Equation (7.10) becomes

$$w^2 = \sum_{i=1}^2 \sum_{j=1}^2 \frac{(\mu_{ij} - \phi_i \theta_j)^2}{\phi_i \theta_j} = \delta^2 \sum_i \sum_j \frac{1}{\phi_i \theta_j} = \frac{\delta^2}{\phi_1 \theta_1 \phi_2 \theta_2}, \quad (7.14)$$

Solving for δ in (7.13), and plugging into the expression for signal size (7.14) yields Relation (7.11).

The other three cases ($1/2 \leq \theta_1, \phi_1 \leq 1$; $0 < \theta_1 \leq 1/2 \leq \phi_1 \leq 1$; and $0 \leq \phi_1 \leq 1/2 \leq \theta_1 \leq 1$) may be obtained similarly, or by appealing to the symmetry of the problem. \square

To understand Proposition 7.1, we illustrate Relation (7.11) for selected values of marginals θ_1 and ϕ_1 in Fig. 7.1. Observe in the figure that an odds ratio further away from one corresponds to stronger statistical signal per sample, ceteris paribus. However, this “valley” pattern is in general not symmetric around 1, except for balanced marginal distributions ($\phi_1 = 1/2$ or $\theta_1 = 1/2$). While the odds ratio R can be arbitrarily close to 0 or diverge to $+\infty$ for any marginal distribution, the signal sizes w^2 are bounded from above by constants that depend only on the marginals.

Corollary 7.1 *The signal size as a function of the odds ratio $w^2(R)$ is decreasing on $(0, 1)$ and increasing on $(1, \infty)$, with limits*

$$\lim_{R \rightarrow 0_+} w^2(R) = \min \left\{ \frac{\phi_1 \theta_1}{\phi_2 \theta_2}, \frac{\phi_2 \theta_2}{\phi_1 \theta_1} \right\}, \quad (7.15)$$



258 and

$$259 \quad \lim_{R \rightarrow +\infty} w^2(R) = \min \left\{ \frac{\phi_1 \theta_2}{\phi_2 \theta_1}, \frac{\phi_2 \theta_1}{\phi_1 \theta_2} \right\}. \quad (7.16)$$

Proof As in the proof of Proposition 7.1, we examine the case where $0 < \theta_1, \phi_1 \leq 1/2$, and leave the other three cases an exercise. Take the first derivative of the expression for w^2 in Eq. (7.14) with respect to δ , it is evident that $w^2(\delta)$ is decreasing on $[\delta_{\min}, 0]$, increasing on $(0, \delta_{\max}]$, with limits

$$\lim_{d \rightarrow \delta_{\min}} w^2(d) = \frac{\phi_1 \theta_1}{\phi_2 \theta_2}, \quad \text{and} \quad \lim_{d \rightarrow \delta_{\max}} w^2(d) = \min \left\{ \frac{\phi_1 \theta_2}{\phi_2 \theta_1}, \frac{\phi_2 \theta_1}{\phi_1 \theta_2} \right\}.$$

260

□

261 Corollary 7.1 immediately implies that balanced designs with roughly equal number
262 of cases and controls are not necessarily the most informative.

263 For example, in a study where a third of the recruited subjects carry the genetic
264 variant positively correlated with the trait (i.e., $\theta_1 = 1/3$), an unbalanced design with
265 $\phi_1 = 1/3$ would maximize w^2 at large odds ratios. This unbalanced design is much
266 more efficient compared to, say, a balanced design with $\phi_1 = 1/2$. In the first case,
267 we have $w^2 \rightarrow 1$ as $R \rightarrow \infty$; whereas in the second design, $w^2 < 1/2$ no matter how
268 large R is. This difference can also be read by comparing the dashed curve ($\theta_1 = 1/3$,
269 $\phi_1 = 1/2$) in the left panel of Fig. 7.1, with the solid curve ($\theta_1 = 1/3$, $\phi_1 = 1/3$) in
270 the right panel of Fig. 7.1.

271 7.3 Optimal Study Designs and Rare Variants

272 For a study with a fixed budget, i.e., a fixed total number of subjects n , the researcher
273 is free to choose the fraction of cases ϕ_1 to be included in the study. A natural question
274 is how this budget should be allocated to maximize the statistical power of discovery,
275 or equivalently, the signal sizes $\lambda = nw^2$.

276 In principal, Relation (7.11) can be optimized with respect to the fraction of cases
277 ϕ_1 in order to find optimal designs, if θ_1 is known and held constant. In practice,
278 this is not the case. While the fraction of cases can be controlled, the distributions
279 of genotypes *in the study* are often unknown prior to data collection, and can change
280 with the case-to-control ratio.

281 Fortunately, the conditional distributions of genotypes in the healthy control
282 groups are often estimated by existing studies, and are made available by consortia
283 such as the NHGRI-EBI GWAS catalog (MacArthur et al. 2016). We denote the
284 conditional frequency of the first genetic variant in the control group as $(f, 1 - f)$,
285 where

$$286 \quad f := \mu_{21}/\phi_2 = \mu_{21}/(1 - \phi_1). \quad (7.17)$$

287 The multinomial probability is fully parametrized by the new trio: (f, ϕ_1, R) .

288 Probabilities	Genotype		Total by phenotype
	Variant 1	Variant 2	
	Cases	$\frac{\phi_1 f R}{f R + 1 - f}$	
Controls	$f(1 - \phi_1)$	$(1 - f)(1 - \phi_1)$	$1 - \phi_1$

289 Proposition 7.1 may also be re-stated in terms of the new parametrization.

290 **Corollary 7.2** *In the 2-by-2 multinomial distribution with marginals $(\phi_1, \phi_2 = 291 1 - \phi_1)$, and conditional distribution of the variants in the control group 292 $(f, 1 - f)$, Relation (7.11) holds with $\theta_1 = \phi_1 f R / (f R + 1 - f) + f(1 - \phi_1)$ and 293 $\theta_2 = 1 - \theta_1$.*

294 The choice of ϕ_1 now has a practical solution.

295 **Corollary 7.3** *In the context of Corollary 7.2, the optimal design (ϕ_1^*, ϕ_2^*) that maximizes the signal size per sample w^2 is prescribed by*

$$297 \phi_1^* = \frac{f R + 1 - f}{f R + 1 - f + \sqrt{R}}, \quad \text{and} \quad \phi_2^* = 1 - \phi_1^*. \quad (7.18)$$

298 **Proof** Using the parametrization in (7.12), we solve for δ in (7.13) to obtain

$$299 \delta = \frac{\phi_1 f R}{f R + 1 - f} - \left(\frac{\phi_1 f R}{f R + 1 - f} + f(1 - \phi_1) \right) \phi_1 \\ 300 = \frac{f(1 - f)\phi_1(1 - \phi_1)(R - 1)}{f R + 1 - f}. \quad (7.19)$$

302 Substituting (7.19) into the expression (7.14), after some simplification, yields

$$303 w^2 = \frac{f(1 - f)\phi_1(1 - \phi_1)(R - 1)^2}{[\phi_1 R + (1 - \phi_1)D][\phi_1 + (1 - \phi_1)D]}, \quad (7.20)$$

304 where $D = f R + 1 - f > 0$. Therefore, the derivative of (7.20) with respect to ϕ_1 305 is

$$306 \frac{dw^2}{d\phi_1} = \frac{f(1 - f)(R - 1)^2}{[\phi_1 R + (1 - \phi_1)D]^2 [\phi_1 + (1 - \phi_1)D]^2} [(D^2 - R)\phi_1^2 - 2D^2\phi_1 + D^2]. \quad (7.21)$$

307 Further, we obtain the second derivative with respect to ϕ_1 ,

$$308 \frac{d^2w^2}{d\phi_1^2} = h(R, f) [(\phi_1 - 1)D^2 - \phi_1 R], \quad (7.22)$$

309 where h is some function of (R, f) taking on strictly positive values.



Since $[(\phi_1 - 1)D^2 - \phi_1 R] < 0$, the second derivative (7.22) must be strictly negative on $[0, 1]$. This implies that the first derivative (7.21) is strictly decreasing on $[0, 1]$. Since the first derivative (7.21) is strictly positive at $\phi_1 = 0$, and strictly negative at $\phi_1 = 1$, it must have a unique zero between 0 and 1, and hence, the solution to $(D^2 - R)\phi_1^2 - 2D^2\phi_1 + D^2 = 0$ in the interval of $[0, 1]$ must be the maximizer of (7.20)—when $D^2 - R > 0$, the smaller of the two roots maximizes (7.20), and when $D^2 - R < 0$, it is the larger of the two. They share the same expression $D/(D + \sqrt{R})$, which coincides with (7.18). Finally, when $D^2 = R$, the only root $\phi_1^* = 1/2$, which also coincides with (7.18), is the maximizer of (7.20). \square

Of particular interest in the genetics literature are genetic variants with very low allele frequencies in the control group (i.e., $f \approx 0$), known as rare variants. In such cases, Eq. (7.18) can be approximated using the Taylor expansion,

$$\phi_1^* = \frac{1}{1 + \sqrt{R}} + \frac{(R - \sqrt{R})f}{1 + \sqrt{R}} + O(f^2). \quad (7.23)$$

To illustrate, for rare and adversarial factors ($f \approx 0$ and $R > 1$), the optimal ϕ_1^* is less than $1/2$. Therefore, for studies under a fixed budget, controls should constitute the majority of the subjects, in order to maximize power. On the other hand, for rare and protective factors ($f \approx 0$ and $R < 1$), the optimal ϕ_1^* is greater than $1/2$, and cases should be the majority.

7.4 Phase Transitions in Large-Scale Association Screening Studies

Returning to the problem of *high-dimensional* marginal screenings for categorical covariates, we explore the manifestation of the phase transition in the exact–approximate support recovery problem in the genetic context.

Recall Theorem 7.2 predicts that FWER and FNR can be simultaneously controlled in large dimensions if and only if

$$r = \frac{\lambda}{2 \log p} = \frac{w^2 n}{2 \log p} > 1. \quad (7.24)$$

Therefore, if we were to apply FWER-controlling procedures at low nominal levels (say, 5%), then the FNR would experience a phase transition in the following sense. If the signal size is strong enough, i.e.,

$$r > 1 \iff w^2 > \frac{2 \log p}{n}, \quad (7.25)$$

then the FNR can be close to 0; otherwise, FNR must be close to 1.

Using the parametric relationship described in Corollary 7.2 (and Proposition 7.1), the inequalities in (7.25) implicitly define regions of (f, R) where associations are discoverable with high power, for a given ϕ_1 . Further, the boundary of such discoverable regions sharpens as dimensionality diverges. We illustrate this phase transition through a numerical example next.

Example 7.1 Consider association tests on 2×2 contingency tables at p locations as introduced in Sect. 1.2, where the counts follow a multinomial distribution parametrized by (f, R, ϕ_1) as in Sect. 7.3. Assume that the phenotype marginals are fixed at $\phi_1 = \phi_2 = 1/2$. Applying Bonferroni's procedure with nominal FWER at $\alpha = 5\%$ level, we can approximate the marginal power of association tests by

$$\mathbb{P}[\chi_1^2(\lambda) > \chi_{1,\alpha/p}^2], \quad (7.26)$$

where $\chi_{1,\alpha/p}^2$ is the upper (α/p) -quantile of a central chi-squared distribution with 1 degree of freedom. We calculate this marginal power as a function of the parameters (f, R) in three scenarios:

- $p = 4, n = 3 \times 10^4$
- $p = 10^2, n = 1 \times 10^5$
- $p = 10^6, n = 3 \times 10^6$

and visualize the results as heatmaps² (referred to as OR-RAF diagrams) in Fig. 7.2. These parameter values are chosen so that $\log(p)/n$ are roughly constant (around 4.6×10^{-5}).

We also overlay “equi-signal” curves, i.e., functions implicitly defined by the equations $r = c$ for a range of c (dashed curves), and highlight the predicted boundary of phase transition for the exact–approximate support recovery problem $r = 1$ (red curves). The change in marginal power clearly sharpens around the predicted boundary $r = 1$ as dimensionality diverges.

Remark 7.2 In an attempt to find empirical evidence of our theoretical predictions, we chart the genetic variants associated with breast cancer, discovered in a 2017 study by Michailidou et al. (2017) in an OR-RAF diagram. The estimated risk allele frequencies (f) and odds ratios (R) are taken from the NHGRI-EBI GWAS catalog MacArthur et al. (2016), and plotted against a power heatmap calculated according to the reported sample sizes. See lower-right panel of Fig. 7.2.

It is tempting to believe, on careless inspection, that roughly *all* discovered associations fall inside the high power region of the diagram, therefore demonstrating the phase transition in statistical power. Unfortunately, the estimates here are subject to survival bias—the study in fact uses the same dataset for *both* support estimation and parameter estimation, without adjusting the latter for the selection process.

² Since genetic variants can always be relabelled such that Variant 1 is positively associated with cases, we only produce part of the diagram where $R > 1$. Sample sizes marked in the figure are adjusted by a factor of 1/2, to reflect the genetic context where a pair of alleles are measured for every individual at every genomic location.

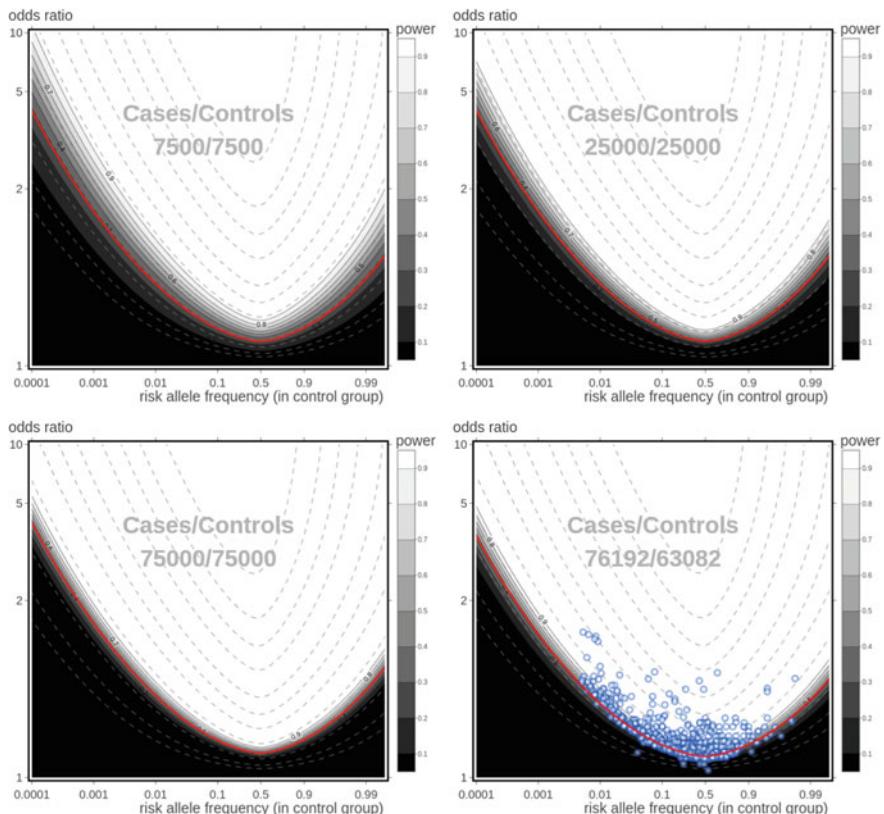


Fig. 7.2 The OR-RAF diagram visualizing the marginal power of discovery in genetic association studies, after applying Bonferroni's procedure with nominal FWER at 5% level. Sample sizes are marked in each panel, and the problem dimensions are, respectively, $p = 4$ (upper-left), $p = 10^2$ (upper-right), and $p = 10^6$ (lower-left), so that $n / \log p$ are roughly constant. Red curves mark the boundaries ($r = 1$) of the phase transition for the exact-approximate support recovery problem; dashed curves are the equi-signal (equi-power) curves. The phase transition in signal sizes λ translates into the phase transition in terms of (f, R) , and sharpens as $p \rightarrow \infty$; see Example 7.1. In the lower right panel, we visualize discovered associations (blue circles) in a recent GWA study (Michailidou et al. 2017); the estimated odds ratios and risk allele frequencies are subject to survival bias and should not be taken at their face values; see Remark 7.2.

377 The seemingly striking agreement between the power calculations and the estimated
 378 effects of reported associations *should not* be taken as evidence for the validity of our
 379 theory. We conjecture, as the theory predicts, that accurate and unbiased parameter
 380 estimates from an independent replication will still place the associations in the high
 381 power region of the diagram.

382 Finally, we demonstrate with an example how results in Sects. 7.1 and 7.2 may
 383 be used for planning prospective association studies.

Example 7.2 In a GWAS with $p = 10^6$ genomic marker locations, researchers wish to locate genetic associations with the trait of interest. Specifically, they wish to maximize power in the region where genetic variants have risk allele frequencies of 0.01 and odds ratios of 1.2. By Corollary 7.3, the optimal design has a fraction of cases $\phi^* = 0.478$, yielding the statistical signal size per sample $w^2 \approx 9.00 \times 10^{-5}$ according to Corollary 7.2.

If we wish to achieve exact–approximate support recovery in the sense of (2.25), Theorem 7.2 predicts that the signal size parameter r has to be at least $f_{EA}(\beta) = 1$. This signal size calls for a sample size of $n = \lambda/w^2 = 2r \log(p)/w^2 \approx 307,011$. In a typical GWAS, a pair of alleles are sequenced for every marker location, bringing the required number of subjects in the study to $n/2 \approx 153,509$.

In comparison, a more accurate power calculation directly using (7.26) predicts that $n/2 = 165,035$ subjects are needed, under the set of parameters ($p = 10^6$, $f = 0.01$, $R = 1.2$) and FWER = 0.05, FNR = 0.5; this is 7% higher than our crude asymptotic approximation. In general, we recommend using the more precise calculations over the back-of-the-envelope asymptotics for planning prospective studies and performing systematic reviews; a user-friendly web application implementing the more precise approximations is provided in Gao et al. (2019). Nevertheless, the theoretical results on phase transitions generate simple, accurate, and powerful insights that cannot be easily derived from numerical calculations.

7.5 Numerical Illustrations of the Phase Transitions in Chi-Square Models

We illustrate with simulations the phase transition phenomena in the chi-square model, and compare numerically the required signal sizes in support recovery problems between the two types of alternatives in the additive error model.

7.5.1 Exact Support Recovery

The sparsity of the signal vectors in the experiments are parametrized as in (7.2). Signal sizes are assumed equal with magnitude $\lambda(i) = 2r \log p$ for $i \in S$. We estimate the support set S using Bonferroni's procedure with nominal FWER level set at $1/(5\log p)$. The nominal FWER levels vanishes slowly, in line with the assumptions in Theorem 7.1. Experiments were repeated 1000 times at each of the 400 sparsity–signal-size combinations, for dimension $p = 10^4$.

The empirical probabilities of exact support recovery under Bonferroni's procedure are shown in Fig. 7.3. The numerical results suggest good accuracy of the predicted boundaries in high-dimensions (left panels of Fig. 7.3).

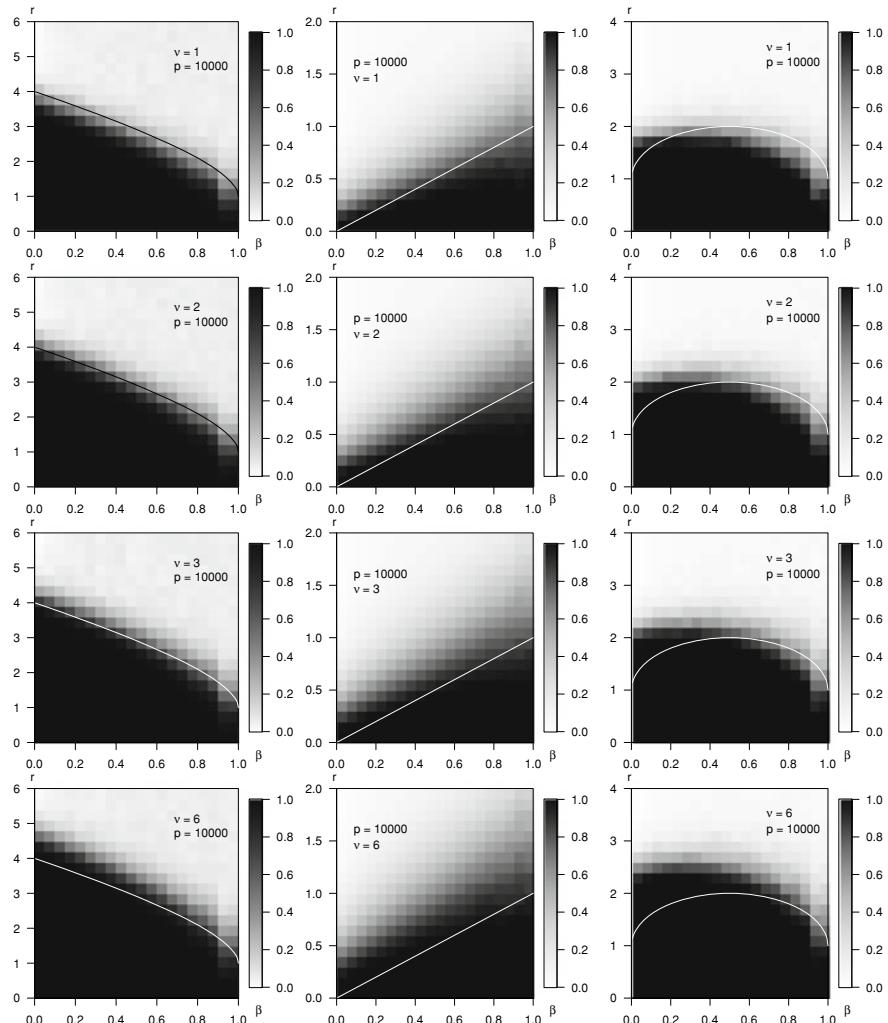


Fig. 7.3 The empirical risks of exact, approximate, and approximate-exact support recovery (left to right) in the chi-squared model (1.3) with Bonferroni's procedure and the Benjamini–Hochberg procedure. We display results as a heatmap for $v = 1, 2, 3, 6$ (first to last row) at dimension $p = 10^4$ (left to right column), for a grid of sparsity levels β and signal sizes r . The experiments were repeated 1000 times for each sparsity-signal size combination; darker color indicates higher probability of exact support recovery. Numerical results are in general agreement with the boundaries described in Theorem 7.1; for large v 's, the phase transitions take place somewhat above the predicted boundaries

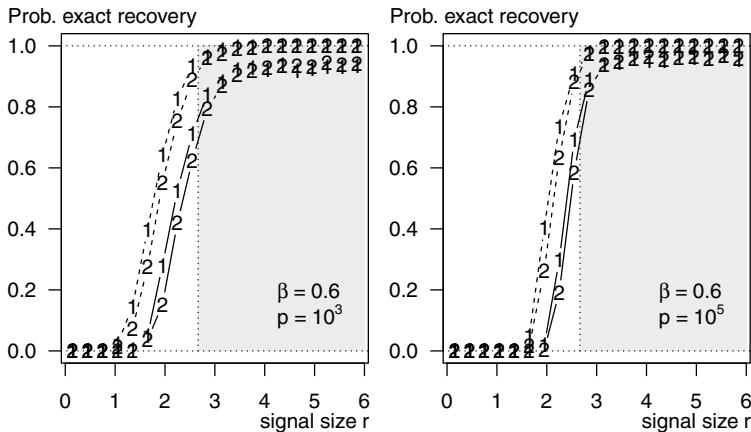


Fig. 7.4 The empirical probability of exact support recovery of Bonferroni's procedure (solid curves) and the oracle procedure (dashed curves) in the chi-squared model with one degree of freedom (marked '2') in the additive Gaussian error model and under one-sided alternatives (marked '1'). We simulate at dimensions $p = 10^2, 10^3, 10^5$ (left to right) for a grid of signal sizes r and sparsity level $\beta = 0.6$. The experiments were repeated 1000 times for each method-model-signal-size combination. Numerical results show evidence of convergence to the 0–1 law as predicted by Theorem 7.1; regions where asymptotically exact support recovery can be achieved are shaded in grey. The difference in power between Bonferroni's procedure and the oracle procedure, as well as in the two types of alternatives both decrease as dimensionality increases

419 We conduct further experiments to examine the optimality claims in Theorem 7.1
 420 by comparing with the oracle procedure with thresholds $t_p = \min_{i \in S} x(i)$. We also
 421 examine the claims in Sect. 7.1.5, and compare the one-sided alternatives in Gaussian
 422 additive models with the two-sided alternatives (or equivalently, the chi-square model
 423 with $\nu = 1$). We apply Bonferroni's procedure and the oracle thresholding procedure
 424 in both settings.

425 The experiments were repeated 1000 times for a grid of signal size values ranging
 426 from $r = 0$ to 6, and for dimensions $10^2, 10^3$, and 10^5 . Results of the experiments,
 427 shown in Fig. 7.4, suggest vanishing difference between difficulties of two-sided vs
 428 one-sided alternatives in the additive error models, as well as vanishing difference
 429 between the powers of Bonferroni's procedures and the oracle procedures as $p \rightarrow \infty$.

430 7.5.2 Approximate, and Approximate–Exact Support 431 Recovery

Similar experiments are conducted to examine the optimality claims in Theorem 7.3, and in Sect. 7.1.5. We define an oracle thresholding procedure for approximate support recovery, where the threshold is chosen to minimize the empirical risk. That is,

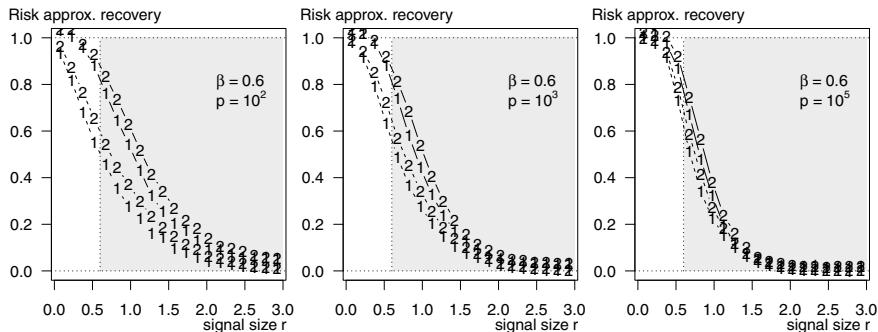


Fig. 7.5 The empirical risk of approximate support recovery of Benjamini–Hochberg’s procedure (solid curves) and the oracle procedure (dashed curves) in the chi-squared model with one degree of freedom (marked ‘2’) and in the additive Gaussian error model under one-sided alternatives (marked ‘1’). We simulate at dimensions $p = 10^2, 10^3, 10^5$ (left to right) for a grid of signal sizes r and sparsity level $\beta = 0.6$. The experiments were repeated 1000 times for each method-model-signal-size combination. Numerical results show evidence of convergence to the 0–1 law as predicted by Theorem 7.3; regions where asymptotically approximate support recovery can be achieved are shaded in grey. The difference in risks between Benjamini–Hochberg’s procedure and the oracle procedure, as well as in the two types of alternatives, both decrease as dimensionality increases

$$t_p(x, S) \in \arg \min_{t \in \mathbb{R}} \frac{|\widehat{S}(t) \setminus S|}{\max\{|\widehat{S}(t)|, 1\}} + \frac{|S \setminus \widehat{S}(t)|}{\max\{|S|, 1\}},$$

432 where $\widehat{S}(t) = \{i \mid x(i) \geq t\}$; in implementation, we only need to scan the values of
 433 observations $t \in \{x(1), \dots, x(p)\}$. The nominal FDR level for the BH procedure
 434 is set at $1/(5\log p)$, therefore slowly vanishing, in line with the assumptions in
 435 Theorem 7.3; all other parameters are identical to that in the experiments for exact
 436 support recovery in Sect. 7.5.1. The results of the experiments are shown in Fig. 7.5
 437 and in the middle column of Fig. 7.3.

We also examine the boundary described in Theorem 7.2. Experimental settings are identical to that in the experiments for approximate support recovery. We compare the performance of the BH procedure with an oracle procedure with threshold

$$t_p(x, S) \in \min_{i \in S} x(i),$$

438 and visualize the results of the experiments in the right column of Fig. 7.3. Notice that
 439 the BH procedure sets its threshold somewhat higher than the oracle, especially for
 440 small β ’s. The empirical risk of the oracle procedure (not shown here in the interest
 441 of space) follows much more closely the predicted boundary (7.7).

Author Queries

Chapter 7

Query Refs.	Details Required	Author's response
AQ1	Please note that the mismatch has been found between author Tex and PDF regarding the usage of Table citations. Hence, we have followed author Tex. Kindly, check and confirm.	

Appendix A

Additional Proofs

- 0 We review some properties of the chi-square distributions in Sect. A.1, before presenting
 1 the proofs of the main theorems on phase transitions in Sects. A.2, A.3,
 2 and A.4.

A.1 Auxiliary Facts of Chi-Square Distributions

- 4 We shall recall, and establish, some auxiliary facts about chi-square distributions.
 5 These facts will be used in the proofs of Theorems 7.1 and 7.3.

6 **Lemma A.1** (Rapid variation of chi-square distribution tails) *The central chi-square
 7 distribution with v degrees of freedom has rapidly varying tails. That is,*

$$8 \quad \lim_{x \rightarrow \infty} \frac{\mathbb{P}[\chi_v^2(0) > tx]}{\mathbb{P}[\chi_v^2(0) > x]} = \begin{cases} 0, & t > 1 \\ 1, & t = 1 \\ \infty, & 0 < t < 1 \end{cases}, \quad (\text{A.1})$$

- 9 where we overloaded the notation $\chi_v^2(0)$ to represent a random variable with the
 10 chi-square distribution.

Proof (Lemma A.1) When $v = 1$, the chi-square distribution reduces to a squared Normal, and (A.1) follows from the rapid variation of the standard Normal distribution. For $v \geq 2$, we recall the following bound on tail probabilities (see, e.g., Inglot 2010),

$$\frac{1}{2}\mathcal{E}_v(x) \leq \mathbb{P}[\chi_v^2(0) > x] \leq \frac{x}{(x - v + 2)\sqrt{\pi}}\mathcal{E}_v(x), \quad v \geq 2, \quad x > v - 2,$$

where $\mathcal{E}_v(x) = \exp\left\{-\frac{1}{2}[(x - v - (v - 2)\log(x/v) + \log v]\right\}$. Therefore, we have

$$\frac{(x - v + 2)\sqrt{\pi}}{2x} \frac{\mathcal{E}_v(tx)}{\mathcal{E}_v(x)} \leq \frac{\mathbb{P}[\chi_v^2(0) > tx]}{\mathbb{P}[\chi_v^2(0) > x]} \leq \frac{2tx}{(tx - v + 2)\sqrt{\pi}} \frac{\mathcal{E}_v(tx)}{\mathcal{E}_v(x)},$$

where $\mathcal{E}_v(tx)/\mathcal{E}_v(x) = \exp\{-\frac{1}{2}[(t - 1)x - (v - 2)\log t]\}$ converges to 0 or ∞ depending on whether $t > 1$ or $0 < t < 1$. The case where $t = 1$ is trivial. \square

Lemma A.1 and Proposition 2.2 yield the following Corollary.

Corollary A.1 *Maxima of independent observations from central chi-square distributions with v degrees of freedom are relatively stable. Specifically, let $\epsilon_p = (\epsilon_p(i))_{i=1}^p$ be independently and identically distributed (iid) $\chi_v^2(0)$ random variables. Then the triangular array $\mathcal{E} = \{\epsilon_p, p \in \mathbb{N}\}$ has relatively stable (RS) maxima in the sense of (2.38).*

Lemma A.2 (Stochastic monotonicity) *The non-central chi-square distribution is stochastically monotone in its non-centrality parameter. Specifically, for two non-central chi-square distributions both with v degrees of freedom, and non-centrality parameters $\lambda_1 \leq \lambda_2$, we have $\chi_v^2(\lambda_1) \stackrel{d}{\leq} \chi_v^2(\lambda_2)$. That is,*

$$\mathbb{P}[\chi_v^2(\lambda_1) \leq t] \geq \mathbb{P}[\chi_v^2(\lambda_2) \leq t], \quad \text{for any } t \geq 0. \quad (\text{A.2})$$

where we overloaded the notation $\chi_v^2(\lambda)$ to represent a random variable with the chi-square distribution with non-centrality parameter λ and degree-of-freedom parameter v .

Proof (Lemma A.2) Recall that non-central chi-square distributions can be written as sums of $v - 1$ standard normal random variables and a non-central normal random variable with mean $\sqrt{\lambda}$ and variance 1,

$$\chi_v^2(\lambda) \stackrel{d}{=} Z_1^2 + \dots + Z_{v-1}^2 + (Z_v + \sqrt{\lambda})^2.$$

Therefore, it suffices to show that $\mathbb{P}[(Z + \sqrt{\lambda})^2 \leq t]$ is non-increasing in λ for any $t \geq 0$, where Z is a standard normal random variable. We rewrite this expression in terms of standard normal probability function Φ ,

$$\begin{aligned} \mathbb{P}[(Z + \sqrt{\lambda})^2 \leq t] &= \mathbb{P}[-\sqrt{\lambda} - \sqrt{t} \leq Z \leq -\sqrt{\lambda} + \sqrt{t}] \\ &= \Phi(-\sqrt{\lambda} + \sqrt{t}) - \Phi(-\sqrt{\lambda} - \sqrt{t}). \end{aligned} \quad (\text{A.3})$$

The derivative of the last expression (with respect to λ) is

$$\frac{1}{2\sqrt{\lambda}} \left(\phi(\sqrt{\lambda} + \sqrt{t}) - \phi(\sqrt{\lambda} - \sqrt{t}) \right) = \frac{1}{2\sqrt{\lambda}} \left(\phi(\sqrt{\lambda} + \sqrt{t}) - \phi(\sqrt{t} - \sqrt{\lambda}) \right), \quad (\text{A.4})$$



39 where ϕ is the density of the standard normal distribution. Notice that we have used
 40 the symmetry of ϕ around 0 in the last expression.

41 Since $0 \leq \max\{\sqrt{\lambda} - \sqrt{t}, \sqrt{t} - \sqrt{\lambda}\} < \sqrt{t} + \sqrt{\lambda}$ when $t > 0$, by monotonicity
 42 of the normal density on $(0, \infty)$, we conclude that the derivative (A.4) is indeed
 43 negative. Therefore, (A.3) is decreasing in λ , and (A.2) follows for $t > 0$. For $t = 0$,
 44 equality holds in (A.2) with both probabilities being 0. \square

45 Finally, we derive asymptotic expressions for chi-square quantiles.

46 **Lemma A.3** (Chi-square quantiles) *Let F be the central chi-square distributions
 47 with v degrees of freedom, and let $u(y)$ be the $(1 - y)$ -th generalized quantile of F ,
 48 i.e.,*

$$49 \quad u(y) = F^{\leftarrow}(1 - y). \quad (\text{A.5})$$

50 Then

$$51 \quad u(y) \sim 2 \log(1/y), \quad \text{as } y \rightarrow 0. \quad (\text{A.6})$$

Proof (Lemma A.3) The case where $v = 1$ follows from the well-known formula for Normal quantiles (see, e.g., Proposition 1.1 in Gao and Stoev 2020)

$$F^{\leftarrow}(1 - y) = \Phi^{\leftarrow}(1 - y/2) \sim \sqrt{2 \log(2/y)} \sim \sqrt{2 \log(1/y)}.$$

52 The case where $v \geq 2$ follows from the following estimates of high quantiles of
 53 chi-square distributions (see, e.g., Inglot 2010),

$$\nu + 2 \log(1/y) - 5/2 \leq u(y) \leq \nu + 2 \log(1/y) + 2\sqrt{\nu \log(1/y)}, \quad \text{for all } y \leq 0.17,$$

54 where both the lower and upper bound are asymptotic to $2 \log(1/y)$. \square

55 A.2 Proof of Theorem 7.1

56 **Proof** (Theorem 7.1) We first prove the sufficient condition. The Bonferroni pro-
 57 cedure sets the threshold at $t_p = F^{\leftarrow}(1 - \alpha/p)$, which, by Lemma A.3, is asymp-
 58 totic to $2 \log p - 2 \log \alpha$. By the assumption on α in (3.17), for any $\delta > 0$, we have
 59 $p^{-\delta} = o(\alpha)$. Therefore, we have $-\log \alpha \leq \delta \log p$ for large p , and

$$60 \quad 1 \leq \limsup_{p \rightarrow \infty} \frac{2 \log p - 2 \log \alpha}{2 \log p} \leq 1 + \delta,$$

61 for any $\delta > 0$. Hence, $t_p \sim 2 \log p$.

62 The condition $r > f_E(\beta)$ implies, after some algebraic manipulation, $\sqrt{r} -$
 63 $\sqrt{1 - \beta} > 1$. Therefore, we can pick $q > 1$ such that



$$64 \quad \sqrt{r} - \sqrt{1 - \beta} > \sqrt{q} > 1. \quad (\text{A.7})$$

65 Setting the $t^* = t_p^* = 2q \log p$, we have $t_p < t_p^*$ for large p .

66 On the one hand, $\text{FWER} = 1 - \mathbb{P}[\widehat{S}_p \subseteq S_p]$ vanishes under the Bonferroni pro-
67 cedure with $\alpha \rightarrow 0$. On the other hand, for large p , the probability of no missed
68 detection is bounded from below by

$$69 \quad \mathbb{P}[\widehat{S}_p \supseteq S_p] = \mathbb{P}[\min_{i \in S} x(i) \geq t_p] \geq \mathbb{P}[\min_{i \in S} x(i) \geq t^*] \geq 1 - p^{1-\beta} \mathbb{P}[\chi_v^2(\Delta) < t^*], \quad (\text{A.8})$$

where we have used the fact that signal sizes are bounded below by Δ , and the stochastic monotonicity of chi-square distributions (Lemma A.2) in the last inequality. Writing

$$\chi_v^2(\Delta) \stackrel{\text{d}}{=} Z_1^2 + \dots + Z_{v-1}^2 + (Z_v + \sqrt{\Delta})^2$$

70 where Z_i 's are iid standard normal variables, we have

$$71 \quad \mathbb{P}[\chi_v^2(\Delta) < t^*] \leq \mathbb{P}[(Z_v + \sqrt{\Delta})^2 < t^*] = \mathbb{P}[|Z_v + \sqrt{\Delta}| < \sqrt{t^*}] \\ 72 \quad \leq \mathbb{P}\left[Z_v < -\sqrt{\Delta} + \sqrt{t^*}\right] \\ 73 \quad = \mathbb{P}\left[Z_v < \sqrt{2 \log p} (\sqrt{q} - \sqrt{r})\right]. \quad (\text{A.9})$$

75 By our choice of q in (A.7), the last probability in (A.9) can be bounded from above
76 by

$$77 \quad \mathbb{P}\left[Z_v < -\sqrt{2(1-\beta) \log p}\right] \sim \frac{\phi\left(-\sqrt{2(1-\beta) \log p}\right)}{\sqrt{2(1-\beta) \log p}} \\ 78 \quad = \frac{1}{\sqrt{2(1-\beta) \log p}} p^{-(1-\beta)},$$

80 where the first line uses Mill's ratio for Gaussian distributions (see Sect. 2.7 and
81 Relation (2.45)). This, combined with (A.8), completes the proof of the sufficient
82 condition for the Bonferroni's procedure.

83 Under the assumption of independence, Sidák's, Holm's, and Hochberg's pro-
84 cedures are strictly more powerful than Bonferroni's procedure, while controlling
85 FWER at the nominal levels. Therefore, the risks of exact support recovery for these
86 procedures also vanishes. This completes the proof for the first part of Theorem 7.1.

87 We now show the necessary condition. We first normalize the maxima by the
88 chi-square quantiles $u_p = F^{-1}(1 - 1/p)$, where F is the distribution of a (central)
89 chi-square random variable,



90

$$\mathbb{P}[\widehat{S}_p = S_p] \leq \mathbb{P}[M_{S^c} < t_p \leq m_S] \leq \mathbb{P}\left[\frac{M_{S^c}}{u_p} < \frac{m_S}{u_p}\right], \quad (\text{A.10})$$

where $M_{S^c} = \max_{i \in S^c} x(i)$ and $m_S = \min_{i \in S} x(i)$. By the relative stability of chi-square random variables (Corollary A.1), we know that $M_{S^c}/u_{|S^c|} \rightarrow 1$ in probability. Further, using the expression for u_p (Lemma A.3), we obtain

$$\frac{u_{p-p^{1-\beta}}}{u_p} \sim \frac{2 \log(p - p^{1-\beta})}{2 \log p} = \frac{\log p + \log(1 - p^{-\beta})}{\log p} \sim 1.$$

91 Therefore, the left-hand-side of the last probability in (A.10) converges to 1,

92

$$\frac{M_{S^c}}{u_p} = \frac{M_{S^c}}{u_{p-p^{1-\beta}}} \frac{u_{p-p^{1-\beta}}}{u_p} \xrightarrow{\mathbb{P}} 1. \quad (\text{A.11})$$

93 Meanwhile, for any $i \in S$, by Lemma A.2 and the fact that signal sizes are bounded
94 above by $\bar{\Delta}$, we have,

95

$$\chi_v^2(\lambda(i)) \stackrel{d}{\leq} \chi_v^2(\bar{\Delta}) \stackrel{d}{=} Z_1^2 + \dots + Z_{v-1}^2 + (Z_v + \sqrt{\bar{\Delta}})^2.$$

96 Dividing through by u_p , and taking minimum over S , we obtain

97

$$\frac{m_S}{u_p} = \min_{i \in S} \frac{\chi_v^2(\lambda(i))}{u_p} \stackrel{d}{\leq} \min_{i \in S} \left\{ \frac{Z_1^2(i) + \dots + Z_{v-1}^2(i)}{u_p} + \frac{(Z_v(i) + \sqrt{\bar{\Delta}})^2}{u_p} \right\}. \quad (\text{A.12})$$

98 Let $i^\dagger = i_p^\dagger$ be the index minimizing the second term in (A.12), i.e.,

99

$$i^\dagger := \arg \min_{i \in S} \frac{(Z_v(i) + \sqrt{\bar{\Delta}})^2}{u_p} = \arg \min_{i \in S} f_p(Z_v(i)), \quad (\text{A.13})$$

100 where $f_p(x) := (x + \sqrt{\bar{\Delta}})^2 / (2 \log p)$. We shall first show that

101

$$\mathbb{P}[f_p(Z_v(i^\dagger)) < 1 - \delta] \rightarrow 1, \quad (\text{A.14})$$

102 for some small $\delta > 0$. On the one hand, we know (by solving a quadratic inequality)
103 that

104

$$f_p(x) < 1 - \delta \iff \frac{x}{\sqrt{2 \log p}} \in (-(\sqrt{r} + \sqrt{1 - \delta}), -(\sqrt{r} - \sqrt{1 - \delta})). \quad (\text{A.15})$$

105 On the other hand, we know (by the relative stability of iid Gaussians, recall Sect. 2.7)
106 that

107

$$\frac{\min_{i \in S} Z_v(i)}{\sqrt{2 \log p}} \rightarrow -\sqrt{1 - \beta} \quad \text{in probability.} \quad (\text{A.16})$$



¹⁰⁸ Further, by the assumption on the signal sizes $\bar{r} < (1 + \sqrt{1 - \beta})^2$, we have,

$$\text{109} \quad -(\sqrt{\bar{r}} + 1) < -1 < -\sqrt{1 - \beta} < -(\sqrt{\bar{r}} - 1).$$

¹¹⁰ Therefore we can picking a small $\delta > 0$ such that

$$\text{111} \quad -(\sqrt{\bar{r}} + 1) < -(\sqrt{\bar{r}} + \sqrt{1 - \delta}) < -\sqrt{1 - \beta} < -(\sqrt{\bar{r}} - \sqrt{1 - \delta}) < -(\sqrt{\bar{r}} - 1). \quad (\text{A.17})$$

¹¹² Combining (A.15), (A.16), and (A.17), we obtain

$$\begin{aligned} \text{113} \quad \mathbb{P}\left[\min_{i \in S} f_p(Z_v(i)) < 1 - \delta\right] &= \mathbb{P}\left[f_p(Z_v(i^\dagger)) < 1 - \delta\right] \\ \text{114} \quad &\geq \mathbb{P}\left[f_p\left(\min_{i \in S} Z_v(i)\right) < 1 - \delta\right] \rightarrow 1, \\ \text{115} \end{aligned}$$

¹¹⁶ and we arrive at (A.14). As a corollary, since $u_p \sim 2 \log p$, it follows that

$$\text{117} \quad \mathbb{P}\left[\min_{i \in S} \frac{(Z_v(i) + \sqrt{\Delta})^2}{u_p} < 1 - \delta\right] \rightarrow 1. \quad (\text{A.18})$$

Finally, by independence between $Z_1^2(i) + \dots + Z_{v-1}^2(i)$ and $(Z_v(i) + \sqrt{\Delta})^2$, and the fact that i^\dagger is a function of only the latter, we have

$$Z_1^2(i^\dagger) + \dots + Z_{v-1}^2(i^\dagger) \stackrel{d}{=} Z_1^2(i) + \dots + Z_{v-1}^2(i) \quad \text{for all } i \in S.$$

¹¹⁸ Therefore, $Z_1^2(i^\dagger) + \dots + Z_{v-1}^2(i^\dagger) = O_{\mathbb{P}}(1)$, and

$$\text{119} \quad \frac{Z_1^2(i^\dagger) + \dots + Z_{v-1}^2(i^\dagger)}{u_p} \rightarrow 0 \quad \text{in probability.} \quad (\text{A.19})$$

¹²⁰ Together, (A.18) and (A.19) imply that

$$\begin{aligned} \text{121} \quad \mathbb{P}\left[\frac{m_S}{u_p} < 1 - \delta\right] &\geq \mathbb{P}\left[\min_{i \in S} \left\{ \frac{Z_1^2(i) + \dots + Z_{v-1}^2(i)}{u_p} + \frac{(Z_v(i) + \sqrt{\Delta})^2}{u_p} \right\} < 1 - \delta\right] \\ \text{122} \quad &\geq \mathbb{P}\left[\frac{Z_1^2(i^\dagger) + \dots + Z_{v-1}^2(i^\dagger)}{u_p} + \frac{(Z_v(i^\dagger) + \sqrt{\Delta})^2}{u_p} < 1 - \delta\right] \rightarrow 1. \\ \text{123} \end{aligned} \quad (\text{A.20})$$

¹²⁴ In view of (A.10), (A.11), and (A.20), we conclude that exact recovery cannot succeed with any positive probability. The proof of the necessary condition is complete. \square



126 **A.3 Proof of Theorem 7.3**

127 We first show the necessary condition. That is, when $\bar{r} < \beta$, no thresholding procedure
128 is able to achieve approximate support recovery.

129 The proof follows the ideas in Arias-Castro and Chen (2017), and is very similar
130 to the proof of Theorem 3.3. One could in principle obtain the proofs in this section
131 by referencing arguments that have appeared in Chap. 3. We choose to present the
132 proof here in full for completeness.

133 **Proof (Necessary condition in Theorem 7.3)** Denote the distributions of $\chi_v^2(0)$,
134 $\chi_v^2(\Delta)$ and $\chi_v^2(\bar{\Delta})$ as F_0 , $F_{\underline{a}}$, and $F_{\bar{a}}$ respectively.

Recall that thresholding procedures are of the form

$$\widehat{S}_p = \{i \mid x(i) > t_p(x)\}.$$

135 Denote $\widehat{S} := \{i \mid x(i) > t_p(x)\}$, and $\widehat{S}(u) := \{i \mid x(i) > u\}$. For any threshold $u \geq t_p$
136 we must have $\widehat{S}(u) \subseteq \widehat{S}$, and hence

$$\text{FDP} := \frac{|\widehat{S} \setminus S|}{|\widehat{S}|} \geq \frac{|\widehat{S} \setminus S|}{|\widehat{S} \cup S|} = \frac{|\widehat{S} \setminus S|}{|\widehat{S} \setminus S| + |S|} \geq \frac{|\widehat{S}(u) \setminus S|}{|\widehat{S}(u) \setminus S| + |S|}. \quad (\text{A.21})$$

138 On the other hand, for any threshold $u \leq t_p$ we must have $\widehat{S}(u) \supseteq \widehat{S}$, and hence

$$\text{NDP} := \frac{|S \setminus \widehat{S}|}{|S|} \geq \frac{|S \setminus \widehat{S}(u)|}{|S|}. \quad (\text{A.22})$$

140 Since either $u \geq t_p$ or $u \leq t_p$ must take place, putting (A.21) and (A.22) together,
141 we have

$$\text{FDP} + \text{NDP} \geq \frac{|\widehat{S}(u) \setminus S|}{|\widehat{S}(u) \setminus S| + |S|} \wedge \frac{|S \setminus \widehat{S}(u)|}{|S|}, \quad (\text{A.23})$$

143 for any u . Therefore it suffices to show that for a suitable choice of u , the RHS of
144 (A.23) converges to 1 in probability; the desired conclusion on FDR and FNR follows
145 by the dominated convergence theorem.

146 Let $t^* = 2q \log p$ for some fixed q , we obtain an estimate of the tail probability

$$\begin{aligned} \overline{F}_0(t^*) &= \mathbb{P}[\chi_v^2(0) > t^*] = \frac{2^{-v/2}}{\Gamma(v/2)} \int_{2q \log p}^{\infty} x^{v/2-1} e^{-x/2} dx \\ &\sim \frac{2^{-v/2}}{\Gamma(v/2)} 2 (2q \log p)^{v/2-1} p^{-q}. \end{aligned} \quad (\text{A.24})$$

where $a_p \sim b_p$ is taken to mean $a_p/b_p \rightarrow 1$; this tail estimate was also obtained
 in Donoho and Jin (2004). Observe that $|\widehat{S}(t^*) \setminus S|$ has distribution $\text{Binom}(p - s, \overline{F}_0(t^*))$ where $s = |S|$, denote $X = X_p := |\widehat{S}(t^*) \setminus S|/|S|$, and we have



$$\mu := \mathbb{E}[X] = \frac{(p-s)\overline{F}_0(t^*)}{s}, \quad \text{and} \quad \text{Var}(X) = \frac{(p-s)\overline{F}_0(t^*)F_0(t^*)}{s^2} \leq \mu/s.$$

150 Therefore for any $M > 0$, we have, by Chebyshev's inequality,

$$151 \quad \mathbb{P}[X < M] \leq \mathbb{P}[|X - \mu| > \mu - M] \leq \frac{\mu/s}{(\mu - M)^2} = \frac{1/(\mu s)}{(1 - M/\mu)^2}. \quad (\text{A.25})$$

Now, from the expression of $\overline{F}_0(t^*)$ in (A.24), we obtain

$$\mu = (p^\beta - 1)\overline{F}_0(t^*) \sim \frac{2^{1-\nu/2}}{\Gamma(\nu/2)} (2q \log p)^{\nu/2-1} p^{\beta-q}.$$

152 Since $\bar{r} < \beta$, we can pick q such that $\bar{r} < q < \beta$. In turn, we have $\mu \rightarrow \infty$, as
153 $p \rightarrow \infty$. Therefore the last expression in (A.25) converges to 0, and we conclude
154 that $X \rightarrow \infty$ in probability, and hence

$$155 \quad \frac{|\widehat{S}(t^*) \setminus S|}{|\widehat{S}(t^*) \setminus S| + |S|} = \frac{X}{X+1} \rightarrow 1 \quad \text{in probability.} \quad (\text{A.26})$$

156 On the other hand, we show that with the same choice of $u = t^*$,

$$157 \quad \frac{|S \setminus \widehat{S}(t^*)|}{|S|} \rightarrow 1 \quad \text{in probability.} \quad (\text{A.27})$$

158 By the stochastic monotonicity of chi-square distributions (Lemma A.2), the prob-
159 ability of missed detection for each signal is lower bounded by $\mathbb{P}[\chi_v^2(\lambda_i) \leq t^*] \geq$
160 $F_{\bar{a}}(t^*)$. Therefore, $|S \setminus \widehat{S}(t^*)| \stackrel{d}{\geq} \text{Binom}(s, F_{\bar{a}}(t^*))$, and it suffices to show that $F_{\bar{a}}(t^*)$
161 converges to 1. This is indeed the case, since

$$162 \quad F_{\bar{a}}(t^*) = \mathbb{P}[Z_1^2 + \dots + Z_v^2 + 2\sqrt{2\bar{r} \log p} Z_v + 2\bar{r} \log p \leq 2q \log p] \\ 163 \quad \geq \mathbb{P}[Z_1^2 + \dots + Z_v^2 \leq (q - \bar{r}) \log p, 2\sqrt{2\bar{r} \log p} Z_v \leq (q - \bar{r}) \log p],$$

165 and both events in the last line have probability going to 1 as $p \rightarrow \infty$. The necessary
166 condition is shown. \square

167 We now turn to the sufficient condition. That is, when $r > \beta$, the Benjamini–
168 Hochberg procedure with slowly vanishing FDR levels achieves asymptotic approx-
169 imate support recovery. The structure for the proof of sufficient condition follows
170 that of Theorem 2 in Arias-Castro and Chen (2017).

171 **Proof (Sufficient condition in Theorem 7.3)** The FDR vanishes by our choice of α
172 and the FDR-controlling property of the BH procedure. It only remains to show that
173 FNR also vanishes.



To do so we compare the FNR under the alternative specified in Theorem 7.3 to one with all of the signal sizes equal to Δ . Let $x(i)$ be vectors of independent observations with $p - s$ nulls having $\chi_v^2(0)$ distributions, and s signals having $\chi_v^2(\Delta)$ distributions. By Lemma 3.2, it suffices to show that the FNR under the BH procedure in this setting vanishes.

Let \widehat{G} denote the empirical survival function as in (3.36). Define the empirical survival functions for the null part and signal part

$$\widehat{W}_{\text{null}}(t) = \frac{1}{p-s} \sum_{i \notin S} \mathbb{1}\{x(i) \geq t\}, \quad \widehat{W}_{\text{signal}}(t) = \frac{1}{s} \sum_{i \in S} \mathbb{1}\{x(i) \geq t\}, \quad (\text{A.28})$$

where $s = |S|$, so that

$$\widehat{G}(t) = \frac{p-s}{p} \widehat{W}_{\text{null}}(t) + \frac{s}{p} \widehat{W}_{\text{signal}}(t).$$

Apply Lemma 3.1 to the two summands in \widehat{G} , we obtain $\widehat{G}(t) = G(t) + \widehat{R}(t)$. where

$$G(t) = \frac{p-s}{p} \overline{F}_0(t) + \frac{s}{p} \overline{F}_a(t), \quad (\text{A.29})$$

where \overline{F}_0 and \overline{F}_a are the survival functions of $\chi_v^2(0)$ and $\chi_v^2(\Delta)$ respectively, and

$$\widehat{R}(t) = O_{\mathbb{P}} \left(\xi_p \sqrt{\overline{F}_0(t) F_0(t)} + \frac{s}{p} \xi_s \sqrt{\overline{F}_a(t) F_a(t)} \right), \quad (\text{A.30})$$

uniformly in t .

Recall (see proof of Lemma 3.2) that the BH procedure is the thresholding procedure with threshold set at τ (defined in (3.37)). The NDP may also be re-written as

$$\text{NDP} = \frac{|S \setminus \widehat{S}|}{|S|} = \frac{1}{s} \sum_{i \in S} \mathbb{1}\{x(i) < \tau\} = 1 - \widehat{W}_{\text{signal}}(\tau),$$

so that it suffices to show that

$$\widehat{W}_{\text{signal}}(\tau) \rightarrow 1 \quad (\text{A.31})$$

in probability. Applying Lemma 3.1 to $\widehat{W}_{\text{signal}}$, we know that

$$\widehat{W}_{\text{signal}}(\tau) = \overline{F}_a(\tau) + O_{\mathbb{P}} \left(\xi_s \sqrt{\overline{F}_a(\tau) F_a(\tau)} \right) = \overline{F}_a(\tau) + o_{\mathbb{P}}(1).$$

So it suffices to show that $F_a(\tau) \rightarrow 0$ in probability. Now let $t^* = 2q \log(p)$ for some q such that $\beta < q < \underline{r}$. We have



$$\begin{aligned}
F_a(t^*) &= \mathbb{P}[\chi_v^2(\underline{\Delta}) \leq t^*] \leq \mathbb{P}\left[2\sqrt{\underline{\Delta}}Z_v \leq t^* - \underline{\Delta}\right] \\
&= \mathbb{P}\left[Z_v \leq \frac{t^*}{2\sqrt{\underline{\Delta}}} - \frac{\sqrt{\underline{\Delta}}}{2}\right] = \mathbb{P}\left[Z_v \leq \frac{q-r}{2\sqrt{r}}\sqrt{2\log p}\right] \rightarrow 0. \quad (\text{A.32})
\end{aligned}$$

Hence in order to show (A.31), it suffices to show

$$\mathbb{P}[\tau \leq t^*] \rightarrow 1. \quad (\text{A.33})$$

By (A.29), the mean of the empirical process \widehat{G} evaluated at t^* is

$$G(t^*) = \frac{p-s}{p}\overline{F_0}(t^*) + \frac{s}{p}\overline{F_a}(t^*). \quad (\text{A.34})$$

The first term, using Relation (A.24), is asymptotic to $p^{-q}L(p)$, where $L(p)$ is the logarithmic term in p . The second term, since $\overline{F_a}(t^*) \rightarrow 1$ by Relation (A.32), is asymptotic to $p^{-\beta}$. Therefore, $G(t^*) \sim p^{-q}L(p) + p^{-\beta} \sim p^{-\beta}$, since $p^{\beta-q}L(p) \rightarrow 0$ where $q > \beta$.

The fluctuation of the empirical process at t^* , by Relation (A.30), is

$$\begin{aligned}
\widehat{R}(t^*) &= O_{\mathbb{P}}\left(\xi_p\sqrt{\overline{F_0}(t^*)F_0(t^*)} + \frac{s}{p}\xi_s\sqrt{\overline{F_a}(t^*)F_a(t^*)}\right) \\
&= O_{\mathbb{P}}\left(\xi_p\sqrt{\overline{F_0}(t^*)}\right) + o_{\mathbb{P}}(p^{-\beta}).
\end{aligned}$$

By (A.24) and the expression for ξ_p , the first term is $O_{\mathbb{P}}(p^{-(q+1)/2}L(p))$ where $L(p)$ is a poly-logarithmic term in p . Since $\beta < \min\{q, 1\}$, we have $\beta < (q+1)/2$, and hence $\widehat{R}(t^*) = o_{\mathbb{P}}(p^{-\beta})$.

Putting the mean and the fluctuation of $\widehat{G}(t^*)$ together, we obtain

$$\widehat{G}(t^*) = G(t^*) + \widehat{R}(t^*) \sim_{\mathbb{P}} G(t^*) \sim p^{-\beta},$$

and therefore, together with (A.24), we have

$$\overline{F_0}(t^*)/\widehat{G}(t^*) = p^{\beta-q}L(p)(1 + o_{\mathbb{P}}(1)),$$

which is eventually smaller than the FDR level α by the assumption (3.17) and the fact that $\beta < q$. That is,

$$\mathbb{P}[\overline{F_0}(t^*)/\widehat{G}(t^*) < \alpha] \rightarrow 1.$$

By definition of τ (recall (3.37)), this implies that $\tau \leq t^*$ with probability tending to 1, and (A.33) is shown. The proof for the sufficient condition is complete. \square



212 **A.4 Proof of Theorems 7.2 and 7.4**

213 As with the proof of Theorem 7.3, one could shorten the presentations in this section
214 by referencing arguments in Chap. 3.

215 **Proof (Theorem 7.2)** We first show the sufficient condition. Similar to the proof of
216 Theorem 7.3, it suffices to show that

217
$$\text{NDP} = 1 - \widehat{W}_{\text{signal}}(t_p) \rightarrow 0, \quad (\text{A.35})$$

218 where t_p is the threshold of Bonferroni's procedure.

Since $\underline{r} > f_{\text{EA}}(\beta) = 1$, we can pick q such that $1 < q < \underline{r}$. Let $t^* = 2q \log p$, we have $t_p < t^*$ for large p as in the proof of Theorem 7.1. Therefore for large p , we have

$$\widehat{W}_{\text{signal}}(t_p) \geq \widehat{W}_{\text{signal}}(t^*) \geq \overline{F_a}(t^*) + o_{\mathbb{P}}(1),$$

219 where the last inequality follows from the stochastic monotonicity of the chi-square
220 family (Lemma A.2), and Lemma 3.1. Indeed, $F_a(t^*) \rightarrow 0$ by (A.32) and our choice
221 of $q < \underline{r}$. The proof of the sufficient condition is complete.

222 Proof of the necessary condition follows a similar structure to that of Theorem 7.3.
223 That is, we show that FWER + FNR has liminf at least 1 by working with the lower
224 bound

225
$$\text{FWER}(\mathcal{R}) + \text{FNR}(\mathcal{R}) \geq \mathbb{P} \left[\max_{i \in S^c} x(i) > u \right] \wedge \mathbb{E} \left[\frac{|S \setminus \widehat{S}(u)|}{|S|} \right], \quad (\text{A.36})$$

226 which holds for any thresholding procedure \mathcal{R} and for arbitrary $u \in \mathbb{R}$. By the
227 assumption that $\bar{r} < f_{\text{EA}}(\beta) = 1$, we can pick q such that $\bar{r} < q < 1$ and let $u =$
228 $t^* = 2q \log p$. By relative stability of chi-squared random variables (Lemma A.1),
229 we have

230
$$\mathbb{P} \left[\frac{\max_{i \in S^c} x(i)}{2 \log p} > \frac{t^*}{2 \log p} \right] \rightarrow 1. \quad (\text{A.37})$$

231 where the first fraction in (A.37) converges to 1, while the second converges to $q < 1$.
232 On the other hand, by our choice of $q > \bar{r}$, the second term in (A.36) also converges
233 to 1 as in (A.27). This completes the proof of the necessary condition. \square

234 **Proof (Theorem 7.4)** We first show the sufficient condition. Since FDR control is
235 guaranteed by the BH procedure, we only need to show that the FWNR also vanishes,
236 that is,

237
$$\mathbb{P} \left[\min_{i \in S} x(i) \geq \tau \right] \rightarrow 1, \quad (\text{A.38})$$

238 where τ is the threshold for the BH procedure.

239 By the assumption that $\underline{r} > f_{\text{AE}}(\beta) = (\sqrt{\beta} + \sqrt{1-\beta})^2$, we have
240 $\sqrt{\underline{r}} - \sqrt{1-\beta} > \sqrt{\beta}$, so we can pick $q > 0$, such that



$$\sqrt{r} - \sqrt{1 - \beta} > \sqrt{q} > \sqrt{\beta}. \quad (\text{A.39})$$

Let $t^* = 2q \log p$, we claim that

$$\mathbb{P}[\tau \leq t^*] \rightarrow 1. \quad (\text{A.40})$$

Indeed, by our choice of $q > \beta$, (A.40) follows in the same way that (A.33) did.

With this t^* , we have

$$\mathbb{P}\left[\min_{i \in S} x(i) \geq \tau\right] \geq \mathbb{P}\left[\min_{i \in S} x(i) \geq t^*, t^* \geq \tau\right]. \quad (\text{A.41})$$

However, by our choice of $\sqrt{q} < \sqrt{r} - \sqrt{1 - \beta}$, the probability of the first event on the right-hand side of (A.41) also goes to 1 according to (A.8) and (A.9). Together with (A.40), this proves (A.38), and completes proof of the sufficient condition.

The necessary condition follows from the lower bound

$$\text{FDR}(\mathcal{R}) + \text{FWNR}(\mathcal{R}) \geq \mathbb{E}\left[\frac{|\widehat{S}(u) \setminus S|}{|\widehat{S}(u) \setminus S| + |S|}\right] \wedge \mathbb{P}\left[\min_{i \in S} x(i) < u\right], \quad (\text{A.42})$$

which holds for any thresholding procedure \mathcal{R} and for arbitrary $u \in \mathbb{R}$.

By the assumption that $\bar{r} < f_{AE}(\beta) = (\sqrt{\beta} + \sqrt{1 - \beta})^2$, we can pick a constant $q > 0$, such that

$$\sqrt{r} - \sqrt{1 - \beta} < \sqrt{q} < \sqrt{\beta}. \quad (\text{A.43})$$

Let also $u = t^* = 2q \log p$. By our choice of $q < \beta$, we know from (A.26) that the first term on the right-hand-side of (A.42) converges to 1. It remains to show that the second term in (A.42) also converges to 1.

For the second term in (A.42), dividing through by $2 \log p$, we obtain

$$\mathbb{P}\left[\min_{i \in S} x(i) < t^*\right] = \mathbb{P}\left[\frac{m_S}{2 \log p} < q\right]. \quad (\text{A.44})$$

Similar to (A.12), we have

$$\frac{m_S}{2 \log p} \stackrel{d}{\leq} \min_{i \in S} \frac{Z_1^2(i) + \dots + Z_{v-1}^2(i)}{2 \log p} + \frac{(Z_v(i) + \sqrt{\Delta})^2}{2 \log p}. \quad (\text{A.45})$$

Define $i^\dagger = i_p^\dagger$ to be the index minimizing the second term in (A.45), i.e.,

$$i^\dagger := \arg \min_{i \in S} f_p(Z_v(i)), \quad (\text{A.46})$$

where $f_p(x) := (x + \sqrt{\Delta})^2 / (2 \log p)$.



Since $\sqrt{q} > \sqrt{\bar{r}} - \sqrt{1-\beta}$ and $q > 0$, we have $\frac{\sqrt{\bar{r}} - \sqrt{q}}{\sqrt{1-\beta}} < 1$. Also, since

$$\frac{\sqrt{\bar{r}} + \sqrt{q}}{\sqrt{1-\beta}} > 0, \quad \text{and} \quad \frac{\sqrt{\bar{r}} - \sqrt{q}}{\sqrt{1-\beta}} < \frac{\sqrt{\bar{r}} + \sqrt{q}}{\sqrt{1-\beta}},$$

we can further pick a constant $\beta_0 \in (0, 1]$ such that

$$\frac{\sqrt{\bar{r}} - \sqrt{q}}{\sqrt{1-\beta}} < \sqrt{\beta_0} < \frac{\sqrt{\bar{r}} + \sqrt{q}}{\sqrt{1-\beta}}. \quad (\text{A.47})$$

Let $Z_{[1]} \leq Z_{[2]} \leq \dots \leq Z_{[s]}$ be the order statistics of $\{Z_v(i)\}_{i \in S}$ and define $k = \lfloor s^{1-\beta_0} \rfloor$. Applying Lemma A.4 (stated below), we obtain

$$\frac{Z_{[k]}}{\sqrt{2 \log p}} = \frac{Z_{[k]}}{\sqrt{2 \log s}} \frac{\sqrt{2 \log s}}{\sqrt{2 \log p}} \rightarrow -\sqrt{\beta_0(1-\beta)} \quad \text{in probability.} \quad (\text{A.48})$$

Since we know (by solving a quadratic inequality) that

$$f_p(x) < q \iff \frac{x}{\sqrt{2 \log p}} \in \left(-(\sqrt{\bar{r}} + \sqrt{q}), -(\sqrt{\bar{r}} - \sqrt{q}) \right), \quad (\text{A.49})$$

combining (A.47), (A.48), and (A.49), it follows that

$$\mathbb{P}[f_p(Z_v(i^\dagger)) < q] \geq \mathbb{P}[f_p(Z_{[k]}) < q] \rightarrow 1.$$

Finally, using (A.19), we conclude that

$$\mathbb{P}\left[\min_{i \in S} x(i) < t^*\right] = \mathbb{P}\left[\frac{m_S}{2 \log p} < q\right] \geq \mathbb{P}[o_{\mathbb{P}}(1) + f_p(Z_v(i^\dagger)) < q] \rightarrow 1.$$

Therefore, the two terms on the right-hand-side of (A.42) both converge 1. This completes the proof of the necessary condition. \square

It only remains to justify (A.48).

Lemma A.4 (Relative stability of order statistics) *Let $Z_{[1]} \leq \dots \leq Z_{[s]}$ be the order statistics of s iid standard Gaussian random variables. Let $\beta_0 \in (0, 1]$ and define $k = \lfloor s^{1-\beta_0} \rfloor$, then we have*

$$\frac{Z_{[k]}}{\sqrt{2 \log s}} \rightarrow -\sqrt{\beta_0} \quad \text{in probability.} \quad (\text{A.50})$$

Proof (Lemma A.4) Using the Renyi representation for order statistics, we write

$$Z_{[i]} = \Phi^{-1}(U_{[i]}), \quad (\text{A.51})$$



where $U_{[i]}$ is the i^{th} (smallest) order statistic of s independent uniform random variables over $(0, 1)$. Since $U_{[i]}$ has a $\text{Beta}(i, s+1-i)$ distribution, with mean and standard deviation,

$$\mathbb{E}[U_{[k]}] = k/(s+1) \sim s^{-\beta_0}, \quad \text{and} \quad \text{sd}(U_{[k]}) = \frac{1}{s+1} \sqrt{\frac{k(s+1-k)}{s+2}} \sim s^{-\frac{1+\beta_0}{2}},$$

we obtain by Chebyshev's inequality

$$\mathbb{P}[s^{-\beta_0}(1-\epsilon) < U_{[k]} < s^{-\beta_0}(1+\epsilon)] \rightarrow 1,$$

284 where ϵ is an arbitrary positive constant. This implies, by representation (A.51),

285
$$\mathbb{P}[\Phi^{-}(s^{-\beta_0}(1-\epsilon)) < Z_{[k]} < \Phi^{-}(s^{-\beta_0}(1+\epsilon))] \rightarrow 1. \quad (\text{A.52})$$

286 Using the expression for standard Gaussian quantiles (see, e.g., Proposition 1.1. in
287 Gao and Stoev 2020), we know that

288
$$\begin{aligned} \Phi^{-}(s^{-\beta_0}(1-\epsilon)) &\sim -\sqrt{2 \log(s^{\beta_0}/(1-\epsilon))} \\ 289 &= -\sqrt{2(\beta_0 \log s - \log(1-\epsilon))} \sim -\sqrt{2\beta_0 \log s}, \end{aligned}$$

291 and similarly $\Phi^{-}(s^{-\beta_0}(1+\epsilon)) \sim -\sqrt{2\beta_0 \log s}$. Since both ends of the interval in
292 (A.52) are asymptotic to $-\sqrt{2\beta_0 \log s}$, the desired conclusion follows. \square



0 Appendix B

1 Exact Support Recovery in Non AGG Models

2 B.1 Strong Classification Boundaries in Other Light-Tailed 3 Error Models

4 The strong classification boundaries extend beyond the AGG models. As our analysis
 5 in Chap. 4 suggests, all additive error models where the errors have URS maxima
 6 exhibit this phase transition phenomenon under appropriate parametrization of the
 7 sparsity and signal sizes. We derive explicit boundaries for two additional classes of
 8 models under the general form of the additive noise models (1.1) with *heavier* and
 9 *lighter* tails than the AGG models, respectively.

We would like to point out that the sparsity and signal sizes can be re-parametrized for the boundaries to have different shapes. For example in the case of Gaussian errors, if we re-parametrize sparsity s with $\tilde{\beta} = 2 - (1 + \sqrt{1 - \beta})^2$ where $\tilde{\beta} \in (0, 1)$, then the signal sparsity would have a slightly more complicated form:

$$|S_p| = \lfloor p^{1-\beta} \rfloor = \left\lfloor p^{(\sqrt{2-\tilde{\beta}}-1)^2} \right\rfloor,$$

10 while the strong classification boundary would take on the simpler form:

$$f_E(\beta) = \tilde{f}_E(\tilde{\beta}) = 2 - \tilde{\beta}. \quad (\text{B.1})$$

11 In the next two classes of models we will adopt parametrizations such that the bound-
 12 aries are of the form \tilde{g} in (B.1).

¹⁴ **B.1.1 Additive Error Models with Heavier-Than-AGG Tails**

¹⁵ Distributions such as the log-normal have heavier tails than the AGG model, yet
¹⁶ the tails are nevertheless rapidly-varying. Therefore, Proposition 2.2 applies, and
¹⁷ we expect to see phase-transition-type results when the additive errors have these
¹⁸ heavier-than-AGG tails.

¹⁹ **Example B.1 (Heavier than AGG)** Let $\gamma > 1$, $c > 0$, and suppose that

$$\log \bar{F}(x) = -(\log x)^\gamma (c + M(x)), \quad (\text{B.2})$$

²¹ where $\lim_{x \rightarrow \infty} M(x) \log^\gamma x = 0$. Then, Relation (2.39) holds under model (B.2).

²² Further, if the entries in the array are independent, the maxima are relatively stable.

²³ The behavior of the quantiles u_p in this model is as follows. As $p \rightarrow \infty$,

$$\text{u}_p \sim \exp \left\{ (c^{-1} \log p)^{1/\gamma} \right\} \iff c (\log u_p)^\gamma + o(1) = \log(p) = -\log \bar{F}(u_p).$$

²⁵ since u_p diverges, and $M(u_p)$ is $o((\log^\gamma u_p)^{-1})$.

²⁶ Following Example B.1, assume that the errors in Model (1.1) have rapidly varying
²⁷ right tails

$$\log \bar{F}(x) = -(\log x)^\gamma (c + M(x)), \quad (\text{B.3})$$

²⁹ as $x \rightarrow \infty$, and left tails

$$\log F(x) = -(\log(-x))^\gamma (c + M(-x)), \quad (\text{B.4})$$

³¹ as $x \rightarrow -\infty$.

Theorem B.1 Suppose the marginals F follows (B.3) and (B.4). Let

$$k(\beta) = \log p - ((\log p)^{1/\gamma} + \log(1 - \beta))^\gamma,$$

and let the signal μ have

$$|S_p| = \lfloor p e^{-k(\beta)} \rfloor$$

non-zero entries. Assume the magnitudes of non-zero signal entries are in the range between

$$\underline{\Delta} = \exp \{(\log p)^{1/\gamma}\} \underline{r} \quad \text{and} \quad \overline{\Delta} = \exp \{(\log p)^{1/\gamma}\} \overline{r}.$$

³² If $\underline{r} > \tilde{f}_{\text{E}}(\beta) = 2 - \beta$, then Bonferroni's procedure \hat{S}_p (defined in (2.21)) with appro-
³³ priately calibrated FWER $\alpha \rightarrow 0$ achieves asymptotic perfect support recovery,
³⁴ under arbitrary dependence of the errors.

³⁵ On the other hand, when the errors are uniformly relatively stable, if $\overline{r} < \tilde{f}_{\text{E}}(\beta) =$
³⁶ $2 - \beta$, then no thresholding procedure can achieve asymptotic perfect support recov-
³⁷ ery with positive probability.



³⁸ **B.1.2 Additive Error Models with Lighter-Than-AGG Tails**

³⁹ Similar to how Proposition 2.2 applies to models with heavier-than-AGG tails, it also
⁴⁰ to error models with lighter tails than the AGG class.

⁴¹ **Example B.2 (Lighter than AGG)** With $v > 0$, and $L(x)$ a slowly varying function,
⁴² the class of distributions

⁴³
$$\log \bar{F}(x) = -\exp \{x^v L(x)\}, \quad (\text{B.5})$$

⁴⁴ is rapidly varying. The quantiles can be derived explicitly in a subclass of (B.5)
⁴⁵ where $L(x) \rightarrow 1$, or equivalently, when $\log |\log \bar{F}(x)| \sim x^v$,

⁴⁶
$$u_p \sim (\log \log p)^{1/v} \iff \exp \{u_p^v (1 + o(1))\} = \log(p) = -\log \bar{F}(u_p).$$

⁴⁷ Following Example B.2, assume that errors in Model (1.1) has rapidly varying
⁴⁸ right tails

⁴⁹
$$\log \bar{F}(x) = -\exp \{x^v L(x)\}, \quad (\text{B.6})$$

⁵⁰ where $L(x)$ is a slowly varying function, as $x \rightarrow \infty$, and left tails

⁵¹
$$\log \bar{F}(x) = -\exp \{-x^v L(-x)\}, \quad (\text{B.7})$$

⁵² as $x \rightarrow -\infty$.

⁵³ The phase transition results in multiple testing problems under such tail assumptions
⁵⁴ is characterizes as follows.

Theorem B.2 Suppose marginals F follow (B.6) and (B.7). Let

$$k(\beta) = \log p - (\log(p))^{(1-\beta)v},$$

and let the signal μ have

$$|S_p| = \lfloor p e^{-k(\beta)} \rfloor$$

non-zero entries. Assume the magnitudes of non-zero signal entries are in the range between

$$\underline{\Delta} = \log \log p^{1/v} r_- \text{ and } \bar{\Delta} = \log \log p^{1/v} \bar{r}_.$$

⁵⁵ If $r > \tilde{f}_E(\beta) = 2 - \beta$, then Bonferroni's procedure \hat{S}_p (defined in (2.21)) with appro-
⁵⁶ priately calibrated FWER $\alpha \rightarrow 0$ achieves asymptotic perfect support recovery,
⁵⁷ under arbitrary dependence of the errors.

⁵⁸ On the other hand, when the errors are uniformly relatively stable, if $\bar{r} < \tilde{f}_E(\beta) =$
⁵⁹ $2 - \beta$, then no thresholding procedure can achieve asymptotic perfect support recov-
⁶⁰ ery with positive probability.



61 B.2 Thresholding Procedures Under Heavy-Tailed Errors

62 We analyze the performance of thresholding estimators under heavy-tailed models
 63 in this section, and illustrate its lack of phase transition. Suppose we have iid errors
 64 with Pareto tails in Model (1.1), that is, $\epsilon(i)$'s have common marginal distribution F
 65 where

$$66 \quad \bar{F}(x) \sim x^{-\alpha} \quad \text{and} \quad F(-x) \sim x^{-\alpha}, \quad (\text{B.8})$$

67 as $x \rightarrow \infty$. It is well-known (see, e.g., Theorem 1.6.2 of Leadbetter et al. 1983) that
 68 the maxima of iid Pareto random variables have Frechet-type limits. Specifically, we
 69 have

$$70 \quad \frac{\max_{i \in \{1, \dots, p\}} \epsilon(i)}{u_p} \implies Y, \quad (\text{B.9})$$

71 in distribution, where $u_p = F^\leftarrow(1 - 1/p) \sim p^{1/\alpha}$, and Y is a standard α -Frechet
 72 random variable, i.e.,

$$73 \quad \mathbb{P}[Y \leq t] = \exp\{-t^{-\alpha}\}, \quad t > 0.$$

74 By symmetry in our assumptions, the same argument applies to the minima as well.

75 **Theorem B.3** *Let errors in Model (1.1) be as described in Relation (B.8). Let the
 76 signal have $s = |S| = fp$ non-zero entries, with magnitude $\Delta = rp^{1/\alpha}$, where both
 77 $f \in (0, 1)$ and $r \in (0, +\infty)$ may depend on p , so that no generality is lost.*

78 *Under these assumptions, the necessary condition for thresholding procedures \widehat{S}
 79 to achieve exact support recovery ($\mathbb{P}[\widehat{S} = S] \rightarrow 1$) is*

$$80 \quad \liminf_{p \rightarrow \infty} r = \infty. \quad (\text{B.10})$$

81 *Condition (B.10) is also sufficient for the oracle thresholding procedure to succeed
 82 in the exact support recovery problem.*

On the other hand, the necessary and sufficient condition for all thresholding
 procedures to fail exact support recovery ($\mathbb{P}[\widehat{S} = S] \rightarrow 0$) is

$$\limsup_{p \rightarrow \infty} r = 0.$$

83 In other words, Theorem B.3 states that there does not exist a non-trivial phase
 84 transition for thresholding procedures when errors have (two-sided) α -Pareto tails.

85 **Proof (Theorem B.3)** Recall the oracle thresholding procedure $\widehat{S}^* = \{i : x(i) \geq
 86 x_{[s]}\}$, and the set of all thresholding procedures, denoted \mathcal{S} (see Definition 2.20).
 87 The probability of exact support recovery by any thresholding procedure $\widehat{S} \in \mathcal{S}$ is
 88 bounded above by that of \widehat{S}^* , that is,



$$\begin{aligned}
99 & \max_{\widehat{S} \in \mathcal{S}} \mathbb{P}[\widehat{S} = S] = \mathbb{P}[\widehat{S}^* = S] = \mathbb{P}\left[\max_{i \in S^c} x(i) \leq \min_{i \in S} x(i)\right] \\
90 & = \mathbb{P}\left[\frac{\max_{i \in S^c} x(i)}{u_p} \leq \frac{\min_{i \in S} x(i)}{u_p}\right] \\
91 & = \mathbb{P}\left[\frac{M_{S^c}}{u_p} \leq \frac{m_S}{u_p} + r_p\right], \tag{B.11}
92
\end{aligned}$$

93 where $M_{S^c} = \max_{i \in S^c} \epsilon(i)$ and $m_S = \min_{i \in S} \epsilon(i)$. For any $\alpha > 0$, the following
94 elementary relations hold,

95 $0 < L \leq (1-f)^{1/\alpha} + f^{1/\alpha} \leq U < \infty$, for all $f \in (0, 1)$,

96 where $L = \min\{1, 2(1/2)^{1/\alpha}\}$ and $U = \max\{1, 2(1/2)^{1/\alpha}\}$. Therefore we have,

97 $U \max\left\{\frac{M_{S^c}}{u_p}, -\frac{m_S}{u_p}\right\} < r_p \implies (1-f)^{1/\alpha} \frac{M_{S^c}}{u_p} - f^{1/\alpha} \frac{m_S}{u_p} < r_p, \tag{B.12}$

98 and

99 $L \min\left\{\frac{M_{S^c}}{u_p}, -\frac{m_S}{u_p}\right\} < r_p \iff (1-f)^{1/\alpha} \frac{M_{S^c}}{u_p} - f^{1/\alpha} \frac{m_S}{u_p} < r_p. \tag{B.13}$

100 Putting together (B.11), (B.12), and (B.13), we have

101 $\mathbb{P}\left[\max\left\{\frac{M_{S^c}}{u_p}, -\frac{m_S}{u_p}\right\} < r_p/U\right] \leq \mathbb{P}[\widehat{S}^* = S] \leq \mathbb{P}\left[\min\left\{\frac{M_{S^c}}{u_p}, -\frac{m_S}{u_p}\right\} < r_p/L\right]. \tag{B.14}$

102 We know from the weak convergence result (B.9) that for any $\epsilon > 0$ there is a constant
103 N such that for all $p > N$ we have

104 $\mathbb{P}\left[\max\left\{\frac{M_{S^c}}{u_p}, -\frac{m_S}{u_p}\right\} < r_p/U\right] \geq \mathbb{P}\left[\max\{Y^{(1)}, Y^{(2)}\} < r_p/U\right] - \epsilon, \tag{B.15}$

105 where $Y^{(1)}$ and $Y^{(2)}$ are independent α -Frechet random variables with scale coefficients
106 $(1-f)^{1/\alpha}$ and $f^{1/\alpha}$ respectively. That is,

107 $\mathbb{P}[Y^{(1)} \leq t] = \exp\{-(1-f)/t^\alpha\}, \text{ and } \mathbb{P}[Y^{(2)} \leq t] = \exp\{-f/t^\alpha\}.$

Since the distributional limit in (B.15) has a density (with respect to the Lebesgue measure), we know that density is bounded above by a finite constant, say, K . For the same choice of ϵ as before, we can find a further constant N' such that for all $p > \max\{N, N'\}$ we have

$$\liminf r_p < \epsilon/K + r_p,$$



108 so that the right hand side of (B.15) is bounded by

$$\mathbb{P}\left[\max\{Y^{(1)}, Y^{(2)}\} < r_p/U\right] - \epsilon \geq \mathbb{P}\left[\max\{Y^{(1)}, Y^{(2)}\} < \frac{\liminf r_p}{U}\right] - 2\epsilon. \quad (\text{B.16})$$

109 110 By the arbitrariness in the choice of ϵ , we conclude from (B.15) and (B.16) that

$$\liminf \mathbb{P}\left[\max\left\{\frac{M_{S^c}}{u_p}, -\frac{m_S}{u_p}\right\} < r_p/U\right] \geq \mathbb{P}\left[\max\{Y^{(1)}, Y^{(2)}\} < \frac{\liminf r_p}{U}\right]. \quad (\text{B.17})$$

111 Combining Relations (B.14) and (B.17), we know that if $\liminf r_p = \infty$, we must have

$$\liminf \mathbb{P}[\widehat{S}^* = S] \geq \mathbb{P}\left[\max\{Y^{(1)}, Y^{(2)}\} < \frac{\liminf r_p}{U}\right] = 1.$$

112 Conversely, if $\liminf \mathbb{P}[\widehat{S}^* = S] < 1$, we must have $\liminf r_p < \infty$.

113 114 Similarly, we can obtain the upper bound of exact support recovery probability for the optimal thresholding procedure,

$$\limsup \mathbb{P}\left[\min\left\{\frac{M_{S^c}}{u_p}, -\frac{m_S}{u_p}\right\} < r_p/L\right] \leq \mathbb{P}\left[\min\{Y^{(1)}, Y^{(2)}\} < \frac{\limsup r_p}{L}\right]. \quad (\text{B.18})$$

115 116 117 The conclusions of the second part of Theorem B.3 follow from (B.14) and (B.18). \square

118 119 120 The probability of exact recovery can be approximated if the parameters r and f converge. The next result follows from a small modification of the arguments in the proof of Theorem B.3.

Corollary B.1 *Under the assumptions in Theorem B.3, if $\lim r = r^*$, and $\lim f = f^*$, for some constant $r^* \geq 0$ and $f^* \in [0, 1]$, then*

$$\lim \mathbb{P}[\widehat{S}^* = S] = \mathbb{P}\left[(1 - f^*)^{1/\alpha} Z_1 + (f^*)^{1/\alpha} Z_2 < r^*\right].$$

121 122 where Z_1 and Z_2 are independent standard α -Frechet random variables, i.e., $\mathbb{P}[Z_i \leq x] = \exp\{-x^{-\alpha}\}$, $x > 0$.

123 124 **Remark B.1** Of course one might wonder if it would be meaningful to derive a “phase transition” under a different parametrization of the signal sizes, say

$$\Delta = p^{r/\alpha}. \quad (\text{B.19})$$

126 127 128 In this case, Theorem B.3 suggests that a “phase transition” takes place at $r = 1$. However, this non-multiplicative parametrization of the signal sizes would make power analysis (like in Example 3.1) dimension-dependent.

129 130 To illustrate, in the case of Gaussian errors with variance 1, if we were interested in small signals of size $\sqrt{2r \log p}$, where $r < 1$ is below the boundary (4.5), then



131 we only need $n > 2/r$ samples to guarantee discovery of their support. In the Pareto
132 case with parametrization (B.19), however, if we were interested in small signals of
133 size $p^{r/\alpha}$, where $r < 1$, then the “boundary” says that we will need $n > p^{2(1-r)/\alpha}$
134 samples, which is exponential in the dimension p and quickly diverges. Recall that
135 the “boundary” is really an asymptotic result in p . Such an approximation in finite
136 dimensions becomes invalid.

0

Bibliography

2

- 3 Adler, R.J., Taylor, J.E.: Random Fields and Geometry. Springer Science & Business Media, Berlin
4 (2009)
- 5 Agresti, A.: An Introduction to Categorical Data Analysis. Wiley, Hoboken (2018)
- 6 Anderson, T.W., Darling, D.A.: Asymptotic theory of certain “goodness of fit” criteria based on
7 stochastic processes. Ann. Math. Stat. 193–212 (1952)
- 8 Arias-Castro, E., Chen, S.: Distribution-free multiple testing. Electron. J. Stat. **11**(1), 1983–2001
9 (2017)
- 10 Arias-Castro, E., Wang, M.: Distribution-free tests for sparse heterogeneous mixtures. Test **26**(1),
11 71–94 (2017)
- 12 Barber, R.F., Candès, E.J.: Controlling the false discovery rate via knockoffs. Ann. Stat. **43**(5),
13 2055–2085 (2015)
- 14 Barndorff-Nielsen, O.: On the limit behaviour of extreme order statistics. Ann. Math. Stat. **34**(3),
15 992–1002 (1963)
- 16 Benjamini, Y., Hochberg, Y.: Controlling the false discovery rate: a practical and powerful approach
17 to multiple testing. J. R. Stat. Society. Ser. B (Methodol.) 289–300 (1995)
- 18 Benjamini, Y., Yekutieli, D.: The control of the false discovery rate in multiple testing under depen-
19 dency. Ann. Stat. 1165–1188 (2001)
- 20 Berman, S.M.: Limit theorems for the maximum term in stationary sequences. Ann. Math. Stat.
21 502–516 (1964)
- 22 Bingham, N.H., Goldie, C.M., Teugels, J.L.: Regular Variation. Cambridge University Press, Cam-
23 bridge (1987)
- 24 Bogdan, M., Chakrabarti, A., Frommlet, F., Ghosh, J.K.: Asymptotic Bayes-optimality under spar-
25 sity of some multiple testing procedures. Ann. Stat. **39**(3), 1551–1579 (2011)
- 26 Boucheron, S., Lugosi, G., Massart, P.: Concentration Inequalities. A Nonasymptotic Theory of
27 Independence. With a foreword by Michel Ledoux. Oxford University Press, Oxford (2013)
- 28 Bush, W.S., Moore, J.H.: Genome-wide association studies. PLoS Comput. Biol. **8**(12), e1002822
29 (2012)
- 30 Butucea, C., Ndaoud, M., Stepanova, N.A., Tsybakov, A.B.: Variable selection with Hamming loss.
31 Ann. Stat. **46**(5), 1837–1875 (2018)
- 32 Cai, T.T., Jeng, X.J., Jin, J.: Optimal detection of heterogeneous and heteroscedastic mixtures. J. R.
33 Stat. Soc.: Ser. B (Stat. Methodol.) **73**(5), 629–662 (2011)

- 34 Cai, T.T., Jin, J., Low, M.G.: Estimation and confidence sets for sparse normal mixtures. *Ann. Stat.*
 35 **35**(6), 2421–2449 (2007)
- 36 Cai, T.T., Wu, Y.: Optimal detection of sparse mixtures against a given null distribution. *IEEE Trans.*
 37 *Inf. Theory* **60**(4), 2217–2232 (2014)
- 38 Candès, E.J.: Lecture 3: global testing, chi-square test, optimality of chi-square test for distributed
 39 mild effects. In: Stats 300C: Theory of Statistics (Spring 2018). Stanford Lecture Notes (2018).
 40 <https://statweb.stanford.edu/~candes/teaching/stats300c/index.html>
- 41 Chatterjee, S.: Superconcentration and Related Topics, vol. 15. Springer, Berlin (2014)
- 42 Comminges, L., Dalalyan, A.S.: Tight conditions for consistency of variable selection in the context
 43 of high dimensionality. *Ann. Stat.* **40**(5), 2667–2696 (2012)
- 44 Conlon, D., Fox, J., Sudakov, B.: Recent developments in graph Ramsey theory. In: Surveys in
 45 Combinatorics 2015. London Mathematical Society Lecture Note Series, vol. 424, pp. 49–118.
 46 Cambridge University Press, Cambridge (2015)
- 47 Cramér, H.: On the composition of elementary errors: first paper: mathematical deductions. *Scand.*
 48 *Actuar. J.* **1928**(1), 13–74 (1928)
- 49 De Haan, L., Ferreira, A.: Extreme Value Theory: An Introduction. Springer Science & Business
 50 Media, Berlin (2007)
- 51 Dedecker, J., Doukhan, P., Lang, G., León, J.R., Louhichi, S., Prieur, C.: Weak Dependence: With
 52 Examples and Applications. Lecture Notes in Statistics, vol. 190. Springer, New York (2007)
- 53 Domingos, P., Pazzani, M.: On the optimality of the simple Bayesian classifier under zero-one loss.
 54 *Mach. Learn.* **29**(2–3), 103–130 (1997)
- 55 Donoghue, W.F.: Distributions and Fourier Transforms, vol. 32. Academic, Cambridge (2014)
- 56 Donoho, D., Jin, J.: Higher criticism for detecting sparse heterogeneous mixtures. *Ann. Stat.* **32**(3),
 57 962–994 (2004)
- 58 Donoho, D., Jin, J.: Special invited paper: higher criticism for large-scale inference, especially for
 59 rare and weak effects. *Stat. Sci.* 1–25 (2015)
- 60 Dudoit, S., Shaffer, J.P., Boldrick, J.C.: Multiple hypothesis testing in microarray experiments. *Stat.*
 61 *Sci.* 71–103 (2003)
- 62 Dunn, O.J.: Multiple comparisons among means. *J. Am. Stat. Assoc.* **56**(293), 52–64 (1961)
- 63 Efron, B.: Large-scale simultaneous hypothesis testing: the choice of a null hypothesis. *J. Am. Stat.*
 64 *Assoc.* **99**(465), 96–104 (2004)
- 65 Eicker, F.: The asymptotic distribution of the suprema of the standardized empirical processes. *Ann.*
 66 *Stat.* 116–138 (1979)
- 67 Embrechts, P., Klüppelberg, C., Mikosch, T.: Modelling Extremal Events: For Insurance and
 68 Finance, vol. 33. Springer Science & Business Media, Berlin (2013)
- 69 Erdős, P., Szekeres, G.: A combinatorial problem in geometry. *Compos. Math.* **2**, 463–470 (1935)
- 70 Fan, J.: Test of significance based on wavelet thresholding and Neyman's truncation. *J. Am. Stat.*
 71 *Assoc.* **91**(434), 674–688 (1996)
- 72 Ferguson, T.S.: A Course in Large Sample Theory. Routledge, Milton Park (2017)
- 73 Feynman, R.: The Character of Physical Law. MIT Press, Cambridge (2017)
- 74 Fox, J.: Lecture 5: Ramsey Theory. In: MAT 307: Combinatorics (Spring 2009). MIT Lecture Notes
 75 (2009). <http://math.mit.edu/~fox/MAT307.html>
- 76 Gao, Z. (2019). U-PASS: <https://power.stat.lsa.umich.edu/u-pass/>. An R Shiny App accompanying the paper “U-PASS: unified power analysis and forensics for qualitative traits in genetic association studies”
- 77 Gao, Z., Stoev, S.: Fundamental limits of exact support recovery in high dimensions. *Bernoulli*
 78 **26**(4), 2605–2638 (2020)
- 79 Gao, Z., Terhorst, J., Van Hout, C.V., Stoev, S.: U-PASS: unified power analysis and forensics for
 80 qualitative traits in genetic association studies. *Bioinformatics* **36**(3), 974–975 (2019)
- 81 Genovese, C., Wasserman, L.: Operating characteristics and extensions of the false discovery rate
 82 procedure. *J. R. Stat. Soc.: Ser. B (Stat. Methodol.)* **64**(3), 499–517 (2002)
- 83 Genovese, C.R., Jin, J., Wasserman, L., Yao, Z.: A comparison of the lasso and marginal regression.
 84 *J. Mach. Learn. Res.* **13**, 2107–2143 (2012)

- Gnedenko, B.: Sur la distribution limite du terme maximum d'une série aléatoire. *Ann. Math.* **42**–453 (1943)
- Hall, P., Jin, J.: Innovated higher criticism for detecting sparse signals in correlated noise. *Ann. Stat.* **38**(3), 1686–1732 (2010)
- He, Y., Xu, G., Wu, C., Pan, W.: Asymptotically independent U-statistics in high-dimensional testing (2018). arXiv preprint [arXiv:1809.00411](https://arxiv.org/abs/1809.00411)
- Hochberg, Y.: A sharper Bonferroni procedure for multiple tests of significance. *Biometrika* **75**(4), 800–802 (1988)
- Holm, S.: A simple sequentially rejective multiple test procedure. *Scand. J. Stat.* 65–70 (1979)
- Hsing, T.: A note on the asymptotic independence of the sum and maximum of strongly mixing stationary random variables. *Ann. Probab.* 938–947 (1995)
- Inglot, T.: Inequalities for quantiles of the chi-square distribution. *Probab. Math. Stat.* **30**(4), 339–351 (2010)
- Ingster, Y.I.: Minimax detection of a signal for ℓ_n^p -balls. *Math. Methods Stat.* **7**(4), 401–428 (1998)
- Ji, P., Jin, J.: UPS delivers optimal phase diagram in high-dimensional variable selection. *Ann. Stat.* **40**(1), 73–103 (2012)
- Jin, J., Zhang, C.-H., Zhang, Q.: Optimality of graphlet screening in high dimensional variable selection. *J. Mach. Learn. Res.* **15**(1), 2723–2772 (2014)
- Kallitsis, M., Stoev, S.A., Bhattacharya, S., Michailidis, G.: AMON: an open source architecture for online monitoring, statistical analysis, and forensics of multi-gigabit streams. *IEEE J. Sel. Areas Commun.* **34**(6), 1834–1848 (2016)
- Kartsioskas, R., Gao, Z., Stoev, S.: On the rate of concentration of maxima in Gaussian arrays (2019). arXiv preprint [arXiv:1910.04259](https://arxiv.org/abs/1910.04259)
- Klass, M.J.: The minimal growth rate of partial maxima. *Ann. Probab.* **12**(2), 380–389 (1984)
- Leadbetter, M.R., Lindgren, G., Rootzén, H.: Extremes and Related Properties of Random Sequences and Processes. Springer Series in Statistics. Springer, Berlin (1983)
- Li, X., Fithian, W.: Optimality of the max test for detecting sparse signals with Gaussian or heavier tail (2020). arXiv:2006.12489
- MacArthur, J., Bowler, E., Cerezo, M., Gil, L., Hall, P., Hastings, E., Junkins, H., McMahon, A., Milano, A., Morales, J.: The new NHGRI-EBI catalog of published genome-wide association studies (GWAS catalog). *Nucl. Acids Res.* **45**(D1), D896–D901 (2016)
- Mccormick, W., Mittal, Y.: On weak convergence of the maximum. Department of Statistics, Stanford University (1976)
- Meinshausen, N., Bühlmann, P.: High-dimensional graphs and variable selection with the lasso. *Ann. Stat.* **34**(3), 1436–1462 (2006)
- Michailidou, K., Lindström, S., Dennis, J., Beesley, J., Hui, S., Kar, S., Lemaçon, A., Soucy, P., Glubb, D., Rostamianfar, A.: Association analysis identifies 65 new breast cancer risk loci. *Nature* **551**(7678), 92 (2017)
- Naveau, P.: Almost sure relative stability of the maximum of a stationary sequence. *Adv. Appl. Probab.* **35**(3), 721–736 (2003)
- Neuvial, P., Roquain, E.: On false discovery rate thresholding for classification under sparsity. *Ann. Stat.* **40**(5), 2572–2600 (2012)
- Nichols, T., Hayasaka, S.: Controlling the familywise error rate in functional neuroimaging: a comparative review. *Stat. Methods Med. Res.* **12**(5), 419–446 (2003)
- Pipiras, V., Taqqu, M.S.: Long-Range Dependence and Self-similarity. Cambridge Series in Statistical and Probabilistic Mathematics, [45]. Cambridge University Press, Cambridge (2017)
- Ramsey, F.P.: On a problem of formal logic. In: Classic Papers in Combinatorics, pp. 1–24. Springer, Berlin (2009)
- Resnick, S.I.: Extreme Values, Regular Variation and Point Processes. Springer, Berlin (2013)
- Resnick, S.I.: A probability path. Modern Birkhäuser Classics. Birkhäuser/Springer, New York (2014). Reprint of the fifth (2005) printing of the 1999 original [MR1664717]
- Resnick, S.I., Tomkins, R.: Almost sure stability of maxima. *J. Appl. Probab.* **10**(2), 387–401 (1973)
- Robbins, H.: A remark on Stirling's formula. *Am. Math. Mon.* **62**(1), 26–29 (1955)

- 140 Sah, A.: Diagonal Ramsey via effective quasirandomness (2020)
- 141 Šidák, Z.: Rectangular confidence regions for the means of multivariate normal distributions. *J.
Am. Stat. Assoc.* **62**(318), 626–633 (1967)
- 142 Skorokhod, A.V.: Limit theorems for stochastic processes. *Theory Probab. Its Appl.* **1**(3), 261–290
(1956)
- 143 Slepian, D.: The one-sided barrier problem for Gaussian noise. *Bell Labs Tech. J.* **41**(2), 463–501
(1962)
- 144 Smirnov, N.: Table for estimating the goodness of fit of empirical distributions. *Ann. Math. Stat.*
19(2), 279–281 (1948)
- 145 Storey, J.D.: The optimal discovery procedure: a new approach to simultaneous significance testing.
J. R. Stat. Soc.: Ser. B (Stat. Methodol.) **69**(3), 347–368 (2007)
- 146 Sun, W., Cai, T.T.: Oracle and adaptive compound decision rules for false discovery rate control. *J.
Am. Stat. Assoc.* **102**(479), 901–912 (2007)
- 147 Tanguy, K.: Some superconcentration inequalities for extrema of stationary Gaussian processes.
Stat. Probab. Lett. **106**, 239–246 (2015a)
- 148 Tanguy, K.: Some superconcentration inequalities for extrema of stationary Gaussian processes.
Stat. Probab. Lett. **106**, 239–246 (2015b)
- 149 Taqqu, M.S.: Fractional Brownian motion and long-range dependence. In: Doukhan, P., Oppen-
heim, G., Taqqu, M.S. (eds.) *Theory and Applications of Long-Range Dependence*, pp. 5–38.
Birkhäuser, Basel (2003)
- 150 Tsagris, M., Beneki, C., Hassani, H.: On the folded normal distribution. *Mathematics* **2**(1), 12–28
(2014)
- 151 Tukey, J.W.: T13N: the higher criticism. In: *Statistics 411*. Princeton University Lecture Notes
(1976)
- 152 Wainwright, M.J.: Information-theoretic limits on sparsity recovery in the high-dimensional and
noisy setting. *IEEE Trans. Inf. Theory* **55**(12), 5728–5741 (2009a)
- 153 Wainwright, M.J.: Sharp thresholds for high-dimensional and noisy sparsity recovery using ℓ_1 -
constrained quadratic programming (Lasso). *IEEE Trans. Inf. Theory* **55**(5), 2183–2202 (2009b)
- 154 Wainwright, M.J.: *High-Dimensional Statistics: A Non-asymptotic Viewpoint*. Cambridge Univer-
sity Press, Cambridge (2019)
- 155 Wasserman, L., Roeder, K.: High dimensional variable selection. *Ann. Stat.* **37**(5A), 2178 (2009)
- 156 Wu, C., Xu, G., Pan, W.: An adaptive test on high-dimensional parameters in generalized linear
models. *Stat. Sin.* **29**, 2163–2186 (2019)
- 157 Xu, G., Lin, L., Wei, P., Pan, W.: An adaptive two-sample test for high-dimensional means.
Biometrika **103**(3), 609–624 (2016)
- 158 Zhang, J.: Powerful goodness-of-fit tests based on the likelihood ratio. *J. R. Stat. Soc.: Ser. B (Stat.
Methodol.)* **64**(2), 281–294 (2002)
- 159 Zhao, P., Yu, B.: On model selection consistency of Lasso. *J. Mach. Learn. Res.* **7**, 2541–2563
(2006)
- 160 Zhong, P.-S., Chen, S.X., Xu, M.: Tests alternative to higher criticism for high-dimensional means
under sparsity and column-wise dependence. *Ann. Stat.* **41**(6), 2820–2851 (2013)

MARKED PROOF

Please correct and return this set

Please use the proof correction marks shown below for all alterations and corrections. If you wish to return your proof by fax you should ensure that all amendments are written clearly in dark ink and are made well within the page margins.

Instruction to printer	Textual mark	Marginal mark
Leave unchanged	... under matter to remain	✓
Insert in text the matter indicated in the margin	↪	New matter followed by ↪ or ↪ ^②
Delete	/ through single character, rule or underline ———— through all characters to be deleted	⊖ or ⊖ ^②
Substitute character or substitute part of one or more word(s)	/ through letter or ———— through characters	new character / or new characters /
Change to italics	— under matter to be changed	←
Change to capitals	≡ under matter to be changed	≡
Change to small capitals	= under matter to be changed	=
Change to bold type	~ under matter to be changed	~
Change to bold italic	~~ under matter to be changed	~~
Change to lower case	Encircle matter to be changed	≠
Change italic to upright type	(As above)	+
Change bold to non-bold type	(As above)	—
Insert 'superior' character	/ through character or ↪ where required	Y or X under character e.g. ³ or ³
Insert 'inferior' character	(As above)	↪ over character e.g. ²
Insert full stop	(As above)	◎
Insert comma	(As above)	,
Insert single quotation marks	(As above)	‘ or ’ and/or ‘ or ’
Insert double quotation marks	(As above)	“ or ” and/or “ or ”
Insert hyphen	(As above)	—
Start new paragraph	—	—
No new paragraph	≈	≈
Transpose	□	□
Close up	linking () characters	()
Insert or substitute space between characters or words	/ through character or ↪ where required	Y
Reduce space between characters or words	between characters or words affected	T